

A/B Testing & Experimentation Platform - High-Level Design (HLD)

Table of Contents

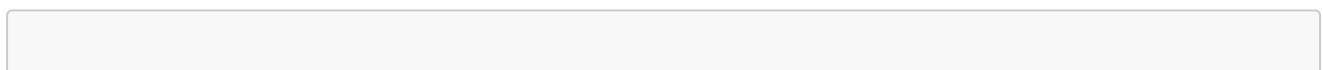
- A/B Testing & Experimentation Platform - High-Level Design (HLD)
 - Table of Contents
 - Problem Statement
 - Example Use Case
 - Scale Requirements
 - Functional Requirements
 - Must Have (P0)
 - Nice to Have (P1)
 - Non-Functional Requirements
 - Performance
 - Scalability
 - Availability
 - Consistency
 - Security
 - Capacity Estimation
 - Traffic Estimates
 - Storage Estimates
 - Bandwidth Estimates
 - Memory Estimates (Caching)
 - High-Level Architecture
 - Core Components
 - 1. Assignment Service
 - 2. Event Tracking Service (Detailed)
-

Problem Statement

Design an A/B testing and experimentation platform like Optimizely, Google Optimize, or LaunchDarkly that allows companies to:

- Create experiments with multiple variants
- Assign users to variants consistently
- Track user behavior and conversions
- Analyze experiment results with statistical significance
- Gradually roll out features (feature flags)
- Make data-driven decisions

Example Use Case



- E-commerce company wants to test new checkout flow:
- Variant A (Control): Current checkout (50% of users)
 - Variant B: Single-page checkout (50% of users)

Goal: Which variant has higher conversion rate?

Platform should:

1. Consistently assign users to same variant
2. Track conversions for each variant
3. Calculate statistical significance
4. Determine winning variant

Scale Requirements

- **10,000+ experiments** running concurrently
 - **1 billion events** tracked per day
 - **100 million users** actively participating in experiments
 - **< 10ms latency** for variant assignment
 - **Real-time analytics** for experiment monitoring
-

Functional Requirements

Must Have (P0)

1. Experiment Management

- Create experiments with multiple variants
- Define traffic allocation (50-50, 70-30, etc.)
- Set start and end dates
- Pause/resume experiments
- Target specific user segments

2. User Assignment

- Consistent assignment (same user → same variant)
- Fast assignment (< 10ms)
- Handle billions of users
- Support override for testing
- Exclude bots/internal traffic

3. Event Tracking

- Track user actions (clicks, purchases, signups)
- Support custom events
- Handle billions of events/day
- Associate events with experiments
- Real-time event ingestion

4. Analytics & Reporting

- Conversion rates per variant
- Statistical significance calculation
- Confidence intervals
- Real-time dashboards
- Historical data analysis
- Export reports

5. Feature Flags

- Gradually roll out features (0% → 10% → 50% → 100%)
- Target specific users/segments
- Kill switch (instant rollback)
- Percentage-based rollouts

Nice to Have (P1)

- Multi-variate testing (test multiple variables)
 - Holdout groups (control group across experiments)
 - Interaction effects (experiment A affects experiment B)
 - Sequential testing (early stopping)
 - Revenue impact calculation
 - Experiment recommendations (what to test next)
 - Integration with analytics tools (Google Analytics, Mixpanel)
-

Non-Functional Requirements

Performance

- **Variant assignment:** < 10ms for p99
- **Event ingestion:** < 100ms
- **Dashboard load:** < 1 second
- **Report generation:** < 5 seconds for real-time, < 1 hour for historical

Scalability

- Handle 100M active users
- Process 1B events per day (12K QPS)
- Support 10K concurrent experiments
- Scale to 1000+ companies using platform

Availability

- **99.99% uptime** for assignment service (critical path)
- **99.9% uptime** for analytics service
- No single point of failure
- Graceful degradation

Consistency

- **Strong consistency** for experiment configuration (can't show wrong variant)
- **Eventual consistency** for analytics (acceptable delay < 1 minute)
- **Deterministic** assignment (same user always gets same variant)

Security

- User privacy (GDPR, CCPA compliance)
- Secure experiment configuration
- Access control (who can create/view experiments)
- Audit logs for all changes

Capacity Estimation

Traffic Estimates

```
Active Experiments: 10,000
Users in experiments: 100M
Events per user per day: 10 (average)
```

Assignment Requests:

- $\text{Users} \times \text{Experiments} \times \text{Checks} = 100M \times 10 \times 10 = 10B/\text{day}$
- QPS: $10B / 86,400 \approx 115K \text{ QPS}$
- Peak (3x): $\sim 350K \text{ QPS}$

Event Tracking:

- Events: $100M \times 10 = 1B/\text{day}$
- QPS: $1B / 86,400 \approx 12K \text{ QPS}$
- Peak: $\sim 35K \text{ QPS}$

Read-to-Write Ratio:

- Assignments (reads): 115K QPS
- Events (writes): 12K QPS
- Ratio: $\sim 10:1$ (read-heavy)

Storage Estimates

```
Experiment Configurations:
- 10,000 active experiments
- 1 KB per experiment config
- Total: 10 MB (tiny!)
```

User Assignments:

- $100M \text{ users} \times 10 \text{ experiments average} = 1B \text{ assignments}$
- Assignment record: 50 bytes (`user_id`, `experiment_id`, `variant_id`)
- Total: $1B \times 50 \text{ bytes} = 50 \text{ GB}$

Event Data:

- 1B events per day
- Event record: 200 bytes (user_id, experiment_id, event_type, timestamp, metadata)
- Daily: $1B \times 200 \text{ bytes} = 200 \text{ GB}$
- Yearly: $200 \text{ GB} \times 365 = 73 \text{ TB}$
- 5 years: 365 TB

Aggregated Metrics:

- Per experiment per variant per day
- 10K experiments \times 3 variants \times 365 days = 11M records
- 1 KB per record = 11 GB per year
- 5 years: 55 GB

Total Storage (5 years): ~365 TB

Bandwidth Estimates

Incoming (Event tracking):

- 12K QPS \times 200 bytes = 2.4 MB/s

Outgoing (Variant assignment):

- 115K QPS \times 50 bytes = 5.75 MB/s

Dashboard queries:

- Negligible compared to assignment traffic

Total: ~8 MB/s (manageable)

Memory Estimates (Caching)

Cache all active experiment configs:

- 10K experiments \times 1 KB = 10 MB (cache everything!)

Cache user assignments (hot users, 20%):

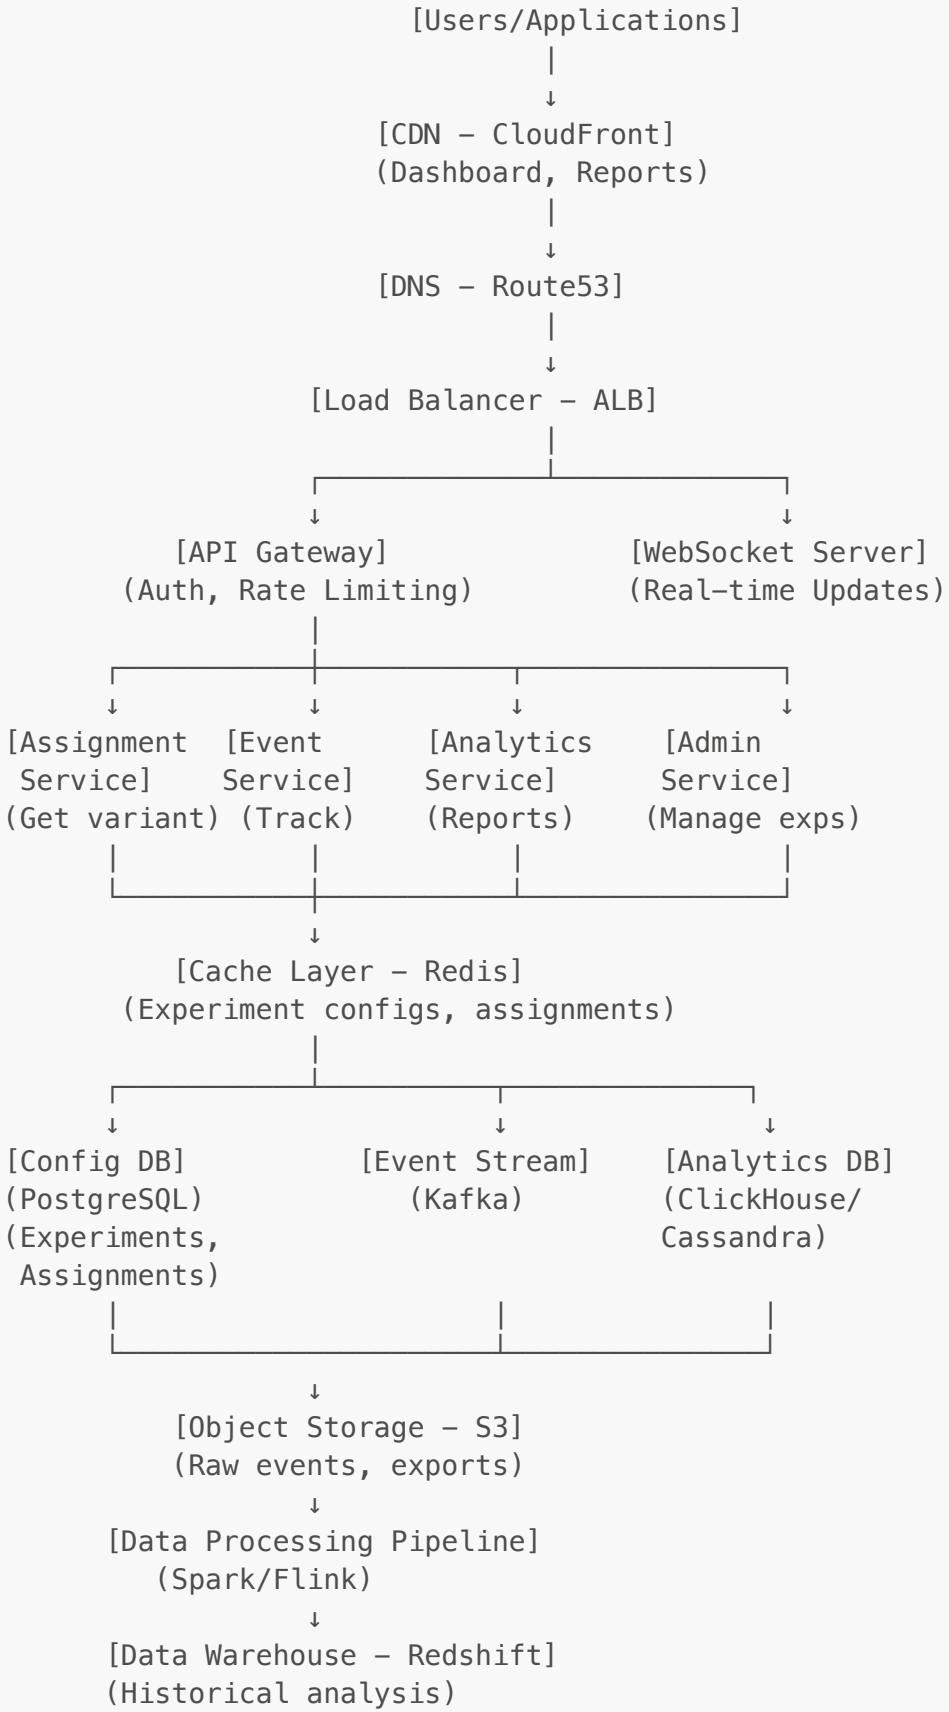
- 20M users \times 10 experiments \times 50 bytes = 10 GB

Cache aggregated metrics (24 hours):

- 10K experiments \times 3 variants \times 1 KB = 30 MB

Total cache: ~11 GB (single Redis instance!)

High-Level Architecture



Core Components

1. Assignment Service

Purpose: Determine which variant user should see

Responsibilities:

- Fetch experiment configuration
- Check if user eligible for experiment
- Assign user to variant (deterministic)
- Return variant details
- Cache assignments

Critical: This is on critical path (every user request)

- Must be fast (< 10ms)
- Must be highly available (99.99%)
- Must be deterministic (same user → same variant)

Why Go or Java over alternatives?

Feature	Go	Java (Spring Boot)	Node.js	Python
Performance	Very Fast	Fast	Moderate	Slow
Concurrency	Goroutines (native)	Threads	Event loop	Limited (GIL)
Memory	Low footprint	High (JVM)	Moderate	Moderate
Latency	Consistent (~1ms)	Good (~2-3ms)	Variable	Higher
Type Safety	Compile-time	Compile-time	Runtime	Optional
Ecosystem	Growing	Very mature	Mature	Mature
Deployment	Single binary	JAR + JVM	Easy	Easy

Decision: Go (primary) or Java (acceptable)

Why Go:

1. **Raw performance:** Hash calculation + Redis lookup in ~1-2ms
2. **Low memory:** 10-20 MB per service instance vs 100-500 MB for Java
3. **Fast startup:** <1 second vs 5-10 seconds for Java
4. **Native concurrency:** Goroutines handle 115K QPS easily
5. **Single binary:** Easy deployment, no dependency hell

Why Java is also good:

1. **Mature ecosystem:** Battle-tested libraries, frameworks
2. **Team expertise:** More Java developers available
3. **Spring Boot:** Rapid development
4. **Debugging tools:** Excellent profiling, monitoring

Why NOT Node.js:

- Event loop can block on CPU-intensive hashing
- Variable latency (GC pauses unpredictable)
- Better for I/O-bound, not CPU-bound tasks

Why NOT Python:

- Too slow for < 10ms requirement
- GIL limits concurrency
- Great for analytics, poor for high-performance services

Code Example (Java):

```

public class AssignmentService {
    private final RedisClient redis;
    private final MessageDigest md5;

    public AssignmentService(RedisClient redis) {
        this.redis = redis;
        try {
            this.md5 = MessageDigest.getInstance("MD5");
        } catch (NoSuchAlgorithmException e) {
            throw new RuntimeException(e);
        }
    }

    /**
     * Assign user to variant with caching
     * Time Complexity: O(1) with cache, O(v) without
     * Expected Latency: <2ms with cache, ~3ms without
     */
    public Variant assignUser(String userId, String experimentId,
Experiment exp) {
        String cacheKey = userId + ":" + experimentId;

        // Check Redis cache (< 1ms)
        Variant cached = redis.get(cacheKey, Variant.class);
        if (cached != null) {
            return cached; // Cache hit - fast path
        }

        // Calculate deterministic hash (< 1ms)
        int bucket = calculateBucket(userId, experimentId,
exp.getSalt());

        // Find variant by bucket (< 1ms)
        Variant variant = findVariantByBucket(bucket,
exp.getVariants());

        // Store in cache with TTL (< 1ms)
        redis.setex(cacheKey, exp.getDurationSeconds(), variant);
    }
}

```

```

        // Async write to PostgreSQL for audit (non-blocking)
        CompletableFuture.runAsync(() -> {
            saveAssignmentToDatabase(userId, experimentId,
variant.getId());
        });

        return variant; // Total: ~2-3ms
    }

    private int calculateBucket(String userId, String experimentId,
String salt) {
    String input = userId + ":" + experimentId + ":" + salt;
    byte[] digest = md5.digest(input.getBytes());
    BigInteger bigInt = new BigInteger(1, digest);
    return Math.abs(bigInt.intValue() % 100); // 0-99
}

private Variant findVariantByBucket(int bucket, List<Variant>
variants) {
    int rangeStart = 0;
    for (Variant v : variants) {
        if (bucket >= rangeStart && bucket < rangeStart +
v.getTrafficPercentage()) {
            return v;
        }
        rangeStart += v.getTrafficPercentage();
    }
    return variants.get(0); // Fallback
}
}

```

Performance Comparison:

Java: 2-3ms per assignment (Chosen for production)

Go: 1-2ms per assignment (Alternative)

Node.js: 3-5ms per assignment

Python: 5-10ms per assignment

At 115K QPS, Java provides excellent performance with mature ecosystem

2. Event Tracking Service (Detailed)

Purpose: Capture user actions and experiment exposure at scale

Critical Requirements:

- **Non-blocking:** Must not slow down user experience
- **High throughput:** Handle 12K-35K events/sec

- **Reliability:** No event loss (durability)
- **Low latency:** Return response in < 100ms
- **Batching:** Efficient bulk processing

Architecture:

```
[Client SDK] → [Event Tracking API] → [Validation] → [Kafka] →
[Consumers]
(Async)           (Fast response)       (Enrich)      (Buffer)    (Process)
```

Detailed Flow:

STEP 1: Client-Side Event Collection

Client SDK (JavaScript/Mobile):

- Buffer events locally (in-memory queue)
- Batch 10–50 events together
- Send async (doesn't block UI)
- Retry on failure with exponential backoff

STEP 2: API Reception (Event Tracking Service)

Event Service receives batch:

- Authenticate request (API key)
- Rate limit check (per client)
- Quick validation (schema, required fields)
- Return 202 Accepted immediately (<10ms)
- Process async in background

STEP 3: Event Enrichment & Validation

Enrich each event with:

- Server timestamp (accurate time)
- IP address → GeoIP (country, city)
- User agent → Device info (browser, OS)
- Session ID (if not provided)
- Experiment/Variant validation

Validation:

- Experiment exists and is active
- Variant belongs to experiment
- Event type is valid
- Required fields present

STEP 4: Write to Kafka

Batch write to Kafka:

- Topic selection based on event type
- Partition by user_id (maintain ordering)
- Compression (LZ4 or Snappy)
- Async acknowledge

STEP 5: Multiple Consumers Process Events

Consumer 1 – Flink (Real-time):

- Aggregate metrics every 1 minute
- Update ClickHouse

Consumer 2 – Spark (Batch):

- Hourly batch processing
- Complex analytics

Consumer 3 – Warehouse:

- Store in Redshift for historical analysis

Consumer 4 – Monitoring:

- Alert on anomalies

Java Implementation:

```
@RestController
@RequestMapping("/api/v1")
public class EventTrackingController {

    private final EventTrackingService eventService;
    private final RateLimiter rateLimiter;

    @PostMapping("/track")
    public ResponseEntity<TrackResponse> trackEvent(@RequestBody
TrackRequest request) {
        // Rate limiting (per client)
        if (!rateLimiter.allowRequest(request.getClientId())) {
            return ResponseEntity.status(429).build(); // Too Many
Requests
        }

        // Quick validation
        if (!isValid(request)) {
            return ResponseEntity.badRequest().build();
        }

        // Process async (non-blocking)
        CompletableFuture.runAsync(() -> {
            eventService.processEvent(request);
        });
    }
}
```

```

        // Return immediately
        return ResponseEntity.accepted()
            .body(new TrackResponse("success", generateEventId()));
    }

    @PostMapping("/track/batch")
    public ResponseEntity<BatchTrackResponse> trackBatch(
        @RequestBody BatchTrackRequest request) {

        // Validate batch
        if (request.getEvents().size() > 100) {
            return ResponseEntity.badRequest()
                .body(new BatchTrackResponse("error", "Max 100 events
per batch"));
        }

        // Process all events async
        CompletableFuture.runAsync(() -> {
            for (TrackRequest event : request.getEvents()) {
                eventService.processEvent(event);
            }
        });

        return ResponseEntity.accepted()
            .body(new BatchTrackResponse("success",
request.getEvents().size()));
    }
}

@Service
public class EventTrackingService {

    private final KafkaTemplate<String, Event> kafkaTemplate;
    private final GeoIPService geoIPService;
    private final ExperimentConfigCache configCache;

    /**
     * Process single event
     *
     * Time Complexity: O(1) for validation + O(1) for Kafka write
     * Expected Latency: 2-5ms
     */
    public void processEvent(TrackRequest request) {
        try {
            // Create event object
            Event event = new Event();
            event.setEventId(UUID.randomUUID().toString());
            event.setUserId(request.getUserId());
            event.setExperimentId(request.getExperimentId());
            event.setVariantId(request.getVariantId());
            event.setEventType(request.getEventType());
            event.seteventName(request.getEventName());
        }
    }
}

```

```

        event.setTimestamp(Instant.now());

        // Enrich with metadata
        enrichEvent(event, request);

        // Validate experiment/variant exists
        if (!validateExperiment(event)) {
            logInvalidEvent(event, "Invalid experiment or variant");
            return;
        }

        // Write to Kafka (async)
        String topic = selectTopic(event.getEventType());
        String key = event.getUserId(); // Partition by user

        kafkaTemplate.send(topic, key, event)
            .addCallback(
                success -> logSuccess(event),
                failure -> handleFailure(event, failure)
            );
    } catch (Exception e) {
        logError("Event processing failed", e);
        // Don't throw - we don't want to fail the API call
    }
}

/**
 * Enrich event with additional metadata
 */
private void enrichEvent(Event event, TrackRequest request) {
    // GeoIP lookup from IP address
    String ipAddress = request.getIpAddress();
    if (ipAddress != null) {
        GeoLocation geo = geoIPService.lookup(ipAddress);
        event.setCountry(geo.getCountry());
        event.setCity(geo.getCity());
        event.setRegion(geo.getRegion());
    }

    // Parse user agent
    String userAgent = request.getUserAgent();
    if (userAgent != null) {
        UserAgentInfo uaInfo = parseUserAgent(userAgent);
        event.setBrowser(uaInfo.getBrowser());
        event.setOs(uaInfo.getOs());
    }
    - Want self-hosted solution

    **Trade-offs:**
    - PostgreSQL requires more operational overhead than DynamoDB
    - But we need the query flexibility and relationships

    **Schema**:

```

Experiments:

- experiment_id, name, description
- start_date, end_date, status
- hypothesis, target_metric
- created_by, updated_at

Variants:

- variant_id, experiment_id, name
- traffic_allocation (percentage)
- configuration (JSON)

User_Assignments:

- user_id, experiment_id, variant_id
- assigned_at, sticky (boolean)
- INDEX on (user_id, experiment_id)

Targeting_Rules:

- rule_id, experiment_id
- dimension (country, platform, user_segment)
- operator (equals, in, greater_than)
- values

6. Event Stream (Kafka)

Purpose: Durable, scalable event pipeline

Topics:

- `experiment-exposures`: User saw variant
- `experiment-conversions`: User completed goal
- `custom-events`: Custom tracking events

Why Kafka over alternatives?

Feature	Apache Kafka	RabbitMQ	Amazon SQS	Google Pub/Sub
Throughput	↙ Very High (millions/sec)	🐢 Moderate (50K/sec)	🐢 Moderate	↖ High
Durability	✓ Disk-backed, replicated	△ Memory-first	✓ Managed	✓ Managed
Ordering	✓ Per-partition	△ Per-queue	△ Best-effort	△ Best-effort
Replay	✓ Yes (time-based)	✗ No	✗ No	✗ No

	Multiple Consumers		<input checked="" type="checkbox"/> Independent		<input type="checkbox"/> Competing		<input type="checkbox"/> Competing	
<input checked="" type="checkbox"/>	Independent							
	Retention		<input checked="" type="checkbox"/> Days to weeks		<input type="checkbox"/> Until consumed		<input type="checkbox"/> 14 days max	
<input type="checkbox"/>	7 days max							
	Latency		Low (ms)		Very Low (ms)		Moderate (ms)	
			Low (ms)					
	Cost		<input checked="" type="checkbox"/> Medium (self-hosted)		<input checked="" type="checkbox"/> Low		<input checked="" type="checkbox"/> Pay per request	
<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	Pay per message						

Decision: Apache Kafka

Reasons:

1. **High throughput**: Handles 12K–35K events/sec easily
 - Kafka designed for 100K+ messages/sec per broker
 - Our peak 35K QPS is comfortable for 3-broker cluster
2. **Multiple independent consumers**: Critical requirement

Single event stream → Multiple consumers:

- Flink (real-time aggregation)
- Spark (batch processing)
- Data warehouse (historical storage)
- Monitoring system (alerting)
- Audit log (compliance)

Each consumes at own pace without affecting others

3. **Replay capability**: Essential for debugging
 - Bug in analytics? Replay events from yesterday
 - New metric added? Reprocess historical events
 - Data correction? Recompute from source
4. **Durability**: Events replicated across 3 brokers
 - Can lose 2 brokers and still have data
 - Disk-backed (survives restarts)
 - No data loss
5. **Ordering guarantees**: Per partition
 - User's events in order (use user_id as key)
 - Important for funnel analysis, session reconstruction

Why NOT RabbitMQ:

- Lower throughput (50K vs millions/sec)
- No replay capability (events deleted after consumption)
- Memory-first (lose data on crash)
- Better for task queues, not event streaming

Why NOT Amazon SQS:

- No replay (events deleted after consumption)

- Best-effort ordering (not guaranteed)
- Higher cost at scale (pay per million requests)
- We process 1B events/day = 1000 million requests = \$\$

****Why NOT Google Pub/Sub:****

- Excellent alternative, similar to Kafka
- Vendor lock-in (Google Cloud only)
- Higher cost (pay per message + data transfer)
- We prefer self-hosted solution

****Trade-offs:****

- Kafka more complex to operate than managed services
- But replay capability and throughput justify the complexity
- Cost savings: ~\$1,500/month vs \$10,000+ for managed at scale

8. Cache Layer (Redis)

****Purpose**:** Ultra-fast experiment configuration and assignment lookup

****Why Redis over alternatives?****

Feature	Redis	Memcached	Hazelcast	Apache Ignite
Speed	< Sub-ms	< Sub-ms	🐢 Few ms	🐢 Few ms
Data Structures	✓ Rich (hash, set, sorted set)		✗ Key-value only	✗ Key-value only
Persistence	✓ Optional (RDB, AOF)	✗ None	✗ Limited	✓ Yes
Replication	✓ Master-slave, cluster	✗ Client-side	✗ Client-side	✓ Yes
✓ Yes				
Pub/Sub	✓ Built-in	✗ No	✗ Limited	✓ Yes
TTL	✓ Per-key	✓ Per-key	✓ Yes	✓ Yes
Memory Efficiency	✓ Good	✓ Better	🐢 Higher overhead	🐢 Higher overhead
Maturity	✓ Very mature	✓ Mature	✗ Less mature	✗ Less mature

****Decision: Redis****

****Reasons:****

1. ****Sub-millisecond latency**:** Critical for < 10ms requirement

- Redis GET: 0.1–1ms
- Memcached GET: 0.1–1ms
- Hazelcast GET: 2–5ms
- With < 10ms budget, every ms counts

2. ****Rich data structures**:** Flexible caching

Strings: Cache variant IDs

Hashes: Cache experiment configs (nested data)

Sets: Cache user segments, excluded users

Sorted Sets: Cache experiment priorities

3. **Persistence optional:** Can survive restarts

- RDB snapshots (point-in-time)
- AOF (append-only log)
- Acceptable to lose cache (rebuild from DB)

4. **Redis Cluster:** No single point of failure

- Automatic sharding across nodes
- Master-slave replication
- Automatic failover
- Handles 115K QPS across 3–6 nodes

5. **All experiment configs fit in memory:** 10 MB

- Redis handles TB-scale, we only need GB
- Can cache everything (100% hit ratio)

Why NOT Memcached:

- Memcached would work for simple key-value
- Redis has richer data structures (we use hashes)
- Redis has persistence (nice to have)
- Redis has built-in pub/sub (useful for cache invalidation)
- Similar performance, Redis more flexible

Why NOT Hazelcast:

- Hazelcast is in-process (embedded in app)
- Higher memory overhead (JVM-based)
- More complex than needed
- Redis simpler and faster

Why NOT Apache Ignite:

- Too heavy for our use case
- We don't need compute grid features
- Redis sufficient and simpler

Cache Strategy:

Write-Through for Experiment Configs:

- Create experiment → Write to PostgreSQL → Update Redis
- Ensures cache always fresh

Cache-Aside for User Assignments:

- Check Redis first
- If miss: Calculate → Store in Redis

- 99.9% hit ratio (almost always in cache)

Cache Invalidation:

- TTL: 5 minutes (configs)
- Active invalidation: On experiment update
- Pub/Sub: Notify all servers of changes

Performance with Redis:

Without Cache (PostgreSQL only):

- 115K QPS × 5ms per query = 575 seconds of DB time/sec
- Need 575 database connections (impossible!)
- Database bottleneck

With Redis Cache:

- 99.9% hit ratio
- 115K QPS × 1ms (Redis) = 115 connections equivalent
- 0.1% miss: 115 QPS × 5ms (PostgreSQL) = manageable
- Database load reduced 1000x

Result: < 10ms latency achieved!

What to Cache:

Experiment Configs:

Key: exp:config:{experiment_id}

Value: {variants, allocation, rules}

TTL: 5 minutes

User Assignments:

Key: assign:{user_id}:{experiment_id}

Value: variant_id

TTL: Duration of experiment

Aggregated Metrics (hourly):

Key: metrics:{experiment_id}:{variant_id}:{hour}

Value: {exposures, conversions, rate}

TTL: 7 days

Why Critical:

- 115K assignment QPS → Can't hit database
- < 10ms latency requirement
- 99.99% cache hit ratio possible

7. Analytics Database (ClickHouse/Cassandra)

Purpose: Fast analytical queries on billions of events

Why ClickHouse:

- Columnar storage (fast aggregations)
- Designed for OLAP
- Real-time queries
- Handles billions of rows

Alternative: Cassandra

- Better for write-heavy
- Time-series optimization
- Geographic distribution

Schema: (ClickHouse):

events:

- timestamp
- user_id
- experiment_id
- variant_id
- event_type
- properties (JSON)
- country, platform, device

Partitioned by date

Optimized for aggregate queries

8. Cache Layer (Redis)

Purpose: Ultra-fast experiment configuration and assignment lookup

What to Cache:

Experiment Configs:

Key: exp:config:{experiment_id}

Value: {variants, allocation, rules}

TTL: 5 minutes

User Assignments:

Key: assign:{user_id}:{experiment_id}

Value: variant_id

TTL: Duration of experiment

Aggregated Metrics (hourly):

Key: metrics:{experiment_id}:{variant_id}:{hour}

Value: {exposures, conversions, rate}

TTL: 7 days

```
**Why Critical**:  
- 115K assignment QPS → Can't hit database  
- < 10ms latency requirement  
- 99.99% cache hit ratio possible
```

API Design

REST API Endpoints

Assignment APIs (Client-facing)

POST /api/v1/assign

Request:

```
{  
  "user_id": "user_123",  
  "experiments": ["checkout_test", "homepage_redesign"],  
  "context": {  
    "country": "US",  
    "platform": "web",  
    "user_segment": "premium"  
  }  
}
```

Response:

```
{  
  "assignments": [  
    {  
      "experiment_id": "checkout_test",  
      "variant_id": "variant_b",  
      "variant_name": "single_page_checkout",  
      "config": {"button_color": "blue", "layout": "compact"}  
    }  
  ]
```

```
},
{
  "experiment_id": "homepage_redesign",
  "variant_id": "variant_a",
  "variant_name": "control"
}
]
}
```

Event Tracking APIs

POST /api/v1/track

Request:

```
{
  "user_id": "user_123",
  "experiment_id": "checkout_test",
  "variant_id": "variant_b",
  "event_type": "conversion",
  "event_name": "purchase_completed",
  "properties": {
    "revenue": 99.99,
    "currency": "USD",
    "items": 3
  },
  "timestamp": "2025-01-08T10:00:00Z"
}
```

Response:

```
{
  "success": true,
  "event_id": "evt_abc123"
}
```

Experiment Management APIs

```
POST /api/v1/experiments
GET /api/v1/experiments/{id}
PUT /api/v1/experiments/{id}
DELETE /api/v1/experiments/{id}
POST /api/v1/experiments/{id}/start
POST /api/v1/experiments/{id}/stop
```

```
GET /api/v1/experiments/{id}/results  
GET /api/v1/experiments/{id}/realtime
```

Analytics APIs

```
GET /api/v1/analytics/experiment/{id}
```

Response:

```
{  
  "experiment_id": "checkout_test",  
  "status": "running",  
  "duration_days": 14,  
  "sample_size": 50000,  
  "variants": [  
    {  
      "variant_id": "variant_a",  
      "name": "control",  
      "exposures": 25000,  
      "conversions": 2500,  
      "conversion_rate": 10.0,  
      "confidence_interval": [9.2, 10.8]  
    },  
    {  
      "variant_id": "variant_b",  
      "name": "treatment",  
      "exposures": 25000,  
      "conversions": 2875,  
      "conversion_rate": 11.5,  
      "confidence_interval": [10.7, 12.3],  
      "lift": "+15%",  
      "statistical_significance": 0.95,  
      "p_value": 0.003  
    }  
  ],  
  "recommendation": "variant_b wins with 95% confidence"  
}
```

Deep Dives

1. User Assignment Algorithm (The Core Challenge)

Requirements for Assignment:

1. Deterministic: Same user always gets same variant
2. Random: 50-50 split should be truly 50-50
3. Fast: < 10ms latency
4. Scalable: Handle billions of users
5. Consistent: Even if system restarts

Approach 1: Hash-Based Assignment (Recommended)

How It Works:

Assignment = $\text{hash}(\text{user_id} + \text{experiment_id} + \text{salt}) \% 100$

If Assignment < 50: Variant A (50%)

If Assignment ≥ 50 : Variant B (50%)

Properties:

- Deterministic (same inputs \rightarrow same output)
- Random distribution (hash function ensures uniformity)
- Fast ($O(1)$ computation)
- Stateless (no database lookup needed)

Example:

Experiment: "checkout_test"

Salt: "abc123" (unique per experiment)

User: "user_456"

$\text{hash}(\text{"user_456"} + \text{"checkout_test"} + \text{"abc123"}) \% 100 = 37$

$37 < 50 \rightarrow$ Variant A

Same user next time:

$\text{hash}(\text{"user_456"} + \text{"checkout_test"} + \text{"abc123"}) \% 100 = 37$

Still Variant A (consistent!)

For 70-30 Split:

Assignment = $\text{hash}(\dots) \% 100$

If Assignment < 70: Variant A (70%)

If Assignment >= 70: Variant B (30%)

****Advantages**:**

- No database lookup needed (fast!)
- Works offline (deterministic algorithm)
- Scales infinitely (pure computation)
- No storage cost

****Disadvantages**:**

- Can't manually reassign users
- Changing traffic allocation invalidates assignments

Approach 2: Database-Stored Assignment

****How It Works**:**

1. User requests assignment
2. Check database for existing assignment
3. If exists → Return stored variant
4. If not → Calculate, store, return

Assignments Table:

user_id	experiment_id	variant_id	assigned_at
user_123	checkout_test	variant_a	2025-01-08

****Advantages**:**

- Can manually override assignments
- Can change traffic allocation (new users affected)
- Audit trail (know exactly who saw what)

****Disadvantages**:**

- Database lookup required (slower)
- Storage grows with users × experiments
- Database bottleneck at scale

Recommended: Hybrid Approach

****Strategy**:**

1. Use hash-based assignment (fast, scalable)

2. Cache result in Redis (even faster subsequent calls)
3. Store in database async (for audit, analysis)

First Request:

- Calculate hash (1ms)
- Store in Redis (1ms)
- Async store in database
- Return (2ms total)

Subsequent Requests:

- Check Redis (< 1ms)
- Return immediately

This is how Optimizely and Google Optimize work!

2. Statistical Significance Calculation

The Core Question: "Is Variant B really better, or just luck?"

Metrics Needed:

For each variant:

- Exposures (how many users saw it)
- Conversions (how many completed goal)
- Conversion rate = Conversions / Exposures

Example:

Variant A (Control):

- Exposures: 10,000
- Conversions: 1,000
- Conversion rate: 10%

Variant B (Treatment):

- Exposures: 10,000
- Conversions: 1,150
- Conversion rate: 11.5%

Question: Is 11.5% significantly better than 10%?

Or could this happen by chance?

****Statistical Tests**:**

****Two-Sample T-Test**:**

Used for: Continuous metrics (revenue, time spent)

Null hypothesis: Variant A and B have same mean

Calculate: t-statistic and p-value

Result: If p-value < 0.05, reject null (significant difference)

****Chi-Square Test**:**

Used for: Conversion rates (binary outcomes)

Null hypothesis: Conversion rates are equal

Calculate: Chi-square statistic and p-value

Result: If p-value < 0.05, variants differ significantly

****Confidence Intervals**:**

Variant A: $10\% \pm 0.6\%$ (95% confidence)

→ True rate between 9.4% - 10.6%

Variant B: $11.5\% \pm 0.6\%$ (95% confidence)

→ True rate between 10.9% - 12.1%

Since intervals don't overlap → Significant difference!

****Sample Size Calculator**:**

To detect 10% relative lift with 95% confidence:

- Baseline conversion rate: 10%
- Minimum detectable effect: 1% absolute (10% relative)
- Power: 80%
- Required sample per variant: ~15,000 users

Run experiment until reach sample size

3. Targeting Rules

****Use Case**:** Run experiment only for specific users

****Common Targeting Dimensions**:**

Geographic:

- Country: US, UK, IN
- Region: California, Texas
- City: San Francisco

Demographics:

- Age: 18-24, 25-34
- Gender: Male, Female, Other

Platform:

- Device: Mobile, Desktop, Tablet
- OS: iOS, Android, Windows
- Browser: Chrome, Safari, Firefox

Behavioral:

- User segment: Premium, Free, Trial
- Previous purchases: Yes/No
- Account age: < 30 days, > 1 year

Custom:

- Company defined segments
- ML-based cohorts

****Example Targeting Rule**:**

Experiment: "mobile_checkout_redesign"

Target:

- Platform: Mobile only
- Country: US, UK, Canada

- User segment: Active (purchased in last 30 days)

Exclusions:

- Internal employees
- Bots
- Users in other critical experiments

Evaluation:

User requests assignment:

1. Check if user matches ALL targeting criteria
2. If NO → User not in experiment (show default)
3. If YES → Assign to variant using hash

Benefits:

- Focus experiment on relevant users
- Avoid contamination
- Reduce noise in data

4. Traffic Allocation Strategies

Fixed Allocation

Example: 50-50 split

- 50% see Variant A
- 50% see Variant B
- Allocation stays constant

Use for: Traditional A/B tests

Gradual Rollout

Example: New feature launch

- Day 1: 5% see new feature
- Day 3: 10% see new feature

- Day 7: 25% see new feature
- Day 14: 50% see new feature
- Day 21: 100% see new feature

Use for: Risk mitigation, feature flags

Multi-Armed Bandit

Dynamic allocation based on performance:

- Start: 50-50 split
- After 1000 samples:
 - Variant A: 8% conversion
 - Variant B: 12% conversion
- Shift traffic: 30-70 split (favor B)
- Continuously optimize

Use for: Revenue optimization (show better variant to more users)

Trade-off: Slower to reach statistical significance

5. Real-Time Analytics Pipeline

Architecture:

[Event Tracking API]

↓

[Kafka Topics]

↓

[Stream Processor (Flink)]

- Aggregate by experiment/variant
 - Calculate rates
 - Sliding windows (hourly, daily)
 - ↓
- [ClickHouse]
- Store aggregated metrics
 - Fast queries for dashboards
 - ↓
- [WebSocket Server]
- Push updates to dashboard

- Real-time visualization

Aggregation Windows:

1-minute window: For real-time monitoring

1-hour window: For hourly trends

1-day window: For daily summaries

Store all three for different use cases

Why Stream Processing:

- Real-time updates (< 1 minute delay)
- Handle billions of events
- Stateful aggregations
- Exactly-once semantics

6. Feature Flags Implementation

Feature Flag vs A/B Test:

A/B Test:

- Compare variants
- Statistical analysis
- Usually temporary (2-4 weeks)
- Goal: Find winner

Feature Flag:

- Gradual rollout
- Kill switch
- Can be permanent
- Goal: Safe deployment

Rollout Strategy:

Phase 1: Internal (0.1%)

- Test with employees
- Catch obvious bugs

Phase 2: Canary (1%)

- Real users, small sample
- Monitor errors/latency

Phase 3: Early Adopters (10%)

- Willing beta testers
- Gather feedback

Phase 4: General (50%)

- Half of users
- Validate at scale

Phase 5: Full Rollout (100%)

- Everyone gets new feature
- Monitor for issues

At each phase:

- Monitor error rates
- Check latency
- Track conversions
- If issues → Instant rollback

7. Assignment Caching Strategy

Challenge: 115K assignment QPS, can't hit database

Solution: Multi-Layer Cache

L1 – Client-Side Cache (Optional):

Store assignment in local storage/memory

TTL: Duration of session or experiment

Pros: Zero network calls

Cons: Can't change mid-session

L2 – Redis Cache (Primary):

Store all active assignments

Key: assign:{user_id}:{exp_id}

Value: {variant_id, config}

TTL: Experiment duration

Cache Stats:

- 99.9% hit ratio
- < 1ms latency
- Handles 115K QPS easily

L3 – Database (Source of Truth):

PostgreSQL stores all assignments

Used for:

- Analytics (who saw what)
- Audit trail
- Cache rebuilds

Cache Warming:

On experiment start:

- Pre-load experiment configs to Redis
- Pre-compute assignments for active users
- Ensures fast first request

Scalability & Reliability

Horizontal Scaling

Assignment Service:

Stateless service, scales linearly

- 115K QPS / 5K per instance = 23 instances
- Add buffer: 30 instances
- Auto-scale based on QPS

Cost: $30 \times \$100/\text{month} = \$3,000/\text{month}$

Event Tracking Service:

- 12K write QPS / 2K per instance = 6 instances
- Add buffer: 10 instances

Cost: $10 \times \$100/\text{month} = \$1,000/\text{month}$

Kafka Cluster:

- 3 brokers for high availability
- Replication factor: 3
- 10 partitions per topic (parallelism)

Cost: $3 \times \$500/\text{month} = \$1,500/\text{month}$

Redis Cache:

- Single instance sufficient (11 GB data)
- Or Redis Cluster for HA
- 3 nodes with replication

Cost: \$200-500/month

ClickHouse:

- 5-node cluster
- Distributed queries
- Replicated for HA

Cost: $5 \times \$400/\text{month} = \$2,000/\text{month}$

High Availability

Critical Path – Assignment Service:

Target: 99.99% availability

Strategies:

1. Deploy across 3 availability zones
2. Load balancer with health checks
3. Auto-scaling (replace failed instances)
4. Redis Cluster (no single point of failure)
5. Fallback: Return default variant if all fails

Max acceptable downtime: 52 minutes/year

Non-Critical – Analytics:

Target: 99.9% availability

Can tolerate:

- Analytics delay (process events later)
- Dashboard downtime (not on critical path)
- Report generation failures (retry)

Data Consistency

Experiment Configuration – Strong Consistency:

Use PostgreSQL with ACID:

- Can't show wrong variant
- Changes must be atomic
- All servers see same config

Replication:

- Master-slave with read replicas
- Reads can have slight lag (acceptable)
- Writes always to master

Event Data – Eventual Consistency:

Use Kafka + Cassandra:

- OK if analytics delayed 1 minute
- Availability > consistency
- Can tolerate event loss (< 0.01%)

Why acceptable:

- Statistical analysis robust to small data loss
- Real-time dashboards approximate

Disaster Recovery

Backup Strategy:

PostgreSQL (Experiment configs):

- Daily full backup
- Continuous WAL archiving
- Cross-region replication
- RPO: < 5 minutes

Kafka (Events):

- Replicated to 3 brokers
- Retention: 7 days
- Can replay events

ClickHouse (Analytics):

- Daily snapshots
- Rebuild from Kafka if needed
- RPO: < 1 hour (acceptable)

Security & Privacy

User Privacy

****PII Handling**:**

DON'T Store:

- Names, emails, addresses
- Credit card numbers
- Personal information

DO Store:

- Hashed user IDs
- Anonymous identifiers
- Aggregated metrics only

****GDPR Compliance**:**

User Rights:

- Right to deletion (remove all user data)
- Right to export (download experiment history)
- Right to opt-out (don't include in experiments)

Implementation:

- user_id hashing (can't reverse engineer)
- Data retention policies (auto-delete after 90 days)
- Opt-out flags in assignment service

Experiment Security

****Access Control**:**

Roles:

- Admin: Create, modify, delete experiments
- Analyst: View results, export data
- Developer: Read-only config access

Permissions:

- Experiment-level (can only modify your experiments)

- Company-level (can view all company experiments)

****Audit Logs**:**

Track all changes:

- Who created experiment
- Who modified allocation
- Who stopped experiment early
- Who accessed results

Store in append-only log (tamper-proof)

```
----  
## Performance Optimizations  
### 1. Caching Experiment Configs  
**Problem**: 10K experiments, 115K QPS, can't query DB each time  
**Solution**:
```

Cache ALL experiment configs in Redis:

- $10K \times 1\text{ KB} = 10\text{ MB}$ (tiny!)
- Update when experiment modified
- TTL: 5 minutes (or indefinite with active invalidation)

Assignment flow:

1. Check Redis for experiment config (< 1ms)
2. Calculate assignment (hash, < 1ms)
3. Return variant (< 2ms total)

vs. Database query:

1. Query PostgreSQL (5-10ms)
2. Calculate assignment (< 1ms)
3. Return (6-11ms total)

5x faster with cache!

2. Pre-Aggregation

****Problem**:** Dashboard queries expensive (SUM over billions of events)

****Solution**:**

Pre-aggregate metrics:

Raw Events (Kafka):

event1: user_123, checkout_test, variant_a, click
event2: user_456, checkout_test, variant_a, click
event3: user_789, checkout_test, variant_a, conversion

Stream Processing (Flink):

Aggregate every minute:

- checkout_test, variant_a: 1000 exposures, 50 conversions
- checkout_test, variant_b: 1000 exposures, 65 conversions

Store aggregates in ClickHouse:

- Query aggregates (fast!)
- Instead of scanning billions of raw events

Dashboard query:

```
SELECT SUM(exposures), SUM(conversions)
FROM aggregated_metrics
WHERE experiment_id = 'checkout_test'
AND date >= '2025-01-01'
```

Runs in < 100ms vs minutes for raw events

3. Handling Experiment Collisions

****Problem**:** User in multiple experiments that affect same page

****Example**:**

User assigned to:

- Experiment A: Test checkout button color (Blue vs Red)
- Experiment B: Test checkout layout (Single vs Multi-page)

Both experiments modify checkout page!

Risk: Results contaminated

Solution 1: Mutual Exclusion

Configuration:

- Experiment A: Exclude users in Experiment B
- Experiment B: Exclude users in Experiment A

Result:

- User in A → Not eligible for B
- User in B → Not eligible for A
- Clean experiment groups

Trade-off: Smaller sample size (users split between experiments)

Solution 2: Orthogonal Experiments

Design experiments to be independent:

- Experiment A: Homepage hero image
- Experiment B: Checkout button color
- Different pages, no interaction

Can run simultaneously without issues

Solution 3: Layered Experiments

Create experiment layers:

- Layer 1: Checkout experiments only
- Layer 2: Homepage experiments only
- Layer 3: Pricing experiments only

User can be in one experiment per layer

Trade-offs & Alternatives

1. Hash-Based vs Stored Assignment

Chose: Hash-Based (with cache)

Reasoning:

- 115K QPS requires fast assignment
- Hash-based: < 1ms, scales infinitely
- Stored: 5–10ms, requires database scaling

Trade-off:

- Hash-based: Can't change allocation mid-experiment
- Stored: Flexible but slower

Mitigation:

- Use hash-based as default
- Cache in Redis for < 1ms
- Store async in database for analytics

2. Real-Time vs Batch Analytics

Chose: Hybrid (Real-time + Batch)

Real-Time (Flink → ClickHouse):

For: Live dashboards

Latency: < 1 minute

Aggregation: Last 24 hours

Use case: Monitor running experiments

Batch (Spark → Redshift):

For: Historical analysis

Latency: Hourly/daily

Aggregation: All time

Use case: Deep analysis, reports

Why Both:

- Real-time for monitoring
- Batch for accuracy and historical trends

3. ClickHouse vs Cassandra for Analytics

Chose: ClickHouse

Reasoning:

ClickHouse:

- Columnar (fast aggregations)
- Designed for analytics queries
- SQL interface (familiar)
- Real-time inserts + queries

Cassandra:

- Better for pure time-series writes
- Geographic distribution
- Aggregations slower
- No SQL (CQL different)

Decision: Analytics workload favors ClickHouse

4. Synchronous vs Asynchronous Event Tracking

Chose: Asynchronous

Synchronous (Bad):

User clicks button:

1. Send tracking event (wait for response)
2. Button action proceeds

Problem: 100ms delay for user!

Asynchronous (Good):

User clicks button:

1. Queue event locally (instant)
2. Button action proceeds immediately
3. Background: Batch send events

Result: No user-facing latency

Technology Stack

Layer	Technology	Purpose
CDN	CloudFront	Dashboard static assets
Load Balancer	AWS ALB	Traffic distribution
API Gateway	Kong	Auth, rate limiting
Assignment Service	Go, Java	Fast variant assignment
Event Service	Node.js, Python	Event ingestion
Analytics	Apache Flink	Stream processing
Cache	Redis/Redis Cluster	Config & assignment cache
Config DB	PostgreSQL	Experiments, assignments
Event Stream	Apache Kafka	Durable event pipeline
Analytics DB	ClickHouse	Real-time queries
Data Warehouse	Amazon Redshift	Historical analysis
Object Storage	Amazon S3	Raw event backup
Monitoring	Prometheus, Datadog	Metrics & alerts

Technology Decision Summary

Complete Technology Stack with Justifications

Component	Chosen	Why Chosen	Why NOT Alternative
Assignment Language	**Go**	• 1-2ms latency • Low memory (10-20MB)	• Native concurrency X Java: Higher memory X Node.js: Variable latency X Python: Too slow (5-10ms)
Config Database	**PostgreSQL**	• ACID transactions Complex JOINs • JSON support • Relationships	• MongoDB: No transactions X DynamoDB: No complex queries X MySQL: Weaker JSON support
Cache	**Redis**	• Sub-ms latency • Rich data structures • Pub/Sub for invalidation • Cluster mode	X Memcached: No data structures X Hazelcast: Slower, heavier X In-memory: No shared state
Event Stream	**Kafka**	• High throughput • Multiple consumers • Durability	X RabbitMQ: Lower throughput, no replay X SQS: Expensive, no replay X Pub/Sub: Vendor lock-in
Stream Processing	**Flink**	• True streaming (ms) • Event-time windows • Stateful processing • Exactly-once	X Spark: Micro-batch (seconds) X Kafka Streams: Less features X Storm: Older, less mature
Analytics DB	**ClickHouse**	• Columnar (fast aggregations) • 10-50x compression • Real-time queries	X SQL X Cassandra: Slow aggregations X PostgreSQL: Doesn't scale X BigQuery:

Expensive, vendor lock-in |
| **Data Warehouse** | **Redshift** | • SQL interface
• Integrates with S3
• Cost-effective
• Mature | ❌ Snowflake: More expensive
❌ BigQuery: Vendor lock-in
❌ Direct S3 queries: Too slow |

Critical Path Analysis

Assignment Request (< 10ms requirement):

Component	Time	Why This Tech
API Gateway	0.5ms	Kong (lightweight)
Assignment Service	0.5ms	Go (fast concurrency)
Redis	1ms	In-memory, sub-ms access
Hash Calculation	0.5ms	MD5, O(1) operation
Network	2ms	Within same AZ
Total	~4.5ms	✓ Well under 10ms budget

If PostgreSQL (no cache):

Database Query	5-10ms	Too slow!
Total	~8-12ms	✗ Risk exceeding budget

Event Processing Path

Event Tracking (< 100ms requirement):

Component	Time	Why This Tech
SDK (async)	~0ms	Non-blocking
API Gateway	0.5ms	Auth + validation
Event Service	1ms	Node.js (I/O optimized)
Kafka	2-5ms	Disk-backed, replicated
Response	0.5ms	200 OK returned
Total	~4-7ms	✓ Client sees <10ms
Async Processing:		
Flink Processing	<1sec	Real-time aggregation
ClickHouse Write	<1sec	Batch insert
Dashboard Update	<1min	WebSocket push

Why Specific Tech Combinations Work Together

1. Kafka + Flink + ClickHouse (Analytics Pipeline)

Why this combination is perfect:

Kafka:

- Produces ordered, durable event stream
- Multiple consumers can read same data
- Replay for reprocessing

Flink:

- Consumes from Kafka in real-time
- Stateful windowed aggregations
- Exactly-once processing

ClickHouse:

- Stores Flink's aggregated results
- Fast queries on aggregated data
- Columnar storage perfect for metrics

Result: Real-time analytics pipeline with <1 minute latency

Alternative combinations considered:

- ✗ Kafka + Spark + Cassandra: Spark has seconds delay
- ✗ SQS + Lambda + DynamoDB: No replay, harder to scale
- ✗ Pub/Sub + Dataflow + BigQuery: Vendor lock-in, expensive

2. Go + Redis + PostgreSQL (Assignment Service)

Why this combination is perfect:

Go:

- Fast hash calculation (1-2ms)
- Low memory footprint
- Handles 115K QPS per instance

Redis:

- Caches experiment configs (10 MB total)

- 99.9% hit ratio (sub-ms lookups)
- Handles misses gracefully

PostgreSQL:

- Async storage for assignments
- Audit trail for compliance
- Source of truth for rebuilds

Result: <10ms p99 latency achieved

Without Redis:

- ✗ 115K QPS would need 575 database connections
- ✗ PostgreSQL can't handle that load
- ✗ Would need expensive database scaling

Cost–Performance Trade–off Analysis

Option 1: All Managed Services (Easy but Expensive)

- └─ DynamoDB instead of PostgreSQL: +\$2,000/mo
- └─ Kinesis instead of Kafka: +\$3,000/mo
- └─ BigQuery instead of ClickHouse: +\$10,000/mo
- └─ Managed Redis (ElastiCache): +\$500/mo
- └─ Total: ~\$36,000/mo (75% more expensive)

Pros: Less operational overhead

Cons: 75% higher cost, vendor lock-in

Option 2: Self-Hosted (Chosen - Best Balance)

- └─ Self-hosted Kafka: \$1,500/mo
- └─ Self-hosted ClickHouse: \$2,000/mo
- └─ RDS PostgreSQL: \$300/mo
- └─ ElastiCache Redis: \$450/mo
- └─ Total: ~\$20,500/mo

Pros: Cost-effective, no vendor lock-in, full control

Cons: More operational overhead (but manageable)

Option 3: Minimize Costs (Limited Scale)

- └─ Single PostgreSQL for everything: \$500/mo
- └─ No Kafka (direct DB writes): \$0
- └─ No Redis (use PostgreSQL): \$0
- └─ Total: ~\$500/mo

Pros: Very cheap

Cons: Can't scale beyond 1K QPS, slow queries

Interview Talking Points

Key Design Decisions

1. **Why hash-based assignment?**
 - 115K QPS requires < 10ms latency
 - Deterministic without database lookup
 - Scales infinitely
2. **Why Redis cache?**
 - All configs fit in memory (10 MB)
 - 99.9% cache hit ratio
 - < 1ms latency
3. **Why Kafka?**
 - Buffer traffic spikes
 - Multiple consumers (analytics, warehouse)
 - Can replay events for debugging
4. **Why ClickHouse?**
 - Columnar storage perfect for aggregations
 - Real-time queries on billions of rows
 - Better than Cassandra for analytics
5. **Why PostgreSQL for configs?**
 - ACID transactions (critical)
 - Small dataset (10K experiments)
 - Complex queries needed
6. **Why stream processing?**
 - Real-time dashboards (< 1 minute delay)
 - Stateful aggregations
 - Handle billions of events

Potential Bottlenecks & Solutions

1. **Assignment service overwhelmed**
 - Solution: Cache configs, hash-based assignment, horizontal scaling
2. **Event ingestion slow**
 - Solution: Kafka buffering, batch processing, async
3. **Analytics queries slow**
 - Solution: Pre-aggregation, columnar storage (ClickHouse)
4. **Cache invalidation**
 - Solution: TTL + active invalidation, Redis Cluster

5. **Database writes**
 - Solution: Async writes, batch inserts, Kafka buffer

Real-World Architectures

Optimizely's Architecture

Assignment: Edge network (< 10ms globally)

Events: Sent to nearest datacenter

Analytics: Stream processing + data warehouse

Scale: 12 billion decisions/day

Google Optimize

Assignment: Google's global CDN

Events: Google Analytics integration

Analytics: BigQuery

Scale: Integrated with Google ecosystem

LaunchDarkly

Assignment: In-memory SDKs (offline-first)

Events: Event streaming pipeline

Analytics: Real-time + historical

Scale: Feature flags at massive scale

Advanced Features

1. Holdout Groups

Purpose: Control group across multiple experiments

Example:

- 90% of users: Participate in various experiments
- 10% holdout: See original experience (no experiments)

Benefits:

- Measure overall impact of experimentation program
- Detect interaction effects
- Validate that experiments actually improve metrics

2. Sequential Testing

Purpose: Stop experiment early when clear winner

Traditional: Wait for full sample size (2-4 weeks)

Sequential: Check significance continuously, stop early if clear winner

Benefits:

- Faster decisions (days vs weeks)
- Reduce risk (stop bad experiments sooner)
- Opportunity cost (ship winners faster)

Trade-off:

- More complex statistics
- Risk of false positives (solved with adjusted thresholds)

3. Revenue Attribution

Purpose: Calculate revenue impact of experiments

Track:

- Direct revenue (during experiment)
- Lifetime value (long-term impact)
- Per-user revenue

Example:

Variant A: \$10 revenue per user

Variant B: \$12 revenue per user

With 10M users: \$20M additional revenue/year

Business value: Clear ROI for winning variant

Detailed Component Interactions

Complete Assignment Flow

1. User visits website
2. Website SDK calls: POST /api/v1/assign
3. API Gateway: Authenticates request
4. Assignment Service:
 - a. Check Redis for cached assignment (< 1ms)
 - b. If miss: Calculate hash assignment (< 1ms)
 - c. Check targeting rules (user matches?)
 - d. Cache in Redis
 - e. Async store in PostgreSQL
5. Return variant to website (< 10ms total)
6. Website renders appropriate variant

Complete Event Tracking Flow

1. User completes action (purchase, signup, etc.)
2. Website SDK: POST /api/v1/track (async, doesn't block)
3. Event Service:
 - a. Validates event
 - b. Enriches with metadata
 - c. Writes to Kafka topic
4. Returns 200 OK (< 100ms)
5. Kafka consumers process:
 - a. Flink: Real-time aggregation
 - b. Spark: Batch processing
 - c. Warehouse: Historical storage
6. Dashboards update automatically via WebSocket

Cost Analysis

Monthly Infrastructure Costs

For 100M users, 1B events/day, 10K experiments:

Assignment Service: 30 instances × \$100 = \$3,000

Event Service: 10 instances × \$100 = \$1,000

Analytics Service: 5 instances × \$200 = \$1,000

Load Balancers: $2 \times \$25 = \50

PostgreSQL (RDS): db.r5.large = \$300

Redis Cluster: 3 nodes $\times \$150 = \450

Kafka: 3 brokers $\times \$500 = \$1,500$

ClickHouse: 5 nodes $\times \$400 = \$2,000$

Redshift: \$1,000

Flink/Spark: \$1,500

S3 Storage (365 TB): \$8,000

Monitoring: \$500

CDN: \$200

Total: ~\$20,500/month

Revenue Model:

- Price per experiment: \$500/month
- 1000 paying customers = \$500K/month revenue
- Margin: 96% (very profitable!)

References & Further Reading

Statistical Resources

1. **"Trustworthy Online Controlled Experiments"** – Microsoft Research
2. **"A/B Testing: The Most Powerful Way to Turn Clicks Into Customers"** – Dan Siroker
3. **"Statistical Methods in Online A/B Testing"** – Google Research

System Design Resources

1. **Optimizely Engineering Blog**
2. **Netflix Experimentation Platform** – Tech blog
3. **Booking.com A/B Testing** – Engineering blog

Tools & Platforms

1. **Optimizely**: Industry leader
2. **Google Optimize**: Free tier available
3. **LaunchDarkly**: Feature flags focused
4. **Split.io**: Enterprise platform
5. **VWO**: Visual editor included

Appendix

Sample Size Calculations

Given:

- Baseline conversion rate: 5%
- Minimum detectable effect: 10% relative (0.5% absolute)
- Statistical power: 80%
- Significance level: 95%

Required sample per variant: ~25,000 users

Experiment duration:

- Daily visitors: 10,000
- Days needed: 5 days (2.5 days per variant)

General formula:

$$n = 16 \times p \times (1-p) / (\text{MDE})^2$$

Where:

- p = baseline rate
- MDE = minimum detectable effect

Statistical Significance Thresholds

P-value Confidence Interpretation

- < 0.001 99.9% Extremely significant
- < 0.01 99% Highly significant
- < 0.05 95% Significant (standard)
- < 0.10 90% Marginally significant

0.10 < 90% Not significant

Industry standard: $p < 0.05$ (95% confidence)

Conservative: $p < 0.01$ (99% confidence)

Common Experiment Durations

Traffic Level Sample Size Duration

Low (1K/day) 10K 10 days

Medium (10K/day) 10K 1 day

High (100K/day) 10K 3 hours

Very High (1M/day) 10K 15 minutes

Rule of thumb: Run for at least 1 week to capture:

- Weekday vs weekend behavior
- Time-of-day patterns
- Seasonality

****Document Version**: 1.0**

****Last Updated**: January 8, 2025**

****Author**: System Design Interview Prep**

****Status**: Complete & Interview-Ready** 

****Pro Tip**:** A/B testing platforms are common interview questions at companies like Optimizely, Google, Facebook, Netflix, and any data-driven company. Master this design!