

Distributed Key-Value Store System Design (Like Redis, DynamoDB, Cassandra)

Table of Contents

- Distributed Key-Value Store System Design (Like Redis, DynamoDB, Cassandra)
 - Table of Contents
 - Problem Statement & Requirements
 - Problem Definition
 - Functional Requirements
 - Non-Functional Requirements
 - Capacity Estimation
 - Assumptions
 - QPS (Queries Per Second)
 - Memory Requirements
 - Network Bandwidth
 - Storage Requirements
 - API Design
 - Basic Operations
 - Batch Operations
 - Advanced Operations
 - Admin Operations
 - Data Model & Storage
 - In-Memory Index Structure
 - On-Disk Storage Format (SSTable)
 - LSM Tree (Log-Structured Merge Tree)
 - Write-Ahead Log (WAL)
 - High-Level Architecture
 - Architecture Layers
 - Core Components
 - 1. Storage Engine
 - 2. Consistent Hashing
 - 3. Replication Coordinator
 - 4. Gossip Protocol
 - 5. Failure Detector
 - 6. Anti-Entropy (Merkle Trees)
 - Data Partitioning (Sharding)
 - Why Partition?
 - Partitioning Strategies
 - 1. Hash-Based Partitioning
 - 2. Range-Based Partitioning
 - 3. Consistent Hashing (Recommended)
 - Partition Management

- Replication Strategy
 - Replication Factor
 - Replica Placement
 - Replication Protocols
 - 1. Synchronous Replication
 - 2. Asynchronous Replication
 - 3. Quorum-Based (Hybrid)
- Consistency Models
 - CAP Theorem
 - Consistency Levels
 - 1. Strong Consistency
 - 2. Eventual Consistency
 - 3. Read-Your-Writes Consistency
 - 4. Monotonic Reads
 - Vector Clocks
- Failure Detection & Recovery
 - Types of Failures
 - Failure Detection Mechanisms
 - 1. Heartbeat Protocol
 - 2. Phi Accrual Failure Detector
 - Recovery Strategies
 - 1. Replica Takeover
 - 2. Hinted Handoff
 - 3. Read Repair
 - 4. Anti-Entropy Repair
- Read/Write Path
 - Write Path (Detailed)
 - Read Path (Detailed)
 - Read Repair Example
- Performance Optimization
 - 1. Caching Strategy
 - 2. Bloom Filters
 - 3. Compaction Strategies
 - 4. Batch Operations
 - 5. Connection Pooling
 - 6. Compression
 - 7. Read-Ahead
- Multi-Datacenter Deployment
 - Architecture
 - Replication Strategies
 - 1. Active-Passive
 - 2. Active-Active (Multi-Master)
 - Geo-Replication
 - Network Partition Handling
- Design Trade-offs

- 1. Consistency vs Availability
 - 2. LSM Tree vs B-Tree
 - 3. Replication Factor
 - 4. Partition Size
 - Real-World Implementations
 - 1. Amazon DynamoDB
 - 2. Apache Cassandra
 - 3. Redis
 - 4. etcd
 - Summary & Key Takeaways
 - Core Design Principles
 - When to Use Key-Value Stores
 - Interview Success Tips
 - Common Pitfalls to Avoid
 - Further Reading
 - Essential Papers
 - Books
 - Online Resources
-

Problem Statement & Requirements

Problem Definition

Design a distributed key-value store that can:

- Store billions of key-value pairs
- Handle millions of operations per second
- Provide high availability (99.99%+)
- Scale horizontally across thousands of nodes
- Support various consistency models
- Replicate data across multiple datacenters

Real-World Examples: Redis, Memcached, DynamoDB, Cassandra, etcd, Consul

Functional Requirements

1. Basic Operations

- **PUT(key, value)**: Store a key-value pair
- **GET(key)**: Retrieve value by key
- **DELETE(key)**: Remove a key-value pair
- **EXISTS(key)**: Check if key exists

2. Advanced Operations

- **LIST(prefix)**: List keys by prefix
- **SCAN(cursor)**: Iterate over keys

- `TTL(key, seconds)`: Set time-to-live
- `BATCH_PUT(keys, values)`: Atomic batch write
- `COMPARE_AND_SWAP(key, old_value, new_value)`: Conditional update

3. Configuration

- Configurable replication factor
 - Tunable consistency levels
 - Custom data retention policies
 - Multi-tenancy support
-

Non-Functional Requirements

1. Performance

- Read latency: P99 < 10ms
- Write latency: P99 < 20ms
- Throughput: 100K+ ops/second per node
- Support 10M+ ops/second cluster-wide

2. Scalability

- Horizontal scaling (add/remove nodes)
- Support 1000+ node clusters
- Store petabytes of data
- Linear scalability with node count

3. Availability

- 99.99% uptime (52 minutes downtime/year)
- No single point of failure
- Automatic failover (< 30 seconds)
- Graceful degradation

4. Consistency

- Support multiple consistency models:
 - Strong consistency
 - Eventual consistency
 - Read-your-writes consistency
- Tunable per operation

5. Durability

- Data persisted to disk
- Replication across nodes/datacenters
- Snapshot and backup support
- 99.999999999% durability (11 nines)

6. Partition Tolerance

- Continue operating during network partitions
 - Automatic partition healing
 - Data reconciliation
-

Capacity Estimation

Assumptions

```
Total Data: 100 TB
Average Key Size: 50 bytes
Average Value Size: 1 KB
Total Keys: 100 TB / 1 KB ≈ 100 Billion keys
Replication Factor: 3
Actual Storage: 100 TB × 3 = 300 TB
```

QPS (Queries Per Second)

```
Read:Write Ratio: 80:20 (read-heavy workload)
Total Operations: 10 Million ops/second

Reads: 8M ops/sec
Writes: 2M ops/sec

Per Node (100 nodes):
- Reads: 80K ops/sec
- Writes: 20K ops/sec
```

Memory Requirements

```
Metadata per key: 100 bytes (key, value pointer, timestamps, version)
Total Keys: 100B

Memory for indexes:
100B keys × 100 bytes = 10 TB metadata

With caching (20% hot data):
Cache: 20 TB of actual data + 2 TB metadata = 22 TB
Per Node (100 nodes): 220 GB memory
```

Network Bandwidth

Average Operation Size: 1 KB
10M ops/sec × 1 KB = 10 GB/sec

With replication (factor 3):
Writes: 2M × 3 replicas = 6M ops/sec
Total Bandwidth: 16 GB/sec

Per Node: 160 MB/sec

Storage Requirements

Data: 300 TB (with 3x replication)
Write-Ahead Log (WAL): 30 TB (10% overhead)
Snapshots: 100 TB (1 full snapshot)
Total: 430 TB

Per Node (100 nodes): 4.3 TB per node
Recommended: 8 TB SSD per node (room for growth)

API Design

Basic Operations

```
# Put – Store a key-value pair
PUT /api/v1/keys/{key}
Headers:
  Content-Type: application/json
  Consistency-Level: QUORUM (optional)
Body: {
  "value": "some_value",
  "ttl": 3600, // optional, in seconds
  "if_not_exists": false // optional, for conditional put
}
Response: {
  "key": "user:123",
  "version": 5,
  "timestamp": 1640995200,
  "status": "success"
}

# Get – Retrieve value by key
GET /api/v1/keys/{key}
Headers:
  Consistency-Level: EVENTUAL (optional)
Response: {
  "key": "user:123",
```

```

    "value": "some_value",
    "version": 5,
    "timestamp": 1640995200,
    "ttl_remaining": 3500
}

# Delete – Remove a key
DELETE /api/v1/keys/{key}
Response: {
    "key": "user:123",
    "deleted": true,
    "version": 6
}

# Exists – Check if key exists
HEAD /api/v1/keys/{key}
Response:
  200 OK (exists)
  404 Not Found (doesn't exist)

```

Batch Operations

```

# Batch Put – Write multiple keys atomically
POST /api/v1/batch/put
Body: {
    "operations": [
        {"key": "user:123", "value": "data1"},
        {"key": "user:456", "value": "data2"}
    ],
    "atomic": true // fail all if any fails
}
Response: {
    "success_count": 2,
    "failed_keys": [],
    "results": [...]
}

# Batch Get – Read multiple keys
POST /api/v1/batch/get
Body: {
    "keys": ["user:123", "user:456", "user:789"]
}
Response: {
    "results": [
        {"key": "user:123", "value": "data1", "found": true},
        {"key": "user:456", "value": "data2", "found": true},
        {"key": "user:789", "found": false}
    ]
}

```

Advanced Operations

```
# Compare and Swap – Atomic conditional update
POST /api/v1/keys/{key}/cas
Body: {
    "expected_version": 5,
    "new_value": "updated_value"
}
Response: {
    "success": true,
    "new_version": 6
}

# List Keys by Prefix
GET /api/v1/keys?prefix=user:&limit=100
Response: {
    "keys": ["user:123", "user:456", ...],
    "cursor": "next_page_token",
    "has_more": true
}

# Scan – Iterate over all keys
GET /api/v1/scan?cursor={token}&count=1000
Response: {
    "keys": [...],
    "next_cursor": "new_token",
    "done": false
}

# Set TTL – Time to live
POST /api/v1/keys/{key}/ttl
Body: {
    "ttl_seconds": 3600
}

# Increment – Atomic increment (for counters)
POST /api/v1/keys/{key}/incr
Body: {
    "delta": 1
}
Response: {
    "key": "counter:views",
    "value": 1001,
    "version": 1002
}
```

Admin Operations

```

# Get Cluster Status
GET /api/v1/admin/cluster/status
Response: {
    "nodes": [
        {
            "node_id": "node-1",
            "address": "10.0.1.5:6379",
            "status": "healthy",
            "role": "master",
            "keys_count": 10000000,
            "memory_used_mb": 5120
        }
    ],
    "total_keys": 100000000,
    "replication_factor": 3
}

# Add Node to Cluster
POST /api/v1/admin/cluster/nodes
Body: {
    "node_address": "10.0.1.10:6379"
}

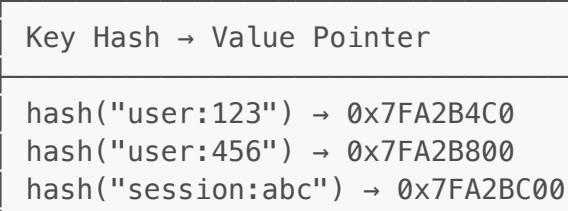
# Rebalance Cluster
POST /api/v1/admin/cluster/rebalance

```

Data Model & Storage

In-Memory Index Structure

Hash Table (Main Index):



Value Pointer Structure:

```

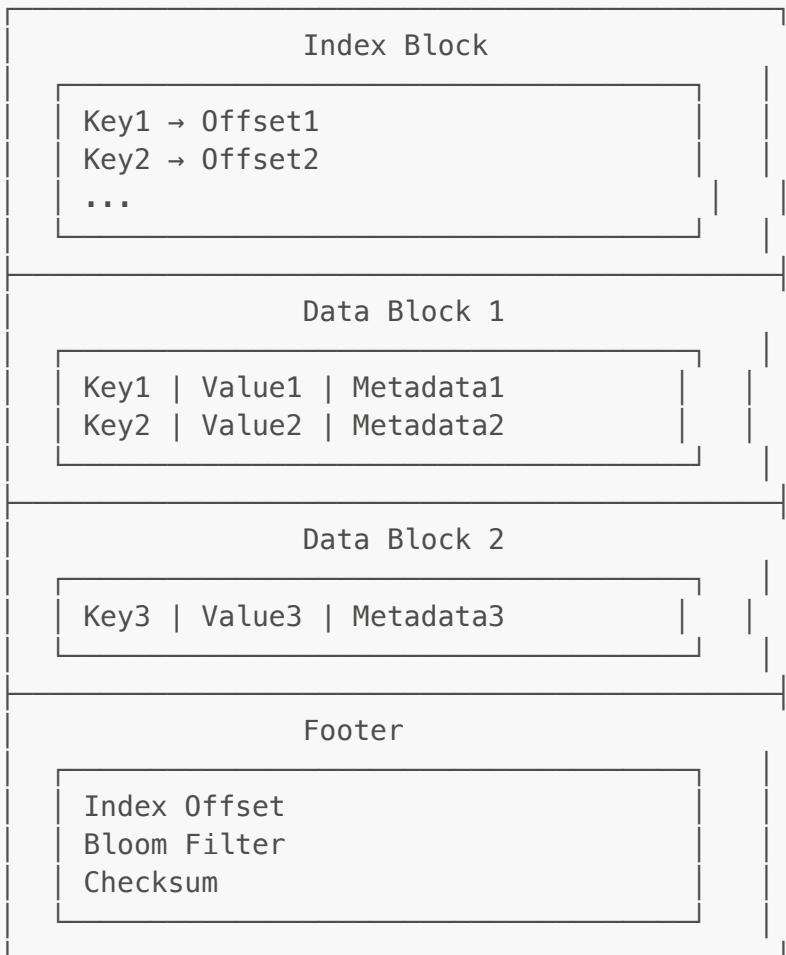
struct ValuePointer {
    char* data;           // pointer to value
    uint32_t size;        // value size
    uint64_t version;     // version number
    uint64_t timestamp;   // last modified
    uint32_t ttl;         // time to live
    uint8_t flags;        // metadata flags
}

```

```
}
```

On-Disk Storage Format (SSTable)

SSTable (Sorted String Table) Layout:



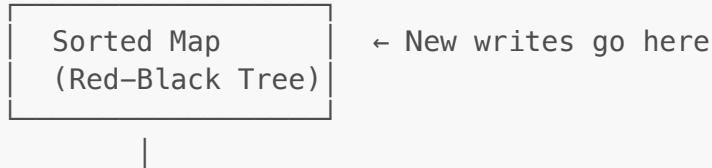
Bloom Filter: Probabilistic data structure

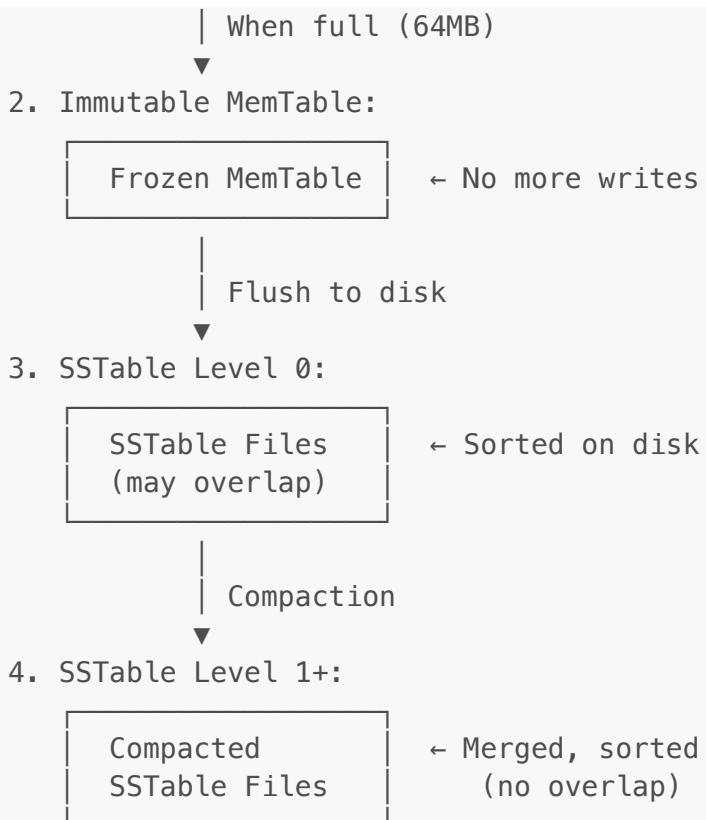
- Check if key exists (no false negatives)
- ~10 bytes per key
- 1% false positive rate

LSM Tree (Log-Structured Merge Tree)

Write Path (LSM Tree):

1. MemTable (In-Memory):





Compaction Strategy:

- Size-Tiered: Merge SSTables of similar size
- Leveled: Each level 10x larger than previous
- Time-Window: Merge based on time ranges

Write-Ahead Log (WAL)

WAL Structure:

```

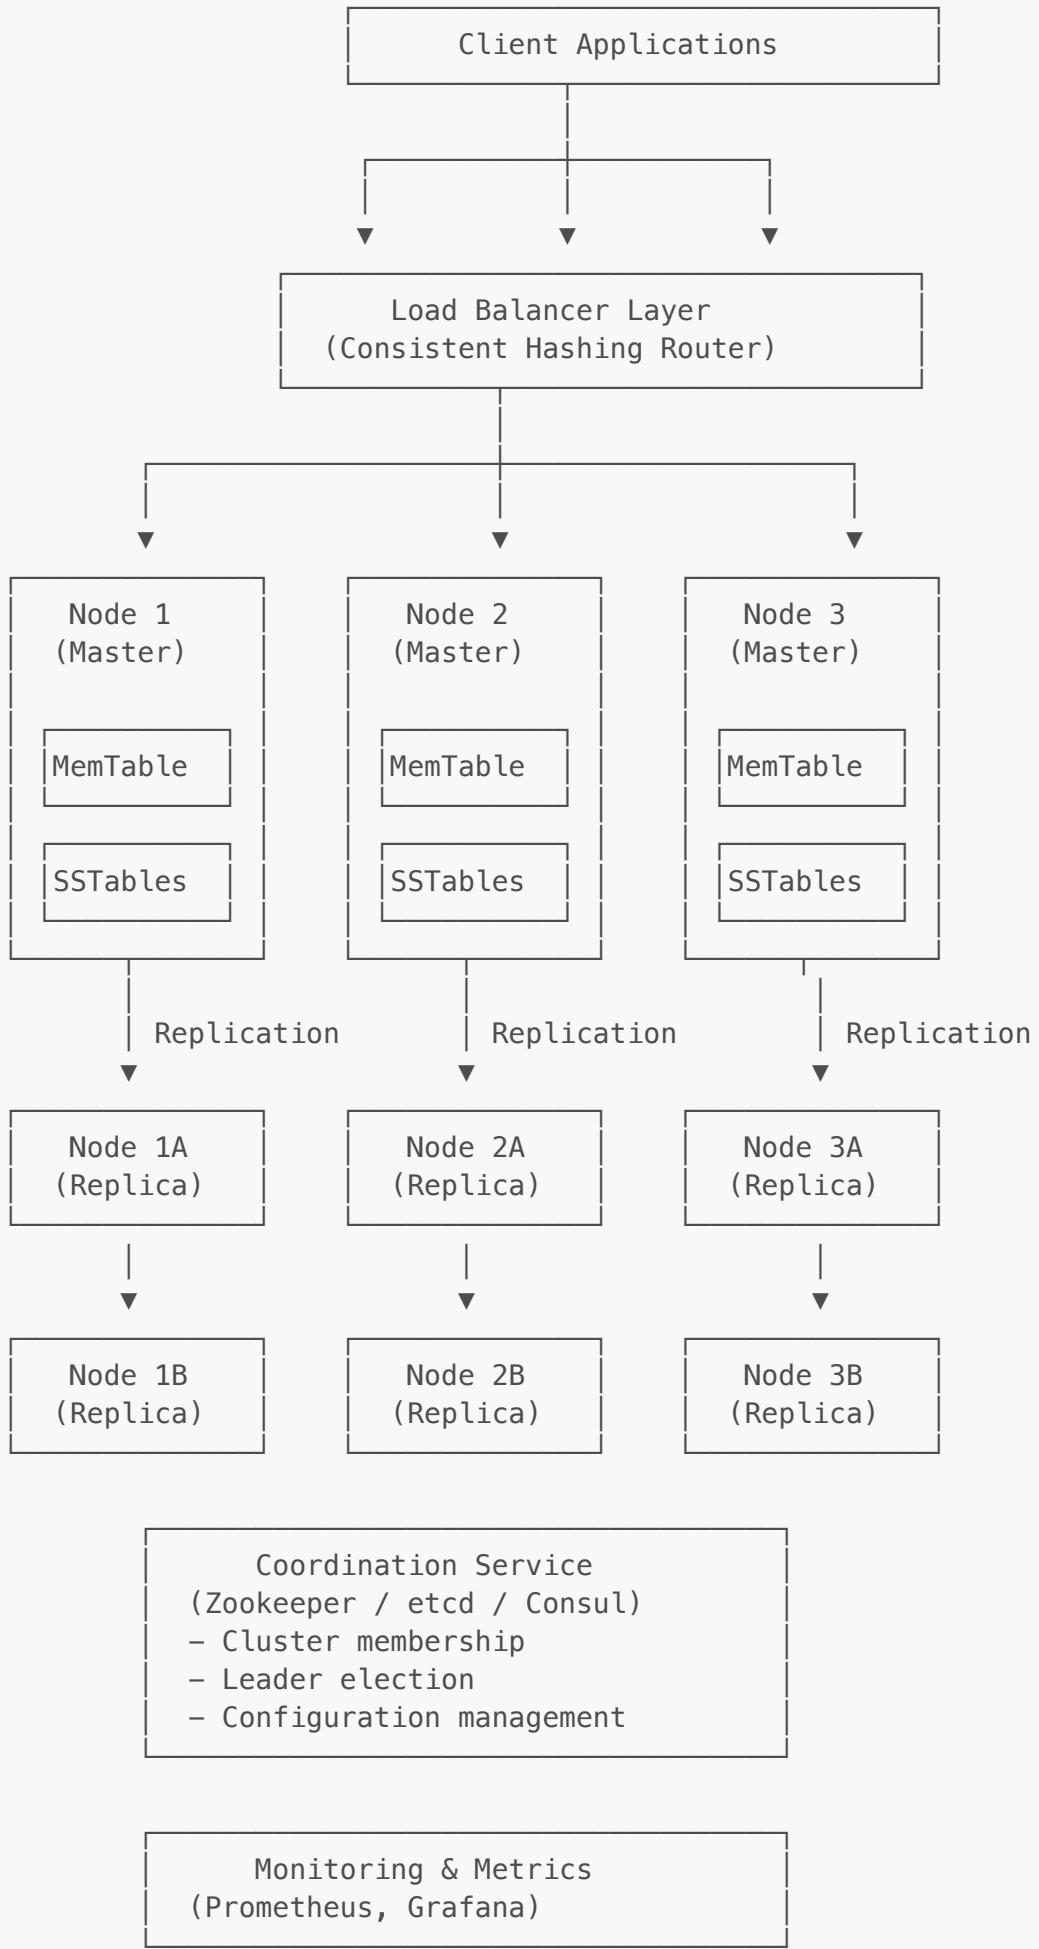
Entry 1: PUT user:123 "value1" v1 ts1
Entry 2: PUT user:456 "value2" v1 ts2
Entry 3: DELETE user:789 v2 ts3
Entry 4: PUT session:abc "data" v1 ts4
...
Checksum

```

Purpose:

- Durability (survive crashes)
- Replay on recovery
- Asynchronous flushing to disk

High-Level Architecture



Architecture Layers

1. Client Layer

- Client libraries (smart clients)
- Connection pooling
- Automatic retry logic
- Consistent hashing

2. Routing Layer

- Load balancers
- Request routing based on key
- Health checks
- Circuit breakers

3. Storage Node Layer

- Master nodes (handle reads/writes)
- Replica nodes (handle reads)
- Data persistence
- Replication coordination

4. Coordination Layer

- Cluster membership
- Leader election
- Configuration sync
- Failure detection

5. Monitoring Layer

- Metrics collection
- Alerting
- Dashboard
- Distributed tracing

Core Components

1. Storage Engine

Responsibilities:

- Store key-value pairs on disk
- Provide efficient reads/writes
- Handle compaction
- Manage cache

Implementation: LSM Tree

Components:

1. MemTable (In-Memory):
 - Sorted data structure (Skip List or Red-Black Tree)
 - Receives all writes
 - Fast writes ($O(\log n)$)
 - Limited size (64–128 MB)
2. Immutable MemTable:
 - Frozen when MemTable full
 - Background flush to disk
 - Read-only during flush
3. SSTable Files:
 - Immutable on-disk files
 - Sorted by key
 - Multiple levels (L0, L1, L2, ...)
 - Bloom filters for fast lookups
4. Write-Ahead Log (WAL):
 - Append-only log
 - Durability guarantee
 - Replayed on crash recovery

Operations:

Write Flow:

1. Append to WAL (sync to disk)
2. Write to MemTable
3. Return success
4. Background: Flush MemTable → SSTable

Read Flow:

1. Check MemTable
2. Check Immutable MemTable
3. Check SSTables (L0 → L1 → L2...)
4. Use Bloom filters to skip files
5. Return result or NOT_FOUND

Compaction:

- Merge overlapping SSTables
- Remove deleted keys
- Reduce read amplification
- Reclaim disk space

Alternative: B-Tree

Pros:

- ✓ Better read performance (lower read amplification)
- ✓ Simpler implementation
- ✓ In-place updates

Cons:

- ✗ Worse write performance (random I/O)
- ✗ Write amplification
- ✗ Fragmentation issues

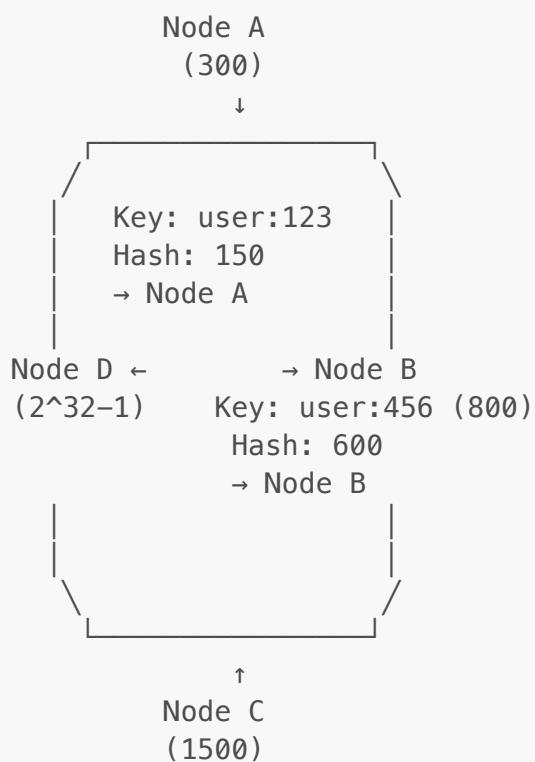
Use When:

- Read-heavy workload
- Limited write throughput
- Strong consistency required

2. Consistent Hashing

Purpose: Distribute keys evenly across nodes

Hash Ring (0 to $2^{32} - 1$):



Algorithm:

1. Hash each node (node_id) → position on ring
2. Hash each key → position on ring
3. Walk clockwise to find first node
4. Key stored on that node (+ replicas)

Virtual Nodes:

- Each physical node → multiple virtual nodes
- Better load distribution
- Easier rebalancing
- Example: 256 virtual nodes per physical node

Key Operations:

```

class ConsistentHash:
    def __init__(self, nodes, virtual_nodes=256):
        self.virtual_nodes = virtual_nodes
        self.ring = {}
        self.sorted_keys = []

    for node in nodes:
        self.add_node(node)

    def add_node(self, node):
        for i in range(self.virtual_nodes):
            virtual_key = f"{node}:{i}"
            hash_value = hash(virtual_key) % (2**32)
            self.ring[hash_value] = node

        self.sorted_keys = sorted(self.ring.keys())

    def get_node(self, key):
        if not self.ring:
            return None

        hash_value = hash(key) % (2**32)

        # Binary search for first node >= hash_value
        idx = bisect.bisect_left(self.sorted_keys, hash_value)

        # Wrap around if necessary
        if idx == len(self.sorted_keys):
            idx = 0

        return self.ring[self.sorted_keys[idx]]

    def get_replicas(self, key, replication_factor=3):
        replicas = []
        hash_value = hash(key) % (2**32)
        idx = bisect.bisect_left(self.sorted_keys, hash_value)

        # Get next N unique nodes
        seen_nodes = set()
        while len(seen_nodes) < replication_factor:
            if idx >= len(self.sorted_keys):
                idx = 0

            node = self.ring[self.sorted_keys[idx]]

```

```

        if node not in seen_nodes:
            seen_nodes.add(node)
            replicas.append(node)

    idx += 1

    return replicas

```

Benefits:

- Minimal data movement when nodes added/removed
 - Even distribution with virtual nodes
 - Automatic load balancing
 - Fault tolerance (replicas on different nodes)
-

3. Replication Coordinator

Responsibilities:

- Coordinate writes to replicas
- Handle consistency levels
- Detect and repair inconsistencies
- Manage replica placement

Replication Strategies:

1. Master-Slave Replication:

Client → Master (Write)
 Master → Replicas (Async)
 Client → Replicas (Read)

Pros: Simple, read scalability
 Cons: Write bottleneck, replication lag

2. Multi-Master Replication:

Client → Any Node (Write)
 Node → Other Nodes (Async)

Pros: Write scalability, no bottleneck
 Cons: Conflict resolution needed

3. Quorum-Based Replication:

R + W > N (R=reads, W=writes, N=replicas)

Example: N=3, W=2, R=2
 – Write to 2 out of 3 replicas

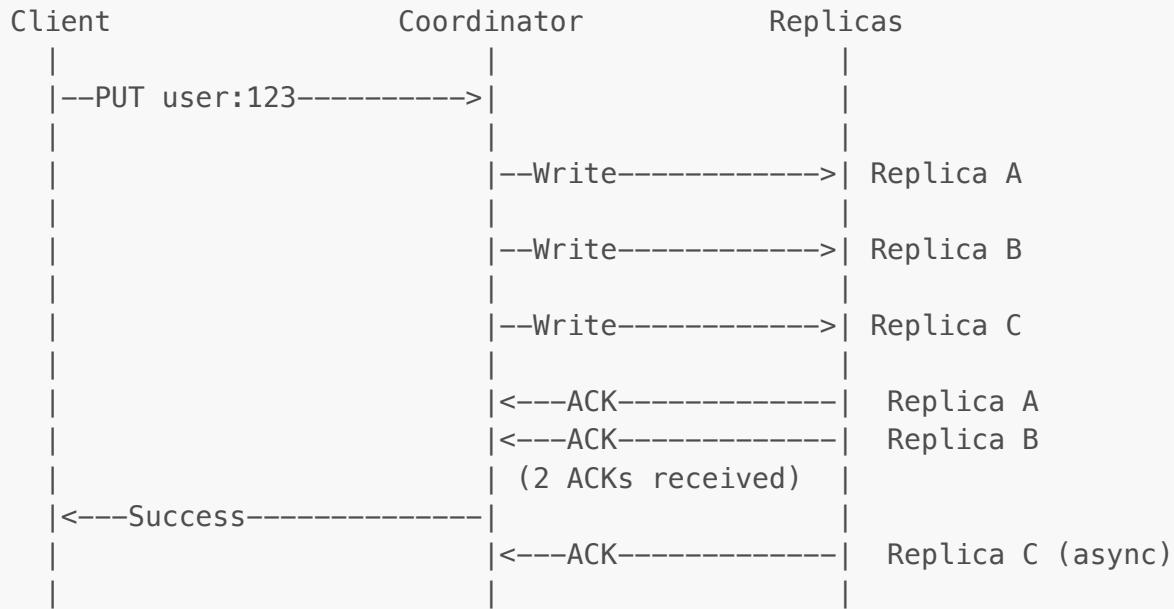
- Read from 2 out of 3 replicas
- Guaranteed to see latest write

Pros: Tunable consistency

Cons: Higher latency

Replication Protocol:

Write with Quorum ($W=2$, $N=3$):



Hinted Handoff (for failed replica):

- Coordinator detects Replica B down
- Writes hint to another node (Replica D)
- When Replica B recovers, replay hints
- Ensures eventual consistency

4. Gossip Protocol

Purpose: Disseminate cluster state without central coordinator

Gossip Protocol Flow:

Every node periodically (1 second):

1. Select random peer
2. Exchange state information
3. Merge states
4. Detect failures

Example:

Node A knows: {A:alive, B:alive, C:unknown}

Node B knows: {A:alive, B:alive, C:dead}

After gossip:

Node A knows: {A:alive, B:alive, C:dead}

Node B knows: {A:alive, B:alive, C:dead}

Information spreads exponentially:

- After $\log(N)$ rounds, all nodes know
- Self-healing on network partition
- No single point of failure

Implementation:

```
class GossipProtocol:  
    def __init__(self, node_id, peers):  
        self.node_id = node_id  
        self.peers = peers  
        self.state = {node_id: {'status': 'alive', 'version': 0}}  
  
    def gossip_round(self):  
        # Select random peer  
        peer = random.choice(self.peers)  
  
        # Send state to peer  
        peer_state = self.send_state(peer, self.state)  
  
        # Merge states  
        self.merge_state(peer_state)  
  
        # Increment local version  
        self.state[self.node_id]['version'] += 1  
  
    def merge_state(self, peer_state):  
        for node, info in peer_state.items():  
            if node not in self.state:  
                self.state[node] = info  
            elif info['version'] > self.state[node]['version']:  
                self.state[node] = info  
  
    def detect_failures(self):  
        current_time = time.time()  
        for node, info in self.state.items():  
            if current_time - info['last_seen'] > 10: # 10 sec timeout  
                info['status'] = 'suspected'  
            if current_time - info['last_seen'] > 30: # 30 sec timeout  
                info['status'] = 'dead'
```

5. Failure Detector

Responsibilities:

- Detect node failures
- Trigger automatic failover
- Prevent false positives (network delays)

Phi Accrual Failure Detector:

Concept:

- Track heartbeat intervals
- Calculate probability of failure (phi)
- Threshold-based decision

Algorithm:

1. Track last N heartbeat arrival times
2. Calculate mean and variance
3. Compute phi value:
 $\phi(t) = -\log_{10}(P(\text{arrival_time} > t))$

4. Decision:

- $\phi < 5$: Node likely alive
- $\phi > 8$: Node likely dead
- $5 < \phi < 8$: Suspicious

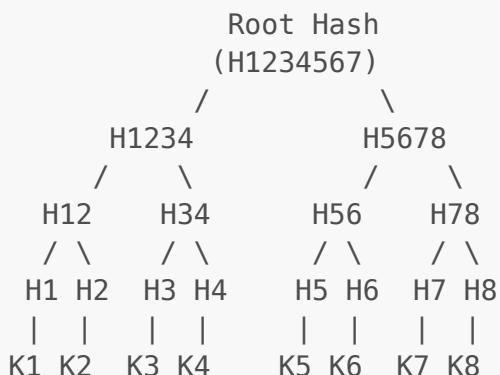
Advantages:

- Adaptive to network conditions
- Configurable threshold
- Low false positive rate

6. Anti-Entropy (Merkle Trees)

Purpose: Detect and repair inconsistencies between replicas

Merkle Tree Structure:



Each leaf = hash of key-value pair
Each parent = hash of children

Comparison Algorithm:

1. Exchange root hashes
2. If equal → replicas in sync
3. If different → recursively compare children
4. Identify differing keys
5. Sync only differing keys

Benefit: $O(\log n)$ comparisons instead of $O(n)$

Anti-Entropy Process:

Periodic background task (every hour):

1. Build Merkle tree for key range
2. Exchange tree with replica
3. Compare trees recursively
4. Identify inconsistent keys
5. Repair inconsistencies (read repair)

Benefits:

- Proactive inconsistency detection
- Efficient comparison ($O(\log n)$)
- Background operation (no latency impact)
- Self-healing system

Data Partitioning (Sharding)

Why Partition?

Problem: Single node limitations

- Storage: Limited disk space (single server ~10TB)
- Memory: Limited RAM for indexes
- CPU: Limited processing power
- Network: Single NIC bottleneck

Solution: Distribute data across multiple nodes

Partitioning Strategies

1. Hash-Based Partitioning

```
Partition = hash(key) % num_partitions
```

Example:

```
hash("user:123") = 8472
```

$8472 \% 10 = 2 \rightarrow$ Partition 2

Pros:

- ✓ Even distribution
- ✓ Simple implementation
- ✓ Good for random access

Cons:

- ✗ Poor range queries
- ✗ All partitions needed for scan

2. Range-Based Partitioning

Partitions by key ranges:

Partition 1: A–F

Partition 2: G–M

Partition 3: N–S

Partition 4: T–Z

Example:

"user:alice" → Partition 1

"user:john" → Partition 2

"user:tom" → Partition 4

Pros:

- ✓ Efficient range queries
- ✓ Prefix searches fast
- ✓ Good for time-series data

Cons:

- ✗ Hot spots (uneven distribution)
- ✗ Rebalancing complex

3. Consistent Hashing (Recommended)

Already covered in Core Components section

Partition Management

Token Assignment:

Each node owns token ranges:

Node 1: tokens 0 – 99

Node 2: tokens 100 – 199

Node 3: tokens 200 – 299

...

```
Token = hash(key) % total_tokens
```

Benefits:

- Fine-grained control
- Easy rebalancing (move tokens)
- Flexible scaling

Rebalancing Strategy:

Add New Node:

Before:

Node A: [0–199]
Node B: [200–399]
Node C: [400–599]

Add Node D:

After:

Node A: [0–149]
Node B: [200–349]
Node C: [400–549]
Node D: [150–199, 350–399, 550–599]

Process:

1. Allocate tokens to new node
2. Stream data from existing nodes
3. Verify checksums
4. Update routing table
5. Drain old data

Replication Strategy

Replication Factor

N = Replication Factor (typically 3)

Each key stored on N nodes:

- Primary node (coordinator)
- N-1 replica nodes

Example: N=3

Key "user:123" → Node A (primary)
→ Node B (replica)
→ Node C (replica)

Replica Placement

Strategy: Cross-Rack and Cross-DC

Rack-Aware Placement:

Datacenter 1:

- Rack 1: Node A (primary)
- Rack 2: Node B (replica)

Datacenter 2:

- Rack 3: Node C (replica)

Benefits:

- Survive rack failure
- Survive datacenter failure
- Local reads in each DC

Replication Protocols

1. Synchronous Replication

Write Flow:

1. Client → Coordinator
2. Coordinator → All replicas (parallel)
3. Wait for all ACKs
4. Coordinator → Client (success)

Pros:

- ✓ Strong consistency
- ✓ No data loss

Cons:

- ✗ High latency (slowest replica)
- ✗ Reduced availability (any failure blocks)

2. Asynchronous Replication

Write Flow:

1. Client → Primary
2. Primary → Client (success immediately)
3. Primary → Replicas (background)

Pros:

- ✓ Low latency

- ✓ High availability

Cons:

- ✗ Eventual consistency
- ✗ Potential data loss

3. Quorum-Based (Hybrid)

Tunable Consistency:

$N = 3$ (total replicas)

$W = 2$ (write quorum)

$R = 2$ (read quorum)

Rule: $R + W > N$

Write Flow:

1. Write to 2 out of 3 replicas
2. Return success when 2 ACK
3. Third replica updated async

Read Flow:

1. Read from 2 out of 3 replicas
2. Return newest version
3. Read repair if inconsistent

Consistency Levels:

- ONE: $R=1, W=1$ (eventual consistency)
- QUORUM: $R=2, W=2$ (balanced)
- ALL: $R=3, W=3$ (strong consistency)

Consistency Models

CAP Theorem

Can only guarantee 2 of 3:

- Consistency
- Availability
- Partition Tolerance

Partition Tolerance

← Always needed (network fails)

Choose one:



▼ ▼ ▼
Consistency Availability Both
(CP) (AP) (CA – impossible in distributed systems)

Examples:

CP: MongoDB, HBase, Redis

AP: Cassandra, DynamoDB, Riak

Consistency Levels

1. Strong Consistency

Definition: All readers see same value immediately after write

Implementation:

- Synchronous replication
- Read from all replicas
- Write to all replicas

Guarantee: $R + W > N$ and $W > N/2$

Example: $N=3$, $W=3$, $R=1$

- Must write to all 3 replicas
- Read from any 1 replica sees latest

Use Cases:

- Financial transactions
- Inventory management
- User authentication

2. Eventual Consistency

Definition: All replicas converge to same value eventually

Implementation:

- Asynchronous replication
- Read from any replica
- Write to any replica

Guarantee: No immediate guarantee

Convergence Time: Typically seconds to minutes

Use Cases:

- Social media feeds
- Product catalogs

- DNS
- CDN caching

3. Read-Your-Writes Consistency

Definition: User sees their own writes immediately

Implementation:

- Route reads to same replica as write
- Use session tokens
- Track write version numbers

Example:

1. User writes comment
2. User refreshes page
3. User sees their comment (even if replicas lag)

Use Cases:

- User profiles
- Shopping carts
- Comment systems

4. Monotonic Reads

Definition: Once you read version V, never see older than V

Implementation:

- Sticky sessions (same replica)
- Version vectors
- Causal consistency

Example:

1. Read version 5
2. Later read version 6 or 5 (never 4)

Use Cases:

- Time-series data
- Audit logs
- Message threads

Vector Clocks

Track causality across replicas:

Vector Clock: {NodeA: 3, NodeB: 2, NodeC: 1}

Comparison:

V1 = {A:2, B:1}

V2 = {A:1, B:2}

→ Concurrent (conflict)

V1 = {A:2, B:1}

V2 = {A:3, B:1}

→ V2 happens after V1

Conflict Resolution:

1. Detect concurrent writes
2. Keep both versions
3. Application resolves (last-write-wins, merge, etc.)

Failure Detection & Recovery

Types of Failures

1. Node Failure:

- Crash (process dies)
- Hardware failure
- OS freeze

2. Network Failure:

- Partition (split brain)
- Packet loss
- High latency

3. Disk Failure:

- Corruption
- Out of space
- I/O errors

Failure Detection Mechanisms

1. Heartbeat Protocol

Each node sends periodic heartbeats:

Node A → [heartbeat] → Node B (every 1 second)

Node B tracks:

- Last heartbeat time
- Missed heartbeat count

Detection:

```
If no heartbeat for 10 seconds:  
    → Mark as suspected  
If no heartbeat for 30 seconds:  
    → Mark as dead  
    → Trigger failover
```

2. Phi Accrual Failure Detector

Already covered in Core Components

Recovery Strategies

1. Replica Takeover

Scenario: Primary node fails

Before:

Primary (Node A) → Replica (Node B)
→ Replica (Node C)

After Detection:

1. Coordinator detects Node A down
2. Promotes Node B to primary
3. Node B serves reads/writes
4. When Node A recovers:
 - Becomes replica
 - Catches up via anti-entropy

Time to Failover: < 30 seconds

2. Hinted Handoff

Scenario: Replica temporarily unavailable

Normal:

Client → Coordinator → Replica A ✓
→ Replica B ✗ (down)
→ Replica C ✓

With Hinted Handoff:

Client → Coordinator → Replica A ✓
→ Replica D ✓ (hint for B)
→ Replica C ✓

When Replica B recovers:

Replica D → Replays hints → Replica B

Benefits:

- Maintains write availability
- Ensures eventual consistency
- No data loss

3. Read Repair

Scenario: Inconsistent replicas detected during read

Read Request (R=2):

Client → Coordinator → Replica A (v5)
→ Replica B (v3)
→ Replica C (not queried)

Process:

1. Return v5 to client (newest)
2. Background: Update B to v5
3. Background: Check C and update if needed

Benefits:

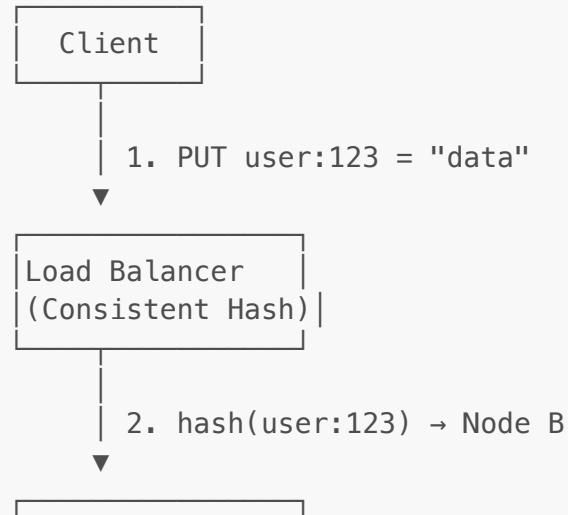
- Passive repair during reads
- Self-healing
- No extra operations

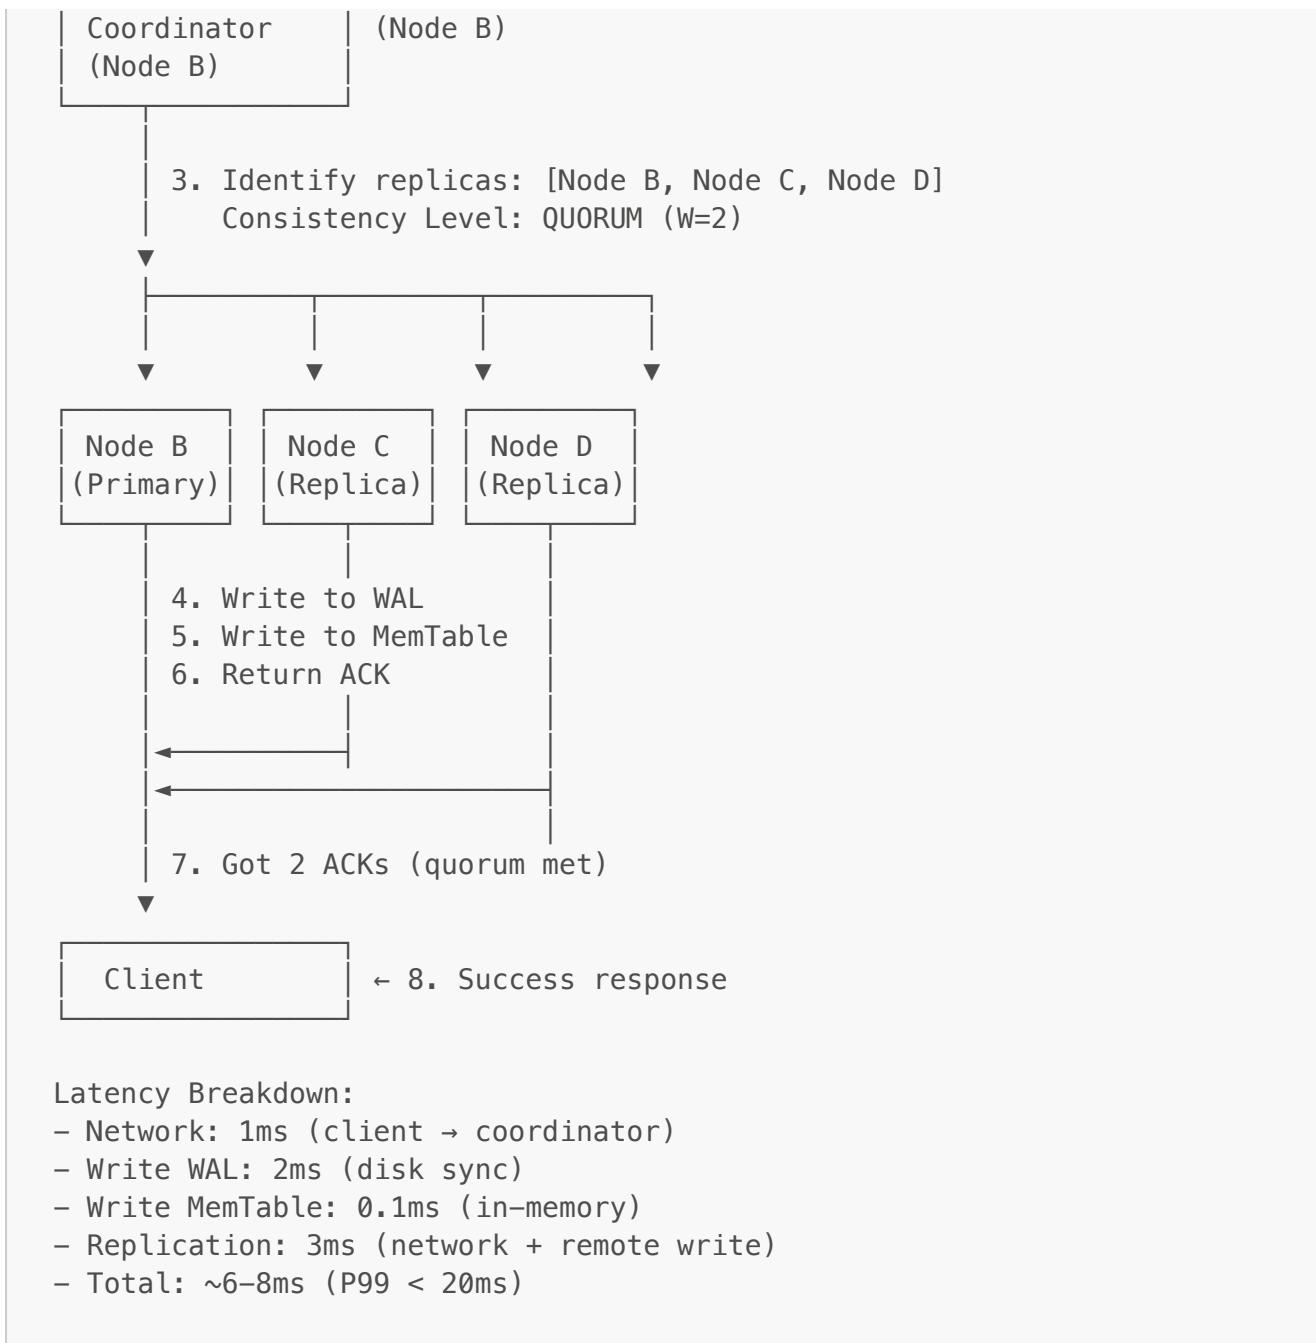
4. Anti-Entropy Repair

Already covered in Core Components (Merkle Trees)

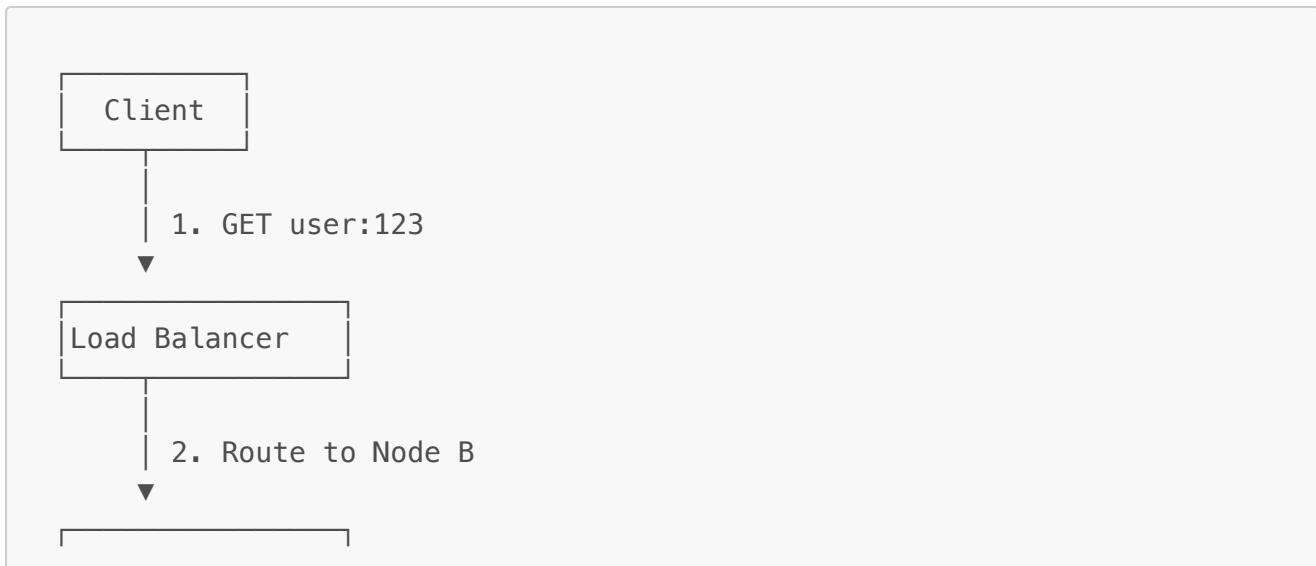
Read/Write Path

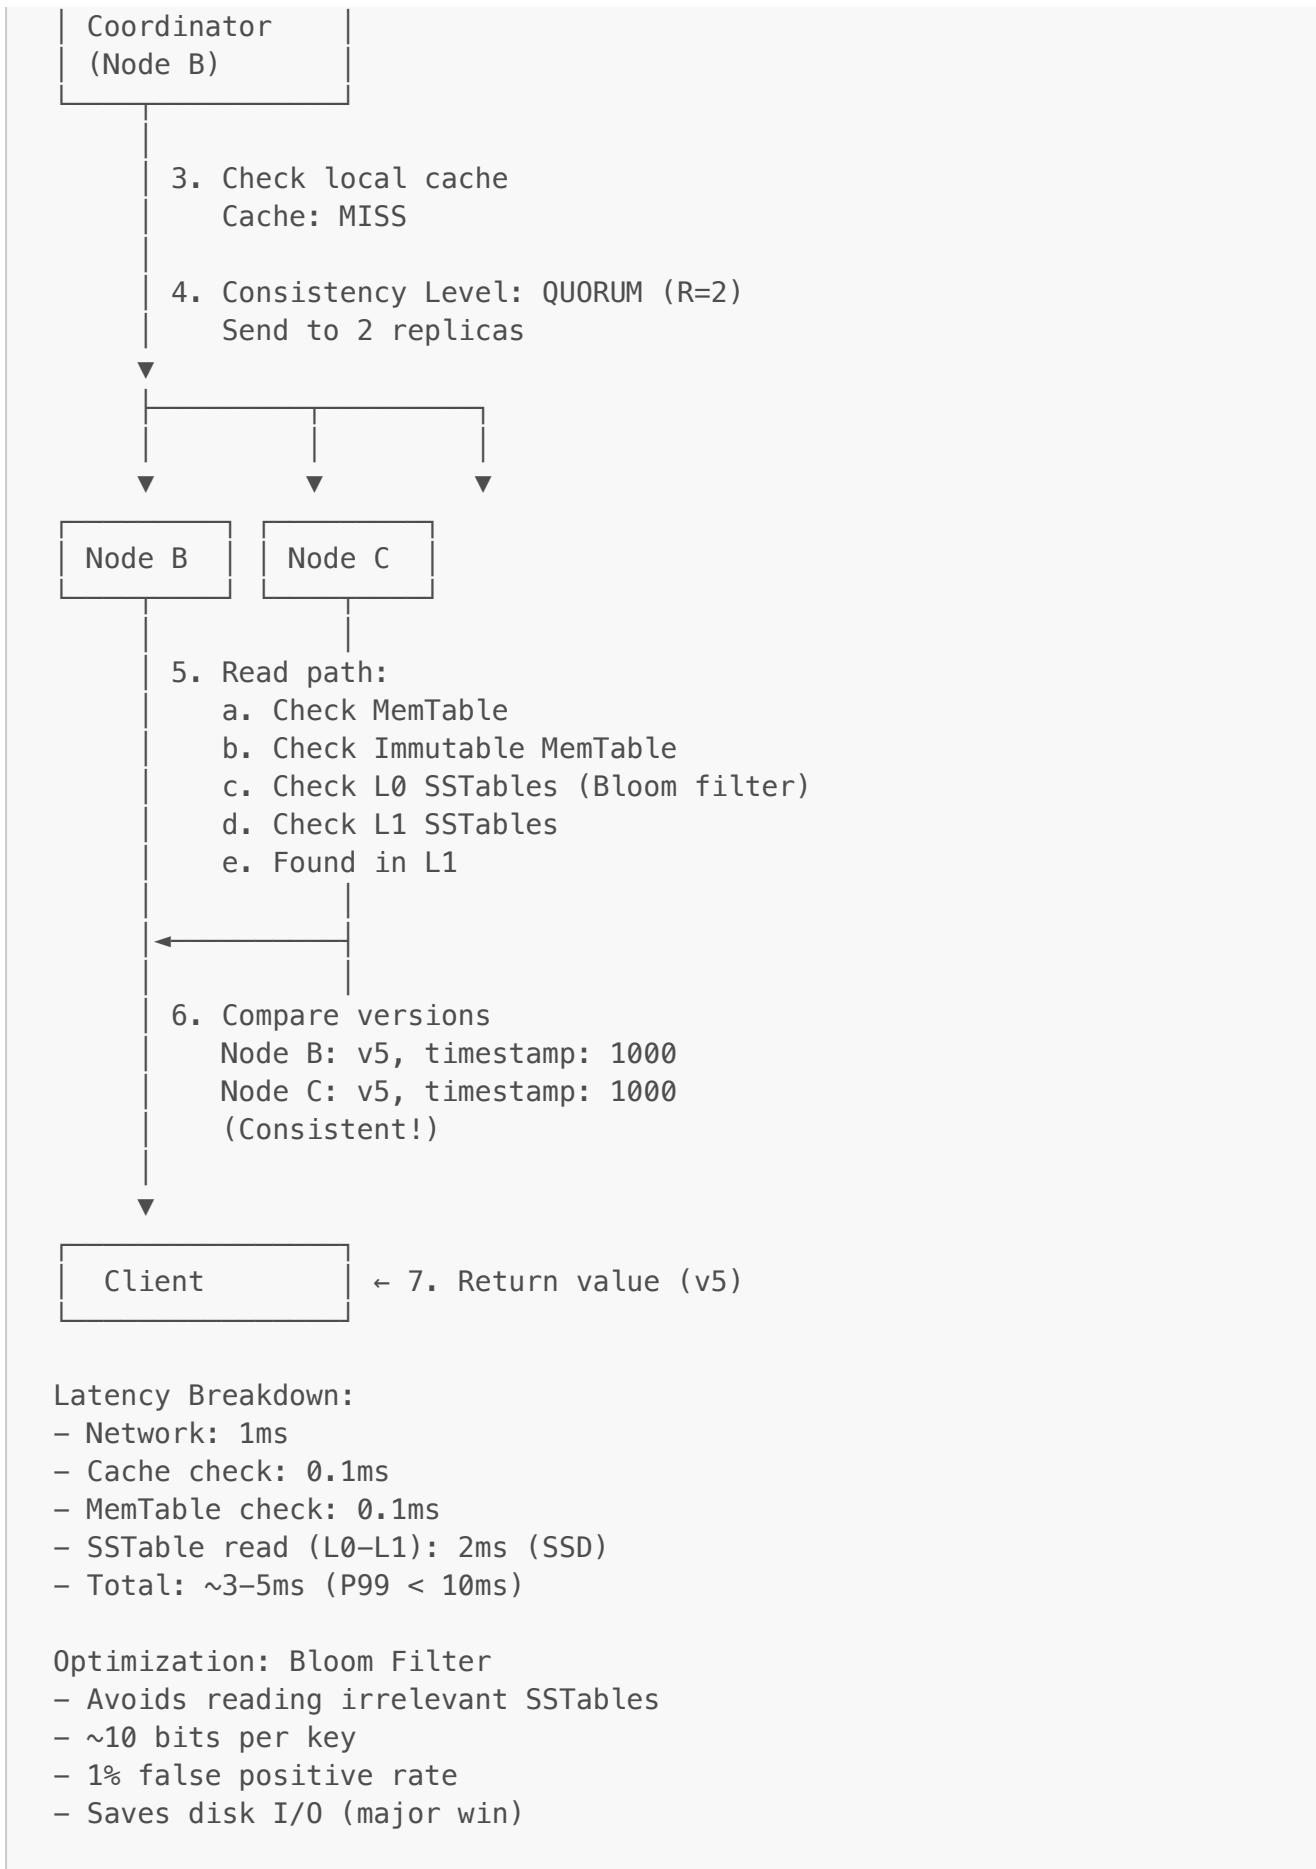
Write Path (Detailed)





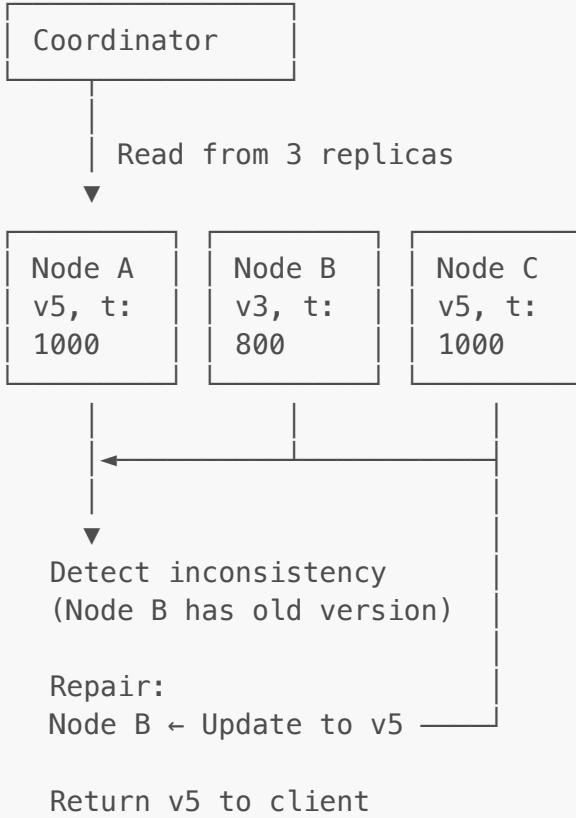
Read Path (Detailed)





Read Repair Example

Read with Inconsistency:



Performance Optimization

1. Caching Strategy

Multi-Layer Cache:

L1: In-Memory Cache (Node)
– Hot keys (20% of data)
– LRU eviction
– Hit rate: 85%
– Latency: 0.1ms

↓ (cache miss)

L2: SSTable Block Cache
– Recently read blocks
– 10GB per node
– Hit rate: 60%
– Latency: 1ms

↓ (cache miss)

- | |
|--|
| <ul style="list-style-type: none"> L3: Disk (SSTables) - All data - SSD - Latency: 2–5ms |
|--|

2. Bloom Filters

Purpose: Avoid unnecessary disk reads

Structure:

- Bit array (10 bits per key)
- Multiple hash functions (3–5)
- Probabilistic (false positives possible)

Check if key exists:

1. Hash key with H1, H2, H3
2. Check bits at positions
3. If all bits = 1 → Maybe exists
4. If any bit = 0 → Definitely NOT exists

Benefit:

- 99% of non-existent keys filtered
- Saves disk I/O
- Small memory overhead (1.2MB per 1M keys)

3. Compaction Strategies

Size-Tiered Compaction:

- Merge SSTables of similar size
- Good for write-heavy workloads
- Higher read amplification

Leveled Compaction:

- Each level 10x previous
- No overlap within level
- Better for read-heavy workloads
- Lower space amplification

Time-Window Compaction:

- Merge based on time ranges
- Good for time-series data
- Easy to expire old data

4. Batch Operations

Benefits:

- Amortize network overhead
- Reduce round trips
- Higher throughput

Example:

Single writes: $100 \text{ ops} \times 5\text{ms} = 500\text{ms}$
Batch write (100): $1 \text{ op} \times 15\text{ms} = 15\text{ms}$

Improvement: 33x faster

5. Connection Pooling

Problem: Creating connections expensive (10–50ms)

Solution:

```
Client Pool
Min: 10 conns
Max: 100 conns
Idle timeout:
60s
```

Benefits:

- Reuse connections
- Lower latency
- Reduced server load

6. Compression

Value Compression:

- Snappy (fast, 2–4x compression)
- LZ4 (faster, 2–3x compression)
- Zstandard (balanced, 3–5x compression)

Trade-off:

- CPU overhead vs storage savings
- Latency vs throughput

Recommended: LZ4 for hot data, Zstd for cold data

7. Read-Ahead

Prefetch adjacent keys:

User requests: user:123

Prefetch: user:124, user:125, user:126

Benefits:

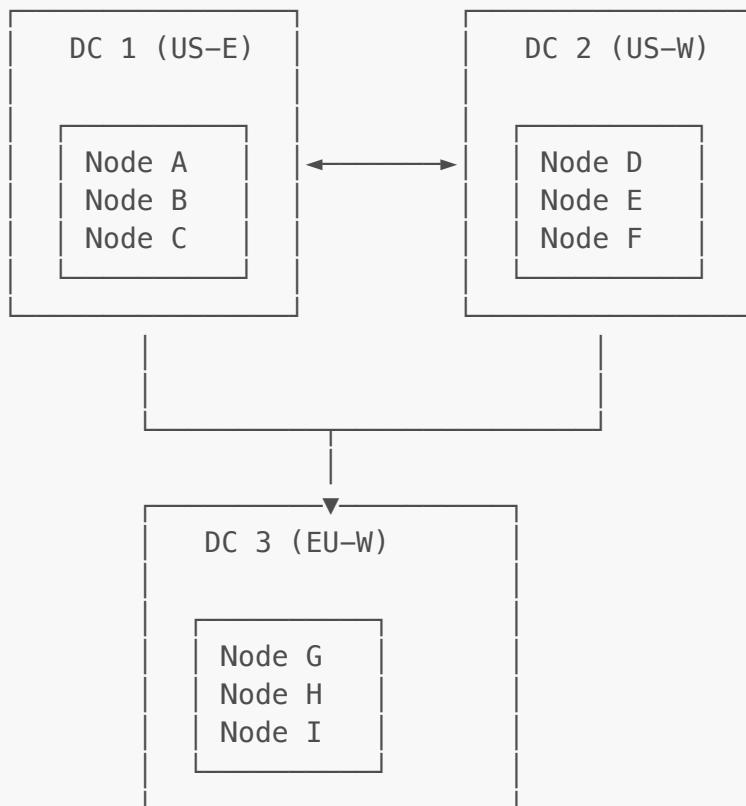
- Faster sequential reads
- Better cache utilization
- Anticipates access patterns

Use Cases:

- Range queries
- Scan operations
- Time-series data

Multi-Datacenter Deployment

Architecture



Cross-DC Replication:

- Async replication (10–100ms lag)
- Conflict resolution with vector clocks
- Per-DC quorum for availability

Replication Strategies

1. Active-Passive

Primary DC (Active):

- Handles all writes
- Replicates to passive DCs

Secondary DCs (Passive):

- Read-only
- Failover destination

Pros:

- ✓ No conflicts
- ✓ Simpler consistency

Cons:

- ✗ Write latency for remote users
- ✗ Underutilized capacity

2. Active-Active (Multi-Master)

All DCs Accept Writes:

- Users write to nearest DC
- Low latency
- Async replication between DCs

Conflict Resolution:

1. Last-Write-Wins (timestamp)
2. Application-specific merge
3. Vector clocks

Pros:

- ✓ Low latency globally
- ✓ Better utilization

Cons:

- ✗ Conflict resolution complexity
- ✗ Eventual consistency

Geo-Replication

Strategy: Place replicas globally

US-East: Primary + 2 replicas

US-West: 2 replicas

Europe: 2 replicas
Asia: 2 replicas

Read Strategy:

- Route to nearest DC
- Latency: <50ms globally

Write Strategy:

- Write to local DC
- Async replicate globally
- Quorum per DC (local quorum)

Network Partition Handling

Scenario: DC1 and DC2 partitioned

Option 1: Majority Quorum

- Require majority of DCs to ack
- DC with majority continues
- Minority becomes read-only

Option 2: Per-DC Quorum

- Each DC operates independently
- Merge conflicts on heal
- Higher availability

Option 3: Designated Tie-Breaker

- Third DC decides in split
- Prevents split-brain
- Single point of failure

Design Trade-offs

1. Consistency vs Availability

Strong Consistency:

Pros:

- ✓ Simple programming model
- ✓ No conflicts
- ✓ Predictable behavior

Cons:

- ✗ Lower availability (CAP)
- ✗ Higher latency
- ✗ Requires synchronous replication

Use When:

- Financial transactions
- Inventory counts
- User authentication

Eventual Consistency:

Pros:

- ✓ High availability
- ✓ Low latency
- ✓ Geo-replication friendly

Cons:

- ✗ Complex programming model
- ✗ Conflict resolution needed
- ✗ Stale reads possible

Use When:

- Social media
- Product catalogs
- Analytics

2. LSM Tree vs B-Tree

LSM Tree:

Pros:

- ✓ Write optimized (sequential I/O)
- ✓ Better write throughput
- ✓ Lower write amplification

Cons:

- ✗ Read amplification (check multiple levels)
- ✗ Compaction overhead
- ✗ Space amplification during compaction

Use When:

- Write-heavy workload
- Time-series data
- Logging systems

B-Tree:

Pros:

- ✓ Read optimized (single lookup)
- ✓ No compaction needed
- ✓ Predictable performance

Cons:

- ✗ Random I/O for writes
- ✗ Write amplification
- ✗ Fragmentation

Use When:

- Read-heavy workload
- Point queries dominant
- Storage efficiency critical

3. Replication Factor

N=2:

- Pros: Lower storage cost, faster writes
 Cons: Lower availability, single failure dangerous

N=3 (Recommended):

- Pros: Good balance, survive 1 failure
 Cons: 3x storage cost

N=5:

- Pros: High durability, survive 2 failures
 Cons: High cost, slower writes

Decision Factors:

- Data criticality
- Cost constraints
- Availability requirements

4. Partition Size

Small Partitions (1GB):

Pros:

- ✓ Fine-grained rebalancing
- ✓ Faster recovery
- ✓ Better load distribution

Cons:

- ✗ More metadata overhead
- ✗ More coordination
- ✗ Higher operational complexity

Large Partitions (100GB):

Pros:

- ✓ Less metadata
- ✓ Simpler operations
- ✓ Fewer network hops

Cons:

- ✗ Coarse rebalancing
- ✗ Slower recovery
- ✗ Hot spots

Real-World Implementations

1. Amazon DynamoDB

Architecture:

- Multi-master, masterless design
- Consistent hashing for partitioning
- Quorum-based replication (N=3)
- Vector clocks for conflict resolution

Key Features:

- Fully managed service
- Automatic scaling
- Global tables (multi-region)
- Point-in-time recovery

Consistency Model:

- Eventual consistency (default)
- Strong consistency (optional)
- Transactional support (ACID)

Scale:

- Trillions of requests/day
- Millisecond latency at any scale
- Petabytes of data

2. Apache Cassandra

Architecture:

- Ring-based peer-to-peer
- No single point of failure
- Tunable consistency (ONE to ALL)
- Gossip protocol for membership

Key Features:

- Linear scalability
- Multi-datacenter support
- CQL (SQL-like query language)
- Time-series optimized

Consistency Model:

- Eventual consistency
- Tunable per-operation
- Lightweight transactions (Paxos)

Use Cases:

- Apple: 75,000+ nodes
- Netflix: Billions of operations/day
- Instagram: Storing user data

3. Redis

Architecture:

- In-memory data store
- Single-threaded (per core)
- Master-replica replication
- Redis Cluster for sharding

Key Features:

- Sub-millisecond latency
- Rich data structures (lists, sets, sorted sets)
- Pub/Sub messaging
- Lua scripting

Persistence:

- RDB (snapshots)
- AOF (append-only file)
- Hybrid (RDB + AOF)

Use Cases:

- Caching layer
- Session store
- Real-time analytics
- Message queuing

4. etcd

Architecture:

- Raft consensus algorithm
- Strong consistency (CP system)
- Leader-based replication
- Key-value with versioning

Key Features:

- Distributed coordination
- Watch mechanism (notifications)
- Lease/TTL support
- Transaction support

Use Cases:

- Kubernetes: Cluster state
 - Service discovery
 - Configuration management
 - Distributed locking
-

Summary & Key Takeaways

Core Design Principles

1. Horizontal Scalability

- Use consistent hashing for partitioning
- Virtual nodes for better distribution
- Stateless coordinators

2. High Availability

- Replication (N=3 recommended)
- No single point of failure
- Automatic failover

3. Fault Tolerance

- Gossip protocol for membership
- Hinted handoff for temporary failures
- Anti-entropy for long-term consistency

4. Performance

- LSM tree for write optimization
- Bloom filters to reduce disk I/O
- Multi-layer caching
- Batch operations

5. Tunability

- Configurable consistency levels
- Flexible replication strategies
- Adjustable partition sizes

When to Use Key-Value Stores

✓ Good Fit:

- High throughput requirements (100K+ ops/sec)
 - Simple data model (key-value pairs)
 - Horizontal scalability needed
 - High availability critical
-

- Global distribution required

X Poor Fit:

- Complex transactions (use RDBMS)
- Complex queries/joins (use RDBMS)
- Strong consistency required everywhere (use RDBMS)
- Small scale (<1TB, single server sufficient)

Interview Success Tips

1. Start with Requirements

- Scale (TB vs PB)
- Consistency needs
- Latency requirements
- Read/write ratio

2. Discuss Trade-offs

- CAP theorem
- Consistency vs availability
- Write vs read optimization
- Cost vs performance

3. Draw Clear Diagrams

- High-level architecture
- Data flow
- Partition distribution

4. Reference Real Systems

- "Similar to DynamoDB's approach..."
- "Cassandra solves this with..."
- Shows awareness of production systems

5. Cover Failure Scenarios

- Node failures
- Network partitions
- Data center outages

6. Mention Monitoring

- Key metrics to track
- Alerting strategies
- Operational concerns

Common Pitfalls to Avoid

- ✖ Over-engineering for day one
 - ✖ Ignoring CAP theorem
 - ✖ Not considering hot keys
 - ✖ Forgetting about operational complexity
 - ✖ Underestimating network effects
 - ✖ Not planning for data migration
 - ✖ Ignoring monitoring and observability
-

Further Reading

Essential Papers

- **Dynamo: Amazon's Highly Available Key-value Store** (2007)
- **Cassandra - A Decentralized Structured Storage System** (2010)
- **Bigtable: A Distributed Storage System for Structured Data** (Google, 2006)
- **The Log-Structured Merge-Tree (LSM-Tree)** (1996)

Books

- **Designing Data-Intensive Applications** - Martin Kleppmann
- **Database Internals** - Alex Petrov
- **NoSQL Distilled** - Pramod Sadalage & Martin Fowler

Online Resources

- DynamoDB Paper: aws.amazon.com/dynamodb/
 - Cassandra Documentation: cassandra.apache.org
 - Redis Documentation: redis.io
 - etcd Documentation: etcd.io
-

Document Version: 1.0 (High-Level Design)

Last Updated: 2025

**Focus