

# Global File Storage & Sharing System - High-Level Design (HLD)

---

---

## Table of Contents

1. Problem Statement
  2. Requirements
  3. Capacity Estimation
  4. Back-of-Envelope Calculations
  5. High-Level Architecture
  6. Core Components
  7. Data Flow
  8. Database Design
  9. Design Decisions & Trade-offs
  10. Scalability Strategy
  11. Reliability & Availability
  12. Security
- 

## Problem Statement

Design a scalable, reliable, globally distributed file storage and sharing system like Google Drive or Dropbox.

## Core Functionality

- Upload/download files of any size (up to 10GB)
- Share files with other users (permissions management)
- Access files from multiple devices
- Real-time synchronization across devices
- Search files by name and content
- Version history (last 30 days)

## Target Scale

- **1 Billion** total users
  - **200 Million** daily active users (DAU)
  - **10 PB** total storage
  - **99.99%** availability
  - Global distribution
- 

## Requirements

### Functional Requirements

1. **File Operations:** Upload, download, delete, update files
2. **Folder Management:** Create folders, organize hierarchy
3. **Sharing:** Share files/folders with users or via public links
4. **Synchronization:** Real-time sync across devices
5. **Search:** Find files by name, content, metadata
6. **Version Control:** Maintain and restore file versions
7. **Permissions:** Owner, editor, viewer access levels

## Non-Functional Requirements

1. **Scalability:** Handle billions of users, petabytes of data
  2. **Availability:** 99.99% uptime (4 nines)
  3. **Consistency:** Strong for metadata, eventual for sync
  4. **Performance:**
    - Upload/Download: < 100ms latency (small files)
    - Sync latency: < 5 seconds
  5. **Durability:** 99.999999999% (11 nines)
  6. **Security:** Encryption at rest and in transit, access control
- 

## Capacity Estimation

### Scale Assumptions

#### Users:

- Total Users: 1 Billion
- Daily Active Users (DAU): 200 Million (20%)
- Average Storage per User: 10 GB

#### Files:

- Average File Size: 2 MB
- Files per User: 5,000
- Total Files: 5 Trillion files

Read/Write Ratio: 3:1 (more downloads than uploads)

## Storage Requirements

Total Storage = 1B users × 10 GB = 10 PB  
With 3x replication = 30 PB

## Traffic Estimation

### Upload Traffic:

- Daily uploads:  $200\text{M users} \times 5 \text{ files} = 1\text{B files/day}$
- Upload data:  $1\text{B} \times 2 \text{ MB} = 2 \text{ PB/day}$
- Upload bandwidth:  $\sim 23 \text{ GB/sec}$

### Download Traffic:

- $3x$  uploads =  $\sim 69 \text{ GB/sec}$

Total Peak Bandwidth:  $\sim 100 \text{ GB/sec}$

## QPS (Queries Per Second)

### Total Daily Requests:

- Uploads:  $1\text{B requests}$
- Downloads:  $3\text{B requests}$
- Metadata operations:  $2\text{B requests}$
- Total:  $6\text{B requests/day}$

Average QPS:  $\sim 70,000$

Peak QPS ( $3x$ ):  $\sim 200,000$

## Memory for Caching

### Metadata Cache:

- 1 KB per file metadata
- Cache 20% hot files =  $1\text{T files}$
- Required cache:  $\sim 1 \text{ TB}$

## Back-of-Envelope Calculations

### Latency Numbers Every Programmer Should Know

L1 cache reference	0.5 ns		
Branch mispredict	5 ns		
L2 cache reference	7 ns	14x L1 cache	
Mutex lock/unlock	25 ns		
Main memory reference	100 ns	20x L2, 200x L1	
Compress 1K with Snappy	10,000 ns	10 µs	
Send 1 KB over 1 Gbps network	10,000 ns	10 µs	
Read 4 KB randomly from SSD	150,000 ns	150 µs	~1GB/sec SSD
Read 1 MB sequentially from memory	250,000 ns	250 µs	
Round trip within datacenter	500,000 ns	500 µs	
Read 1 MB sequentially from SSD	1,000,000 ns	1000 µs	1 ms
Disk seek	10,000,000 ns	10000 µs	10 ms

```
Read 1 MB sequentially from disk 30,000,000 ns 30000 µs 30 ms
Send packet US→Europe→US           150,000,000 ns   150 ms
```

## Powers of Two Table

Power	Exact Value	Approx Value	Bytes
10	1,024	~1 thousand	1 KB
16	65,536	~65 thousand	64 KB
20	1,048,576	~1 million	1 MB
30	1,073,741,824	~1 billion	1 GB
40	1,099,511,627,776	~1 trillion	1 TB
50	~1 quadrillion		1 PB

## Server Capacity Estimation

### Given:

- Upload bandwidth needed: 23 GB/sec
- Assume each server handles: 1 GB/sec network throughput
- Peak load factor: 3x average

### Calculation:

- Average servers =  $23 \text{ GB/sec} \div 1 \text{ GB/sec} = 23 \text{ servers}$
- Peak servers =  $23 \times 3 = 69 \text{ servers}$
- With redundancy ( $N+1$ ) = ~70–80 servers

### For global deployment:

- $3 \text{ regions} \times 80 \text{ servers} = 240 \text{ servers total}$

## Database Sizing

Metadata per file: ~1 KB

- `file_id`: 8 bytes
- `owner_id`: 8 bytes
- `name`: 255 bytes
- `size`: 8 bytes
- `timestamps`: 24 bytes
- `path`: 500 bytes
- other fields: ~200 bytes

Total: ~1 KB

### For 5 Trillion files:

- Storage needed =  $5T \times 1 \text{ KB} = 5 \text{ PB metadata}$
- With indexing (2x) = 10 PB
- Sharded across 1000 DB servers = 10 GB per server

## Network Bandwidth Analysis

### Upload Scenario:

- 200M users upload 5 files/day
- Average file size: 2 MB
- Total data:  $200M \times 5 \times 2 \text{ MB} = 2 \text{ PB/day}$
- Per second:  $2 \text{ PB} / 86,400 \text{ sec} \approx 23 \text{ GB/sec}$
- With peaks (3x):  $\sim 70 \text{ GB/sec}$

### Download Scenario (3:1 read/write ratio):

- Download traffic:  $23 \times 3 = 69 \text{ GB/sec}$
- With peaks:  $\sim 210 \text{ GB/sec}$

### Total bandwidth requirement:

- Average:  $\sim 100 \text{ GB/sec}$
- Peak:  $\sim 300 \text{ GB/sec}$

## Cost Estimation (Rough)

### Storage Costs (AWS S3):

- $10 \text{ PB} \times \$23/\text{TB/month} = \$230,000/\text{month}$
- With replication (3x):  $\sim \$700,000/\text{month}$

### Bandwidth Costs:

- $2 \text{ PB upload/day} \times 30 \text{ days} = 60 \text{ PB/month}$
- $6 \text{ PB download/day} \times 30 \text{ days} = 180 \text{ PB/month}$
- At  $\$0.09/\text{GB}$ :  $\sim \$21\text{M}/\text{month}$  for bandwidth

### Database:

- $1000 \text{ instances} \times \$500/\text{month} = \$500,000/\text{month}$

Total Monthly Cost:  $\sim \$22\text{M}/\text{month}$

Annual Cost:  $\sim \$260\text{M}/\text{year}$

Per User Cost:  $\$260\text{M} \div 1\text{B users} = \$0.26/\text{user/year}$

## Availability Calculation

### Target: 99.99% (Four 9s)

### Downtime allowed per year:

- $365 \text{ days} \times 24 \text{ hours} \times 60 \text{ min} \times (1 - 0.9999)$
- $= 52.56 \text{ minutes/year}$
- $\approx 4.38 \text{ minutes/month}$

Component availability in sequence:

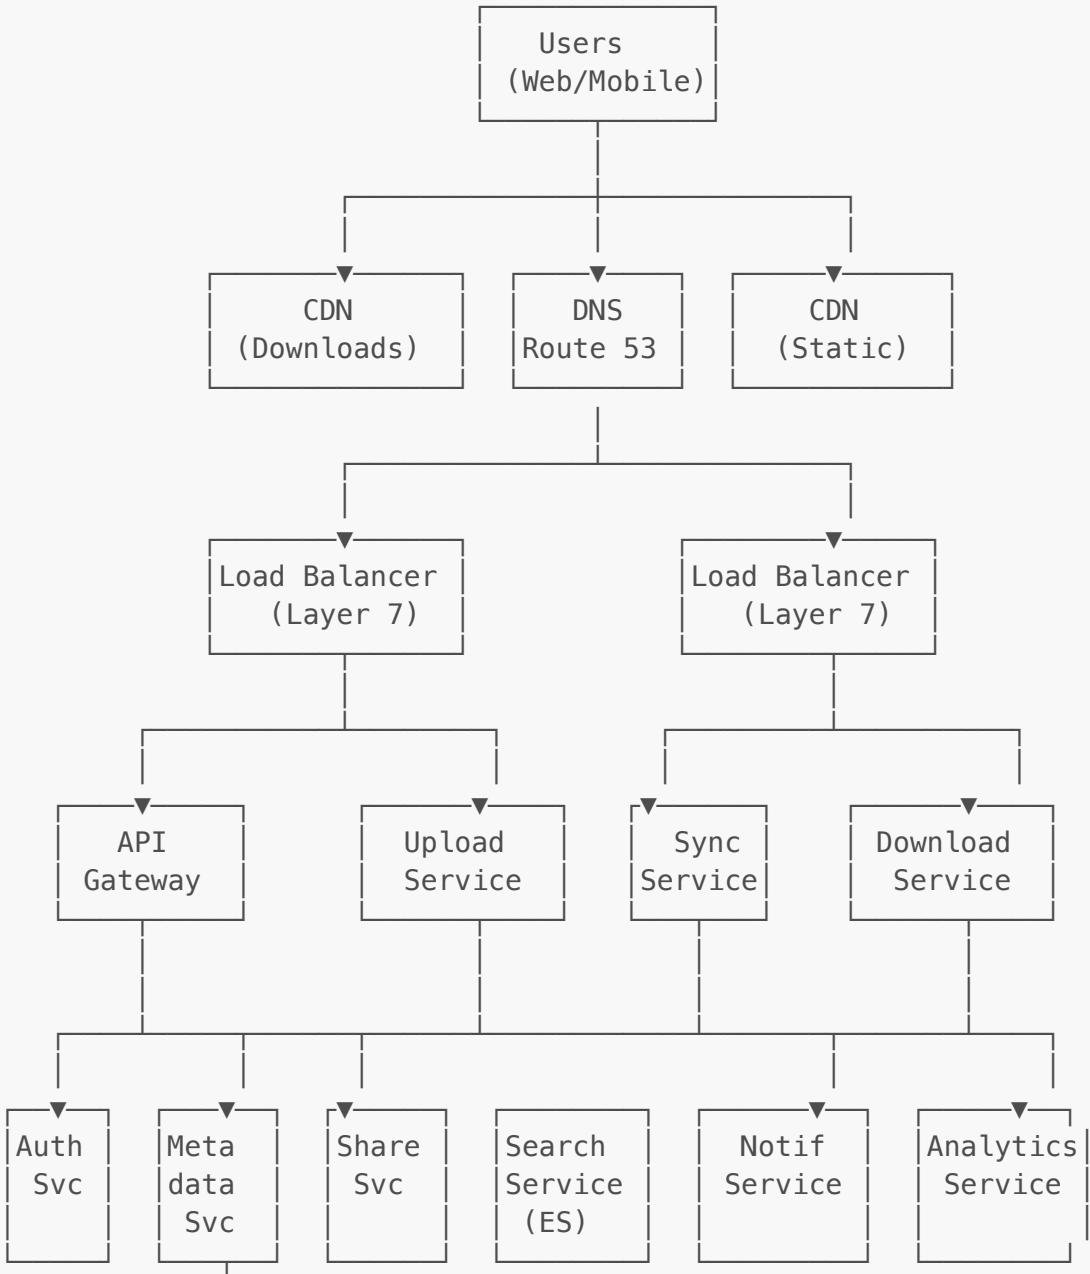
- Load Balancer: 99.99%
- Application: 99.99%
- Database: 99.99%
- Object Storage: 99.99%

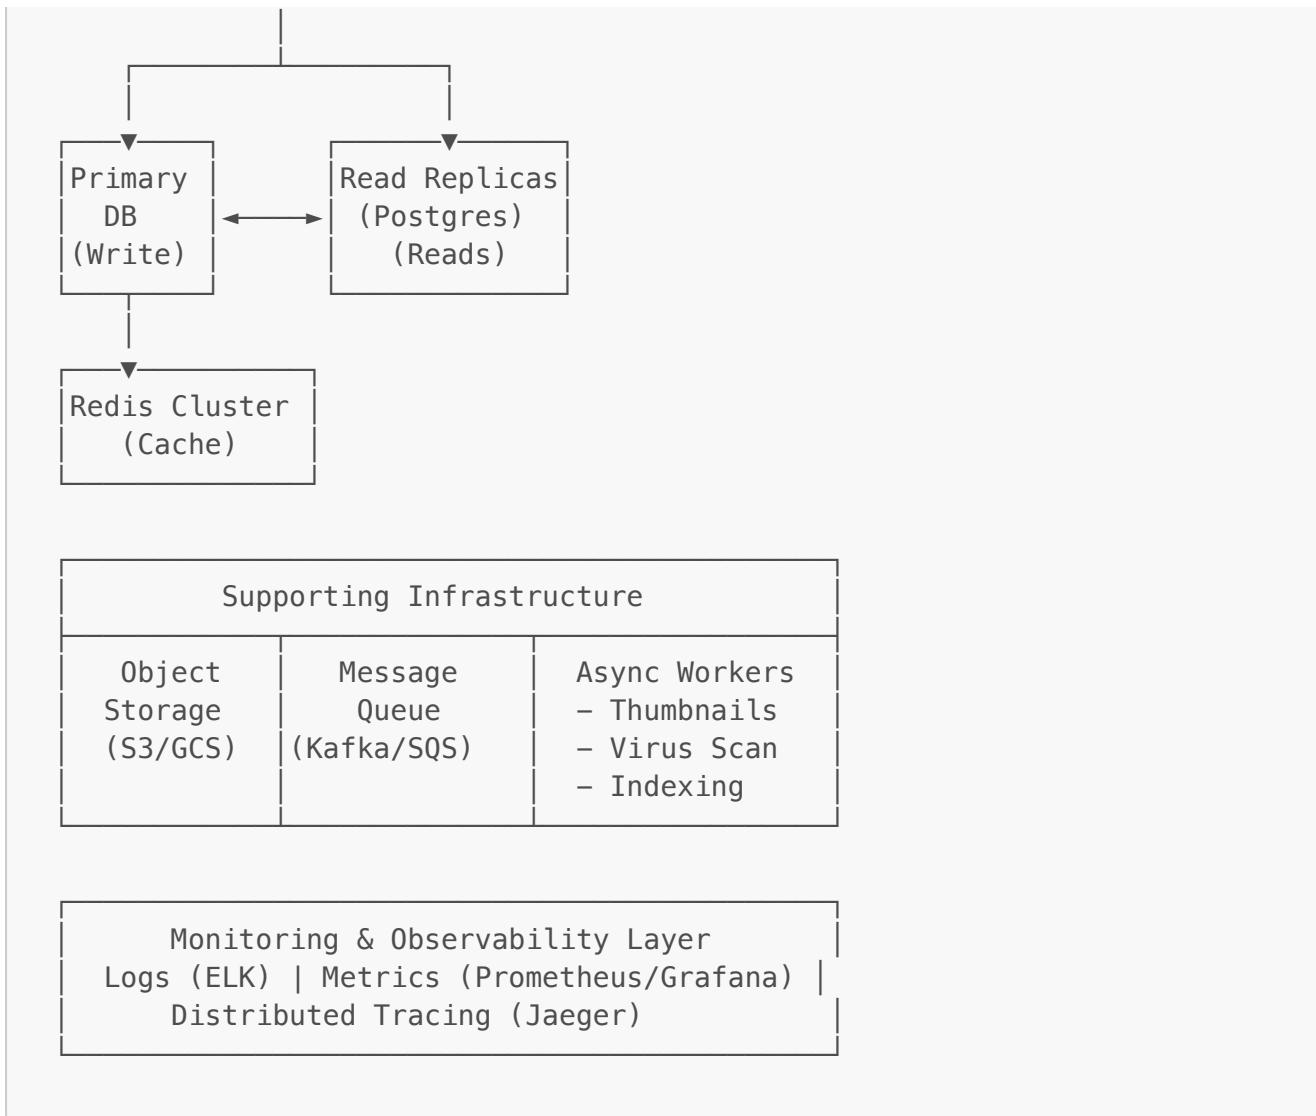
$$\text{Total} = 0.9999^4 = 99.96\%$$

To achieve 99.99%:

- Need redundancy (active-active, failover)
- Multiple availability zones
- Health checks and auto-recovery

## High-Level Architecture





## Core Components Deep Dive

### 1. API Gateway

**Purpose:** Single entry point for all client requests

**Responsibilities:**

- Request routing to appropriate services
- Authentication and authorization (JWT tokens)
- Rate limiting (token bucket algorithm)
- Request/response transformation
- SSL termination
- API versioning

**Technology Options Analysis:**

Technology	Pros	Cons	Use Case

Technology	Pros	Cons	Use Case
Kong	Plugins, extensible, open-source	Complex setup, learning curve	<b>Recommended</b> for flexibility
AWS API Gateway	Managed, integrated with AWS	Vendor lock-in, limited control	Quick MVP
Nginx	Lightweight, fast, proven	Manual scaling, less features	High-performance needs
Envoy	Modern, service mesh ready	Newer, smaller community	Microservices architecture

**Decision:** Kong API Gateway

### Why Kong?

- **Plugin ecosystem:** 50+ plugins for auth, rate limiting, logging
- **Scalability:** Proven at 100K+ RPS (Mashape experience)
- **Flexibility:** Can customize with Lua plugins
- **Multi-protocol:** REST, gRPC, WebSocket support
- **Observability:** Built-in metrics and tracing

### Trade-offs Accepted:

- Best features and flexibility
- Open-source with enterprise option
- Active community
- More complex than managed solution
- Need to manage infrastructure ourselves

### Configuration Strategy:

```
Rate Limiting:
└─ Free tier: 100 req/min per user
└─ Pro tier: 1,000 req/min per user
└─ Enterprise: 10,000 req/min per user
└─ Burst: 2x sustained rate for 10 seconds
```

```
Circuit Breaker:
└─ Failure threshold: 50% error rate
└─ Open duration: 30 seconds
└─ Half-open: Try 1 request
└─ Close when: 10 consecutive successes
```

## 2. Upload Service

**Purpose:** Handle file uploads efficiently at scale

## Challenge: How to upload large files reliably?

### Problem Analysis:

Single-shot upload of 5GB file:

- Network interruption → Complete failure
- 30-minute upload time
- No progress indication
- Cannot resume
- Server memory bottleneck (buffer entire file)

### Solution Comparison:

Approach	Max File Size	Resume	Speed	Complexity	Choice
Single Upload	100 MB	✗	Slow	Simple	✗
Chunked Sequential	10 GB	✓	Slow	Medium	✗
Chunked Parallel	10 GB+	✓	Fast	Medium	✓ Recommended
Streaming	Unlimited	⚠	Medium	High	Special cases

**Decision:** Chunked Parallel Upload

### Architecture:

#### STEP 1: Initiate Upload

Client:

```
└── File: vacation.mp4 (500 MB)
└── Calculate SHA-256: abc123...def
└── POST /api/v1/files/initiate
    Body: {
        filename: "vacation.mp4",
        size: 524288000,
        checksum: "abc123...def",
        parent_folder_id: "folder_456"
    }
```

#### STEP 2: Deduplication Check

Upload Service:

```
└── Query: SELECT file_id FROM files WHERE checksum = 'abc123...def'
└── Result: NULL (file doesn't exist)
└── Decision: Proceed with upload
```

If file existed:

- Create file reference for user
- No actual upload needed
- Response: "Upload complete" (instant)
- Savings: 500 MB bandwidth + storage

### STEP 3: Generate Upload Plan

- Upload Service:
- Calculate chunks:  $500 \text{ MB} \div 5 \text{ MB} = 100 \text{ chunks}$
  - Generate upload\_id: "upload\_789"
  - Create presigned S3 URLs (one per chunk):
    - Chunk 1: PUT s3://bucket/uploads/upload\_789/chunk\_1?sig=...
    - Chunk 2: PUT s3://bucket/uploads/upload\_789/chunk\_2?sig=...
    - ... (100 URLs total)
  - Store in Redis:
    - Key: upload:789:metadata
    - Value: {status: "in\_progress", chunks: 100, completed: 0}
    - TTL: 24 hours
  - Response:
    - upload\_id: "upload\_789",
    - chunk\_size: 5242880,
    - chunks: 100,
    - urls: [...]
  - }

### STEP 4: Upload Chunks (Parallel)

- Client:
- Split file into 100 chunks
  - Upload 10 chunks concurrently (parallel)
  - For each chunk:
    - PUT to presigned URL
    - Verify ETag matches
    - Notify server: POST /api/v1/files/upload\_789/chunks/1
    - Update progress bar
  - Continue until all 100 chunks uploaded

#### Why 10 concurrent chunks?

- Too few (1-5): Slow, underutilizes bandwidth
- Optimal (8-12): Saturates bandwidth, manageable
- Too many (20+): Overhead, diminishing returns

### STEP 5: Track Progress

- Upload Service:
- Each chunk completion:
    - HINCRBY upload:789:metadata completed 1
    - HSET upload:789:chunks chunk\_1 "completed"
  - WebSocket: Push progress to client
    - {uploaded: 45, total: 100, percent: 45}

- └ Client shows: "Uploading... 45%"
- Failure Handling:
  - └ Chunk 42 fails (network timeout)
  - └ Client retries chunk 42 (3 attempts)
  - └ Success on retry
  - └ Continue with remaining chunks

- Resume from Failure:
  - └ Client crashes at 60% complete
  - └ Client restarts
  - └ GET /api/v1/files/upload\_789/status
  - └ Response: {completed\_chunks: [1-60], pending: [61-100]}
  - └ Resume uploading chunks 61-100 only

#### STEP 6: Assemble File

- Client:
  - └ POST /api/v1/files/upload\_789/complete
  - Body: {checksum: "abc123...def"}

- Upload Service:
  - └ Verify all 100 chunks received
  - └ S3 Multipart Complete API:
    - └ Combines all chunks into single file
  - └ Calculate final checksum
  - └ Verify matches client checksum
  - └ Atomic operation (all-or-nothing)

- If checksum mismatch:
  - └ File corrupted during upload
  - └ Delete all chunks
  - └ Response: 400 Bad Request
  - └ Client retries entire upload

#### STEP 7: Finalize & Trigger Jobs

- Upload Service:
  - └ Move file to permanent location:
    - └ s3://bucket/users/user\_123/files/file\_789
  - └ Update metadata database:
    - └ INSERT INTO files (...)
    - └ UPDATE users SET storage\_used += 500MB
  - └ Delete temporary chunks & Redis state
  - └ Publish to message queue:
    - └ Event: file\_uploaded
    - └ Payload: {file\_id, user\_id, size, mime\_type}
    - └ Consumers:
      - └ Thumbnail generator (if image/video)
      - └ Virus scanner (all files)
      - └ Content indexer (if document)

└— Sync notifier (notify other devices)

Total Latency Breakdown:

- └— Deduplication check: 5ms
- └— Generate URLs: 10ms
- └— Upload chunks (parallel): 30s (depends on bandwidth)
- └— Assemble file: 2s
- └— Update metadata: 50ms
- └— Trigger jobs: 10ms

## Key Optimizations:

### 1. Direct to S3 Upload

Without presigned URLs:

- Client → Upload Service → S3
- Upload Service buffers data
- Memory bottleneck
- Double bandwidth usage

With presigned URLs:

- Client → S3 (direct)
- Zero server load
- Single network hop
- Upload Service only coordinates

### 2. Chunk Size Selection: Why 5MB?

Too Small (1 MB):

- └— 500 chunks for 500 MB file
- └— High HTTP overhead
- └— More S3 API calls (\$\$\$)
- └— ✗ Not optimal

Optimal (5 MB):

- └— 100 chunks for 500 MB file
- └— Balanced overhead
- └— Good for resume
- └— ✅ \*\*Recommended\*\*

Too Large (50 MB):

- └— Only 10 chunks
- └— Less granular progress
- └— Longer retry on failure
- └— ✗ Not optimal

### 3. Deduplication Savings

Scenario: 1000 users upload same 100MB video

Without deduplication:

└ Store:  $1000 \times 100 \text{ MB} = 100 \text{ GB}$

With deduplication:

└ Store: 100 MB (single copy)  
└ 1000 references to same file  
└ Savings: 99.9 GB (99.9%)

Real-world impact:

- Popular files (memes, movies, documents): 20–30% dedup rate
- Backups/system files: 50–70% dedup rate
- Average: 30% storage savings

### 3. Download Service

**Purpose:** Efficient file downloads

**Key Features:**

- **CDN Integration:** Serve files from edge locations
- **Signed URLs:** Time-limited download links (15 min)
- **Range Requests:** Support partial downloads (resume)
- **Permission Validation:** Check access before generating URL
- **Streaming:** Support for large files

### 4. Metadata Service

**Purpose:** Manage file and folder metadata at global scale

**Challenge: How to store 5 trillion file metadata records?**

**Problem Analysis:**

Single PostgreSQL instance limits:

└ Max connections: ~10,000  
└ Max storage: ~16 TB (practical limit)  
└ QPS capacity: ~10,000 reads, ~1,000 writes  
└ Our need: 70,000 QPS, 10 PB metadata

Conclusion: Single database cannot handle the scale

**Database Technology Comparison:**

Database	Consistency	Queries	Scale	Ops	Choice
PostgreSQL	Strong (ACID)	Complex (JOINs)	Shard required	Mature	<input checked="" type="checkbox"/> Primary
MySQL	Strong (ACID)	Complex	Shard required	Very mature	Alternative
MongoDB	Eventual	Simple	Auto-shard	Easy	<input checked="" type="checkbox"/> Eventual consistency issue
Cassandra	Eventual	Limited	Excellent	Complex	<input checked="" type="checkbox"/> No JOINs
DynamoDB	Eventual	Limited	Excellent	Managed	<input checked="" type="checkbox"/> No JOINs for folders

**Decision:** PostgreSQL with Sharding

### Why PostgreSQL over alternatives?

#### 1. Strong Consistency Needed:

Scenario: User deletes file while another device reads it

With Strong Consistency (PostgreSQL):

- Device A: DELETE file
- Transaction commits
- Device B: SELECT file
- File not found (correct)

With Eventual Consistency (NoSQL):

- Device A: DELETE file
- Device B: SELECT file (different replica)
- File still exists (stale read)
- Conflict! Device B might re-upload deleted file

#### 2. Complex Queries Required:

Query: "Get all files in folder and subfolders shared with user X"

```
SELECT f.*  
FROM files f  
JOIN folders fol ON f.parent_folder_id = fol.folder_id  
JOIN shares s ON (s.resource_id = f.file_id OR s.resource_id =  
fol.folder_id)  
WHERE s.shared_with_user_id = 'user_X'  
    AND fol.path LIKE '/Documents/%'  
ORDER BY f.modified_at DESC;
```

This is trivial in SQL, very complex in NoSQL

### 3. ACID Transactions Required:

Operation: Move file to different folder

```
BEGIN TRANSACTION;  
    UPDATE files SET parent_folder_id = 'new_folder' WHERE file_id =  
    'file_123';  
    UPDATE users SET storage_used = storage_used WHERE user_id =  
    'user_123';  
    INSERT INTO sync_events (...);  
COMMIT;
```

If any fails, all rollback (atomic)  
NoSQL databases struggle with multi-document transactions

### Sharding Strategy: Why shard by user\_id?

#### Alternatives Considered:

Option 1: Shard by file\_id (Hash Sharding)

Pros:

- └ Even distribution
- └ No hot shards

Cons:

- └ User's files scattered across shards
- └ Cannot query "all files for user" on single shard
- └ Cross-shard JOIN for folder hierarchy
- └ X Rejected

Option 2: Shard by geography

Pros:

- └ Low latency (data near users)
- └ Regulatory compliance (data residency)

Cons:

- └ Users travel → need cross-region queries
- └ Uneven distribution (US > Antarctica)
- └ ▲ Consider for global deployment

Option 3: Shard by user\_id (Range Sharding)

Pros:

- └ All user data on same shard ✓
- └ Single-shard queries for user operations ✓
- └ Folder hierarchy queries efficient ✓
- └ ✓ \*\*Selected\*\*

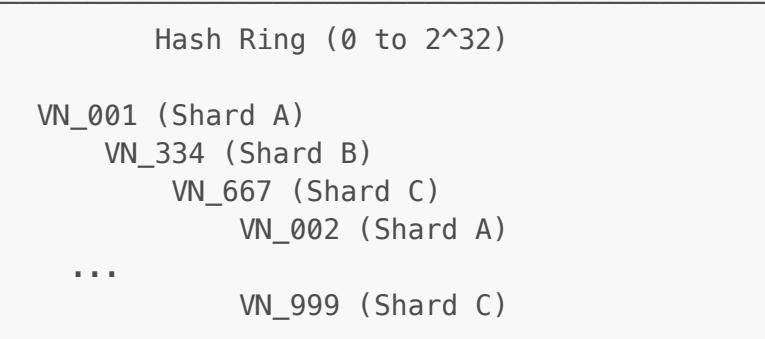
Cons:

- └ Celebrity users create hot shards

- └ Need rebalancing as users grow
  - └ Mitigation: Virtual nodes + monitoring

**Implementation:** Consistent Hashing with Virtual Nodes

## Hash Ring with 1000 virtual nodes:



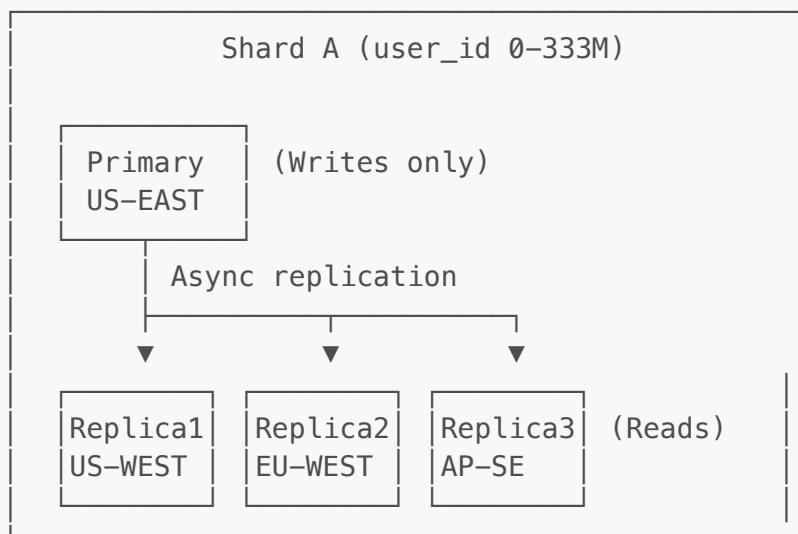
### User assignment:

- ```
|__ hash(user_12345) → 445329847
|__ Find next virtual node: VN_334
|__ VN_334 → Shard B
|__ Store all user_12345 data in Shard B
```

## Benefits:

- └─ Adding shard: Only  $\sim 1/N$  data moves
  - └─ Removing shard: Evenly redistributed
  - └─ Hot users: Can assign dedicated shard

## **Read Replica Strategy:**



## Read distribution:

- 80% reads go to replicas (distributed)

- └─ 20% writes go to primary
- └─ Replication lag: <100ms (acceptable)

- Failover on primary failure:
- └─ Promote Replica1 to primary (<30s)
  - └─ Reconfigure replication topology
  - └─ Resume operations

## Caching Strategy: Multi-Layer

### L1: Application Cache (In-Memory)

- └─ Store: Recently accessed file metadata
- └─ Size: 1 GB per API gateway instance
- └─ TTL: 5 minutes
- └─ Hit rate: 40%
- └─ Latency: 0.1ms

### L2: Redis Cluster

- └─ Store: Hot file metadata, user quotas
- └─ Size: 100 GB total
- └─ TTL: 1 hour (metadata), 5 min (quotas)
- └─ Hit rate: 90% (of L1 misses)
- └─ Latency: 1ms

### L3: Database

- └─ Store: All metadata (source of truth)
- └─ Size: 10 PB (5T files × 1KB × 2 for indexes)
- └─ Hit rate: 10% (cache misses)
- └─ Latency: 10–50ms

Combined Hit Rate =  $40\% + (60\% \times 90\%) = 94\%$

Average Latency =  $(0.4 \times 0.1\text{ms}) + (0.54 \times 1\text{ms}) + (0.06 \times 30\text{ms}) = 2.4\text{ms}$

## Cache Invalidation Strategy:

Problem: File metadata updated, caches become stale

Solutions Compared:

### Option 1: Write-Through Cache

- └─ Update cache + database simultaneously
  - └─ Pros: Always consistent
  - └─ Cons: Complex, all writes slower

### Option 2: Cache Invalidation on Write

- └─ Delete cache key on update
  - └─ Pros: Simple, cache always fresh on next read
  - └─ Cons: Next read slower (cache miss)

\*\*Selected\*\*

#### Option 3: TTL-based Expiration

- └ Let cache expire naturally
  - ├ Pros: Zero overhead
  - └ Cons: Stale data for TTL duration
- △ Use with invalidation

#### Implementation:

- └ On file update:
  - ├ UPDATE database
  - ├ DEL cache key in Redis
  - └ Publish event to invalidate L1 cache
- └ Next read:
  - ├ Cache miss
  - ├ Query database
  - └ Populate cache

## Quota Enforcement:

Challenge: Prevent storage quota exceeded

Naive approach (doesn't work):

GET user quota → Check limit → Upload file → Update quota

Problem: Race condition if concurrent uploads

Solution: Atomic check-and-reserve

```
└ Use database transaction:  
  BEGIN;  
    SELECT storage_used, storage_quota FROM users WHERE user_id = X  
  FOR UPDATE;  
    -- Row locked, no other transaction can modify  
    IF storage_used + file_size <= storage_quota THEN  
      UPDATE users SET storage_used = storage_used + file_size;  
      -- Reserve space  
      COMMIT;  
    ELSE  
      ROLLBACK;  
      RETURN 'Quota exceeded';  
    END IF;  
  COMMIT;
```

Alternative: Pessimistic locking ensures atomicity

## 5. Synchronization Service

**Purpose:** Keep files in sync across devices in real-time

## Challenge: How to sync files across billions of devices in real-time?

### Problem Statement:

Scenario: User has 4 devices

- MacBook (editing document)
- iPhone (wants to view)
- iPad (offline)
- Work PC (wants to edit)

#### Requirements:

- MacBook saves → iPhone sees update in <5 seconds
- iPad offline → Syncs when online
- Work PC edits same file → Conflict resolution
- Scale: 200M users × 3 devices = 600M active connections

### Synchronization Approach Comparison:

| Approach           | Latency | Scale | Offline | Conflicts | Bandwidth | Choice                                          |
|--------------------|---------|-------|---------|-----------|-----------|-------------------------------------------------|
| Polling (5s)       | 5s avg  | ✓     | ✓       | Manual    | High      | <span style="color: red;">✗</span> Outdated     |
| Long Polling       | 1-5s    | ✓✓    | ✓       | Manual    | Medium    | <span style="color: blue;">△</span> Fallback    |
| WebSocket          | <1s     | ✓✓✓   | ✓       | Auto      | Low       | <span style="color: green;">✓</span> Primary    |
| Server-Sent Events | <1s     | ✓✓    | ✓       | Auto      | Low       | <span style="color: blue;">△</span> Alternative |

**Decision:** WebSocket with Long Polling fallback

### Why WebSocket?

#### 1. Real-Time Bidirectional Communication

Polling (old way):

Every 5 seconds:

Client → Server: "Any changes?"

Server → Client: "No" (99% of the time)

Waste: 99% of requests return "no changes"

Latency: Average 2.5 seconds to detect change

WebSocket (new way):

One persistent connection

Server pushes when change occurs

Efficiency: 100× less network traffic

Latency: Milliseconds to detect change

## 2. Scalability Analysis

Polling approach:

- 200M users × 12 polls/min = 2.4B requests/min = 40M RPS
- Each poll: Network roundtrip, server processing
- Cost: Extremely high server load

WebSocket approach:

- 200M persistent connections
- Each connection idle most of the time (epoll)
- One server handles 100K connections (C10K problem solved)
- Need: 2,000 servers for connections  
(Much better than 400,000 servers for polling)

## 3. Battery Efficiency (Mobile)

Polling on mobile:

- Wake up every 5 seconds
- Establish TCP connection
- Send HTTP request
- Battery drain: Significant

WebSocket:

- Single persistent connection
- Server pushes data
- No repeated wake-ups
- Battery drain: Minimal

For Mobile: Use native push notifications (APNs/FCM)

## Architecture in Detail:

SYNC SCENARIO: User edits file on MacBook

Step 1: Detect Local Change (MacBook)

File System Watcher (inotify/FSEvents)  
— Monitors: ~/Drive directory  
— Detects: file.txt modified  
— Debounce: Wait 1s for more changes

↓

Step 2: Calculate Delta

Sync Client

- Get file's last sync version
- Calculate diff (rsync algorithm)
- Compression: gzip the diff
- Result: 500 KB → 50 KB (90% smaller)

↓

### Step 3: Upload Change

```
POST /api/v1-sync/upload
Body: {
  file_id: "file_123",
  device_id: "device_mac",
  version: 42,
  delta: <compressed_binary>,
  checksum: "xyz789"
}
```

↓

### Step 4: Store & Version

```
Upload Service
- Store delta in S3
- Update metadata:
  UPDATE files
  SET version = 43,
      modified_at = NOW()
  WHERE file_id = 'file_123'
```

↓

### Step 5: Publish Sync Event

```
Message Queue (Kafka)
Topic: file_changes
Event: {
  type: "file_modified",
  file_id: "file_123",
  user_id: "user_456",
  version: 43,
  device_id: "device_mac",
  timestamp: 1640995200
}
```

↓

### Step 6: Sync Service Processes Event

```
Sync Service Consumer
- Read from Kafka
- Query: Get all devices for user_456
- Filter: Exclude device_mac (originator)
- Result: [device_iphone, device_ipad,
           device_workpc]
```

↓

## Step 7: Notify Devices

```
Device iPhone (WebSocket – Active)
- Connection alive
- Push: {file_id, version, delta_url}
- Latency: 50ms
```

```
Device iPad (Offline)
- No active connection
- Store in pending_sync queue
- Deliver when device comes online
```

```
Device WorkPC (WebSocket – Active)
- Currently editing same file!
- Conflict detected!
- Handle conflict resolution
```

↓

## Step 8: Apply Changes

```
iPhone Sync Client
- Download delta from URL
- Apply patch to local file
- Verify checksum
- Update local version to 43
- Notify user: "file.txt updated"
```

## Conflict Resolution: The Hard Problem

### CONFLICT SCENARIO:

Time 10:00:00 – Both devices start with file version 42  
Time 10:00:30 – MacBook saves "Hello World" (version 43)  
Time 10:00:31 – WorkPC saves "Goodbye World" (version 43)

Both think they're at version 43!

```
Conflict Detection Algorithm
```

### Server receives MacBook update:

```
— Current server version: 42
— MacBook claims base version: 42 ✓
— No conflict
— Accept, increment to version 43
```

Server receives WorkPC update (1 second later):

- └─ Current server version: 43 (MacBook just updated)
- └─ WorkPC claims base version: 42 ✗
- └─ CONFLICT DETECTED!
- └─ Need resolution strategy

### Conflict Resolution Strategies Compared:

Strategy 1: Last-Write-Wins (LWW)

- └─ Compare timestamps
- └─ MacBook: 10:00:30
- └─ WorkPC: 10:00:31 → Winner
- └─ WorkPC version becomes version 44
- └─ MacBook version discarded
- └─ Problem: Data loss! (MacBook's edits lost)

Strategy 2: Merge Attempts (Operational Transform)

- └─ Analyze both edits
- └─ Try to merge intelligently
- └─ "Hello World" + "Goodbye World" = ???
- └─ Problem: Very complex, doesn't always work

Strategy 3: Last-Write-Wins + Conflict Copy ✓

- └─ WorkPC version becomes version 44 (latest timestamp)
- └─ Save MacBook version as conflict copy:  
    "file (conflicted copy from MacBook 10:00:30).txt"
- └─ Notify user on both devices
- └─ User manually resolves
- Pros: No data loss, simple, always works
- Cons: Manual intervention needed

### Implementation of LWW + Conflict Copy:

#### Conflict Resolution Flow

Server detects conflict:

- └─ Base version: 42
- └─ Current version: 43 (MacBook)
- └─ Incoming version: 43 (WorkPC claims 42 base)
- └─ Timestamp comparison:
  - └─ MacBook: 10:00:30
  - └─ WorkPC: 10:00:31 (winner)

Resolution:

- └─ Accept WorkPC as version 44
- └─ Create conflict file for MacBook:

```

    └── Name: "file (MacBook 10:00:30 conflicted).txt"
    └── Content: MacBook's version
    └── Metadata: marked as conflict
└── INSERT INTO files (conflict file)
└── Publish sync events:
    └── MacBook: Download conflict file
    └── WorkPC: Download winning version
└── Notify both users:
    "Conflicted copy created"

```

User Experience:

MacBook sees:

```

    └── file.txt (WorkPC version)
    └── file (MacBook 10:00:30 conflicted).txt (their version)
    └── Notification: "Conflict detected, review changes"

```

WorkPC sees:

```

    └── file.txt (their version, now official)
    └── file (MacBook 10:00:30 conflicted).txt (other version)
    └── Notification: "Another device edited simultaneously"

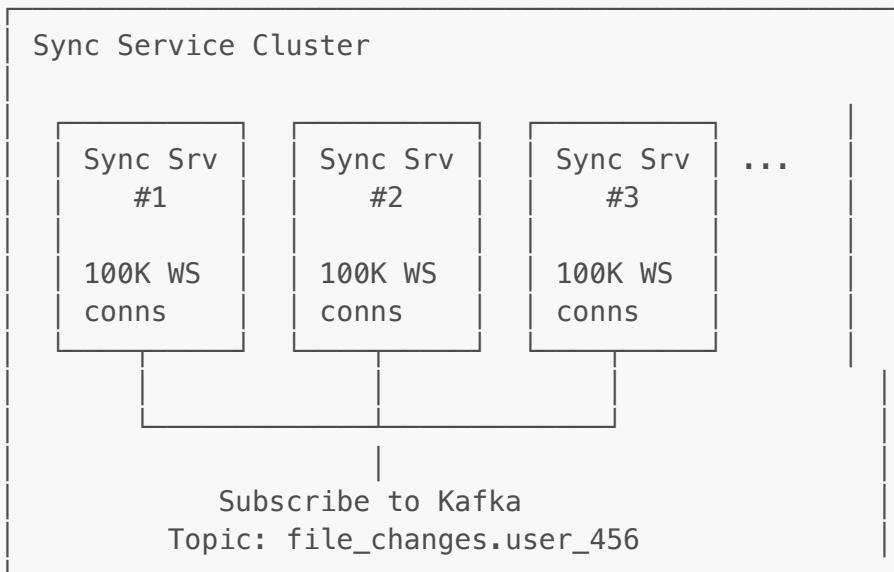
```

### **WebSocket Connection Management:**

Challenge: Maintain 600M WebSocket connections

Solution: Connection Pooling + Load Distribution

Architecture:



Distribution Strategy:

- └── Each server: 100,000 connections
- └── Total servers:  $600M \div 100K = 6,000$  servers
- └── Load balancer: Consistent hashing by user\_id
- └── Benefit: All user's devices on same server  
(Easier to broadcast changes)

#### Connection stickiness:

- └─ user\_456 → Always routes to Sync Server #42
- └─ All user\_456's devices on same server
- └─ Broadcast to 3–5 devices instead of searching 6000 servers

#### Cost optimization:

- └─ WebSocket keep-alive: 30 seconds
- └─ Idle connections use ~4 KB memory
- └─ 100K connections = 400 MB memory
- └─ Acceptable on modern servers (128 GB RAM)

### Fallback to Long Polling:

Why need fallback?

#### WebSocket failures:

- └─ Corporate firewalls block WS
- └─ Proxy servers don't support WS
- └─ Old browsers (IE 9)
- └─ Mobile apps occasionally

#### Long Polling mechanism:

1. Client: GET /sync/poll?cursor=123 |  
Server: Hold connection (60s) |
2. If change occurs → Return immediately |  
If 60s timeout → Return "no changes" |
3. Client immediately reconnects  
(Keeps connection quasi-persistent) |

#### Comparison:

- └─ WebSocket: 1 connection, push messages
- └─ Long Poll: Reconnect every 60s, pull messages
- └─ WebSocket preferred, Long Poll acceptable fallback

### Sync Protocol Design:

#### Delta Sync vs Full Sync:

Small change (edited 1 line in 10 MB document):

##### Full Sync:

- └─ Download entire 10 MB file
- └─ Bandwidth: 10 MB

└ Time: 10 seconds

#### Delta Sync:

- └ Download only changed bytes
- └ Bandwidth: 5 KB (0.05%)
- └ Time: 50ms (200x faster)
- └  \*\*Use for files < 100 MB\*\*

#### Algorithm: rsync-like binary diff

- └ Divide file into blocks (4 KB each)
- └ Calculate rolling checksum per block
- └ Compare checksums
- └ Transfer only changed blocks
- └ Reconstruct file from blocks + deltas

#### When to use Full Sync:

- └ New file (no previous version)
- └ File completely rewritten
- └ Delta larger than original (rare)
- └ First sync for device

### Sync State Management:

#### Per-Device State Tracking:

##### Redis storage:

Key: sync:state:user\_456:device\_iphone

Value: {  
    last\_sync\_timestamp: 1640995200,  
    last\_event\_id: "event\_789",  
    pending\_downloads: ["file\_123", "file\_456"],  
    pending\_uploads: [],  
    sync\_status: "synced"  
}

##### Sync cursor:

- └ Tracks last successfully processed event
- └ On reconnect: Send all events after cursor
- └ Ensures no events missed
- └ Even if device offline for days

### Offline Support:

#### iPad goes offline for 6 hours:

##### During offline:

- └ Local changes queued (SQLite queue)
- └ 10 files edited

```
└── 5 files created  
└── Queue stored persistently
```

When online:

1. Reconnect to Sync Service  
GET /sync/changes?cursor=last\_event

↓

2. Download Remote Changes
  - 15 files changed by other devices
  - Download deltas
  - Apply to local files

↓

3. Upload Local Changes
  - 10 edited files → Check conflicts
  - 5 new files → Upload
  - 2 conflicts detected → Create copies

↓

4. Resolve Conflicts
  - file.txt has conflict
  - Create: file (iPad conflicted).txt
  - Notify user

Conflict Detection:

```
└── Remote version: 44 (edited at 10:05)  
└── Local version: 43 (edited at 10:03 while offline)  
└── Both diverged from version 42  
└── CONFLICT! (cannot auto-merge)
```

**Connection Recovery:**

Challenge: WebSocket connections drop frequently

Causes:

```
└── Network switches (WiFi → Cellular)  
└── Network congestion  
└── Server restarts  
└── Load balancer timeout  
└── Laptop sleep/wake
```

Recovery Strategy:

Detection:

```
└── Client: Heartbeat every 30s
```

- └─ Server: Heartbeat every 30s
- └─ No heartbeat for 60s → Connection dead
- └─ Reconnect immediately

**Reconnection:**

- └─ Exponential backoff: 1s, 2s, 4s, 8s, 16s, 30s (max)
- └─ Include cursor in reconnect request
- └─ Server sends missed events
- └─ Resume normal operation

**State preservation:**

- └─ Server: Keep connection state for 5 minutes after disconnect
- └─ Client: Keep pending changes in queue
- └─ No data loss even with frequent reconnections

**Bandwidth Optimization:**

**Smart Sync Features:**

**1. Selective Sync:**

User marks folders:

- └─ "Always available offline" → Full sync
- └─ "Available on demand" → Metadata only
- └─ "Online only" → No local copy

Bandwidth savings: 50–70% for typical user

**2. Sync Scheduling:**

Mobile devices:

- └─ WiFi: Sync immediately
- └─ Cellular: Sync metadata only
- └─ Low battery: Pause sync
- └─ Charging + WiFi: Full sync

Battery & data plan friendly

**3. Compression:**

- └─ Text files: gzip (70% savings)
- └─ Images: Already compressed (skip)
- └─ Videos: Already compressed (skip)
- └─ Applied automatically per file type

**4. Throttling:**

- └─ Detect rapid changes (save every second)
- └─ Debounce: Wait 5 seconds for more changes
- └─ Batch upload: Send one delta with all changes
- └─ Reduces sync events by 80%

## 6. Sharing Service

**Purpose:** Manage file/folder sharing

**Features:**

- **User Sharing:** Share with specific users
- **Public Links:** Generate shareable URLs
- **Permission Levels:** Owner, Editor, Commenter, Viewer
- **Expiration:** Time-limited shares
- **Access Logs:** Track who accessed what

**Security:**

- Cryptographically secure tokens
  - Store hashed tokens (bcrypt)
  - Password protection for public links
  - Audit trail for all access
- 

## 7. Search Service

**Purpose:** Fast file search across content and metadata

**Technology:** Elasticsearch

**Search Capabilities:**

- Full-text search in document content
- Metadata search (name, type, owner, date)
- Filters and facets
- Relevance ranking

**Indexing Strategy:**

- Asynchronous indexing via background workers
  - Text extraction from PDFs, Office docs
  - OCR for images (if enabled)
  - Real-time index updates
- 

## 8. Storage Layer (Object Storage)

**Purpose:** Store actual file data

**Technology:** AWS S3 / Google Cloud Storage

**Features:**

- **Versioning:** Keep multiple versions of files
  - **Lifecycle Policies:** Auto-tier to cold storage
-

- **Cross-Region Replication:** For disaster recovery
- **Encryption:** At rest (AES-256)

#### Storage Classes:

- **Hot:** Frequently accessed (S3 Standard)
  - **Warm:** Infrequent access (S3 IA)
  - **Cold:** Archive (S3 Glacier)
- 

## 9. Message Queue

**Purpose:** Asynchronous communication between services

**Technology:** Apache Kafka, RabbitMQ, AWS SQS

#### Event Types:

- File upload complete → Trigger virus scan, thumbnail generation
- File modified → Trigger sync notification
- File shared → Send notification
- File accessed → Update analytics

#### Benefits:

- Decouple services
  - Handle traffic spikes
  - Guarantee delivery
  - Replay capability
- 

## 10. Async Workers

**Purpose:** Process background jobs

#### Worker Types:

1. **Thumbnail Generator:** Create image/video previews
2. **Virus Scanner:** Scan files with ClamAV
3. **Content Indexer:** Extract text, update search index
4. **Analytics Processor:** Track usage metrics
5. **Notification Sender:** Email/push notifications

#### Processing:

- Pull jobs from message queue
  - Process independently
  - Retry with exponential backoff
  - Dead letter queue for failed jobs
-

## 11. Notification Service

**Purpose:** Real-time notifications to users

**Channels:**

- **WebSocket/SSE:** Real-time web notifications
- **Push Notifications:** Mobile (FCM/APNS)
- **Email:** Via SendGrid/SES
- **In-App:** Notification center

**Events:**

- File shared with you
- Comment on your file
- Upload complete
- Storage quota warning

## 12. Cache Layer (Redis)

**Purpose:** Reduce database load, improve performance

**Cached Data:**

- User sessions (7 day TTL)
- File metadata (1 hour TTL)
- User quotas (5 min TTL)
- Recent files list (1 hour TTL)
- Sync cursors (24 hour TTL)

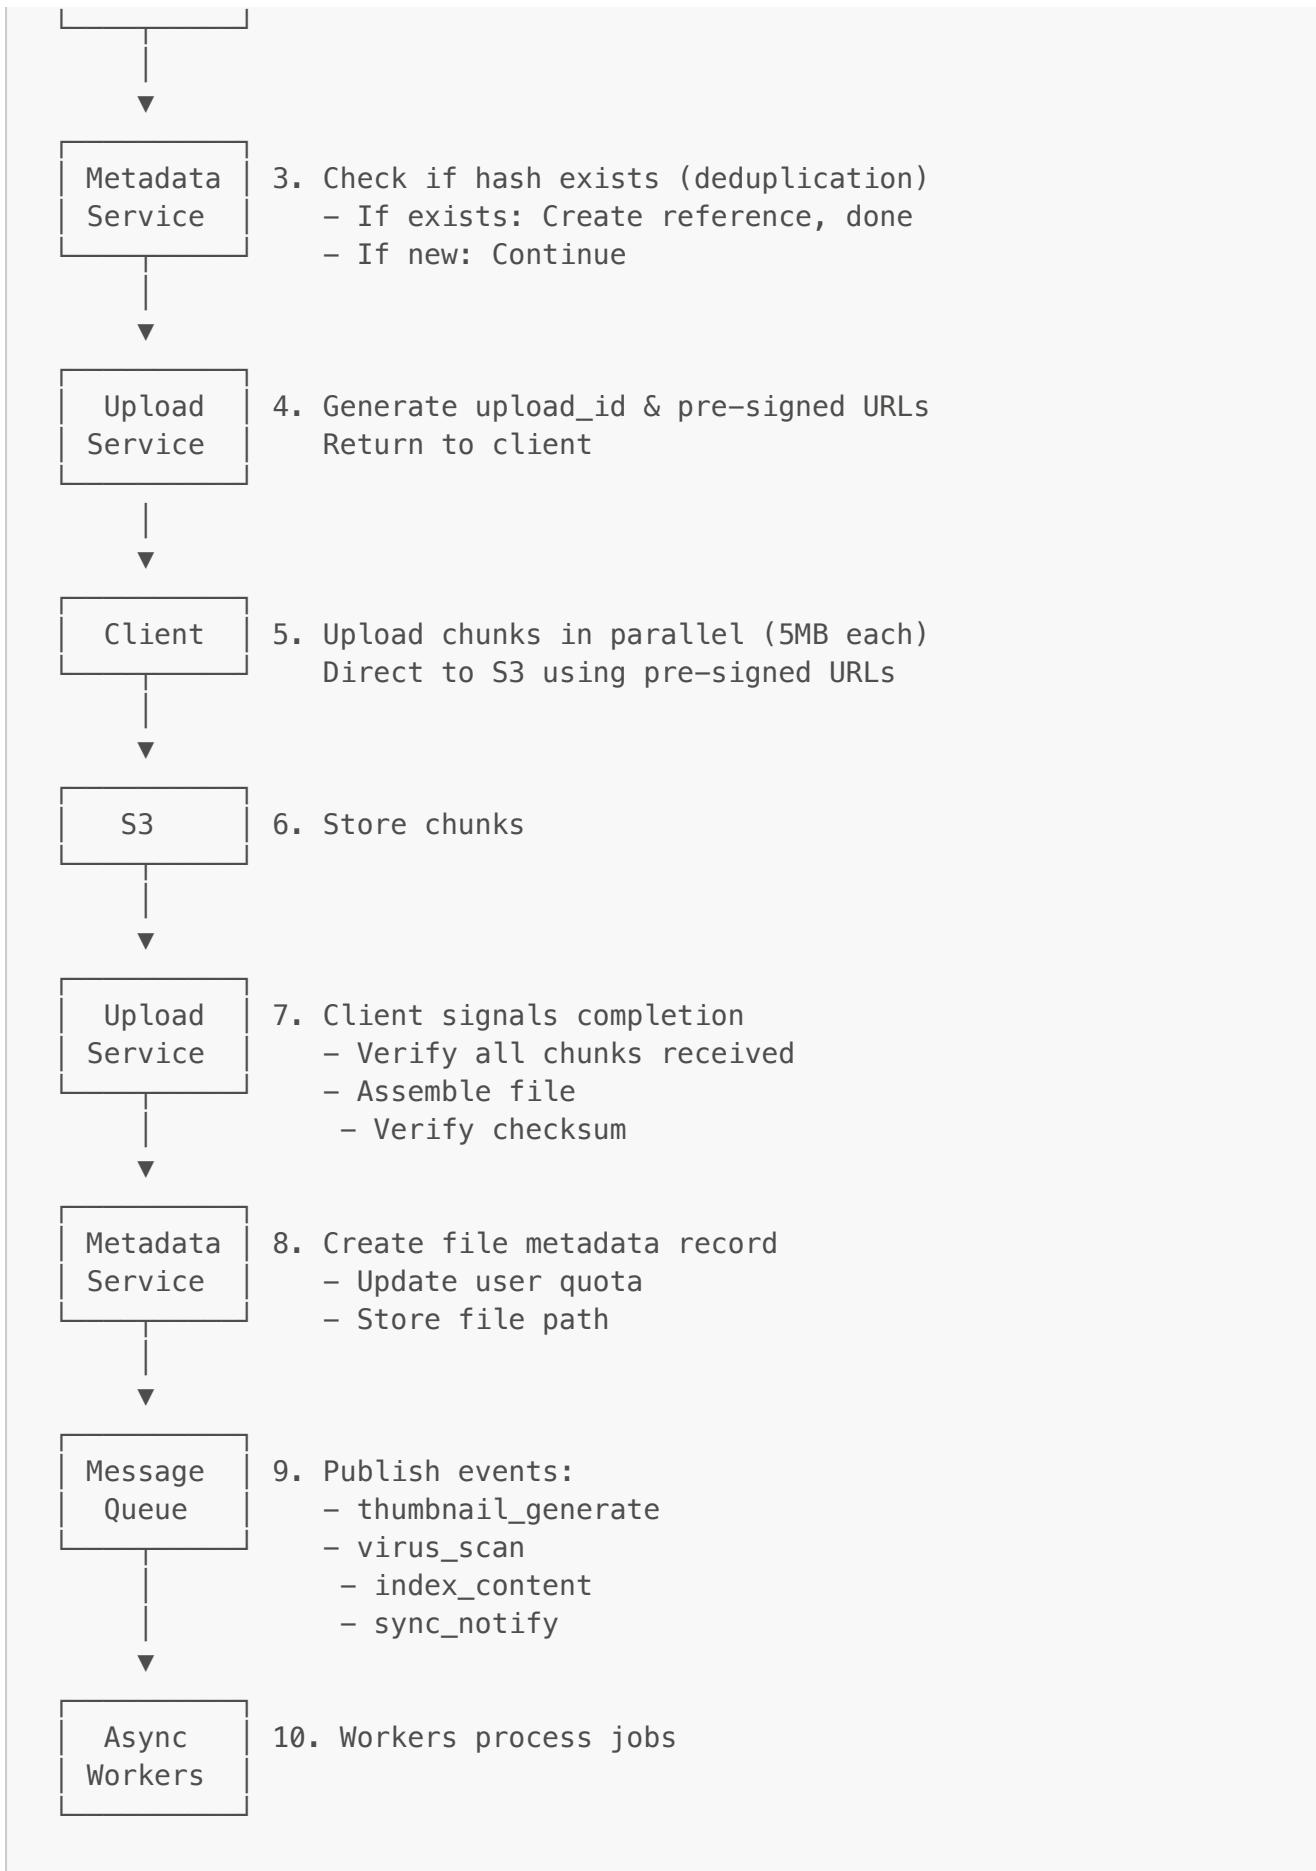
**Strategy:**

- Cache-aside pattern
- Write-through for critical data
- Cache invalidation on updates

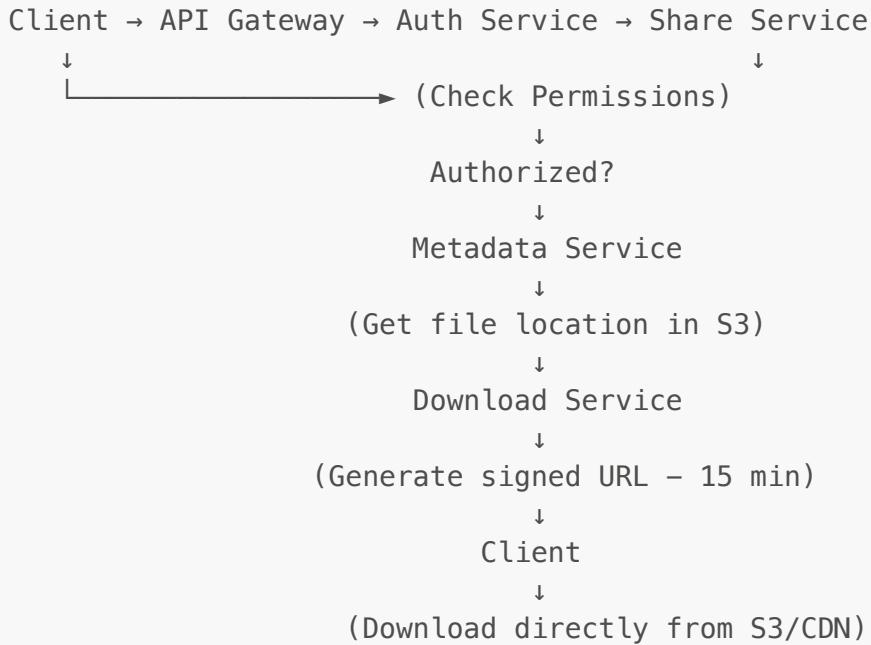
## Data Flow

### 1. File Upload Flow





## 2. File Download Flow



### 3. Real-Time Sync Flow

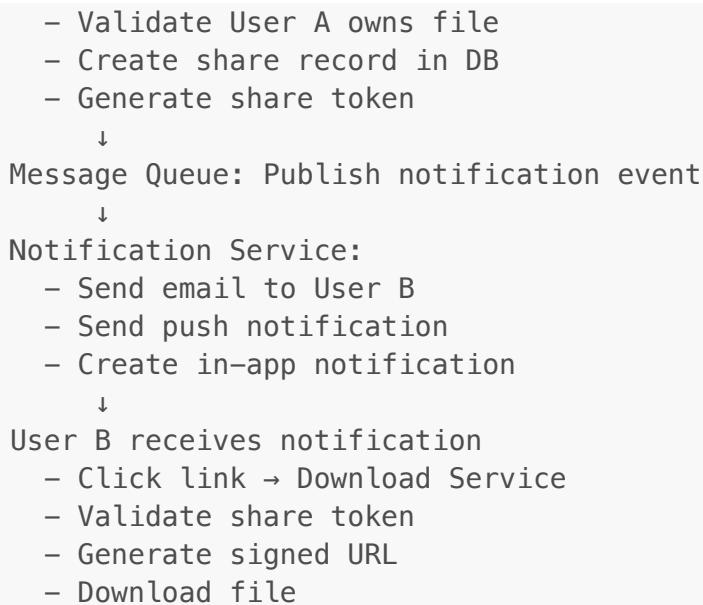
```

Device A: User edits file
    ↓
Local File System Watcher detects change
    ↓
Upload Service: Store new version
    ↓
Metadata Service: Update metadata
    ↓
Message Queue: Publish sync event
    ↓
Sync Service: Consume event
    ↓
Identify affected devices (Device B, C, D)
    ↓
Push notification via WebSocket/Long Polling
    ↓
Device B, C, D: Receive update
    ↓
Download new version and apply
  
```

### 4. File Sharing Flow

```

User A shares file with User B
    ↓
API Gateway → Auth Service
    ↓
Share Service:
  
```



## Database Design

Metadata Database (PostgreSQL)

### Core Entities:

1. **Users**: user\_id, email, storage\_quota, storage\_used
2. **Files**: file\_id, owner\_id, name, size, checksum, parent\_folder\_id, version
3. **Folders**: folder\_id, owner\_id, name, parent\_folder\_id, path
4. **File\_Versions**: version\_id, file\_id, version\_number, storage\_path
5. **Shares**: share\_id, resource\_id, owner\_id, shared\_with\_user\_id, permission
6. **Devices**: device\_id, user\_id, device\_name, last\_sync\_at
7. **Sync\_Events**: event\_id, user\_id, resource\_id, event\_type, timestamp

### Key Indexes:

- owner\_id + parent\_folder\_id (files/folders)
- checksum (for deduplication)
- share\_token (for public links)
- user\_id + timestamp (for sync events)

### Sharding Strategy:

- Shard by user\_id
- Consistent hashing
- Virtual nodes for better distribution

Document Store (MongoDB/DynamoDB)

### Collections:

1. **File\_Chunks**: Track chunks during upload
  2. **Activity\_Logs**: User actions audit trail
  3. **Analytics\_Events**: Usage metrics
- 

## Search Index (Elasticsearch)

### Document Structure:

- file\_id, name, content (extracted text)
  - owner, tags, created\_at, modified\_at
  - path, mime\_type, size
- 

## Cache (Redis)

### Key Patterns:

- session:{session\_id} → User session
  - file:metadata:{file\_id} → File metadata
  - user:quota:{user\_id} → Storage usage
  - user:recent:{user\_id} → Recent files list
  - sync:cursor:{user\_id}:{device\_id} → Sync state
- 

## Design Decisions & Trade-offs

### 1. Metadata: SQL vs NoSQL

**Decision:** PostgreSQL (SQL)

#### Reasoning:

- Strong consistency needed for ownership, permissions
- Complex queries (folder hierarchy, shared files)
- ACID transactions required
- Mature ecosystem

#### Trade-offs:

- Strong consistency
  - Relational integrity
  - Complex queries
  - Harder horizontal scaling (needs sharding)
- 

### 2. Storage: Object Storage (S3/GCS)

**Decision:** Object Storage

#### Reasoning:

- Designed for petabyte scale
- Built-in 11 9's durability
- Cost-effective
- Versioning support
- Global CDN integration

**Trade-offs:**

- Infinite scalability
  - High durability
  - Low cost per GB
  - Higher latency than block storage
- 

### 3. Consistency Model: Hybrid

**Decision:** Strong for metadata, Eventual for sync

**Reasoning:**

- Metadata (ownership, permissions) must be consistent
- Sync can tolerate eventual consistency
- Best balance of consistency and availability

**Trade-offs:**

- Guarantees where needed
  - High availability
  - More complex implementation
- 

### 4. Upload: Chunked Multipart

**Decision:** 5MB chunks, parallel upload

**Reasoning:**

- Support large files (>5GB)
- Faster (parallel upload)
- Resumable after failures
- Better progress tracking

**Trade-offs:**

- Fast for large files
  - Resumable
  - Better UX
  - More complex
- 

### 5. Deduplication: Client-side hash

**Decision:** Client calculates hash, server verifies

**Reasoning:**

- Save bandwidth (don't upload duplicates)
- Save storage (single copy)
- Fast for users

**Trade-offs:**

- 30% bandwidth savings
  - Storage savings
  - Privacy concerns (hash reveals content)
  - Reference counting complexity
- 

## 6. Sync: WebSocket + Push Notifications

**Decision:** WebSocket for web/desktop, FCM/APNS for mobile

**Reasoning:**

- Real-time updates (low latency)
- Single persistent connection
- Battery efficient on mobile

**Trade-offs:**

- Real-time (milliseconds)
  - Efficient
  - Connection management complexity
  - Firewall issues
- 

## 7. Sharding: By user\_id

**Decision:** Shard metadata by user\_id

**Reasoning:**

- User's data co-located
- No cross-shard queries for user operations
- Simple to implement

**Trade-offs:**

- Query efficiency
  - Data locality
  - Hot shards (power users)
  - Rebalancing complexity
-

## 8. Caching: Multi-layer (CDN + Redis + DB)

**Decision:** 3-tier caching

**Reasoning:**

- CDN for downloads (edge caching)
- Redis for metadata (application caching)
- DB buffer for query results

**Trade-offs:**

- 80%+ cache hit rate
  - Reduced DB load
  - Lower latency
  - Cache invalidation complexity
  - Additional cost
- 

## 9. Conflict Resolution: Last-Write-Wins

**Decision:** LWW with conflict copies

**Reasoning:**

- Simple to implement
- Works for 99% of cases
- Preserves all data

**Flow:**

1. Detect conflict
2. Latest timestamp wins
3. Create conflict copy
4. User reviews

**Trade-offs:**

- Simple
  - No data loss
  - User intervention needed
- 

## 10. Database Replication: Primary-Replica

**Decision:** 1 Primary + 3 Read Replicas

**Reasoning:**

- Scale reads (most operations)
  - High availability
-

- Geographic distribution

#### **Trade-offs:**

- Read scalability (3x)
  - HA
  - Replication lag
  - Failover complexity
- 

## Scalability Strategy

### Horizontal Scaling

#### **Application Layer:**

- Stateless services behind load balancers
- Auto-scaling based on CPU/memory
- Can add/remove instances dynamically

#### **Database Layer:**

- Shard by user\_id
- Add more shards as needed
- Read replicas for read scaling
- Eventually migrate to distributed SQL (CockroachDB, YugabyteDB)

#### **Cache Layer:**

- Redis cluster with consistent hashing
- Add nodes to handle more load
- Partition hot keys

#### **Storage Layer:**

- Object storage scales automatically
- No manual intervention needed

### Vertical Scaling (When Needed)

- Upgrade database instance types
- More memory for cache
- Faster disks for database

### Geographic Distribution

- Multi-region deployment
  - Route users to nearest region
  - Cross-region replication for DR
-

# Reliability & Availability

## High Availability Strategies

### 1. Redundancy:

- Multiple availability zones
- Load balancers in active-active
- Database primary-replica setup
- 3x replication for object storage

### 2. Failover:

- Automatic failover for database
- Health checks for all services
- Circuit breakers to prevent cascade failures

### 3. Graceful Degradation:

- Serve stale cache if DB down
- Disable features if dependencies fail
- Show user-friendly error messages

## Disaster Recovery

### Backup Strategy:

- Database: Continuous WAL archiving, daily full backups
- Object Storage: Cross-region replication, versioning
- Metadata: Daily exports to S3

### Recovery:

- RPO (Recovery Point Objective): < 1 hour
  - RTO (Recovery Time Objective): < 4 hours
- 

# Security Considerations

## 1. Authentication & Authorization

- JWT tokens with short expiry (15 min)
- Refresh tokens for session management
- OAuth 2.0 for third-party integrations
- Multi-factor authentication (TOTP, SMS)

## 2. Encryption

- **At Rest:** AES-256 encryption in S3
  - **In Transit:** TLS 1.3 for all connections
  - **End-to-End** (optional): Client-side encryption
-

### 3. Access Control

- Role-based permissions (Owner, Editor, Viewer)
- Least privilege principle
- Regular permission audits

### 4. Network Security

- VPC with private subnets
- Security groups and NACLs
- DDoS protection (CloudFlare, AWS Shield)
- Rate limiting at API Gateway

### 5. Data Protection

- Regular backups
  - Data lifecycle policies
  - Compliance (GDPR, HIPAA)
  - Audit logs for all access
- 

## Summary & Key Points

### Architecture Highlights

1. **Microservices:** Separate concerns (upload, download, metadata, sync, search)
2. **Horizontal Scalability:** Shard databases, object storage, read replicas
3. **High Availability:** Multi-region, replication, failover
4. **Strong Security:** Encryption, access control, audit logging
5. **Optimized Performance:** Caching, CDN, chunking, async processing

### Key Design Principles

- **Scalability First:** Built to handle billions of users
- **Reliability:** Multiple layers of redundancy
- **Performance:** Caching and CDN for low latency
- **Security:** Defense in depth
- **Cost Efficiency:** Lifecycle policies, deduplication

### Technology Choices

- **Metadata:** PostgreSQL (sharded)
- **Storage:** AWS S3 / GCS
- **Cache:** Redis Cluster
- **Message Queue:** Kafka / SQS
- **Search:** Elasticsearch
- **CDN:** CloudFront / Cloudflare

**This design can handle billions of users and petabytes of data while maintaining 99.99% availability and strong security! 🚀**