

Leaderboard System Design - High-Level Design Focus

Table of Contents

1. Problem Statement & Requirements
 2. High-Level Architecture
 3. Core Components
 4. Data Model & Storage Design
 5. Real-time Updates & Data Flow
 6. Ranking Algorithms
 7. Distributed System Design
 8. Scalability & Performance
 9. Caching Strategy
 10. Fault Tolerance & Reliability
 11. Design Trade-offs
 12. Real-World Case Studies
-

Problem Statement & Requirements

Problem Definition

Design a scalable leaderboard system that can:

- Handle millions of users with real-time score updates
- Support multiple leaderboard types (global, regional, friends)
- Provide fast rank lookups (< 100ms)
- Handle 50,000+ score updates/second
- Display top N players efficiently
- Support time-based leaderboards (daily, weekly, monthly, all-time)

Functional Requirements

1. Core Features

- **Score Updates:** Users can update their scores in real-time
- **Rank Queries:** Get a user's current rank instantly
- **Top N Queries:** Retrieve top N users (e.g., top 100)
- **Range Queries:** Get users ranked between M and N
- **User's Neighborhood:** Show users around a specific user's rank
- **Multiple Leaderboards:** Global, friends, regional, by game level

2. Leaderboard Types

Time-based:

- └─ Real-time (continuous)
- └─ Daily (resets at midnight)
- └─ Weekly (resets Monday 00:00)
- └─ Monthly (resets 1st of month)
- └─ All-time (never resets)

Scope-based:

- └─ Global (all users)
- └─ Regional (by country/continent)
- └─ Friends (social graph)
- └─ Guild/Team (group-based)
- └─ Game Level (per level/stage)

3. Advanced Features

- **Historical Snapshots:** View past leaderboard states
- **Percentile Ranks:** "You're in top 5% of players"
- **Achievements:** Milestone-based badges
- **Decay System:** Scores decay over time (for engagement)
- **Tie-breaking:** Handle identical scores fairly

Non-Functional Requirements

1. Performance

- **Score Update Latency:** < 50ms (P95)
- **Rank Query Latency:** < 100ms (P95)
- **Top N Query Latency:** < 200ms (P95)
- **Throughput:** Support 50K+ updates/sec

2. Scalability

- **Users:** 100M+ registered users
- **Active Users:** 10M+ DAU
- **Concurrent Updates:** 50K+ simultaneous score updates
- **Read-Heavy:** 100:1 read-to-write ratio

3. Availability & Reliability

- **Uptime:** 99.9% availability
- **Data Durability:** No score loss
- **Consistency:** Eventual consistency acceptable (< 1 second lag)
- **Fault Tolerance:** Handle node failures gracefully

4. Cost Efficiency

- **Storage:** Optimize for large user base
- **Compute:** Balance real-time vs batch processing
- **Network:** Minimize cross-region data transfer

Capacity Estimation

User Base:

- Total Users: 100M
- Daily Active Users (DAU): 10M
- Concurrent Users (Peak): 1M
- Games per User per Day: 5

Score Updates:

- Daily Updates: $10M \times 5 = 50M$
- Average Updates/Sec: $50M / 86400 \approx 580$
- Peak Updates/Sec: $580 \times 10 = 5,800$
- Safety Margin (10x): 58,000 updates/sec capacity

Read Operations:

- Rank Checks per Game: 2
- Top 100 Views: 1 per game
- Daily Reads: $50M \times 3 = 150M$
- Average Reads/Sec: $150M / 86400 \approx 1,736$
- Peak Reads/Sec: 17,360
- Safety Margin: 173,600 reads/sec capacity

Storage Requirements:

- Per User Record: 50 bytes (user_id: 16, score: 8, timestamp: 8, metadata: 18)
- Global Leaderboard: $100M \times 50 = 5GB$
- Daily/Weekly/Monthly: $3 \times 5GB = 15GB$
- Regional (10 regions): $10 \times 5GB = 50GB$
- Friends (avg 100 friends): Complex, stored in social graph
- Historical Snapshots: $365 \times 5GB = 1.825TB/year$
- Total with replicas (3x): ~6TB active + ~5.5TB archive

Server Requirements:

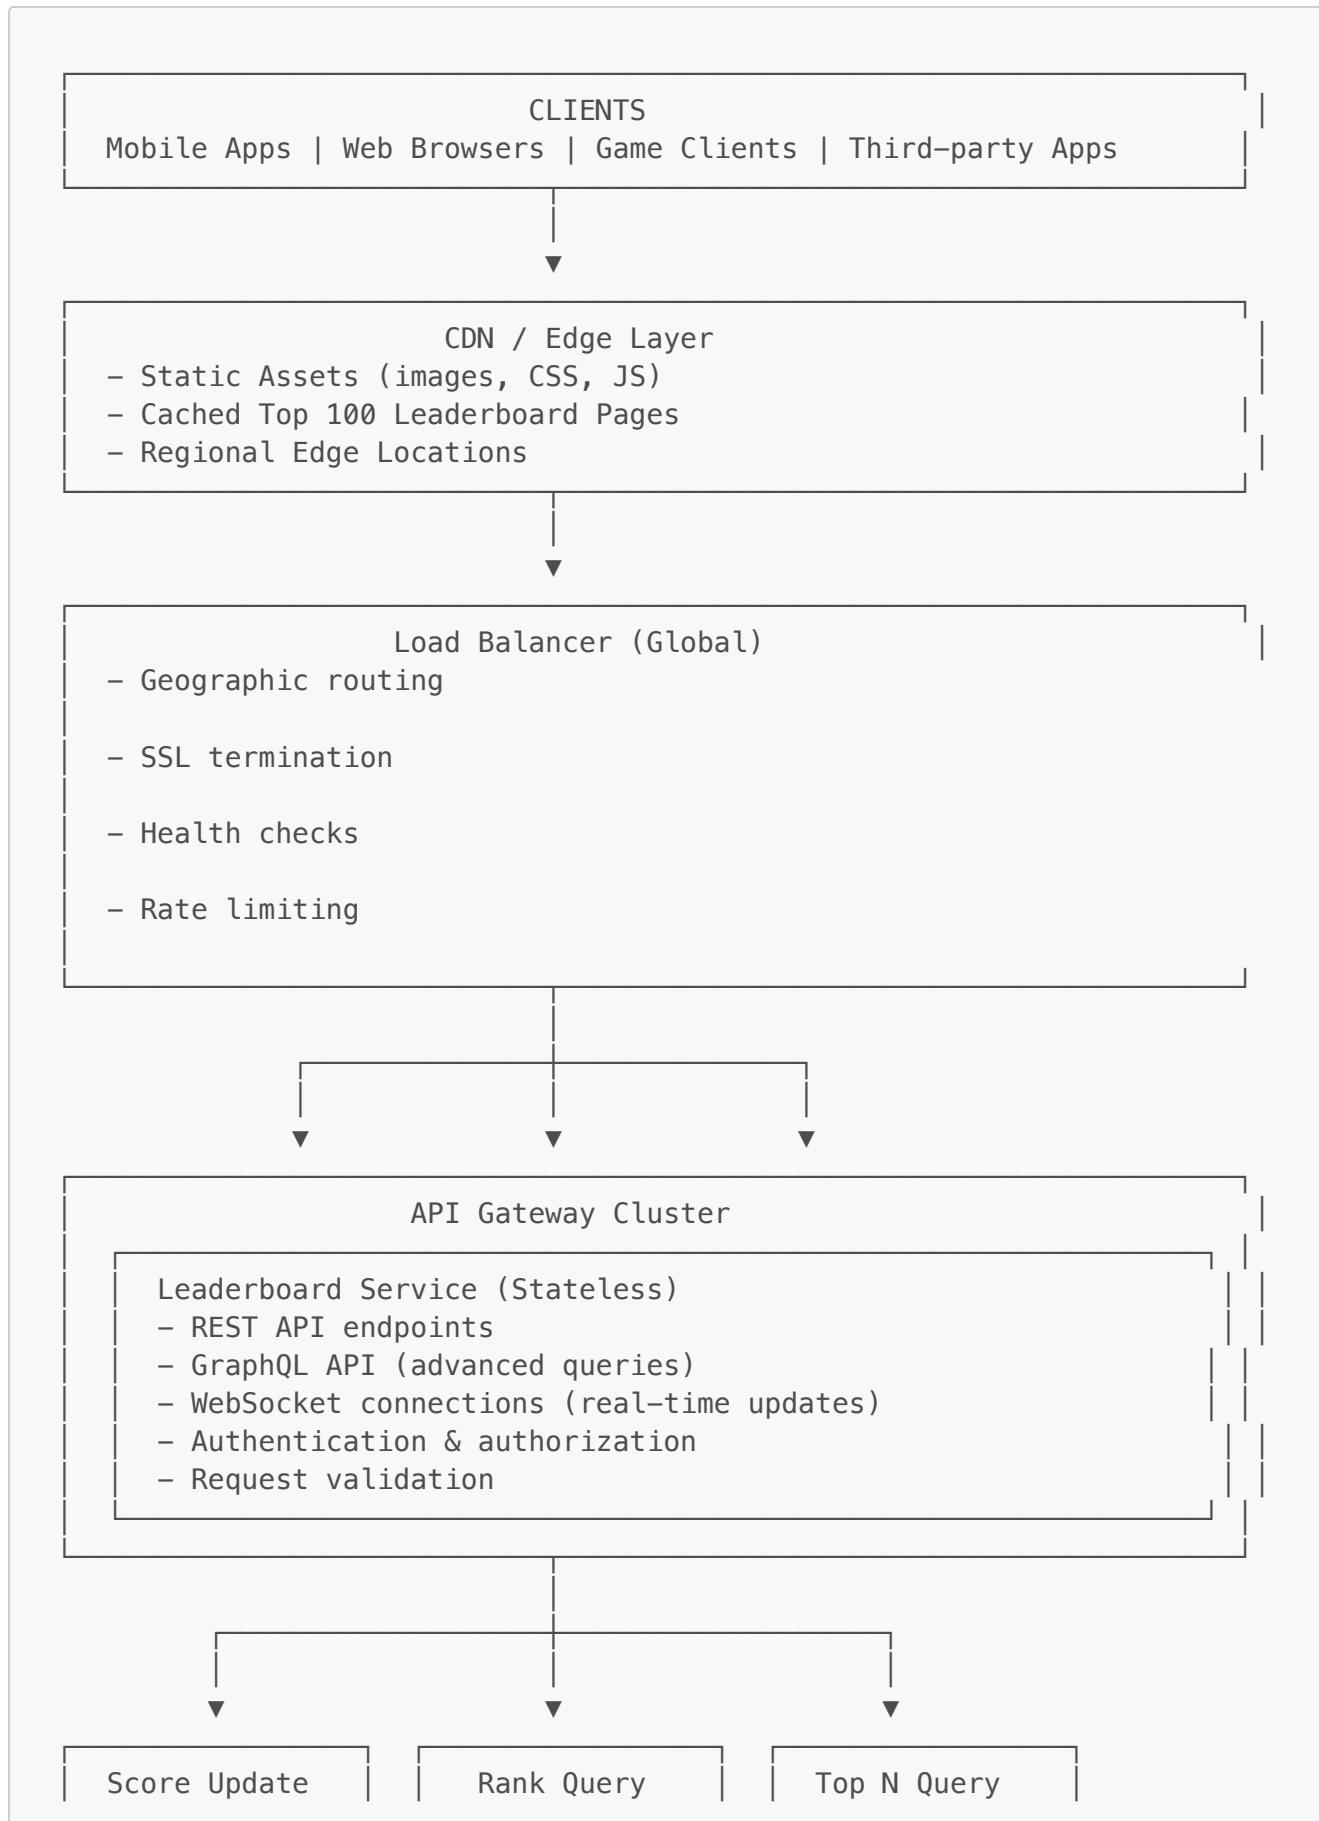
- Redis Nodes (score storage): $20 \text{ nodes} \times 64GB = 1.28TB \text{ memory}$
- Application Servers: 50 servers for processing
- PostgreSQL (metadata): 3 servers with replication
- Message Queue (Kafka): 5 brokers

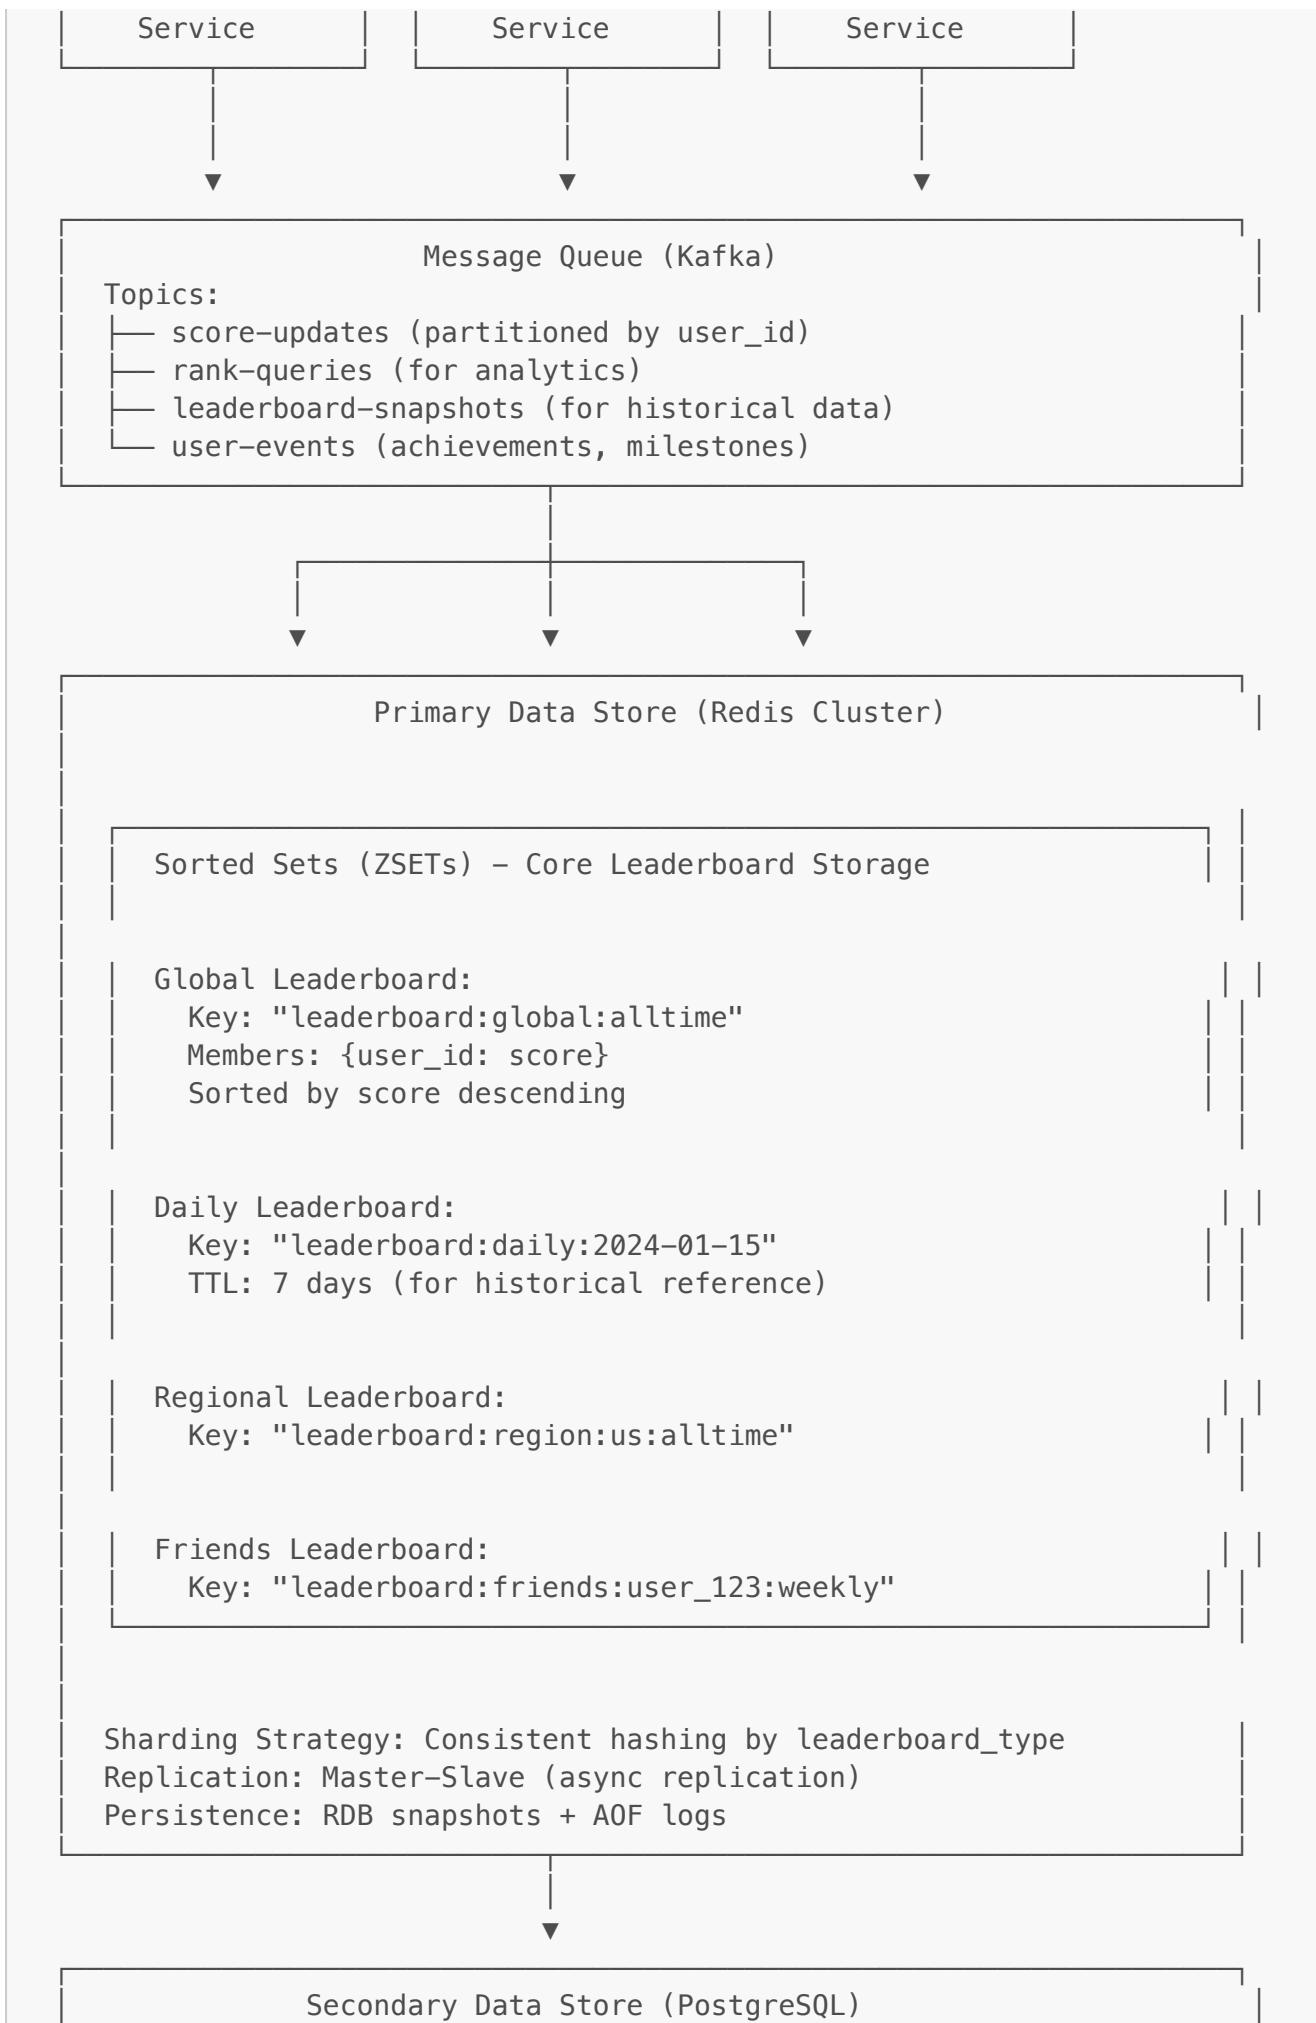
Monthly Cost (AWS):

- Redis (ElastiCache): ~\$10,000
- EC2 Application Servers: ~\$5,000
- RDS PostgreSQL: ~\$2,000
- Kafka (MSK): ~\$1,500
- S3 Storage (archives): ~\$500
- Data Transfer: ~\$1,000
- Total: ~\$20,000/month

High-Level Architecture

System Overview





User Metadata:
user_id, username, email, region, created_at

Historical Rankings:
snapshot_id, user_id, rank, score, timestamp
(Partitioned by timestamp)

Leaderboard Configuration:
leaderboard_id, type, reset_schedule, rules

Achievements & Milestones:
achievement_id, user_id, leaderboard_id, earned_at

Primary-Replica Setup (3 nodes)

Read Replicas for Analytics

Background Jobs (Workers)

Scheduled Tasks:
— Daily Leaderboard Reset (00:00 UTC)
— Weekly Leaderboard Reset (Monday 00:00)
— Monthly Snapshot Generation
— Historical Data Archival (move to S3)
— Achievement Detection & Notification
— Inactive User Cleanup
— Leaderboard Analytics Computation

Cache Layer (Redis)

L1 Cache (Application Memory):
— Top 100 cached for 10 seconds
— User's rank cached for 30 seconds
— Leaderboard metadata cached for 5 minutes

L2 Cache (Redis):
— Top 1000 cached for 1 minute
— Regional top 100 cached for 1 minute
— Friends leaderboard cached for 1 minute
— User profile data cached for 10 minutes

Archive Storage (Amazon S3 / GCS)

- Historical leaderboard snapshots (daily exports)
- Analytics data (user behavior, trends)
- Audit logs (compliance)
- Backup data (disaster recovery)

Monitoring & Observability Stack

Metrics (Prometheus + Grafana):

- Score update throughput
- Query latency (P50, P95, P99)
- Cache hit rates
- Redis memory usage
- Error rates by endpoint

Logs (ELK Stack):

- Application logs
- Access logs
- Error logs with stack traces

Tracing (Jaeger / Zipkin):

- Request flow across services
- Performance bottleneck identification

Alerting (PagerDuty / Opsgenie):

- High error rates
- Latency spikes
- Service unavailability
- Redis memory threshold breaches

Key Architectural Decisions

1. Redis Sorted Sets as Primary Storage

- **Why:** Native support for ranked data structures (ZSET)
- **Capabilities:**
 - $O(\log N)$ for score updates
 - $O(\log N)$ for rank lookups
 - $O(\log N + M)$ for range queries
 - Atomic operations (no race conditions)
- **Trade-off:** In-memory only → requires careful capacity planning

2. Separate Read and Write Paths

- **Write Path:** Score updates → Kafka → Redis
 - Asynchronous processing
 - Buffering during spikes
 - Retry mechanisms
- **Read Path:** Direct Redis queries
 - Fast, synchronous responses
 - Heavy caching
- **Trade-off:** Slight write delay (< 1s) for better read performance

3. Multi-Tier Caching

```
Request → L1 (In-Memory) → L2 (Redis Cache) → Redis ZSET → PostgreSQL
          ↓ Hit (90%)      ↓ Hit (9%)       ↓ Hit (1%)     ↓ Rare
          0.1ms latency    1ms latency     10ms latency   50ms latency
```

4. Time-Based Leaderboard Management

Strategy: Separate Redis keys with TTL

```
Daily: "lb:daily:2024-01-15" (TTL: 7 days)
Weekly: "lb:weekly:2024-W03" (TTL: 4 weeks)
Monthly: "lb:monthly:2024-01" (TTL: 12 months)
All-time: "lb:global:alltime" (No TTL)
```

Reset Mechanism:

- New key created at reset time
- Old key kept for historical reference
- Archived to S3 after TTL expiry

5. Horizontal Partitioning Strategy

Partition by Leaderboard Type:

```
└── Shard 1: Global leaderboards
└── Shard 2: Regional leaderboards (Americas)
└── Shard 3: Regional leaderboards (EMEA)
└── Shard 4: Regional leaderboards (APAC)
└── Shard 5-20: Friends & custom leaderboards
```

Benefits:

- ✓ Isolate load from different leaderboard types
- ✓ Scale each type independently
- ✓ Better fault isolation

Core Components

1. Score Update Service

Purpose: Process score updates with validation and conflict resolution

Component Architecture:

```
Score Update API  
POST /leaderboard/update  
{  
  "user_id": "user_123",  
  "leaderboard_id": "global_alltime",  
  "score": 1500,  
  "metadata": {  
    "game_id": "g456",  
    "timestamp": 1640995200,  
    "session_id": "s789"  
  }  
}
```

↓

```
Validation Layer  
└─ Authentication: Valid user token  
└─ Authorization: User owns this score  
└─ Score Range: Within valid bounds  
└─ Rate Limiting: Max updates per user  
└─ Duplicate Detection: Prevent replays  
└─ Game Rule Validation: Score achievable
```

↓

```
Conflict Resolution Logic  
  
Strategy Options:  
└─ HIGHEST: Keep maximum score  
└─ LATEST: Keep most recent score  
└─ SUM: Add scores together  
└─ INCREMENT: Add delta to current  
└─ CUSTOM: Game-specific logic  
  
Example (HIGHEST):  
  Current Score: 1200  
  New Score: 1500  
  Result: 1500 (new score higher, update)
```

↓

```
Publish to Kafka
Topic: "score-updates"
Partition Key: user_id (ensures ordering)
Message:
{
    "user_id": "user_123",
    "leaderboard_id": "global_alltime",
    "score": 1500,
    "previous_score": 1200,
    "timestamp": 1640995200,
    "operation": "UPDATE"
}
```

↓

```
Kafka Consumer (Worker Nodes)
├── Consume in batches (100 messages)
├── Group by leaderboard_id
└── Execute Redis pipeline
└── Commit offset after success
```

↓

Update Redis Sorted Set

Redis Command:

```
ZADD leaderboard:global:alltime 1500 user_123
```

Complexity: $O(\log N)$ where $N = \text{total users}$

Performance: ~0.1ms for 100M users

Atomic Operations (Lua Script):

```
local old_score = redis.call('ZSCORE', key, member)
if old_score == nil or new_score > old_score then
    redis.call('ZADD', key, new_score, member)
    return {old_score, new_score}
end
return {old_score, old_score}
```

↓

Invalidate Related Caches

```
├── User's rank cache
├── Top 100 if user was in top
├── User's neighborhood view
└── Friends leaderboard containing user
```

↓

Trigger Side Effects (Async)

```
├── Achievement Detection: Did user hit milestone?
└── Notification: Send rank change alerts
```

```
└── Analytics: Log score change event  
└── Audit: Record for compliance
```

Error Handling:

Retry Strategy:

```
└── Kafka Consumer Retries: 3 attempts with exponential backoff  
└── Redis Connection Failures: Circuit breaker pattern  
└── Validation Failures: Return 400 Bad Request immediately  
└── Idempotency: Use unique transaction IDs to prevent duplicates
```

2. Rank Query Service

Purpose: Fetch user's current rank efficiently

API Endpoint:

```
GET /leaderboard/rank?user_id=user_123&leaderboard_id=global_alltime
```

Response:

```
{  
  "user_id": "user_123",  
  "rank": 1547,  
  "score": 1500,  
  "total_users": 10000000,  
  "percentile": 99.98,  
  "rank_change": +23 (compared to yesterday)  
}
```

Implementation Flow:

```
Step 1: Check L1 Cache (In-Memory)  
Key: "rank:user_123:global_alltime"  
TTL: 30 seconds  
Hit Rate: ~70%  
Latency: 0.1ms
```

↓ (cache miss)

```
Step 2: Query Redis Sorted Set  
Command: ZREVRANK leaderboard:global:alltime user_123  
Returns: 1546 (0-indexed, so rank is 1547)  
Complexity: O(log N)  
Latency: ~1ms for 100M users
```

↓

Step 3: Get Score
Command: ZSCORE leaderboard:global:alltime user_123
Returns: 1500
Complexity: O(1)
Latency: ~0.5ms

↓

Step 4: Calculate Percentile
Command: ZCARD leaderboard:global:alltime
Returns: 10000000 (total users)
Percentile: $(1 - 1547/10000000) \times 100 = 99.98\%$

↓

Step 5: Get Rank Change (Historical)
Query: PostgreSQL historical_rankings table
SELECT rank FROM historical_rankings
WHERE user_id = 'user_123'
AND leaderboard_id = 'global_alltime'
AND date = CURRENT_DATE - 1
Result: Previous rank = 1570
Change: $1547 - 1570 = -23$ (improved by 23 positions)

↓

Step 6: Cache Result
Store in L1 cache with 30-second TTL
Return to client

Optimization: Batch Rank Queries:

For multiple users (e.g., friends list):

Naive Approach:
For each friend:
ZREVRANK leaderboard friend_id
Total: $N \times O(\log M) = O(N \log M)$

Optimized Approach:
Pipeline all ZREVRANK commands
Execute in single Redis round trip
Total: $O(N \log M)$ but single network call
Speedup: 10–100x depending on network latency

3. Top N Query Service

Purpose: Retrieve top N players efficiently

API Endpoint:

GET /leaderboard/top?leaderboard_id=global_alltime&limit=100&offset=0

Response:

```
{  
  "leaderboard_id": "global_alltime",  
  "updated_at": 1640995200,  
  "total_users": 10000000,  
  "entries": [  
    {  
      "rank": 1,  
      "user_id": "user_999",  
      "username": "ProGamer",  
      "score": 9999,  
      "avatar_url": "...",  
      "country": "US"  
    },  
    ...  
  ]  
}
```

Implementation Flow:

Step 1: Check CDN Cache

URL: /leaderboard/top?lb=global&limit=100

TTL: 10 seconds

Hit Rate: ~95% (very cacheable)

Latency: <10ms (edge location)

↓ (cache miss)

Step 2: Query Redis (Top N)

Command:

ZREVRANGE leaderboard:global:alltime 0 99 WITHSCORES

Returns:

[user_999, 9999, user_888, 8888, ..., user_100, 1000]

Complexity: O(log N + M) where M = 100

Latency: ~1-2ms

↓

Step 3: Enrich with User Data (Batch)

Fetch user profiles for all 100 users in parallel:

- Username, avatar, country from PostgreSQL

- Use connection pooling
 - Cache results for 5 minutes
- Latency: ~20ms (parallel queries)

↓

Step 4: Format Response & Cache
 Combine Redis scores + PostgreSQL user data
 Cache in CDN for 10 seconds
 Return to client
 Total latency: ~30–50ms

Pagination Strategy:

For large result sets (e.g., top 1000):

Cursor-based:

- GET /leaderboard/top?cursor=user_500&limit=100
- More consistent with changing data
- Better for infinite scroll
- Use ZRANK to find position

Offset-based:

- GET /leaderboard/top?offset=500&limit=100
- Simpler to implement
- Better for page numbers
- May skip entries if data changes

4. User Neighborhood Service

Purpose: Show users around a specific user's rank

API Endpoint:

GET /leaderboard/neighbors?user_id=user_123&range=10

Response:

```
{
  "user_id": "user_123",
  "rank": 1547,
  "neighbors": {
    "above": [
      {"rank": 1537, "user_id": "user_abc", "score": 1510},
      ...
      {"rank": 1546, "user_id": "user_xyz", "score": 1501}
    ],
    "current": {
      "rank": 1547,
```

```

        "user_id": "user_123",
        "score": 1500
    },
    "below": [
        {"rank": 1548, "user_id": "user_def", "score": 1499},
        ...
        {"rank": 1557, "user_id": "user_ghi", "score": 1490}
    ]
}

```

Implementation:

Redis Command:

```

# Get user's rank first
ZREVRANK leaderboard:global:alltime user_123 → 1546

# Get users around this rank
ZREVRANGE leaderboard:global:alltime 1536 1556 WITHSCORES

Returns 21 users (10 above + user + 10 below)
Complexity: O(log N + M) where M = 21

```

Data Model & Storage Design

Redis Data Structures

1. Sorted Set (ZSET) - Core Leaderboard

Key: "leaderboard:{type}:{scope}:{period}"
Member: user_id
Score: user's score (numeric)

Examples:

Key: "leaderboard:global:alltime"

```

user_999 → 9999
user_888 → 8888
user_777 → 7777
...
user_123 → 1500
...
user_001 → 100

```

Total: 100M members, ~5GB memory

Operations:

- └── ZADD: $O(\log N)$ – Add/update score
- └── ZREVRANK: $O(\log N)$ – Get rank
- └── ZREVRANGE: $O(\log N + M)$ – Get top M
- └── ZSCORE: $O(1)$ – Get score
- └── ZCARD: $O(1)$ – Total members
- └── ZINCRBY: $O(\log N)$ – Increment score

2. String - User Rank Cache

Key: "rank:{user_id}:{leaderboard_id}"
Value: JSON {rank: 1547, score: 1500, updated_at: 1640995200}
TTL: 30 seconds

Purpose: Cache frequently queried ranks
Memory: ~100 bytes per cached entry
Cache size: 10M active users \times 100B = 1GB

3. List - Recent Score Updates

Key: "recent_updates:{leaderboard_id}"
Value: List of recent score changes
Structure: LPUSH new updates, LTRIM to keep last 1000

Purpose: Real-time activity feed, leaderboard changes
Memory: ~50KB per leaderboard

4. Set - Active Users

Key: "active_users:{leaderboard_id}:{date}"
Value: Set of user_ids who scored today
TTL: 48 hours

Purpose: Track daily/weekly active users
Memory: 10M users \times 16 bytes = 160MB

PostgreSQL Schema

```
-- Users table (10M–100M rows)
CREATE TABLE users (
    user_id UUID PRIMARY KEY,
```

```

username VARCHAR(50) UNIQUE NOT NULL,
email VARCHAR(255) UNIQUE NOT NULL,
country_code CHAR(2),
region VARCHAR(50),
avatar_url TEXT,
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
last_active_at TIMESTAMP,
INDEX idx_username (username),
INDEX idx_country (country_code),
INDEX idx_last_active (last_active_at)
);

-- Leaderboard configurations (hundreds of rows)
CREATE TABLE leaderboards (
    leaderboard_id VARCHAR(100) PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    description TEXT,
    type VARCHAR(50), -- 'global', 'regional', 'friends', 'custom'
    scope VARCHAR(50), -- 'alltime', 'daily', 'weekly', 'monthly'
    reset_schedule VARCHAR(50), -- cron expression
    score_update_strategy VARCHAR(20), -- 'highest', 'latest', 'sum'
    tie_breaker VARCHAR(20) DEFAULT 'timestamp', -- 'timestamp',
    user_id
    is_active BOOLEAN DEFAULT true,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Historical rankings (billions of rows, partitioned)
CREATE TABLE historical_rankings (
    id BIGSERIAL,
    snapshot_id VARCHAR(50),
    leaderboard_id VARCHAR(100),
    user_id UUID,
    rank INTEGER,
    score BIGINT,
    snapshot_date DATE,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    PRIMARY KEY (id, snapshot_date),
    INDEX idx_user_leaderboard_date (user_id, leaderboard_id,
snapshot_date),
    INDEX idx_leaderboard_date (leaderboard_id, snapshot_date)
) PARTITION BY RANGE (snapshot_date);

-- Create partitions for each month
CREATE TABLE historical_rankings_2024_01 PARTITION OF
historical_rankings
FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');

-- Achievements (millions of rows)
CREATE TABLE achievements (
    achievement_id VARCHAR(100) PRIMARY KEY,
    name VARCHAR(255),
    description TEXT,

```

```

threshold_type VARCHAR(50), -- 'rank', 'score', 'streak'
threshold_value INTEGER,
icon_url TEXT,
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

CREATE TABLE user_achievements (
    id BIGSERIAL PRIMARY KEY,
    user_id UUID NOT NULL,
    achievement_id VARCHAR(100) NOT NULL,
    leaderboard_id VARCHAR(100),
    earned_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    metadata JSONB,
    INDEX idx_user_achievements (user_id),
    INDEX idx_achievement (achievement_id),
    UNIQUE (user_id, achievement_id, leaderboard_id)
);

-- Score audit log (for compliance, billions of rows)
CREATE TABLE score_audit_log (
    id BIGSERIAL PRIMARY KEY,
    user_id UUID NOT NULL,
    leaderboard_id VARCHAR(100),
    old_score BIGINT,
    new_score BIGINT,
    operation VARCHAR(20), -- 'UPDATE', 'DELETE', 'RESET'
    ip_address INET,
    user_agent TEXT,
    session_id VARCHAR(100),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
) PARTITION BY RANGE (created_at);

```

Data Retention Policy

Retention Strategy:

Hot Data (Redis):

- └─ Current leaderboards: Keep in memory
- └─ Daily leaderboards: TTL 7 days
- └─ Weekly leaderboards: TTL 4 weeks
- └─ Monthly leaderboards: TTL 12 months
- └─ Cache entries: TTL 30 seconds – 5 minutes

Warm Data (PostgreSQL):

- └─ Historical rankings: Keep 2 years
- └─ Partition by month
- └─ Archive older data to S3
- └─ Use read replicas for analytics

Cold Data (S3):

- Archived historical snapshots: Keep 5 years
- Compressed (gzip)
- Parquet format for analytics
- Glacier for long-term storage (7+ years)

Lifecycle:

Redis (live) → PostgreSQL (2 years) → S3 (5 years) → Glacier (permanent)

Real-time Updates & Data Flow

Write Path (Score Update)

PHASE 1: Client Submission

- Client submits score
- POST /leaderboard/update
- Latency: 0ms (start)

↓

PHASE 2: API Gateway Processing

- Authentication (JWT validation): 5ms
 - Request validation: 1ms
 - Rate limiting check: 1ms
- Total: 7ms

↓

PHASE 3: Score Update Service

- Game rules validation: 3ms
 - Conflict resolution logic: 2ms
 - Publish to Kafka: 5ms
 - Return 202 Accepted to client: 10ms
- Total: 20ms (client sees response here)

↓

PHASE 4: Asynchronous Processing (Kafka)

- Message queued in Kafka: +1ms
- Consumer picks up message: +100–500ms (batching)
- Process batch (100 messages): 50ms

↓

PHASE 5: Redis Update

- Execute ZADD command: 0.5ms
 - Replicate to slaves: 1–5ms (async)
 - Update multiple leaderboards: 2ms
- Total: 3–8ms

↓

PHASE 6: Side Effects (Parallel)

- Cache invalidation: 1ms
 - Achievement detection: 10ms
 - Notification service: 20ms
 - Analytics logging: 5ms
- (All happen in background, don't block)

Total Perceived Latency (client): ~20ms

Total Actual Processing Time: ~600ms (async)

Read Path (Rank Query)

PHASE 1: Request Arrives

GET /leaderboard/rank?user_id=user_123
Latency: 0ms (start)

↓

PHASE 2: L1 Cache Check (In-Memory)

- Lookup in local cache: 0.1ms
 - Hit rate: 70%
 - If HIT: Return immediately
- Total: 0.1ms (if hit)

↓ (if miss)

PHASE 3: Redis Query

- Connection from pool: 0.1ms
 - ZREVRANK command: 1ms
 - ZSCORE command: 0.5ms
 - Network roundtrip: 0.5ms
- Total: 2ms

↓

PHASE 4: Historical Data (Optional)

- Query PostgreSQL for rank change: 10ms
 - Use read replica
 - Cached query plan
- Total: 10ms

↓

PHASE 5: Response Assembly

- Format JSON: 0.5ms

```
| - Cache result in L1: 0.1ms  
| - Return to client: 1ms  
Total: 1.5ms
```

Total Latency:

```
|-- Cache hit: 0.1ms  
|-- Redis only: 2-3ms  
|-- With historical: 12-15ms
```

Target: <100ms (P95)

Actual: <15ms (P99) ✓

Real-Time Update Notifications (WebSocket)

Architecture:

```
Client ←→ WebSocket → API Gateway ←→ Pub/Sub → Redis  
                                ↓  
      Subscribed to:  
      - User's rank changes  
      - Friends' score updates  
      - Top 10 changes
```

Implementation:

```
When score updates happen:  
1. Redis ZADD updates sorted set  
2. Lua script publishes to Redis Pub/Sub:  
   PUBLISH "leaderboard:updates:global" "user_123:1500"  
3. API Gateway subscribed to channel  
4. Forwards to connected WebSocket clients  
5. Client UI updates in real-time
```

Message Format:

```
{  
  "type": "rank_change",  
  "user_id": "user_123",  
  "old_rank": 1570,  
  "new_rank": 1547,  
  "score": 1500,  
  "timestamp": 1640995200  
}
```

Scalability:

```
|-- Use Redis Pub/Sub for broadcasting  
|-- Sticky sessions for WebSocket connections  
|-- Connection pooling per server  
|-- Limit subscriptions per user (max 10 leaderboards)
```

Ranking Algorithms

1. Standard Ranking (Descending by Score)

Algorithm: Rank by score, highest first

Example:

Rank	User ID	Score	Notes
1	user_999	9999	Highest
2	user_888	8888	
3	user_777	7777	
4	user_666	6666	
5	user_555	5555	

Redis Implementation:

```
ZREVRANGE leaderboard:global:alltime 0 -1 WITHSCORES  
(Sorted automatically by score descending)
```

Complexity: $O(\log N)$ for insert, $O(\log N)$ for rank lookup

2. Tie-Breaking Strategies

Problem: Multiple users with same score

Scenario:

Rank	User ID	Score	Joined Time
?	user_aaa	1000	10:00:00
?	user_bbb	1000	10:05:00
?	user_ccc	1000	10:10:00

How to rank them?

Strategy A: First Come, First Served

User who achieved score first gets higher rank

Implementation:

```
Store composite score: score + (1 / timestamp)
```

Example:

```
user_aaa: 1000.0000001 (achieved at 10:00:00)
user_bbb: 1000.0000000.8 (achieved at 10:05:00)
user_ccc: 1000.0000000.6 (achieved at 10:10:00)
```

Result:

```
Rank 1: user_aaa
Rank 2: user_bbb
Rank 3: user_ccc
```

Pros: Fair, rewards early achievement

Cons: Precision issues with large timestamps

Strategy B: Lexicographic (User ID)

Break ties by user ID (alphabetically)

Implementation:

```
Use Redis ZSET with same score
Redis maintains insertion order for same scores
Then sort by user_id for display
```

Result:

```
Rank 1: user_aaa (alphabetically first)
Rank 2: user_bbb
Rank 3: user_ccc
```

Pros: Simple, deterministic, no precision issues

Cons: Arbitrary (user ID doesn't mean much to users)

Strategy C: Multiple Criteria (Recommended)

Use secondary & tertiary criteria

Primary: Score (highest wins)

Secondary: Timestamp (earlier wins)

Tertiary: User ID (alphabetically)

Implementation:

```
Store in Redis: score as primary
Fetch tied users, sort by timestamp in application
Final sort by user_id if still tied
```

Pros: Most fair and intuitive

Cons: Slightly more complex

3. Time-Decay Ranking

Purpose: Keep leaderboard fresh, reward recent activity

Algorithm: Scores decay over time

Formula:

$$\text{effective_score} = \text{base_score} \times \text{decay_factor}^{\text{days_since_activity}}$$

Example (decay_factor = 0.9, decay per day):

User	Base Score	Days	Decay Factor	Effective
user_active	1000	0	$0.9^0 = 1.0$	1000
user_week	1200	7	$0.9^7 = 0.478$	574
user_month	1500	30	$0.9^{30} = 0.042$	63

Result:

user_active (1000) ranks higher than user_month (63)
Despite user_month having higher base score!

Implementation:

Daily cron job:

```
FOR each user IN leaderboard:  
    days_inactive = TODAY - last_activity_date  
    effective_score = base_score × (0.9 ^ days_inactive)  
    UPDATE Redis ZSET with effective_score
```

Pros: Encourages ongoing engagement

Cons: Computational overhead, may frustrate casual players

4. Percentile Ranking

Purpose: Show relative performance ("You're in top 5%")

Calculation:

$$\text{percentile} = (1 - \text{rank} / \text{total_users}) \times 100$$

Example:

Rank	Total Users	%ile	Interpretation
100	10,000,000	99.999	Top 0.001%
1000	10,000,000	99.99	Top 0.01%
10K	10,000,000	99.9	Top 0.1%
100K	10,000,000	99	Top 1%
1M	10,000,000	90	Top 10%
5M	10,000,000	50	Median

```
Display to User:  
"You're ranked 10,543 out of 10M players (top 0.1%)"  
"You're better than 99.9% of players!"
```

```
Redis Commands:  
ZREVRANK leaderboard user_id → rank  
ZCARD leaderboard → total_users  
Calculate: (1 - rank/total) × 100
```

Distributed System Design

Challenge 1: Handling Concurrent Updates

Problem: Multiple score updates for same user simultaneously

Scenario:

Time T0:

```
Game Server 1: user_123 scores 1000  
Game Server 2: user_123 scores 1200  
Both send updates simultaneously
```

Without synchronization:

```
Server 1 writes 1000 to Redis  
Server 2 writes 1200 to Redis  
OR vice versa → race condition
```

Desired behavior:

Keep highest score (1200)

Solution 1: Atomic Lua Scripts (Recommended)

```
-- Store in Redis as a Lua script  
-- Ensures atomic compare-and-set operation

local key = KEYS[1]
local member = ARGV[1]
local new_score = tonumber(ARGV[2])
local strategy = ARGV[3] -- 'highest', 'latest', 'sum'

local current_score = redis.call('ZSCORE', key, member)

if current_score == false then
    -- User not in leaderboard, add them
    redis.call('ZADD', key, new_score, member)
    return {nil, new_score, 'added'}
end
```

```

current_score = tonumber(current_score)

if strategy == 'highest' then
    if new_score > current_score then
        redis.call('ZADD', key, new_score, member)
        return {current_score, new_score, 'updated'}
    else
        return {current_score, current_score, 'unchanged'}
    end
elseif strategy == 'latest' then
    redis.call('ZADD', key, new_score, member)
    return {current_score, new_score, 'updated'}
elseif strategy == 'sum' then
    local sum = current_score + new_score
    redis.call('ZADD', key, sum, member)
    return {current_score, sum, 'updated'}
end

```

Usage:

```
EVALSHA <script_sha> 1 leaderboard:global:alltime user_123 1200 highest
```

Returns: [1000, 1200, "updated"]

- Previous score was 1000
- New score is 1200
- Status: updated

Atomic guarantee: No race conditions possible

Solution 2: Optimistic Locking

Alternative for complex logic:

1. Read current score + version number
 GET leaderboard:user_123:version → v5
 ZSCORE leaderboard user_123 → 1000
2. Perform logic in application

$$\text{new_score} = \max(1000, 1200) = 1200$$
3. Conditional write with version check
 WATCH leaderboard:user_123:version
 IF version == v5:
 ZADD leaderboard 1200 user_123
 INCR leaderboard:user_123:version
 ELSE:
 RETRY

Pros: Supports complex logic
Cons: Retries needed, higher latency

Challenge 2: Redis Cluster Partitioning

Problem: 100M users don't fit in single Redis instance

Solution: Consistent Hashing with Redis Cluster

Redis Cluster Setup:

```
16,384 hash slots distributed across nodes  
  
Node 1 (Master): Slots 0-5460      → Replica 1  
Node 2 (Master): Slots 5461-10922  → Replica 2  
Node 3 (Master): Slots 10923-16383 → Replica 3
```

Hash Slot Calculation:

```
slot = CRC16(key) % 16384
```

Example:

```
Key: "leaderboard:global:alltime"  
Hash: CRC16("leaderboard:global:alltime") = 12345  
Slot: 12345 % 16384 = 12345  
Node: Node 3 (slots 10923-16383)
```

Routing:

```
Client → Redis Cluster  
Cluster returns: MOVED 12345 node3:6379  
Client connects to node3 directly
```

Alternative: Application-Level Sharding

Strategy: Partition by leaderboard type

Shard Assignment:

```
Shard 1 (Redis Instance): Global leaderboards  
- leaderboard:global:alltime  
- leaderboard:global:daily  
- leaderboard:global:weekly  
  
Shard 2-5: Regional leaderboards  
- leaderboard:region:us:*
```

Shard 6-20: Friends leaderboards - Partition by user_id hash - leaderboard:friends:user_*

Routing Logic (Application):

```
def get_shard(leaderboard_id):
    if leaderboard_id.startswith('global'):
        return shard_1
    elif leaderboard_id.startswith('region'):
        region = extract_region(leaderboard_id)
        return region_shard_map[region]
    else: # friends
        user_id = extract_user_id(leaderboard_id)
        return hash(user_id) % num_friend_shards + 6
```

Pros:

- ✓ Isolate load by leaderboard type
- ✓ Scale each type independently
- ✓ Better debuggability

Cons:

- ✗ Application logic complexity
- ✗ Can't use Redis Cluster features

Challenge 3: Cross-Region Consistency

Problem: Users in different regions see different ranks

Solution: Multi-Region Architecture

Option 1: Regional Independence (Recommended)

Each region has own Redis cluster

US-WEST: leaderboard:global:us-west
US-EAST: leaderboard:global:us-east
EU-WEST: leaderboard:global:eu-west

Async Replication: Every 5 minutes
1. Export top 10K from each region
2. Merge in central aggregator
3. Redistribute merged result

Trade-off:

Slightly stale cross-region ranks (5 min lag)
Fast local queries (< 10ms)

Option 2: Global Redis (High Consistency)

Single Redis cluster for global leaderboard

Primary: US-EAST (central location)

Replicas: US-