

# Logging and Metrics System Design - High-Level Design Focus

---

## Table of Contents

1. Problem Statement & Requirements
  2. High-Level Architecture
  3. Core Components
  4. Storage Architecture
  5. Data Flow & Ingestion Pipeline
  6. Distributed System Design
  7. Scalability & Performance
  8. Fault Tolerance & Reliability
  9. Multi-Region Architecture
  10. Design Trade-offs
  11. Real-World Case Studies
- 

## Problem Statement & Requirements

### Problem Definition

Design a distributed observability system that can:

- Collect and aggregate logs from 10,000+ servers
- Store and query metrics with millisecond precision
- Process 1TB+ of logs per day
- Support real-time alerting and dashboards
- Retain data for compliance (90 days hot, 1 year cold)

### Use Cases

#### 1. Logging System (ELK-style)

- Application error tracking and debugging
- Security audit logs
- User activity tracking
- System performance logs
- Business analytics

#### 2. Metrics System (Prometheus-style)

- Real-time monitoring (CPU, memory, disk)
- Application performance (latency, throughput, error rate)
- Business metrics (orders/sec, revenue, user signups)
- Custom metrics from applications

- Infrastructure health

## Functional Requirements

### Logging

1. **Collection:** Ingest logs from multiple sources (apps, systems, containers)
2. **Parsing:** Extract structured data from unstructured logs
3. **Search:** Full-text search with millisecond latency
4. **Aggregation:** Group and aggregate log data
5. **Alerting:** Trigger alerts on log patterns

### Metrics

1. **Collection:** Scrape metrics from endpoints at configurable intervals
2. **Storage:** Efficient time-series storage
3. **Querying:** Query language for analysis and aggregation
4. **Visualization:** Real-time dashboards and graphs
5. **Alerting:** Alert on metric thresholds and anomalies

## Non-Functional Requirements

### 1. Performance

- Log ingestion: 1M events/second
- Metrics ingestion: 10M samples/second
- Query latency: <1 second for recent data
- Dashboard refresh: <5 seconds

### 2. Scalability

- Handle 10,000+ hosts
- Store 1TB+ logs/day
- Store 100M+ time-series
- Support 1,000+ concurrent users

### 3. Availability

- 99.9% uptime for ingestion
- 99.95% uptime for queries
- No data loss during failures
- Graceful degradation

### 4. Retention

- Hot storage: 7-90 days (fast queries)
- Cold storage: 1-7 years (slower, cheaper)
- Configurable per data source

## Capacity Estimation

#### Logging System:

- Servers: 10,000
- Logs per server: 1,000 lines/minute
- Total log rate: 10M lines/minute = 167K lines/second
- Average log size: 500 bytes
- Ingestion rate: 83 MB/second = 7.2 TB/day
- Storage (90 days): 648 TB

#### Metrics System:

- Servers: 10,000
- Metrics per server: 100 metrics
- Scrape interval: 15 seconds
- Total metrics: 1M metrics
- Samples per day:  $1M \times (86400/15) = 5.76B$  samples
- Sample size: 16 bytes (timestamp + value)
- Daily storage: 92 GB
- Storage (90 days): 8.3 TB

Total Storage: ~656 TB (hot) + archives (cold)

#### Infrastructure:

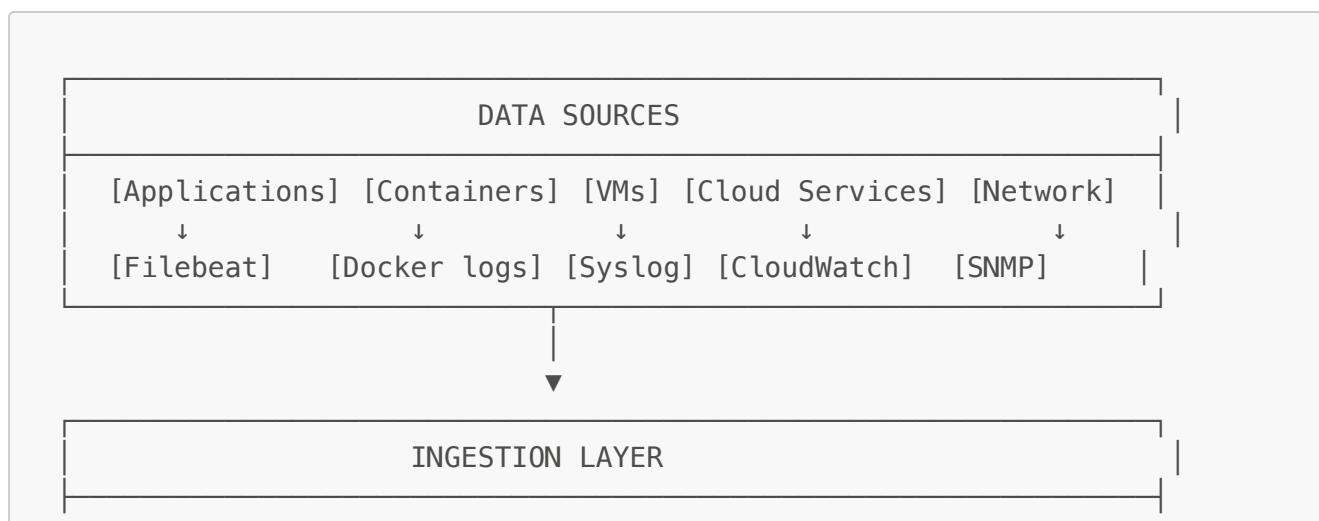
- Ingestion nodes: 20–30 (handle peaks)
- Storage nodes: 50–100 (distributed)
- Query nodes: 10–20 (dedicated)
- Cache nodes: 5–10 (Redis for hot data)

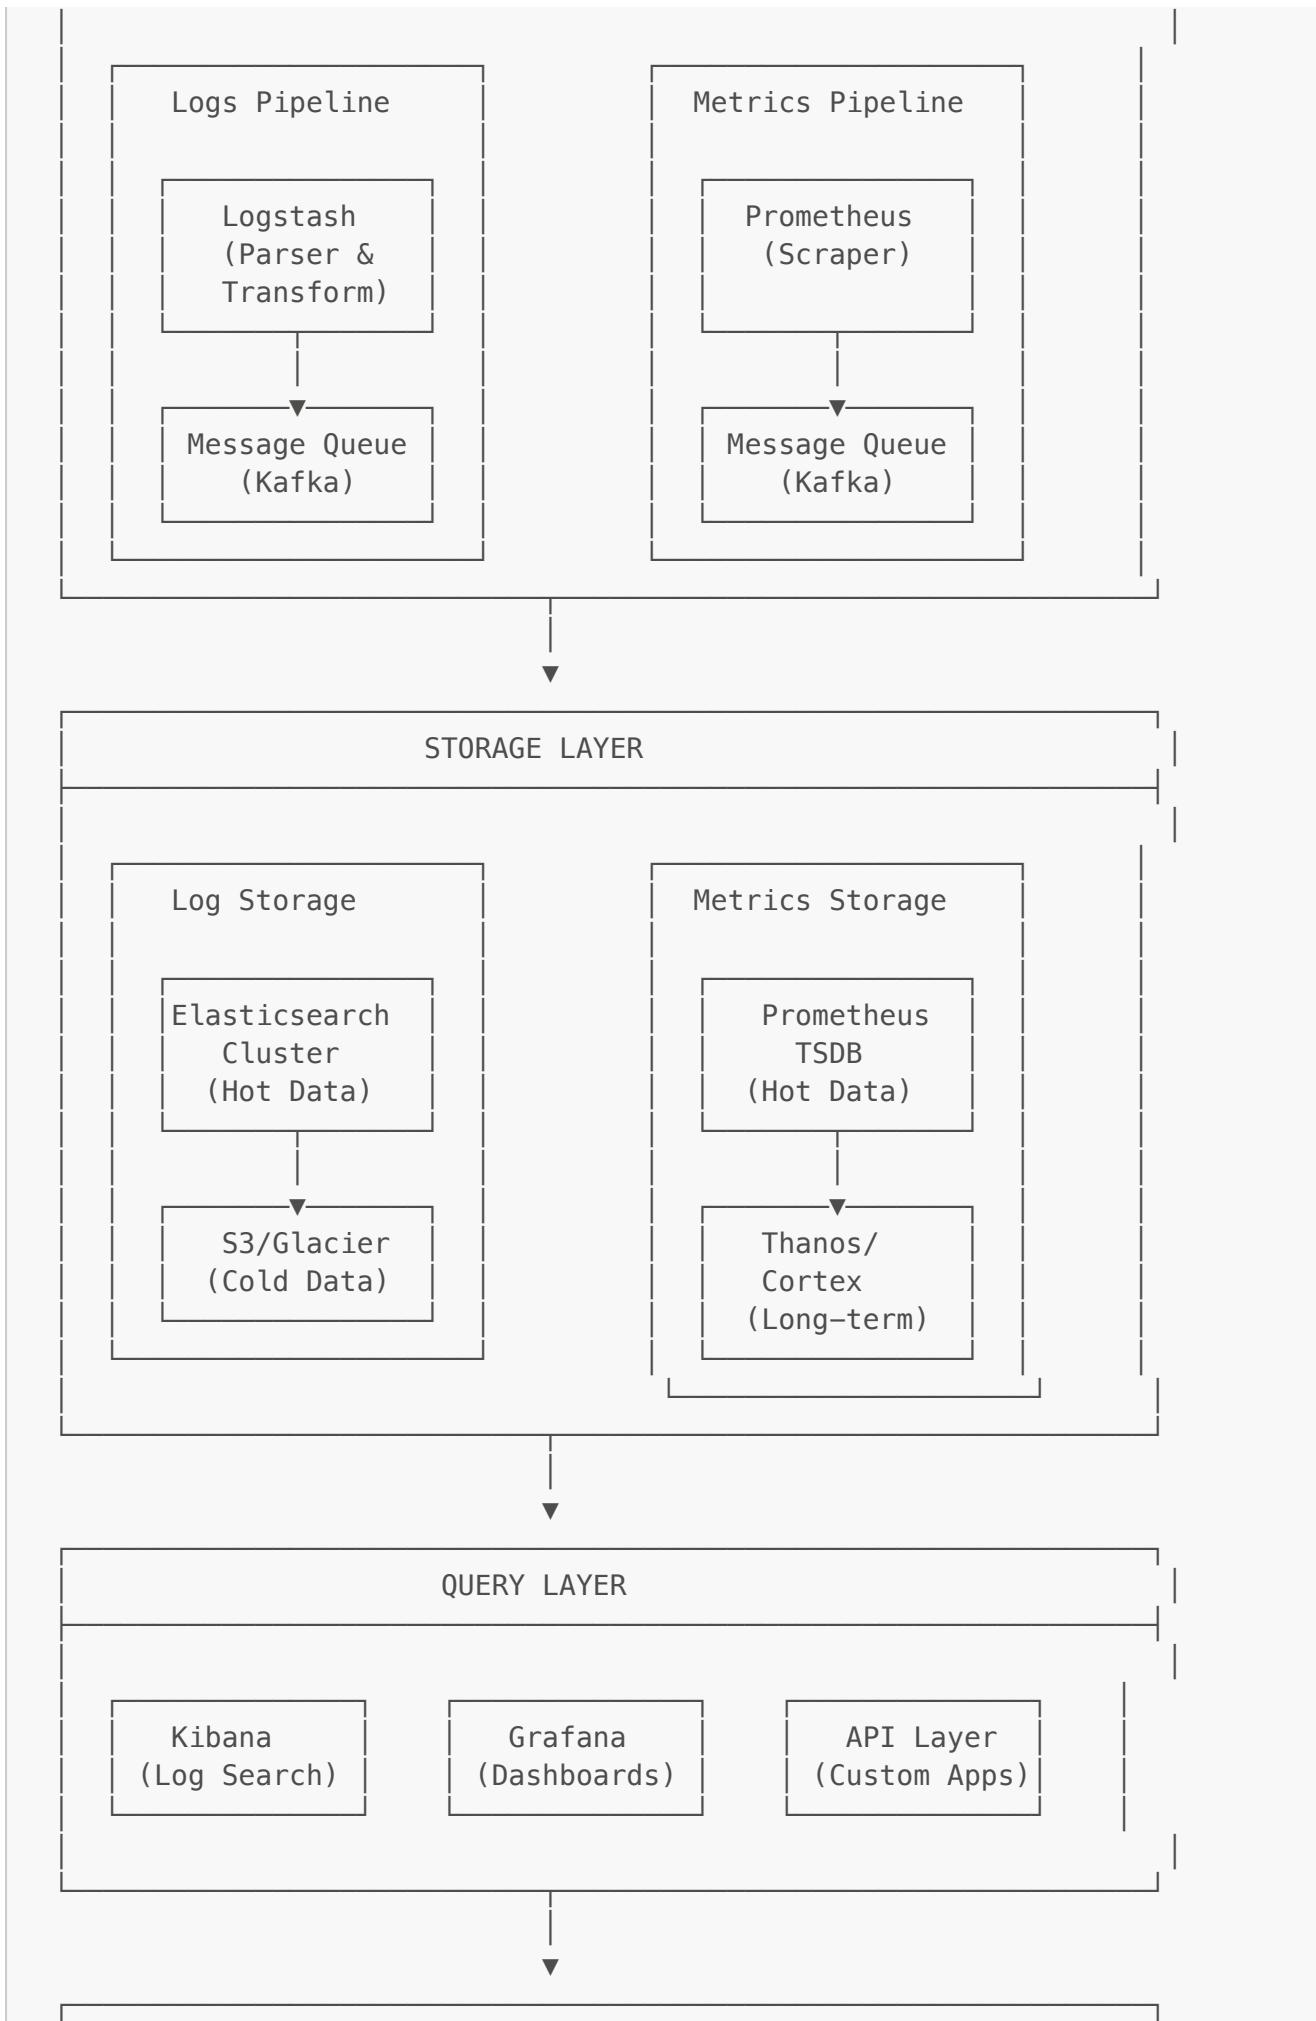
#### Cost Estimate (AWS):

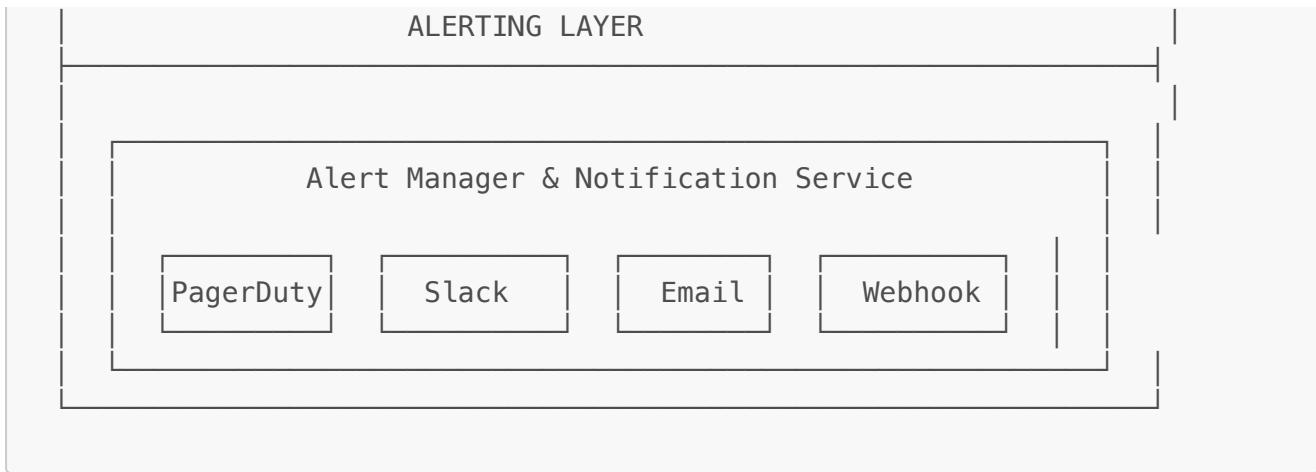
- Compute (EC2): \$15,000/month
- Storage (EBS/S3): \$25,000/month
- Network: \$5,000/month
- Total: ~\$45,000/month

## High-Level Architecture

### System Overview







## Architectural Patterns

### 1. Lambda Architecture (Batch + Stream)

#### Real-time Layer (Stream):

- └── Kafka → Stream Processor → Hot Storage
- └── Latency: <1 second
- └── Retention: 7–90 days
- └── Use: Dashboards, real-time alerts

#### Batch Layer:

- └── S3 → Spark Jobs → Data Warehouse
- └── Latency: Hours
- └── Retention: Years
- └── Use: Analytics, compliance, ML training

#### Serving Layer:

- └── Combines results from both
- └── Provides unified query interface
- └── Handles data quality and corrections

### 2. Kappa Architecture (Stream Only)

#### Single Pipeline:

- └── Kafka → Stream Processor → Storage
- └── Reprocess historical data if needed
- └── Simpler than Lambda
- └── Preferred for most use cases

#### Benefits:

- ✓ Single codebase
- ✓ Easier operations
- ✓ Sufficient for most requirements

# Key Design Decisions

## 1. Push vs Pull Model

### Logs: Push Model (Recommended)

Why:

- Logs are generated continuously
- Applications can't be "polled"
- Need immediate delivery for errors
- Buffering handles backpressure

Implementation:

App → Agent (Filebeat) → Logstash → Kafka → Storage

### Metrics: Pull Model (Prometheus)

Why:

- Service discovery easier
- Target health visible
- Prevents overwhelming targets
- Metrics collected at consistent intervals

Implementation:

Prometheus → Scrape → App /metrics endpoint

### Metrics: Push Model (Alternative)

Why:

- Short-lived jobs (cron, Lambda)
- Behind firewalls
- Dynamic infrastructure

Implementation:

App → Push Gateway → Prometheus

## 2. Hot vs Cold Storage Tiering

Hot Tier (Fast, Expensive):

- └ Recent data (7–90 days)
- └ SSD storage
- └ High IOPS
- └ Low latency queries (<100ms)
- └ Examples: Elasticsearch, Prometheus

#### Warm Tier (Medium):

- └ Less recent (90–180 days)
- └ Slower SSD or fast HDD
- └ Medium IOPS
- └ Acceptable latency (1–5s)
- └ Examples: S3 Standard, Clickhouse

#### Cold Tier (Slow, Cheap):

- └ Archives (>180 days)
- └ Object storage
- └ Low IOPS
- └ Higher latency (minutes)
- └ Examples: S3 Glacier, Google Coldline

#### Benefits:

- ✓ 80% cost reduction
- ✓ Meets compliance requirements
- ✓ Transparent to users (federated queries)

### 3. Centralized vs Distributed

#### Centralized Collection:

All logs → Central cluster

#### Pros:

- ✓ Simple operations
- ✓ Easier to query
- ✓ Lower complexity

#### Cons:

- ✗ Single point of failure
- ✗ Network bandwidth
- ✗ Compliance issues (data locality)

#### Distributed (Recommended):

Regional clusters → Federated queries

#### Pros:

- ✓ Data locality
- ✓ Reduced network costs
- ✓ Failure isolation
- ✓ Compliance friendly

#### Cons:

- ✗ Complex operations
- ✗ Harder to query across regions

# Core Components

## 1. Log Collection Agents

**Purpose:** Collect and ship logs from sources

### Filebeat (Elastic Stack)

```
Features:  
└── Lightweight (Go-based)  
└── Tails log files  
└── Handles log rotation  
└── Built-in modules (nginx, mysql, etc.)  
└── Multiline support  
└── Back-pressure handling  
  
Configuration:  
filebeat.inputs:  
  - type: log  
    paths:  
      - /var/log/app/*.log  
    fields:  
      app: my-app  
      env: production  
  
output.logstash:  
  hosts: ["logstash-1:5044", "logstash-2:5044"]  
  loadbalance: true
```

### Fluentd (CNCF)

```
Features:  
└── Plugin ecosystem (500+ plugins)  
└── JSON structured logging  
└── Memory and file buffering  
└── Tag-based routing  
└── Kubernetes native
```

```
Configuration:  
<source>  
  @type tail  
  path /var/log/app/*.log  
  tag app.logs  
  <parse>  
    @type json  
  </parse>  
</source>
```

```
<match app.logs>
  @type kafka2
  brokers kafka1:9092,kafka2:9092
  topic application-logs
</match>
```

## 2. Log Processing Pipeline

**Purpose:** Parse, transform, enrich logs

### Logstash

Pipeline Stages:

Input → Filter → Output

Input:

- └── Beats (Filebeat, Metricbeat)
- └── Kafka
- └── HTTP
- └── S3
- └── Database CDC

Filter:

- └── Grok (pattern matching)
- └── JSON parser
- └── Date parser
- └── GeoIP enrichment
- └── User-Agent parsing
- └── Custom Ruby code

Output:

- └── Elasticsearch
- └── Kafka
- └── S3
- └── Multiple outputs

### Example Pipeline:

```
input {
  beats {
    port => 5044
  }
}

filter {
  # Parse nginx access logs
  grok {
```

```

    match => {
      "message" => '%{IPORHOST:clientip} -- \[%{HTTPDATE:timestamp}\]
      "%{WORD:method} %{URIPATHPARAM:request} HTTP/%{NUMBER:httpversion}" %
      %{NUMBER:response} %{NUMBER:bytes}"
    }
  }

  # Convert response to integer
  mutate {
    convert => { "response" => "integer" }
    convert => { "bytes" => "integer" }
  }

  # Add GeoIP data
  geoip {
    source => "clientip"
  }

  # Parse timestamp
  date {
    match => [ "timestamp", "dd/MMM/YYYY:HH:mm:ss Z" ]
    target => "@timestamp"
  }
}

output {
  elasticsearch {
    hosts => ["es-1:9200", "es-2:9200"]
    index => "nginx-logs-%{+YYYY.MM.dd}"
  }
}

```

### 3. Message Queue (Kafka)

**Purpose:** Buffer and decouple ingestion from storage

#### Why Kafka?

- └── High throughput (millions/sec)
- └── Durability (replication)
- └── Ordered delivery
- └── Replay capability
- └── Backpressure handling
- └── Multiple consumers

#### Topic Design:

- └── application-logs (partitioned by app\_id)
- └── system-metrics (partitioned by host)
- └── error-logs (single partition for ordering)
- └── audit-logs (compacted for deduplication)

```
Configuration:  
└── Replication factor: 3  
└── Retention: 7 days  
└── Compression: LZ4  
└── Batch size: 1MB
```

## 4. Log Storage (Elasticsearch)

**Purpose:** Store and index logs for fast search

```
Architecture:  
└── Master nodes (3): Cluster coordination  
└── Data nodes (20+): Store and query data  
└── Ingest nodes (5): Pre-processing  
└── Coordinating nodes (5): Route requests  
└── Machine learning nodes (3): Anomaly detection
```

```
Index Structure:  
logs-application-2025.01.08  
└── Shards: 5 primary  
└── Replicas: 1 per shard  
└── Refresh interval: 5s (near real-time)  
└── Lifecycle: Hot → Warm → Cold → Delete
```

Index Template:

```
{  
    "index_patterns": ["logs-*"],  
    "settings": {  
        "number_of_shards": 5,  
        "number_of_replicas": 1,  
        "index.lifecycle.name": "logs-policy"  
    },  
    "mappings": {  
        "properties": {  
            "@timestamp": {"type": "date"},  
            "level": {"type": "keyword"},  
            "message": {"type": "text"},  
            "host": {"type": "keyword"},  
            "app_id": {"type": "keyword"}  
        }  
    }  
}
```

### Index Lifecycle Management:

```
Phase 1: Hot (0–7 days)  
└── Fast SSDs  
└── High resource allocation
```

```

    └── All queries hit this tier
        └── Rollover at 50GB or 7 days

Phase 2: Warm (7–30 days)
    ├── Slower SSDs
    ├── Reduced replicas (1 → 0)
    ├── Force merge to reduce segments
    └── Shrink shards (5 → 1)

Phase 3: Cold (30–90 days)
    ├── Object storage (S3)
    ├── Searchable snapshots
    ├── Slower queries (acceptable)
    └── 90% cost reduction

Phase 4: Delete (>90 days)
    ├── Remove from cluster
    ├── Optional archive to Glacier
    └── Meets retention policy

```

## 5. Metrics Collection (Prometheus)

**Purpose:** Scrape and store time-series metrics

```

Architecture:
Prometheus Server
    ├── Scraper (pulls metrics)
    ├── TSDB (time-series database)
    ├── Query engine (PromQL)
    └── Alertmanager integration

Service Discovery:
    ├── Static configs
    ├── Kubernetes (pods, services)
    ├── AWS (EC2, ECS)
    ├── Consul, etcd
    └── DNS-based

Scrape Configuration:
scrape_configs:
  - job_name: 'api-servers'
    scrape_interval: 15s
    kubernetes_sd_configs:
      - role: pod
        namespaces:
          names: ['production']
    relabel_configs:
      - source_labels: [__meta_kubernetes_pod_label_app]
        action: keep
        regex: api-server

```

## Metrics Types:

Counter (always increasing):  
http\_requests\_total 12345

Gauge (can go up/down):  
memory\_usage\_bytes 8589934592

Histogram (distribution):  
http\_request\_duration\_seconds\_bucket{le="0.1"} 100  
http\_request\_duration\_seconds\_bucket{le="0.5"} 450  
http\_request\_duration\_seconds\_bucket{le="1.0"} 900

Summary (percentiles):  
http\_request\_duration\_seconds{quantile="0.5"} 0.23  
http\_request\_duration\_seconds{quantile="0.9"} 0.87  
http\_request\_duration\_seconds{quantile="0.99"} 1.23

## 6. Time-Series Database (Prometheus TSDB)

**Purpose:** Efficient storage for metrics

Storage Format:  
└─ Chunks: 2-hour blocks of compressed data  
└─ Indexes: Fast lookup by labels  
└─ WAL: Write-ahead log for durability  
└─ Compaction: Merge old chunks

Storage Efficiency:  
└─ Compression: ~1.5 bytes per sample  
└─ Sample: 16 bytes → Compressed: 1.5 bytes  
└─ 1M metrics × 15s interval = 240B samples/day  
└─ Storage: 240B × 1.5 bytes = 360 GB/day  
└─ Actual with overhead: ~400 GB/day

Retention:  
└─ Local: 15 days (fast queries)  
└─ Remote: Thanos/Cortex (long-term)  
└─ Downsampling: 5m → 1h → 6h aggregates

## 7. Query Engines

**Elasticsearch Query DSL**

```
{
  "query": {
    "bool": {
      "must": [
        {"term": {"level": "ERROR"}},
        {"range": {"@timestamp": {
          "gte": "now-1h"
        }}}
      ],
      "filter": [
        {"term": {"app_id": "payment-service"}}
      ]
    }
  },
  "aggs": {
    "errors_over_time": {
      "date_histogram": {
        "field": "@timestamp",
        "interval": "5m"
      }
    }
  },
  "size": 100
}
```

## PromQL (Prometheus Query Language)

```
# Request rate (requests/second)
rate(http_requests_total[5m])

# 99th percentile latency
histogram_quantile(0.99,
  rate(http_request_duration_seconds_bucket[5m])
)

# Error rate percentage
sum(rate(http_requests_total{status="500"}[5m]))
/
sum(rate(http_requests_total[5m])) * 100

# Memory usage by pod
container_memory_usage_bytes{pod=~"api-.*"}

# Alert condition
avg(cpu_usage_percent) > 80
```

## 8. Visualization Layer

## Kibana (Elastic Stack)

### Features:

- Discover: Ad-hoc log exploration
- Visualize: Charts, graphs, maps
- Dashboard: Combine visualizations
- Canvas: Pixel-perfect presentations
- Machine Learning: Anomaly detection
- Alerting: Built-in alert rules

### Common Use Cases:

- Application error tracking
- Security event analysis
- Business analytics
- User behavior analysis
- Compliance reporting

## Grafana (Metrics)

### Features:

- Time-series visualizations
- Multi-datasource support
- Template variables
- Alert rules and notifications
- Dashboard sharing
- Plugin ecosystem

### Common Dashboards:

- System metrics (CPU, memory, disk)
- Application performance (RED metrics)
- Business metrics (orders, revenue)
- SLO/SLI tracking
- Cost monitoring

## 9. Alerting System

**Purpose:** Notify on-call team of issues

### Alert Manager (Prometheus)

- Alert routing
- Grouping & deduplication
- Silencing
- Inhibition rules
- Notification channels

### Alert Definition:

```

groups:
  - name: api_alerts
    rules:
      - alert: HighErrorRate
        expr: |
          sum(rate(http_requests_total{status="500"}[5m])) /
          sum(rate(http_requests_total[5m]))
          > 0.01
        for: 5m
        labels:
          severity: critical
          team: backend
        annotations:
          summary: "High error rate on {{ $labels.instance }}"
          description: "Error rate is {{ $value }}%"

```

#### Notification Routing:

```

route:
  receiver: 'default'
  group_by: ['alertname', 'cluster']
  routes:
    - match:
        severity: critical
        receiver: 'pagerduty'
    - match:
        severity: warning
        receiver: 'slack'

```

#### Receivers:

```

- name: 'pagerduty'
  pagerduty_configs:
    - service_key: '<key>'
- name: 'slack'
  slack_configs:
    - channel: '#alerts'

```

## Storage Architecture

### Log Storage Optimization

#### 1. Index Partitioning

```

Time-Based Sharding:
logs-app-2025.01.08
├── Contains: Jan 8, 2025 data
├── Benefit: Easy to delete old data
└── Benefit: Queries often time-bound
    └── Rollover: Daily or at 50GB

```

Service-Based Sharding:

- logs-payment-2025.01
- logs-user-2025.01
- logs-notification-2025.01
  - Benefit: Isolate noisy services
  - Benefit: Different retention policies
  - Trade-off: More indexes to manage

Hybrid Approach (Recommended):

- logs-{service}-{date}
- Best of both worlds
- Example: logs-payment-2025.01.08
- Scales with both dimensions

## 2. Index Templates & Mappings

```
{  
  "index_patterns": ["logs-*"],  
  "settings": {  
    "number_of_shards": 5,  
    "number_of_replicas": 1,  
    "refresh_interval": "5s",  
    "index": {  
      "codec": "best_compression",  
      "sorting": {  
        "fields": ["@timestamp"]  
      }  
    }  
  },  
  "mappings": {  
    "properties": {  
      "@timestamp": {  
        "type": "date",  
        "format": "strict_date_optional_time|epoch_millis"  
      },  
      "level": {  
        "type": "keyword"  
      },  
      "message": {  
        "type": "text",  
        "fields": {  
          "keyword": {  
            "type": "keyword",  
            "ignore_above": 256  
          }  
        }  
      },  
      "host": {  
        "type": "keyword"  
      }  
    }  
  }  
}
```

```

},
"trace_id": {
    "type": "keyword"
},
"numeric_fields": {
    "type": "integer"
}
}
}
}
}

```

## Key Decisions:

- **text** for full-text search (tokenized)
- **keyword** for exact matching, aggregations
- **date** for time-based queries
- Compression for storage savings

## 3. Shard Allocation

### Shard Sizing:

- └─ Target: 30–50 GB per shard
- └─ Too small: Management overhead
- └─ Too large: Slow recovery, rebalancing
- └─ Daily index: Adjust shards based on volume

### Formula:

`shards = ceiling(daily_volume_GB / 40)`

### Example:

- └─ Daily volume: 200 GB
- └─ Shards:  $200 / 40 = 5$  shards
- └─ Per shard: ~40 GB

### Replicas:

- └─ Production: 1 replica (2 copies total)
- └─ Critical: 2 replicas (3 copies)
- └─ Dev: 0 replicas (cost optimization)

## Metrics Storage Optimization

### 1. Label Strategy

#### Good Labels (low cardinality):

- └─ job: "api-server"
- └─ instance: "10.0.1.5:9090"
- └─ region: "us-west-2"

```
└── env: "production"
    └── Total combinations: <10,000

Bad Labels (high cardinality):
x user_id: "user123456"
x request_id: "uuid-1234-5678"
x email: "user@example.com"
└── Causes: Memory explosion, slow queries
```

Rule of Thumb:

- Labels should have <100 unique values
- Label combinations <10,000 per metric
- Use labels for dimensions you'll query by

## 2. Downsampling

```
Raw Data (15s interval):
└── Resolution: 15 seconds
    └── Retention: 15 days
        └── Storage: 100 GB
            └── Use: Recent detailed analysis
```

```
5-Minute Aggregates:
└── Resolution: 5 minutes (20x reduction)
    └── Retention: 90 days
        └── Storage: 15 GB
            └── Use: Weekly/monthly trends
```

```
1-Hour Aggregates:
└── Resolution: 1 hour (240x reduction)
    └── Retention: 1 year
        └── Storage: 5 GB
            └── Use: Yearly trends, capacity planning
```

Thanos Configuration:

```
downsample:
  - from: 0s
    to: 40h
    resolution: raw
  - from: 40h
    to: 90d
    resolution: 5m
  - from: 90d
    to: 1y
    resolution: 1h
```

## 3. Compression

Prometheus TSDB Compression:

- └─ Chunk format: XOR encoding
- └─ Timestamp delta encoding
- └─ Value compression
- └─ Result: 1.5 bytes/sample average

Example:

Uncompressed: 16 bytes/sample

- └─ Timestamp: 8 bytes
- └─ Value: 8 bytes

Compressed: 1.5 bytes/sample

- └─ Timestamp delta: 0.5 bytes
- └─ XOR value: 1.0 bytes

Savings: 90%

Storage Calculation:

- └─ 1M metrics × 15s scrape = 5.76B samples/day
- └─ Uncompressed: 5.76B × 16 bytes = 92 GB/day
- └─ Compressed: 5.76B × 1.5 bytes = 8.6 GB/day
- └─ Savings: 91% storage reduction

## Data Flow & Ingestion Pipeline

### Log Ingestion Flow

#### STEP 1: Log Generation

Application Server:

```
2025-01-08 10:30:45 ERROR [payment-service] Payment failed
user_id=12345 order_id=67890 error="Card declined"
```

#### STEP 2: Agent Collection (Filebeat)

Filebeat:

- └─ Tail /var/log/app/payment.log
- └─ Add metadata (host, app, env)
- └─ Buffer in memory (10MB)
- └─ Send batch to Logstash (100 events)

#### STEP 3: Processing (Logstash)

Logstash Pipeline:

- └─ Parse timestamp

```
    └── Extract level (ERROR)
    └── Extract service name
    └── Parse key-value pairs
    └── Add GeoIP for user location
└── Structured output:
{
    "@timestamp": "2025-01-08T10:30:45Z",
    "level": "ERROR",
    "service": "payment-service",
    "user_id": "12345",
    "order_id": "67890",
    "error": "Card declined",
    "host": "app-server-23",
    "env": "production"
}
```

#### STEP 4: Buffer (Kafka)

##### Kafka:

```
    └── Topic: logs-payment
    └── Partition: hash(user_id) % 10
    └── Replication: 3 copies
    └── Retention: 7 days
```

#### STEP 5: Storage (Elasticsearch)

##### Elasticsearch:

```
    └── Index: logs-payment-2025.01.08
    └── Shard: Route to shard based on doc ID
    └── Indexing: Create inverted indexes
    └── Refresh: Make searchable (5s delay)
    └── Replica: Copy to replica shard
```

#### STEP 6: Query (Kibana)

User searches: "level:ERROR AND service:payment-service"

```
    └── Query coordinator node
    └── Scatter to all relevant shards
    └── Gather results
    └── Sort and rank
    └── Return to Kibana (< 1 second)
```

#### Latency Breakdown:

End-to-End Latency (log written to searchable):

```
    └── Filebeat buffer: 1–5 seconds
    └── Logstash processing: 0.5–2 seconds
    └── Kafka buffering: 0.1–1 second
```

- Elasticsearch indexing: 1–5 seconds (refresh interval)
- Total: 2.6–13 seconds (acceptable for most use cases)

Real-time scenarios (critical errors):

- Filebeat → Direct to Elasticsearch
- Skip Kafka buffering
- Total: ~2–5 seconds

## Metrics Scraping Flow

### STEP 1: Metrics Exposition

Application exposes /metrics endpoint:

```
# HELP http_requests_total Total HTTP requests
# TYPE http_requests_total counter
http_requests_total{method="GET",status="200"} 12345
http_requests_total{method="POST",status="500"} 42
```

### STEP 2: Service Discovery (Prometheus)

Prometheus discovers targets:

- Kubernetes API: List pods with label app=payment-service
- Result: 5 pods found
- Target list: [10.0.1.1:9090, 10.0.1.2:9090, ...]

### STEP 3: Scrape (Every 15 seconds)

Prometheus:

- HTTP GET 10.0.1.1:9090/metrics
- Parse Prometheus format
- Add labels (instance, job)
- Timestamp samples
- Store in TSDB

### STEP 4: Storage (Prometheus TSDB)

TSDB:

- Write to WAL (durability)
- In-memory chunks (2 hours)
- Compress and write to disk
- Build index for fast queries

### STEP 5: Long-term Storage (Thanos/Cortex)

Thanos Sidecar:

- Upload 2-hour blocks to S3
- Downsample older blocks
- Compact blocks for efficiency
- Retention: Indefinite (with downsampling)

#### STEP 6: Query (Grafana)

- User queries: `rate(http_requests_total[5m])`
- Prometheus evaluates PromQL
  - Queries local TSDB
  - Queries Thanos for historical data
  - Aggregates results
  - Returns to Grafana (< 1 second)

#### Latency Breakdown:

- End-to-End Latency (metric emitted to queryable):
- Scrape interval: 0–15 seconds (average 7.5s)
  - TSDB write: < 1 second
  - In-memory indexing: < 1 second
  - Total: ~8.5 seconds average (acceptable)

Real-time requirements:

- Reduce scrape interval to 5s
- Or use push model for critical metrics
- Total: ~6 seconds

#### Backpressure Handling

**Problem:** What if downstream can't keep up?

Scenario: Elasticsearch cluster overloaded

Without Backpressure:

- Filebeat keeps sending
- Logstash buffers in memory
- Memory fills up
- Logstash OOM crash
- Data loss

With Backpressure (Kafka):

- Kafka buffers logs (disk-backed)
- Logstash processes at sustainable rate
- Filebeat slows down (back-pressure)
- Application logs buffered locally if needed
- No data loss, graceful degradation

## **Implementation:**

```
Filebeat Configuration:  
output.kafka:  
  compression: lz4  
  max_message_bytes: 1000000  
  required_acks: 1  
  bulk_max_size: 2048  
  
Kafka Configuration:  
retention.ms: 604800000 # 7 days  
segment.bytes: 1073741824 # 1GB  
compression.type: lz4  
log.retention.check.interval.ms: 300000  
  
Logstash Configuration:  
pipeline.batch.size: 125  
pipeline.batch.delay: 50  
pipeline.workers: 4
```

---

## Distributed System Design

### Challenge 1: Clock Skew

**Problem:** Servers have different clock times

```
Scenario:  
Server A: 10:00:00 (correct)  
Server B: 10:00:05 (5 seconds fast)  
Server C: 9:59:55 (5 seconds slow)
```

```
Impact:  
- Logs appear out of order  
- Metrics timestamps inconsistent  
- Time-based aggregations wrong
```

### Solutions:

#### **Option 1: NTP Synchronization (Recommended)**

```
All servers sync with NTP  
└─ Accuracy: ±50ms typical  
└─ Good enough for logging/metrics  
└─ Industry standard  
└─ Implementation: ntpd or chrony
```

```
Configuration:  
server 0.pool.ntp.org iburst  
server 1.pool.ntp.org iburst  
driftfile /var/lib/ntp/drift
```

## Option 2: Logical Clocks

- Use Lamport timestamps or vector clocks
  - Total ordering without physical time
  - Complex to implement
  - Use when: Causality is critical

## Option 3: Accept Slight Skew

- Most logging systems handle  $\pm 1$  minute skew
  - Elasticsearch accepts old/future docs
  - Prometheus uses sample time, not ingest time
  - Rarely causes issues in practice

## Challenge 2: Data Consistency

**Problem:** Multiple Elasticsearch nodes, eventual consistency

**Scenario:**  
Write to node A → Not immediately visible on node B

**Timeline:**  
T0: Write log to shard 1 on node A  
T1: Refresh interval (5s) – Now searchable on node A  
T2: Replica copies to node B (async)  
T3: Searchable on node B

**Gap:** User might not see recent logs immediately

## Trade-offs:

**Stronger Consistency:**

- Reduce refresh interval (1s)
- Wait for replicas (wait\_for\_active\_shards)
- Cost: Higher latency, more resources

**Eventual Consistency (Default):**

- Refresh every 5s
- Async replication

- Benefit: Better performance, lower cost
- Trade-off: 5-10s delay for new data

**Decision:** Accept eventual consistency for logs

- Logs are inherently time-delayed anyway
- 5-10 second delay acceptable
- Performance benefits significant

### Challenge 3: Split Brain

**Problem:** Network partition creates two master nodes

Elasticsearch Cluster:

Normal:

[Master] ↔ [Data1] ↔ [Data2]

Network Partition:

[Master] ↔ [Data1] ↔ [Data2]  
↑ Elects new master!

Result: Two masters, diverging data

**Solution:** Quorum-based election

Minimum Master Nodes = (total\_masters / 2) + 1

Example with 3 master-eligible nodes:

- Quorum:  $(3 / 2) + 1 = 2$
- Side with 2+ nodes becomes master
- Side with <2 nodes can't elect master
- Prevents split brain

Configuration:

`discovery.zen.minimum_master_nodes: 2`

### Challenge 4: Hot Shards

**Problem:** One shard gets all the traffic

Scenario: logs-payment-{date}

- Payment service is busiest
- Shard for payment service overloaded
- Other shards idle
- Unbalanced cluster

### Symptoms:

- High CPU on one node
- Slow queries
- Cluster appears "slow" despite spare capacity

### Solutions:

#### Option 1: Better Shard Distribution

Instead of routing by service:

Use routing by hash(doc\_id)

Result:

- Even distribution
- All nodes equally loaded
- Better resource utilization

#### Option 2: More Shards

Increase shards for busy indexes:

logs-payment-\*: 10 shards

logs-notification-\*: 3 shards

Trade-off:

- Better distribution
- More management overhead

#### Option 3: Separate Indexes

Critical services → Dedicated indexes

logs-payment-\* → Dedicated hot tier nodes

Benefit:

- Isolation from noisy neighbors
- Dedicated resources
- SLA protection

## Scalability & Performance

### Horizontal Scaling Strategy

#### Elasticsearch Cluster Scaling

Phase 1: Small (0–100 GB/day)

- └─ 3 master nodes
- └─ 5 data nodes
- └─ Cost: \$2,000/month
- └─ Handles: 100 GB logs/day

Phase 2: Medium (100–500 GB/day)

- └─ 3 master nodes
- └─ 20 data nodes (hot tier)
- └─ 10 data nodes (warm tier)
- └─ Cost: \$8,000/month
- └─ Handles: 500 GB logs/day

Phase 3: Large (500–1TB/day)

- └─ 3 master nodes
- └─ 50 data nodes (hot tier)
- └─ 20 data nodes (warm tier)
- └─ S3 for cold tier
- └─ Cost: \$25,000/month
- └─ Handles: 1 TB logs/day

Phase 4: Very Large (1TB+/day)

- └─ Multiple clusters (regional)
- └─ Federated search across clusters
- └─ Object storage for cold data
- └─ Cost: \$50,000+/month
- └─ Handles: Multiple TB/day

## Prometheus Scaling

Single Prometheus (Up to 1M samples/sec):

- └─ 1 Prometheus server
- └─ 32 GB RAM
- └─ Local TSDB (15 days)
- └─ Cost: \$500/month

Prometheus Federation (1M–10M samples/sec):

- └─ Regional Prometheus servers (3)
- └─ Central federation server
- └─ Thanos for long-term storage
- └─ Cost: \$3,000/month

Prometheus + Cortex (10M+ samples/sec):

- └─ Cortex distributed system
- └─ Horizontal write scaling
- └─ S3/GCS for long-term storage
- └─ Query federation
- └─ Cost: \$10,000+/month

# Performance Optimization

## 1. Query Optimization

### Elasticsearch:

```
Slow Query:  
{  
  "query": {  
    "wildcard": {  
      "message": "*error*"  
    }  
  }  
}  
Problem: Full table scan  
Time: 10+ seconds
```

```
Fast Query:  
{  
  "query": {  
    "match": {  
      "message": "error"  
    }  
  },  
  "post_filter": {  
    "range": {  
      "@timestamp": {  
        "gte": "now-1h"  
      }  
    }  
  }  
}  
Improvement: Uses inverted index + time filter  
Time: <100ms
```

### Prometheus PromQL:

```
Slow Query:  
sum(rate(http_requests_total[5m])) by (instance)  
Problem: High cardinality on instance  
  
Fast Query:  
sum(rate(http_requests_total[5m])) by (job)  
Improvement: Lower cardinality on job  
Time reduction: 5x faster
```

## 2. Caching Strategy

## Multi-Layer Cache:

L1: Browser Cache

- └─ Dashboard JSON
- └─ TTL: 30 seconds
- └─ Reduces API calls

L2: CDN Cache

- └─ Static assets (JS, CSS)
- └─ Dashboard definitions
- └─ Reduces server load

L3: Application Cache (Redis)

- └─ Query results
- └─ TTL: 5 minutes
- └─ Hit rate: 60–80%
- └─ Reduces database load

L4: Database Cache

- └─ Elasticsearch query cache
- └─ Prometheus query cache
- └─ Reduces disk I/O

## 3. Batch Processing

Single inserts (Slow):

For each log:

```
INSERT INTO elasticsearch
```

Time: 100 logs/second

Bulk inserts (Fast):

Batch 1000 logs:

```
BULK INSERT INTO elasticsearch
```

Time: 100,000 logs/second

Improvement: 1000x faster

## 4. Index Optimization

Elasticsearch Index Tuning:

Refresh Interval:

- └─ Default: 1s (near real-time)
- └─ Optimized: 30s (for bulk loading)
- └─ Indexing speed: 3x faster

#### Replica Count:

- └─ Indexing: 0 replicas
- └─ After indexing: 1 replica
- └─ Indexing speed: 2x faster

#### Force Merge:

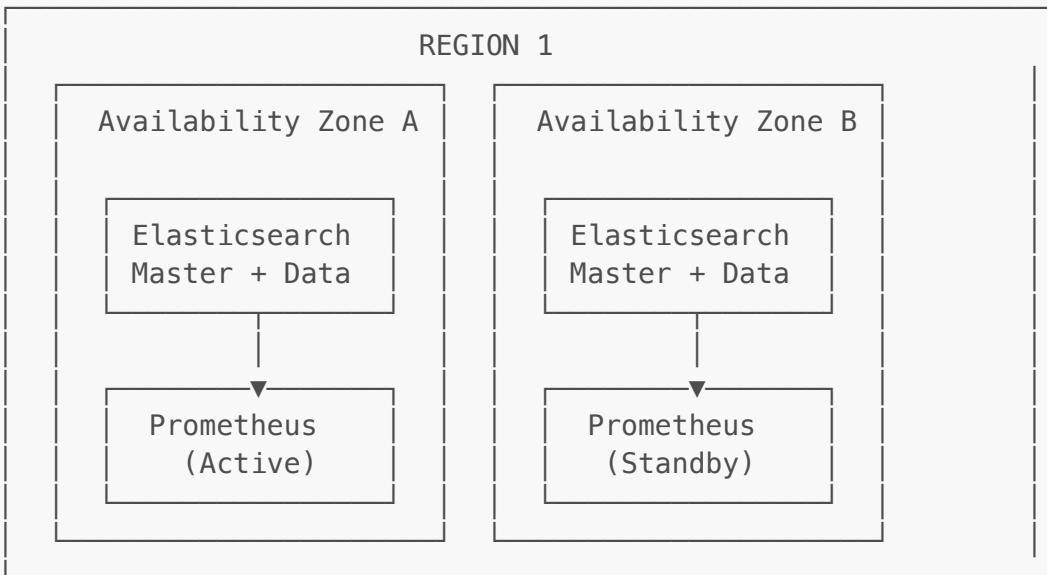
- └─ After rollover to warm tier
- └─ Reduces segment count ( $50 \rightarrow 1$ )
- └─ Query speed: 5x faster

#### Codec:

- └─ Hot tier: default (fast)
- └─ Warm tier: best\_compression
- └─ Storage savings: 30–40%

## Fault Tolerance & Reliability

### High Availability Architecture



#### Elasticsearch:

- └─ 3 master nodes (split across AZs)
- └─ N data nodes (balanced across AZs)
- └─ 1 replica per shard (different AZ)
- └─ Can lose 1 AZ and stay operational

#### Prometheus:

- └─ Active-Standby setup
- └─ Thanos for failover
- └─ Can lose 1 AZ, switch to standby
- └─ Historical data in S3 (survives total loss)

## Failure Scenarios

### Scenario 1: Elasticsearch Node Failure

Normal Operation:

```
[Master] → [Data1] [Data2] [Data3]
      Shard A   Shard B   Shard C
      Replica    Replica    Replica
```

Node2 Fails:

```
[Master] → [Data1] [xData2] [Data3]
      Shard A           Shard C
      Replica B        Replica
```

Recovery:

1. Master detects failure (30s)
2. Promote Replica B to primary (10s)
3. Reallocate Replica A and C (5 min)
4. Cluster back to green (5–10 min)

Impact:

- └ No data loss (replicas available)
- └ Queries continue (other nodes serve)
- └ Slight performance degradation during rebalance
- └ RT0: 30–40 seconds

### Scenario 2: Kafka Broker Failure

Normal Operation:

```
[Producer] → [Broker1] [Broker2] [Broker3]
      Leader     Replica   Replica
```

Broker1 Fails:

```
[Producer] → [xBroker1] [Broker2] [Broker3]
                           ↑ New Leader
```

Recovery:

1. ZooKeeper detects failure (6s)
2. Elect new leader from replicas (2s)
3. Update metadata (2s)
4. Producers/consumers reconnect (automatic)

Impact:

- └ No data loss (replication factor 3)
- └ Brief pause in writes (8–10s)
- └ Automatic recovery
- └ RT0: ~10 seconds

### Scenario 3: Prometheus Server Failure

Active-Standby Setup:

[Active Prometheus] → [Thanos Sidecar] → [S3]  
[Standby Prometheus] → [Thanos Sidecar] → [S3]

Active Fails:

1. Standby starts scraping (manual or auto)
2. Queries redirect to standby
3. Historical data served from S3 (Thanos)

Impact:

- └ Gap in metrics (duration of failover)
- └ Historical data intact
- └ RT0: 1–5 minutes (manual) or 30s (auto)
- └ RPO: Scrape interval (15s)

Mitigation:

Use Cortex for automatic failover:

- └ Multiple Prometheus instances scrape
- └ All write to central Cortex
- └ No single point of failure
- └ RT0: 0 (no downtime)

## Data Durability

### Logs (Elasticsearch):

Durability Mechanisms:

- └ Translog (write-ahead log)
- └ Replication (1+ replicas)
- └ Snapshots to S3 (daily)
- └ Cross-region replication (optional)

Failure Tolerance:

- └ Can lose N-1 replicas
- └ Translog protects unflushed writes
- └ Snapshots for disaster recovery
- └ RPO: < 5 seconds (translog)

### Metrics (Prometheus):

Durability Mechanisms:

- └ WAL (write-ahead log)
- └ 2-hour blocks persisted to disk
- └ Thanos uploads to S3
- └ Downsampled for long-term retention

**Failure Tolerance:**

- └─ WAL protects recent data
- └─ S3 for historical data
- └─ Downsampling preserves trends
- └─ RP0: Scrape interval (15s)

## Disaster Recovery

### **Backup Strategy:**

**Elasticsearch:**

- └─ Daily snapshots to S3
- └─ Retention: 30 days
- └─ Cross-region replication
- └─ Restore time: Hours to TB+
- └─ Test restores monthly

**Prometheus:**

- └─ Continuous upload to S3 (Thanos)
- └─ No backups needed (immutable blocks)
- └─ Downsampling preserves data
- └─ Restore: Point queries at S3

**Configuration:**

- └─ Git repository (versioned)
- └─ Automated deployment
- └─ Restore: Minutes

### **Recovery Procedures:**

#### **Complete Cluster Loss:**

**Elasticsearch:**

1. Provision new cluster
2. Restore latest snapshot from S3
3. Reconfigure applications
4. Time: 2–4 hours

**Prometheus:**

1. Provision new Prometheus
2. Configure Thanos to query S3
3. Historical data immediately available
4. Time: 15–30 minutes

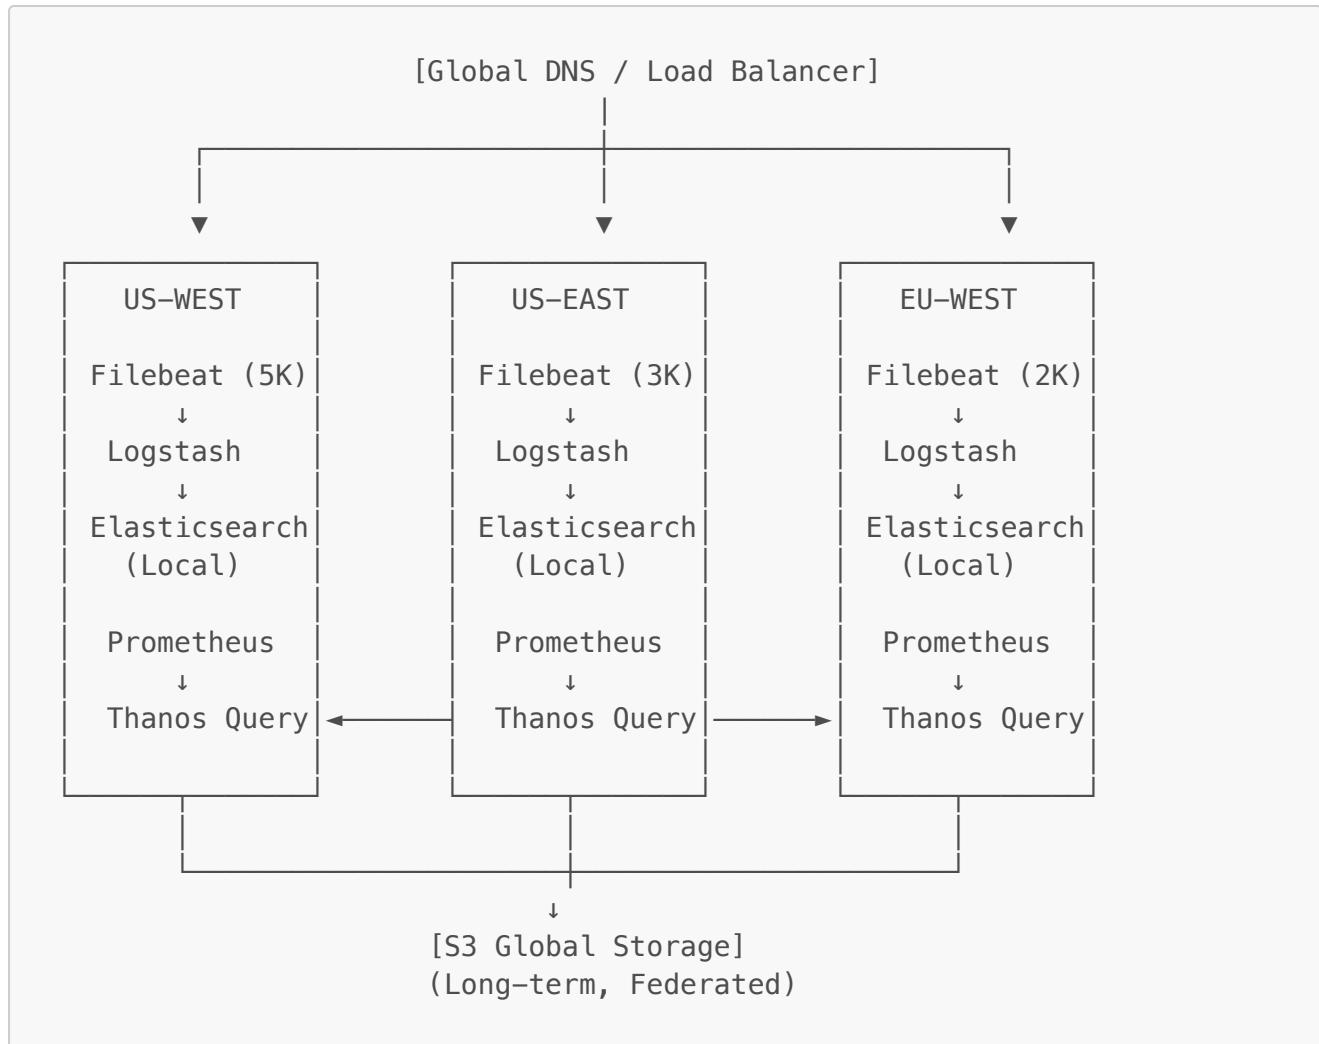
**Prevention:**

- └─ Multi-region deployment

- └ Regular DR drills
- └ Automated runbooks

## Multi-Region Architecture

### Global Deployment



### Regional Independence

**Design:** Each region operates independently

#### Logs:

- └ Stored locally in each region
- └ Cross-region search via federation (optional)
- └ Compliance: Data stays in region
- └ No cross-region latency for queries

#### Metrics:

- └ Scrapped locally by Prometheus
- └ Uploaded to global S3 (Thanos)
- └ Federated queries across regions

## └ Global dashboards via Thanos Query

### Benefits:

- ✓ Data locality (faster, compliance)
- ✓ Regional failures isolated
- ✓ No cross-region bandwidth costs
- ✓ Simple to operate

### Challenges:

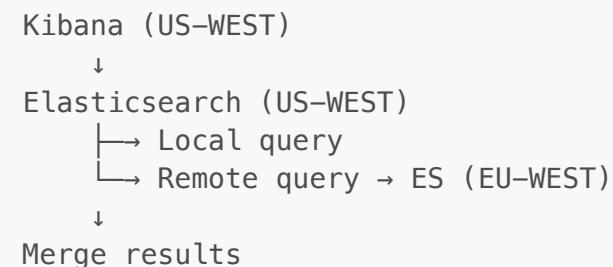
- ✗ Global queries more complex
- ✗ Data scattered across regions

## Federated Queries

### Cross-Region Log Search:

Kibana (US-WEST) needs to search EU logs:

Option 1: Cross-Cluster Search (Elasticsearch)



Latency: +100–200ms for cross-region

Option 2: Separate Queries (Recommended)

- Query each region separately
- User selects region in UI
- Faster, simpler, region-aware

### Global Metrics Dashboard:

#### Thanos Architecture:

Regional Prometheus → Thanos Sidecar → S3  
(US-WEST)   ↓  
   [Upload blocks]

Regional Prometheus → Thanos Sidecar → S3  
(US-EAST)   ↓  
   [Upload blocks]

Regional Prometheus → Thanos Sidecar → S3  
(EU-WEST)

↓  
[Upload blocks]

↓ Query from any region  
[Thanos Query Frontend]  
↓  
Aggregate global metrics

#### Benefits:

- ✓ Single pane of glass
- ✓ Global queries (sum across regions)
- ✓ Historical data centralized
- ✓ No Prometheus federation complexity

## Design Trade-offs

### 1. Push vs Pull

Aspect	Push (Logs)	Pull (Metrics)
Complexity	Higher (need agents)	Lower (server-side)
Flexibility	Can send anytime	Fixed intervals
Firewalls	Works behind firewall	Needs open port
Load	Can overwhelm receiver	Controlled by scraper
Short-lived	Perfect for Lambda/cron	Needs push gateway
Service Discovery	Not needed	Built-in

#### Decision:

- Logs: Push (applications generate continuously)
- Metrics: Pull (Prometheus model, with push gateway for edge cases)

### 2. Structured vs Unstructured Logs

#### Unstructured:

"2025-01-08 10:30:45 ERROR Payment failed for user 12345"

#### Pros:

- ✓ Human readable
- ✓ No schema needed
- ✓ Flexible format

#### Cons:

- ✗ Hard to parse

- ✗ Slow to query
- ✗ Requires grok patterns

Structured (JSON):

```
{
  "timestamp": "2025-01-08T10:30:45Z",
  "level": "ERROR",
  "message": "Payment failed",
  "user_id": 12345
}
```

Pros:

- ✓ Fast to query
- ✓ Easy to aggregate
- ✓ Type-safe

Cons:

- ✗ More verbose
- ✗ Requires schema
- ✗ Less human-readable

**Decision:** Structured (JSON) for new applications

- Better query performance
- Easier aggregation
- Worth the verbosity
- Can always add unstructured "message" field

### 3. Centralized vs Distributed Storage

Centralized (Single Cluster):

All regions → Central Elasticsearch

Pros:

- ✓ Simple operations
- ✓ Easy to query
- ✓ Single source of truth

Cons:

- ✗ Cross-region latency
- ✗ Single point of failure
- ✗ Compliance issues
- ✗ Expensive bandwidth

Distributed (Regional Clusters):

Each region → Local Elasticsearch

Pros:

- ✓ Data locality
- ✓ Failure isolation
- ✓ Compliance friendly

- ✓ Lower latency

Cons:

- ✗ Complex operations
- ✗ Harder to query globally
- ✗ More clusters to manage

**\*\*Decision\*\*: Distributed for large scale**

- Centralized for small deployments (<100 GB/day)
- Distributed for large deployments
- Use federation for cross-region queries

## 4. Retention Policy

Option 1: Single Tier (Simple)

- └─ All data in Elasticsearch
- └─ Retention: 30 days
- └─ Cost: High (all on SSD)
- └─ Query: Fast

Option 2: Multi-Tier (Recommended)

- └─ Hot: 7 days (SSD)
- └─ Warm: 30 days (slower SSD)
- └─ Cold: 90 days (S3)
- └─ Archive: 1 year (Glacier)
- └─ Cost: Medium (tiered)
- └─ Query: Fast for recent, slower for old

Option 3: Immediate Archive

- └─ All data to S3 immediately
- └─ Elasticsearch indexes only recent
- └─ Query: Slow (S3 every time)
- └─ Cost: Low

**\*\*Decision\*\*: Multi-tier**

- Balance of cost and performance
- 90% queries hit hot/warm tier
- 10% queries slower (acceptable)
- 80% cost savings vs single tier

---

## Real-World Case Studies

### Case Study 1: Netflix (1000+ Microservices)

**Scale:**

- 10,000+ EC2 instances
- 100+ TB logs/day

- 2.5 billion metrics per day
- 500+ petabytes historical data

## Architecture:

### Logging Stack:

```
└─ Agent: Vector, Fluentd
└─ Processing: Apache Flink
└─ Storage: Elasticsearch (primary), S3 (archive)
└─ Query: Kibana, custom tools
└─ Retention: 30 days hot, years in S3
```

### Metrics Stack:

```
└─ Collection: Atlas (custom), Prometheus
└─ Storage: Atlas TSDB
└─ Processing: Spark for analytics
└─ Visualization: Grafana, custom dashboards
└─ Alerting: Custom alerting infrastructure
```

## Key Decisions:

### 1. **Atlas (Custom TSDB)**: Built for their specific needs

- Optimized for their workload (dimensional time-series)
- Integration with internal systems
- Better performance than off-the-shelf

### 2. **Elasticsearch for logs**: Industry standard, proven at scale

- Good search performance
- Rich ecosystem
- Multi-tenancy support

### 3. **S3 for archival**: Cost-effective long-term storage

- Compliance requirements
- Rarely queried historical data
- 95% cost savings vs hot storage

## Lessons Learned:

- Build custom tools only when necessary (Atlas for specific needs)
- Use proven solutions where possible (Elasticsearch)
- Invest in automation (alert management, capacity planning)
- Sampling is key at extreme scale (sample 1% of logs for analytics)

## Case Study 2: Uber (Global Ride-Sharing)

### Scale:

- 50,000+ microservices
- Millions of events per second
- Petabyte-scale storage
- Real-time analytics

## **Architecture:**

Logging Stack (M3):

- Collection: Kafka
- Processing: Apache Flink
- Storage: M3DB (custom TSDB)
- Query: Custom query engine
- Visualization: Grafana

Key Innovation: M3DB

- Distributed time-series database
- Built for extreme scale
- Prometheus-compatible
- Global aggregations
- Open-sourced

## **Design Highlights:**

### **1. M3DB Architecture:**

- Namespaces for different retention periods
- Aggressive compression
- Distributed aggregation
- Multi-datacenter replication

### **2. Real-time Processing:**

- Flink for stream processing
- Sub-second latency
- Exactly-once semantics

### **3. Cost Optimization:**

- Downsampling (1s → 1m → 5m)
- Compression (6:1 ratio)
- Tiered storage

## **Lessons Learned:**

- Prometheus doesn't scale to millions of metrics (needs custom solution)
- Real-time anomaly detection is critical for ride-sharing
- Global aggregation necessary for business metrics
- Investment in custom tooling pays off at scale

## Case Study 3: Datadog (SaaS Observability)

### Scale:

- Hundreds of billions of metrics per day
- Petabytes of logs per day
- Thousands of customers
- 99.99% SLA

### Architecture:

#### Multi-Tenant Design:

- Data isolation per customer
- Shared infrastructure
- Usage-based pricing
- Rate limiting per tenant
- SLA guarantees per tier

#### Storage:

- Hot tier: Custom TSDB
- Cold tier: S3
- Real-time: Apache Kafka
- Analytics: Apache Spark
- Retention: Configurable (1 day – 15 months)

#### Querying:

- Custom query language
- Distributed query execution
- Query optimization
- Result caching

### Key Innovations:

1. **Multi-tenancy at scale:** Thousands of isolated customers on shared infrastructure
2. **Intelligent sampling:** Sample to reduce costs while maintaining accuracy
3. **Unified platform:** Metrics, logs, traces in single platform
4. **ML-powered alerts:** Reduce false positives

### Lessons Learned:

- Multi-tenancy requires careful resource isolation
- Query optimization critical for SaaS performance
- Sampling strategies must be transparent to customers
- Rich integrations drive adoption (500+ integrations)

## Case Study 4: LinkedIn (Professional Network)

### Scale:

- 20,000+ hosts
- 1.5+ trillion metrics per day
- 300+ TB logs per day
- Real-time analytics

### Architecture:

Logging Stack (Kafka + Samza):

```
└── Collection: Kafka
└── Processing: Apache Samza
└── Storage: Kafka compacted topics
└── Search: Elasticsearch
└── Analytics: Pinot
```

Metrics Stack (InGraphs):

```
└── Collection: Custom agents
└── Storage: Custom TSDB (Venice)
└── Processing: Spark
└── Query: Custom query engine
└── Visualization: Custom dashboards
```

### Key Innovations:

1. **Kafka as storage:** Use Kafka for both transport and storage
2. **Venice TSDB:** Distributed key-value store optimized for time-series
3. **Pinot:** Real-time OLAP for analytics
4. **Samza:** Stream processing for real-time aggregations

### Lessons Learned:

- Kafka's durability makes it good for log storage
- Custom solutions needed at extreme scale
- Real-time aggregation reduces storage costs
- Query federation across systems for flexibility

## Technology Stack Comparison

### Logging Stacks

Component	ELK Stack	EFK Stack	Splunk	Loki
<b>Agent</b>	Filebeat	Fluentd	Universal Forwarder	Promtail
<b>Processing</b>	Logstash	Fluentd	Splunk Indexers	None
<b>Storage</b>	Elasticsearch	Elasticsearch	Splunk	Object Storage
<b>Query</b>	Kibana	Kibana	Splunk Web	Grafana

Component	ELK Stack	EFK Stack	Splunk	Loki
<b>Cost</b>	Medium	Medium	High	Low
<b>Scale</b>	High	High	Very High	Medium
<b>Ease of Use</b>	Medium	Medium	High	High

#### Recommendations:

- **ELK/EFK**: General purpose, proven at scale
- **Splunk**: Enterprise features, high cost
- **Loki**: Cost-conscious, simple log aggregation

#### Metrics Stacks

Component	Prometheus	Graphite	InfluxDB	Datadog
<b>Collection</b>	Pull	Push	Push	Agent
<b>Storage</b>	Local TSDB	Whisper	TSM Engine	Cloud
<b>Query Language</b>	PromQL	Functions	InfluxQL/Flux	Custom
<b>Scalability</b>	Medium	Low	Medium	Very High
<b>Cost</b>	Free (OSS)	Free (OSS)	Paid	\$\$\$
<b>Long-term</b>	Thanos/Cortex	Carbon Relay	InfluxDB Cloud	Built-in

#### Recommendations:

- **Prometheus**: Cloud-native, Kubernetes, OSS
- **InfluxDB**: IoT, high cardinality
- **Datadog**: SaaS, full observability platform

## Interview Talking Points

### Key Design Decisions

#### 1. Why Elasticsearch for logs?

- Full-text search capability
- Flexible schema (JSON documents)
- Horizontal scalability
- Rich aggregation engine
- Industry standard with large ecosystem

#### 2. Why Prometheus for metrics?

- Pull-based model (better for service discovery)
- Efficient time-series storage (1.5 bytes/sample)

- Powerful query language (PromQL)
- Active community, CNCF graduated
- Kubernetes-native

### 3. Why Kafka as message queue?

- High throughput (millions events/sec)
- Durability (replicated, disk-backed)
- Replay capability for debugging
- Multiple consumers (fanout pattern)
- Decouples producers from consumers

### 4. Why separate hot/warm/cold tiers?

- 80% cost reduction
- 90% queries hit recent data (hot tier)
- Compliance needs long-term retention
- Query performance optimized for common case

### 5. Why not store everything in real-time databases?

- Cost prohibitive (SSD expensive)
- Most queries are recent data
- Historical data queried rarely
- Acceptable to have slower queries for old data

### 6. Why multi-region deployment?

- Data locality (lower latency)
- Compliance (GDPR, data residency)
- Fault isolation (regional failures)
- Reduced cross-region bandwidth costs

## Bottlenecks & Solutions

Bottleneck	Impact	Solution
<b>Elasticsearch indexing</b>	Slow ingestion	Bulk API, tune refresh interval
<b>High cardinality metrics</b>	Memory explosion	Limit label combinations, sampling
<b>Cross-region queries</b>	High latency	Regional clusters, federated search
<b>Hot shards</b>	Uneven load	Better routing, more shards
<b>Query performance</b>	Slow dashboards	Caching, time-range filters, indexing
<b>Storage costs</b>	Budget overrun	Tiered storage, downsampling, compression
<b>Kafka lag</b>	Processing delay	Add consumers, increase partitions
<b>Clock skew</b>	Out-of-order data	NTP sync, logical ordering

## Scale Calculations

### Logs:

- $10,000 \text{ servers} \times 1,000 \text{ lines/min} = 10\text{M lines/min}$
- $10\text{M} \times 500 \text{ bytes} = 5 \text{ GB/min} = 7.2 \text{ TB/day}$
- 90 days retention = 648 TB

### Metrics:

- $10,000 \text{ servers} \times 100 \text{ metrics} = 1\text{M metrics}$
- $1\text{M} \times (86400/15) \text{ samples} = 5.76\text{B samples/day}$
- $5.76\text{B} \times 1.5 \text{ bytes} = 8.6 \text{ GB/day}$
- 90 days retention = 774 GB

### Infrastructure:

- Elasticsearch: 50–100 nodes (10–20 TB each)
- Prometheus: 10–20 servers (500 GB each)
- Kafka: 10 brokers (5 TB each)
- Logstash: 20–30 workers

## Common Interview Questions

### Q: How do you handle 1M log events per second?

A:

- Kafka buffers incoming logs (handles spikes)
- Multiple Logstash workers process in parallel
- Bulk insert to Elasticsearch (1000 docs at once)
- Tune Elasticsearch refresh interval (5s → 30s)
- Use ingest nodes for heavy processing
- Scale horizontally by adding nodes

### Q: How do you prevent high-cardinality metrics from causing issues?

A:

- Design guidelines (labels <100 unique values)
- Automated validation (reject metrics with too many labels)
- Sampling for high-cardinality (1 in 100)
- Use aggregations in recording rules
- Alert on cardinality explosion

### Q: What happens when Elasticsearch cluster is full?

A:

- Index Lifecycle Management (ILM) automatically archives old data
- Disk watermarks prevent complete fill (85%, 90%, 95%)
- At 95%: Block new writes to protect cluster
- Solution: Add nodes or delete old data
- Prevention: Capacity planning, alerts at 70%

**Q: How do you ensure no log data is lost?**

A:

- Filebeat persists position (knows what was sent)
- Kafka replication (3 copies)
- Elasticsearch translog (write-ahead log)
- Elasticsearch replicas (1-2 copies)
- S3 snapshots (disaster recovery)
- Result: Multiple layers of durability

**Q: How do you handle time-series data with high cardinality?**

A:

- Don't use labels for high-cardinality dimensions
- Use relabeling to drop or aggregate labels
- Implement recording rules (pre-aggregation)
- Use Thanos/Cortex for horizontal scaling
- Consider alternative databases (ClickHouse, TimescaleDB)

**Q: How would you reduce costs by 50%?**

A:

- Aggressive tiering (hot → warm → cold)
- Sampling (logs: 10%, metrics: downsampling)
- Compression (best\_compression codec)
- Right-size clusters (don't over-provision)
- Delete unnecessary data
- Use object storage for archives

**Q: How do you ensure 99.99% uptime?**

A:

- Multi-AZ deployment
- Replication (all data has 2+ copies)
- Automatic failover (Elasticsearch, Kafka)
- Health checks and circuit breakers
- Graceful degradation (serve cached data)
- Regular disaster recovery drills

---

## System Design Patterns Used

### 1. Producer-Consumer Pattern

- Applications produce logs/metrics
- Kafka queues decouple producers from consumers
- Multiple consumers process at their own pace
- Enables backpressure handling

## 2. Fan-Out Pattern

- Single log stream → Multiple consumers
- Elasticsearch for search
- S3 for archival
- Spark for analytics
- Independent processing

## 3. CQRS (Command Query Responsibility Segregation)

- Separate write path (ingestion) from read path (query)
- Optimize each independently
- Write: Bulk inserts, batching
- Read: Indexes, caching, replicas

## 4. Circuit Breaker

- Protect downstream services from overload
- Fail fast when service is down
- Automatic recovery when service returns
- Implemented in Filebeat, Logstash

## 5. Bulkhead Pattern

- Isolate resources for different services
- Separate indexes for critical services
- Prevent noisy neighbors
- Dedicated thread pools

## 6. Time-Series Compaction

- Prometheus 2-hour blocks
- Compact older blocks
- Reduce storage and improve query speed
- Immutable blocks after compaction

## 7. Tiered Storage

- Hot tier (SSD) for recent data
- Warm tier (HDD) for older data
- Cold tier (S3) for archives
- Automatic lifecycle management

## 8. Sampling

- Sample high-volume logs (1 in 100)
- Downsample metrics for long-term storage
- Maintain accuracy for rare events
- Reduce costs significantly

# Best Practices & Recommendations

## Log Best Practices

### 1. Structured Logging

```
// Good: Structured JSON
{
  "timestamp": "2025-01-08T10:30:45Z",
  "level": "ERROR",
  "service": "payment-service",
  "trace_id": "abc123",
  "user_id": 12345,
  "error": "Card declined",
  "error_code": "INSUFFICIENT_FUNDS"
}

// Bad: Unstructured string
"2025-01-08 10:30:45 ERROR Payment failed for user 12345 card declined"
```

### 2. Log Levels

```
DEBUG: Detailed diagnostic information
INFO: General informational messages
WARNING: Warning messages (not errors)
ERROR: Error messages (caught exceptions)
FATAL: Critical errors (system shutdown)
```

Best practice:

- Production: INFO and above
- Staging: DEBUG
- Never log sensitive data (passwords, tokens, PII)

### 3. Correlation IDs

```
Add trace_id to all logs in request path:
└─ API Gateway: Generate trace_id
└─ All downstream services: Propagate trace_id
└─ All logs: Include trace_id
└─ Benefit: Track request across services
```

Example:

```
trace_id: "req-abc123"
└─ API Gateway: "Received request"
└─ Auth Service: "User authenticated"
```

- └─ Payment Service: "Payment processed"
- └─ Can correlate entire request flow

## Metrics Best Practices

### 1. Naming Convention

Format: {namespace}\_{subsystem}\_{name}\_{unit}

Good examples:

http\_requests\_total  
http\_request\_duration\_seconds  
node\_cpu\_seconds\_total  
api\_errors\_total

Bad examples:

requests (too generic)  
my\_metric (no context)  
time\_ms (non-standard unit)

### 2. Label Best Practices

Good labels:

http\_requests\_total{method="GET", status="200", endpoint="/api/users"}

Bad labels:

http\_requests\_total{user\_id="12345", request\_id="abc"}  
└─ High cardinality causes problems

Rule:

- Use labels for dimensions you'll query by
- Keep cardinality low (<1000 combinations)
- Document label values

### 3. The Four Golden Signals (Google SRE)

1. Latency: How long requests take  
http\_request\_duration\_seconds

2. Traffic: How much demand  
http\_requests\_total

3. Errors: Rate of failed requests  
http\_requests\_total{status="5xx"}

4. Saturation: How "full" the service is

- CPU usage
- Memory usage
- Disk I/O
- Queue depth

#### 4. RED Method (Requests, Errors, Duration)

```

Requests: Rate of requests
rate(http_requests_total[5m])

Errors: Rate of errors
rate(http_requests_total{status="5xx"} [5m])

Duration: Latency distribution
histogram_quantile(0.99,
    rate(http_request_duration_seconds_bucket[5m])
)

```

### Alert Best Practices

#### 1. Alert on Symptoms, Not Causes

```

Good:
alert: HighErrorRate
expr: error_rate > 0.01
description: "Users experiencing errors"

Bad:
alert: HighCPU
expr: cpu > 80%
description: "CPU is high"
└ CPU high might not affect users

```

#### 2. Alert Fatigue Prevention

```

Strategies:
└ Set appropriate thresholds (not too sensitive)
└ Require duration (for: 5m)
└ Group related alerts
└ Use inhibition rules
└ Meaningful alert names and descriptions
└ Runbooks for each alert

```

```

Example:
alert: DatabaseDown
expr: up{job="postgres"} == 0

```

```
for: 1m # Not just a blip
annotations:
  runbook: "https://wiki.company.com/db-down"
  action: "Check database logs, attempt restart"
```

### 3. Alert Severity Levels

Critical (Page immediately):

- └─ Service is down
- └─ Data loss occurring
- └─ Security breach
- └─ Customer impact

Warning (Notify, don't page):

- └─ Degraded performance
- └─ Approaching capacity
- └─ Non-critical error rate increase
- └─ Can wait for business hours

Info (Log only):

- └─ Successful deployments
- └─ Scaling events
- └─ Routine maintenance

## Cost Optimization Strategies

### Storage Cost Reduction

#### 1. Tiered Storage

Cost comparison (per TB/month):

- └─ SSD (hot): \$100
- └─ HDD (warm): \$40
- └─ S3 Standard: \$23
- └─ S3 IA: \$12.50
- └─ S3 Glacier: \$4

Strategy:

- └─ 7 days on SSD:  $50 \text{ TB} \times \$100 = \$5,000$
- └─ 30 days on HDD:  $200 \text{ TB} \times \$40 = \$8,000$
- └─ 90 days on S3:  $600 \text{ TB} \times \$23 = \$13,800$
- └─ 1 year on Glacier:  $2.5 \text{ PB} \times \$4 = \$10,000$
- └─ Total: \$36,800/month vs \$100,000 (all SSD)

Savings: 63%

## 2. Compression

Elasticsearch:

- └─ Default codec: ~30% compression
- └─ Best\_compression: ~50% compression
- └─ Trade-off: Slightly slower indexing
- └─ Use for warm tier

Prometheus:

- └─ XOR encoding: ~90% compression
- └─ Automatic, no configuration needed
- └─ 92 GB → 8.6 GB per day

## 3. Sampling

Use cases for sampling:

- └─ Debug logs: Sample 10% (still useful)
- └─ Metrics: Downsample after retention period
- └─ Traces: Sample 1–5% of requests
- └─ Analytics: Statistical sampling sufficient

Example:

100 GB/day → Sample 10% → 10 GB/day

Savings: 90% (\$2,300 → \$230/month)

## 4. Retention Optimization

Analyze query patterns:

- └─ 90% queries: Last 7 days
- └─ 9% queries: Last 30 days
- └─ 1% queries: Older than 30 days

Optimize retention:

- └─ Hot (7 days): Fast, expensive
- └─ Warm (30 days): Medium speed/cost
- └─ Cold (90+ days): Slow, cheap
- └─ Delete or archive after 1 year

Result: 70–80% cost reduction

## Compute Cost Reduction

### 1. Right-Sizing

Monitor actual usage:

- └─ CPU utilization: Target 60–70%
- └─ Memory utilization: Target 70–80%
- └─ Disk I/O: Target 60–70%
- └─ Right-size instances accordingly

Example:

- └─ Over-provisioned: c5.4xlarge (16 vCPU, \$0.68/hr)
- └─ Right-sized: c5.2xlarge (8 vCPU, \$0.34/hr)
- └─ Savings: 50%

## 2. Auto-Scaling

Scale based on metrics:

- └─ Scale up: CPU > 70% for 5 minutes
- └─ Scale down: CPU < 30% for 15 minutes
- └─ Min instances: Baseline load
- └─ Max instances: Peak load + buffer
- └─ Savings: 30–40% vs static provisioning

## 3. Spot Instances

Use cases:

- └─ Logstash workers (fault-tolerant)
- └─ Batch processing jobs
- └─ Non-critical workloads
- └─ Savings: 60–80%

Not recommended for:

- ✗ Elasticsearch data nodes (state)
- ✗ Prometheus (metrics loss)
- ✗ Kafka brokers (data loss risk)

# Security & Compliance

## Access Control

### Authentication:

Options:

- └─ LDAP/Active Directory integration
- └─ SAML/OAuth for SSO
- └─ API keys for programmatic access
- └─ JWT tokens for service-to-service

└ MFA for admin access

Implementation:

Kibana → Elasticsearch:

```
xpack.security.enabled: true  
xpack.security.authc.providers:  
  saml.saml1:  
    order: 0  
    realm: saml-realm
```

## Authorization:

Role-Based Access Control (RBAC):

Roles:

- └ Admin: Full access (create, delete, configure)
- └ Developer: Read/write logs and metrics
- └ Viewer: Read-only access
- └ Service: API access for applications
- └ Auditor: Read-only, compliance reports

Implementation:

- Elasticsearch: Index-level permissions
- Prometheus: External authentication proxy
- Kibana: Space-based isolation (multi-tenancy)
- Grafana: Folder and dashboard permissions

## Audit Logging

Track all actions:

- └ Who: User or service account
- └ What: Action performed (read, write, delete)
- └ When: Timestamp
- └ Where: Resource accessed
- └ Result: Success or failure

Implementation:

- Elasticsearch audit logs
- Store in separate index
- Retention: 1+ years (compliance)
- Immutable (append-only)
- Monitored for anomalies

Example audit log:

```
{  
  "user": "john.doe@company.com",  
  "action": "DELETE_INDEX",  
  "resource": "logs-payment-2024.12.01",
```

```
"timestamp": "2025-01-08T10:30:45Z",
"result": "SUCCESS",
"ip": "192.168.1.1"
}
```

## Data Privacy

### PII Handling:

#### Best practices:

- └── Don't log PII (passwords, SSN, credit cards)
- └── Hash sensitive fields (email, user\_id)
- └── Mask sensitive data (show last 4 digits)
- └── Separate indexes for sensitive data
- └── Encryption at rest

#### Implementation:

##### Logstash filter:

```
filter {
    # Mask credit card numbers
    mutate {
        gsub => [
            "message", "\d{4}-\d{4}-\d{4}-(\d{4})",
            "****-****-****-\1"
        ]
    }

    # Hash email addresses
    fingerprint {
        source => "email"
        target => "email_hash"
        method => "SHA256"
    }

    # Remove original email
    mutate {
        remove_field => ["email"]
    }
}
```

### Encryption:

#### In Transit:

- └── TLS 1.3 for all communications
- └── Certificate validation
- └── mTLS for service-to-service
- └── VPN for cross-region

#### At Rest:

- Elasticsearch: Encryption at rest (EBS encryption)
- S3: Server-side encryption (SSE-S3, SSE-KMS)
- Snapshots: Encrypted
- Kafka: Encryption at rest (optional)

## Compliance

### GDPR/Data Residency:

#### Requirements:

- Data must stay in EU for EU users
- Right to access (export user's logs)
- Right to be forgotten (delete user's logs)
- Data processing agreements

#### Implementation:

- Regional clusters (no cross-border transfers)
- User data export API
- Automated deletion workflows
- Index lifecycle management
- Audit logging of all access

### Retention Policies:

#### By data type:

- Application logs: 90 days
- Security logs: 1 year (compliance)
- Audit logs: 7 years (legal requirement)
- Debug logs: 7 days (space-saving)
- Metrics: 15 days raw, 1 year downsampled

#### Automated enforcement:

- ILM policies in Elasticsearch
- Thanos compaction and downsampling
- S3 lifecycle policies
- Regular compliance audits

## Technology Stack Summary

Layer	Logging	Metrics	Purpose
Collection	Filebeat, Fluentd	Prometheus, Telegraf	Gather data
Transport	Kafka	Kafka (optional)	Buffer, decouple

Layer	Logging	Metrics	Purpose
<b>Processing</b>	Logstash, Flink	Recording rules	Transform, enrich
<b>Storage (Hot)</b>	Elasticsearch	Prometheus TSDB	Fast queries
<b>Storage (Cold)</b>	S3, Glacier	Thanos, Cortex	Long-term, cheap
<b>Query</b>	Elasticsearch DSL	PromQL	Retrieve data
<b>Visualization</b>	Kibana	Grafana	Dashboards
<b>Alerting</b>	ElastAlert, Watcher	Alertmanager	Notifications
<b>Analysis</b>	Spark, Flink	Spark	Batch analytics

## Summary & Key Takeaways

### Critical Design Principles

#### 1. Separation of Concerns

- Separate ingestion, storage, and query layers
- Each can scale independently
- Failures isolated

#### 2. Eventual Consistency is Acceptable

- Logs don't need strong consistency
- 5-10 second delay for searchability is fine
- Allows for better performance and cost

#### 3. Push for Logs, Pull for Metrics

- Logs: Continuous generation, need push
- Metrics: Periodic collection, pull works better
- Hybrid approach for edge cases

#### 4. Tiered Storage is Essential

- 80% cost savings
- Query patterns favor recent data
- Compliance needs long-term retention

#### 5. Kafka is the Universal Buffer

- Decouples producers from consumers
- Handles backpressure
- Enables replay for debugging
- Multiple consumers without overhead

#### 6. Sampling Maintains Accuracy

- Statistical sampling (1-10%)
- Maintains trends and patterns
- Massive cost savings
- Use for high-volume, low-value data

## Scaling Strategy

```

Phase 1: Single Server (MVP)
└─ Single Elasticsearch node
└─ Single Prometheus
└─ Filebeat + Logstash on same hosts
└─ Cost: $500/month
└─ Scale: Up to 10 hosts

Phase 2: Small Cluster
└─ 3-node Elasticsearch
└─ Kafka for buffering
└─ Prometheus with local storage
└─ Cost: $2,000/month
└─ Scale: Up to 100 hosts

Phase 3: Regional Deployment
└─ 20+ node Elasticsearch
└─ Kafka cluster (3+ brokers)
└─ Prometheus federation
└─ S3 for archives
└─ Cost: $10,000/month
└─ Scale: Up to 1,000 hosts

Phase 4: Multi-Region
└─ Regional Elasticsearch clusters
└─ Thanos for global metrics
└─ Federated search
└─ Cost: $50,000+/month
└─ Scale: 10,000+ hosts

```

## Common Pitfalls to Avoid

### **✗ Not planning for growth**

- Start with scalable architecture
- Don't paint yourself into corner

### **✗ Logging everything**

- Log what's needed, not everything
- Sampling for high-volume

### **✗ High-cardinality labels**

- Explosion of time-series
- Memory and disk issues

### ✖ No retention policy

- Storage costs spiral
- Set and enforce policies

### ✖ Single point of failure

- Always replicate
- Multi-AZ deployment

### ✖ Ignoring query performance

- Optimize indexes
- Use time range filters
- Cache frequently accessed data

### ✖ No disaster recovery plan

- Regular backups
- Test restores
- Document procedures

### ✖ Alert fatigue

- Too many alerts = ignored alerts
- Focus on actionable alerts
- Meaningful thresholds

---

## Interview Checklist

### Before the Interview

- Understand the three pillars of observability (logs, metrics, traces)
- Know ELK stack architecture and components
- Understand Prometheus architecture and TSDB
- Review time-series compression techniques
- Study distributed systems challenges (clock skew, consistency)
- Understand storage tiering strategies
- Know common query patterns and optimizations
- Review real-world case studies (Netflix, Uber, etc.)

### During the Interview

#### 1. Clarify Requirements (5 min)

- Scale (how many servers, logs/day, metrics)

- Query patterns (real-time vs historical)
- Retention requirements (hot vs cold)
- Budget constraints

## 2. Capacity Estimation (5 min)

- Log volume (lines/sec, bytes/line)
- Metrics volume (samples/sec, cardinality)
- Storage requirements (hot + cold)
- Infrastructure sizing

## 3. High-Level Design (15 min)

- Draw architecture diagram
- Identify major components
- Explain data flow
- Discuss technology choices

## 4. Deep Dives (20 min)

- Ingestion pipeline (buffering, backpressure)
- Storage architecture (sharding, replication)
- Query optimization
- Alerting system
- Cost optimization

## 5. Scalability (5 min)

- Horizontal scaling approach
- Handling growth (10x, 100x)
- Multi-region strategy

## 6. Fault Tolerance (5 min)

- Failure scenarios and handling
- Disaster recovery plan
- Data durability guarantees

## 7. Trade-offs (5 min)

- Discuss alternatives considered
- Identify potential bottlenecks
- Cost vs performance trade-offs

---

## Further Reading

### Essential Resources

## **Observability Fundamentals:**

- "Distributed Systems Observability" - Cindy Sridharan
- "Site Reliability Engineering" - Google SRE Book
- "Observability Engineering" - Charity Majors, Liz Fong-Jones, George Miranda

## **Elastic Stack:**

- Elasticsearch: The Definitive Guide
- ELK Stack Documentation - [elastic.co/guide](http://elastic.co/guide)
- Elastic Blog - [elastic.co/blog](http://elastic.co/blog)

## **Prometheus & Monitoring:**

- "Prometheus: Up & Running" - Brian Brazil
- Prometheus Documentation - [prometheus.io/docs](http://prometheus.io/docs)
- "Monitoring with Prometheus" - James Turnbull

## **Distributed Systems:**

- "Designing Data-Intensive Applications" - Martin Kleppmann
- "Database Internals" - Alex Petrov
- "The Art of Scalability" - Martin Abbott

## Company Engineering Blogs

### **Netflix:**

- Atlas: Custom monitoring system
- Mantis: Real-time event stream processing
- Edgar: Log ingestion and processing

### **Uber:**

- M3: Distributed metrics platform
- Jaeger: Distributed tracing (now CNCF)
- Real-time analytics with Apache Flink

### **LinkedIn:**

- Venice: Distributed key-value store
- Pinot: Real-time OLAP datastore
- Kafka as log storage

### **Datadog:**

- Multi-tenant observability at scale
- Query optimization techniques
- Sampling strategies

## Tools & Technologies

## **Log Management:**

- ELK Stack (Elasticsearch, Logstash, Kibana)
- EFK Stack (Elasticsearch, Fluentd, Kibana)
- Splunk Enterprise
- Grafana Loki
- Graylog

## **Metrics & Monitoring:**

- Prometheus + Grafana
- InfluxDB + Chronograf
- Datadog (SaaS)
- New Relic (SaaS)
- SignalFx (SaaS)

## **Time-Series Databases:**

- Prometheus TSDB
- InfluxDB
- TimescaleDB
- M3DB
- VictoriaMetrics

## **Message Queues:**

- Apache Kafka
- Amazon Kinesis
- RabbitMQ
- Apache Pulsar

---

# **Appendix**

## **Key Metrics Formulas**

### **Storage Calculation:**

$$\begin{aligned}\text{Log Storage} &= \text{Events/sec} \times \text{Avg Size} \times \text{Seconds/day} \times \text{Retention Days} \\ &= 167,000 \times 500 \text{ bytes} \times 86,400 \times 90 \\ &= 648 \text{ TB}\end{aligned}$$

$$\begin{aligned}\text{Metric Storage} &= \text{Metrics} \times \text{Samples/day} \times \text{Sample Size} \times \text{Retention Days} \\ &= 1\text{M} \times 5.76\text{B} \times 1.5 \text{ bytes} \times 90 \\ &= 774 \text{ GB (with compression)}\end{aligned}$$

### **Throughput Calculation:**

Logs:

- └─ 10,000 servers
- └─ 1,000 lines/min per server
- └─ 10M lines/min = 167K lines/sec
- └─ At 500 bytes/line = 83 MB/sec

Metrics:

- └─ 1M metrics
- └─ 15 second scrape interval
- └─ 1M / 15 = 67K samples/sec
- └─ At 16 bytes/sample = 1 MB/sec

## Cluster Sizing:

Elasticsearch nodes = Daily Volume GB / Node Capacity GB  
= 7,200 GB / 150 GB per node  
= 48 data nodes

Prometheus storage = Metrics × Samples/day × Sample Size  
= 1M × 5.76M × 1.5 bytes  
= 8.6 GB/day  
= 129 GB for 15 days

## Common Latency Numbers

L1 cache reference:	0.5 ns
L2 cache reference:	7 ns
Main memory reference:	100 ns
SSD random read:	150 µs
Disk seek:	10 ms
Read 1 MB sequentially from SSD:	1 ms
Read 1 MB from disk:	20 ms
Round trip in datacenter:	0.5 ms
Round trip CA → Netherlands:	150 ms

Observability System Latencies:

- └─ Filebeat to Logstash: 1–5 ms
- └─ Logstash processing: 0.5–2 ms
- └─ Kafka write: 1–10 ms
- └─ Elasticsearch index: 1–5 seconds (refresh)
- └─ Prometheus scrape: 15 seconds (interval)
- └─ Query latency: 10–1000 ms

## Elasticsearch Performance Tuning

#### JVM Heap Size:

- └─ Recommendation: 50% of RAM, max 31 GB
- └─ Example: 64 GB RAM → 31 GB heap
- └─ Reason: Compressed OOPs optimization
- └─ Rest of RAM for OS file cache

#### Shard Best Practices:

- └─ Shard size: 30–50 GB
- └─ Shards per node: <1000
- └─ Replicas: 1 for production
- └─ Over-sharding hurts performance

#### Indexing Performance:

- └─ Bulk API (1000 documents)
- └─ Refresh interval: 30s (vs 1s default)
- └─ Replica: 0 during bulk, then 1
- └─ Result: 10x faster indexing

## Prometheus Best Practices

#### Cardinality Management:

- └─ Monitor cardinality: prometheus\_tsdb\_symbol\_table\_size\_bytes
- └─ Alert on explosion: >10M time-series
- └─ Use recording rules for aggregation
- └─ Drop high-cardinality labels

#### Memory Sizing:

- └─ Formula: (Metrics × Cardinality × Sample Size) / Compression Ratio
- └─ Example: (1000 × 100 × 16 bytes) / 10 = 160 KB
- └─ Add overhead: × 3
- └─ Total: ~500 KB per metric

#### Scrape Interval:

- └─ Default: 15s (good for most cases)
- └─ High-frequency: 5s (critical systems)
- └─ Low-frequency: 60s (slow-changing metrics)
- └─ Trade-off: Granularity vs storage

## Quick Reference

### Elasticsearch Common Queries

```
// Search error logs in last hour
{
  "query": {
    "bool": {
```

```

        "must": [
            {"match": {"level": "ERROR"}},
            {"range": {"@timestamp": {"gte": "now-1h"}}}
        ]
    }
}

// Count errors by service
{
    "size": 0,
    "aggs": {
        "by_service": {
            "terms": {
                "field": "service.keyword",
                "size": 10
            }
        }
    }
}

// Errors over time (time-series)
{
    "size": 0,
    "aggs": {
        "errors_timeline": {
            "date_histogram": {
                "field": "@timestamp",
                "interval": "5m"
            }
        }
    }
}

```

## Prometheus Common Queries

```

# Request rate
rate(http_requests_total[5m])

# Error percentage
100 * sum(rate(http_requests_total{status="500"}[5m]))
    / sum(rate(http_requests_total[5m]))

# 99th percentile latency
histogram_quantile(0.99,
    sum(rate(http_request_duration_seconds_bucket[5m])) by (le))
)

# Memory usage by container
sum(container_memory_usage_bytes) by (container_name)

```

```
# Predict disk full (linear regression)
predict_linear(node_filesystem_free_bytes[1h], 4*3600) < 0
```

## Alert Examples

### Elasticsearch Watcher:

```
{
  "trigger": {
    "schedule": {
      "interval": "5m"
    }
  },
  "input": {
    "search": {
      "request": {
        "indices": ["logs-*"],
        "body": {
          "query": {
            "bool": {
              "must": [
                {"match": {"level": "ERROR"}},
                {"range": {"@timestamp": {"gte": "now-5m"}}}
              ]
            }
          }
        }
      }
    }
  },
  "condition": {
    "compare": {
      "ctx.payload.hits.total": {
        "gt": 100
      }
    }
  },
  "actions": {
    "notify_slack": {
      "webhook": {
        "url": "https://hooks.slack.com/...",
        "body": "High error rate detected"
      }
    }
  }
}
```

### Prometheus Alertmanager:

```

groups:
- name: system_alerts
  interval: 30s
  rules:
    - alert: HighMemoryUsage
      expr: |
        (node_memory_MemTotal_bytes - node_memory_MemAvailable_bytes)
        / node_memory_MemTotal_bytes > 0.85
      for: 5m
      labels:
        severity: warning
      annotations:
        summary: "Memory usage above 85% on {{ $labels.instance }}"

    - alert: DiskWillFillIn4Hours
      expr: |
        predict_linear(node_filesystem_free_bytes[1h], 4*3600) < 0
      for: 5m
      labels:
        severity: critical
      annotations:
        summary: "Disk on {{ $labels.instance }} will fill in 4 hours"

```

**Document Version:** 1.0 (HLD Focus)

**Last Updated:** January 8, 2025

**Author:** System Design Interview Prep

**Focus:** High-Level Design, Architecture, Scalability

**Status:** Complete & Interview-Ready 

### Key Takeaways:

1. **ELK Stack is industry standard** for log management at scale
2. **Prometheus** is the de facto metrics system for cloud-native applications
3. **Kafka provides crucial buffering** and decouples producers from consumers
4. **Tiered storage** saves 70-80% on costs while maintaining performance
5. **Push model for logs**, pull model for metrics (with exceptions)
6. **Structured logging (JSON)** enables faster queries and aggregations
7. **Low-cardinality labels** are critical for metrics performance
8. **Multi-region deployment** provides data locality and fault isolation
9. **Downsampling and compression** enable long-term metric retention
10. **Eventual consistency is acceptable** for observability data

Good luck with your system design interview!  