

Rate Limiter System Design - High-Level Design Focus

Table of Contents

1. Problem Statement & Requirements
 2. High-Level Architecture
 3. Core Components
 4. Algorithm Selection
 5. Data Flow & Request Processing
 6. Distributed System Design
 7. Scalability & Performance
 8. Fault Tolerance & Reliability
 9. Multi-Region Architecture
 10. Design Trade-offs
 11. Real-World Case Studies
-

Problem Statement & Requirements

Problem Definition

Design a distributed rate limiting system that can:

- Limit API requests per user/IP/API key
- Handle 10,000+ requests/second
- Work across multiple servers and regions
- Maintain <10ms latency overhead
- Support multiple rate limit rules and tiers

Functional Requirements

1. **Limit enforcement:** Block requests exceeding configured limits
2. **Multiple dimensions:** Rate limit by user ID, IP, API key, endpoint
3. **Flexible windows:** Support second, minute, hour, day windows
4. **Dynamic rules:** Update limits without service restart
5. **Clear feedback:** Return meaningful error messages with retry information

Non-Functional Requirements

1. **Performance:** P99 latency < 10ms, support 10K+ RPS
2. **Availability:** 99.99% uptime
3. **Scalability:** Horizontal scaling, handle 10x growth
4. **Consistency:** Accept <5% over-limit (eventual consistency)
5. **Durability:** Rules survive restarts

Capacity Estimation

Assumptions:

- 100M total users
- 10M DAU
- 50 requests/user/day
- Peak = 3x average

Daily requests = $10M \times 50 = 500M$

Average RPS = $500M / 86400 \approx 5,800$

Peak RPS = 17,400

Storage (Redis):

- Per user counter: 8 bytes
- 4 time windows: 32 bytes/user
- 10M DAU: 320MB
- With overhead: ~640MB

Servers needed:

- Each server: 1,000 RPS
- Peak: $17,400 / 1,000 = 18$ servers
- With redundancy & growth: 70–100 servers


Cost (AWS):

- API Gateway: ~\$5,000/month
- Redis Cluster: ~\$1,000/month
- Load Balancers: ~\$100/month

Total: ~\$6,000/month


High-Level Architecture

System Overview



- AWS ALB / NGINX
- SSL termination
- Health checks
- Geographic routing

- API Gateway Cluster (Stateless)
- Rate Limiter Middleware
 - Request identification (user, IP, API key)
 - Rule matching & evaluation
 - Counter management



Counter Storage: ratelimit:{user}:{endpoint}:{window}

- Configuration Service (PostgreSQL)
- Rate limit rules
 - User tiers
 - Whitelist/Blacklist
 - Audit logs

Monitoring & Observability
 Metrics: Prometheus | Logs: ELK | Tracing: Jaeger

Key Architectural Decisions

1. Stateless API Gateway

- **Why:** Enables horizontal scaling, no session affinity needed
- **How:** All state in Redis, gateways are compute-only
- **Trade-off:** Network latency to Redis vs local cache

2. Redis as Central Store

- **Why:** Fast (sub-millisecond), atomic operations, proven at scale
- **Alternatives considered:**
 - In-memory (doesn't scale across servers)
 - Cassandra (higher latency ~10ms)
 - DynamoDB (higher latency, eventually consistent)
- **Trade-off:** Single point of failure → mitigated with clustering

3. Separate Configuration Store

- **Why:** Rules change infrequently, don't need sub-ms access
- **How:** PostgreSQL with caching layer in API gateways
- **Trade-off:** Eventual consistency for rule updates (acceptable 5min delay)

4. Multi-Tier Caching

```
L1: In-memory rules cache (5 min TTL)
    ↓ (cache miss)
L2: Redis rules cache (15 min TTL)
    ↓ (cache miss)
L3: PostgreSQL (source of truth)
```

Core Components

1. Request Identifier Service

Purpose: Extract and normalize user identifiers

Input: HTTP Request
Output: Composite Key

Components:

```
User ID: user_12345
IP Address: 192.168.1.1
API Key: sk_live_abc123
Endpoint: /api/payments
Method: POST
```

↓
Generate Key: "ratelimit:user_12345:/api/payments:1640995200"

Design Considerations:

- Hash user IDs for privacy
- Handle missing identifiers (fallback to IP)
- Normalize endpoints (strip query params)
- Consider X-Forwarded-For for true client IP

2. Rules Engine

Purpose: Match requests to rate limit rules with precedence

Rule Priority System:

Priority 1: User-specific + Endpoint-specific
Example: Premium user calling /api/payments
Limit: 10,000 req/hour

↓ (no match)

Priority 2: Endpoint-specific
Example: Any user calling /api/payments
Limit: 1,000 req/hour

↓ (no match)

Priority 3: User tier-specific
Example: Free tier user, any endpoint
Limit: 100 req/hour

↓ (no match)

Priority 4: Global default
Limit: 1,000 req/hour

Key Features:

- Rule caching (5-minute TTL)
- Hot reload without restart
- A/B testing support
- Gradual rollout capability

3. Counter Management Service

Purpose: Track request counts atomically

Architecture:

```
API Gateway
  ↓
[Check Counter] → Redis GET current_window, previous_window
  ↓
[Calculate] → Sliding window estimation
  ↓
[Decision] → Allow / Deny
  ↓
[Update] → Redis INCR current_window (if allowed)
```

Atomicity Guarantees:

- Use Redis Lua scripts for atomic check-and-increment
- Prevents race conditions across servers
- Handles concurrent requests correctly

4. Response Builder

Purpose: Generate consistent API responses

```
Allowed Response:
├── Status: 200 OK
└── Headers:
    ├── X-RateLimit-Limit: 1000
    ├── X-RateLimit-Remaining: 742
    └── X-RateLimit-Reset: 1640995200
└── Body: [Original API Response]

Denied Response:
├── Status: 429 Too Many Requests
└── Headers:
    ├── X-RateLimit-Limit: 1000
    ├── X-RateLimit-Remaining: 0
    ├── X-RateLimit-Reset: 1640995200
    └── Retry-After: 60
└── Body:
    {
        "error": "RATE_LIMIT_EXCEEDED",
        "message": "Try again in 60 seconds",
        "limit": 1000,
        "retry_after": 60
    }
```

Algorithm Selection

Algorithm Comparison Matrix

Criteria	Fixed Window	Sliding Log	Sliding Counter	Token Bucket
Accuracy	★★	★★★★★	★★★★★	★★★★
Memory	★★★★★	★★	★★★★★	★★★★★
Complexity	★★★★★	★★	★★★★	★★★★
Burst Handling	✗	✓	✓	✓ ✓
Boundary Issues	✗	✓	✓	✓
Distributed	✓	⚠	✓	⚠

Recommended: Sliding Window Counter

Why this algorithm?

1. **Good accuracy** (~95-99%, avoids fixed window boundary issue)
2. **Memory efficient** (only 2 counters per user per endpoint)
3. **Distributed-friendly** (works well with Redis)
4. **Handles bursts** (smooth transition between windows)

How it works conceptually:

```
Previous Window: [=====] 80 requests
Current Window: [=      ] 30 requests
                  ↑ Now (50% through window)
```

Estimated count = $30 + (80 \times 50\%) = 30 + 40 = 70$ requests

If limit is 100: ALLOW ($70 < 100$)
If limit is 60: DENY ($70 > 60$)

Architecture:

Redis Keys:

- ratelimit:user_123:/api/data:1640995200 → 30
- ratelimit:user_123:/api/data:1640995140 → 80

Calculation:

1. Get both counters (2 Redis GETs)
2. Calculate weighted sum
3. Compare to limit
4. If allowed: INCR current counter
5. Set TTL on counter (2x window size)

Alternative: Token Bucket (for burst scenarios)

When to use:

- Mobile apps that sync periodically
- Batch operations
- APIs where legitimate bursts are common

Conceptual model:

Bucket Capacity: 100 tokens
Refill Rate: 10 tokens/second

Time 0s: [100 tokens] ← Full
↓ 50 requests
Time 1s: [50 tokens] ← Half full
↓ (refilled 10)
Time 2s: [60 tokens]
↓ 0 requests
Time 3s: [70 tokens] ← Gradually refilling

Data Flow & Request Processing

Detailed Request Flow

STEP 1: Request Arrives

Client Request
↓
POST /api/payments
Headers:
X-API-Key: sk_live_abc123
X-User-ID: user_12345
From IP: 192.168.1.1

STEP 2: Load Balancer Routing

↓
[SSL Termination]
[Health Check]
[Route to API Gateway]
↓
Selected: API Gateway Instance #7

STEP 3: Rate Limiter Middleware Intercepts

```
↓  
Extract Identifiers:  
  user_id: user_12345  
  api_key: sk_live_abc123  
  ip_address: 192.168.1.1  
  endpoint: /api/payments  
↓  
Check Local Rules Cache:  
  Cache Key: "rule:user_12345:/api/payments"  
  Cache TTL: 5 minutes  
  Status: HIT ✓  
↓  
Applicable Rule:  
  limit: 1000 requests  
  window: 3600 seconds (1 hour)  
  algorithm: sliding_window_counter
```

STEP 4: Query Redis for Counters

```
↓  
Current Time: 1640995200  
Current Window: 1640995200  
Previous Window: 1640991600  
↓  
Redis Pipeline (batch operation):  
  GET ratelimit:user_12345:/api/payments:1640995200  
  GET ratelimit:user_12345:/api/payments:1640991600  
↓  
Results:  
  current_count: 45  
  previous_count: 820
```

STEP 5: Calculate Weighted Count

```
↓  
Time elapsed in current window: 1800s (50%)  
Previous window weight: 50%  
↓  
Estimated count = 45 + (820 × 0.5) = 45 + 410 = 455  
↓  
Compare: 455 < 1000 ✓  
Decision: ALLOW
```

STEP 6: Update Counter

```
↓  
Redis Commands:  
  INCR ratelimit:user_12345:/api/payments:1640995200  
  EXPIRE ratelimit:user_12345:/api/payments:1640995200 7200
```

↓
New count: 46

STEP 7: Add Response Headers

↓
X-RateLimit-Limit: 1000
X-RateLimit-Remaining: 545
X-RateLimit-Reset: 1640998800

STEP 8: Forward to Backend Service

↓
[Process Payment]
[Return Response]
↓
Response: 200 OK {"payment_id": "12345"}

STEP 9: Return to Client

↓
Response with rate limit headers

Performance Optimization

Total Latency Breakdown:

Rate Limiter Overhead:
└─ Rule lookup (cache): 0.1ms
└─ Redis GET (2 keys): 0.5ms
└─ Calculation: 0.1ms
└─ Redis INCR: 0.3ms
└─ Total: ~1ms ✓

Target: <10ms (10% of typical API response time)
Actual: ~1ms (well within target)

Distributed System Design

Challenge 1: Maintaining Consistency

Problem: Multiple servers incrementing same counter

Time T0:
Server A reads count: 99
Server B reads count: 99

Time T1:
Server A writes: 100
Server B writes: 100

Result: Lost update! Should be 101

Solution: Atomic Operations

Instead of Read-Then-Write:
 count = GET(key)
 count += 1
 SET(key, count)

Use Atomic Increment:
 INCR(key) ← Atomic operation in Redis

Advanced: Lua Scripts for Complex Logic

```
-- Atomic check-and-increment
local current = redis.call('INCR', key)
if current == 1 then
    redis.call('EXPIRE', key, window)
end

if current > limit then
    redis.call('DECR', key) -- Rollback
    return 0 -- Rejected
else
    return 1 -- Allowed
end
```

Challenge 2: Clock Synchronization

Problem: Server clocks drift

Server A: 10:00:00.000
Server B: 10:00:00.500

Window boundary at 10:00:00:
- Different servers see different windows
- Inconsistent rate limiting

Solutions:

Option 1: Centralized Time Source

- All servers → Query Redis TIME command
- Redis provides synchronized timestamp
- Network latency: +0.5ms
- Trade-off: Slight overhead for consistency

Option 2: NTP Synchronization

- All servers sync with NTP
- Accuracy: ±50ms
- Good enough for rate limiting (not for distributed transactions)
- Zero runtime overhead

Option 3: Accept Slight Inaccuracy

- Most rate limiting can tolerate ±1 second drift
- Simpler architecture
- Focus on approximate fairness

Challenge 3: Redis Failure Scenarios

Scenario A: Master Failure

Normal:

API Gateways → Redis Master → Redis Replica

Failure:

Master crashes

↓

Sentinel detects failure (<5s)

↓

Promotes Replica to Master (<5s)

↓

Updates DNS / Endpoint

↓

API Gateways reconnect

Impact: 5–10s of potential data loss

Mitigation: Fail-open during failover

Scenario B: Network Partition

Datacenter split:
API Gateways (DC1) ↔ Redis (DC2)

Options:

1. Fail Open (Recommended)
 - Allow all requests
 - Log failures
 - Resume when partition heals
2. Fail Closed (High security)
 - Deny all requests
 - Potential service outage
3. Local Fallback
 - Use in-memory counters
 - Less accurate but available
 - Sync when partition heals

Challenge 4: Hot Key Problem

Problem: Celebrity user causes hot key

Taylor Swift's API key: "ts_key"
Rate limit key: "ratelimit:ts_key:*

All requests hit same Redis shard
→ Shard overloaded
→ Latency spikes for everyone

Solutions:

Option 1: Key Splitting

Instead of:
ratelimit:ts_key:endpoint>window

Use:
ratelimit:ts_key:endpoint>window:shard_0
ratelimit:ts_key:endpoint>window:shard_1
ratelimit:ts_key:endpoint>window:shard_2

Sum across shards for total count
Distribute load across Redis shards

Option 2: Dedicated Shard

```
Detect high-traffic users  
Move to dedicated Redis instance  
Isolate from normal traffic
```

Option 3: Client-Side Rate Limiting

```
For known high-volume clients:  
- Implement client-side token bucket  
- Reduce server load  
- Still enforce server-side as backstop
```

Scalability & Performance

Scaling Dimensions

1. Horizontal Scaling (API Gateways)

Auto-Scaling Strategy:

```
Metrics to monitor:  
└─ CPU utilization: Scale at 70%  
└─ Request latency: Scale at P95 > 100ms  
└─ Request rate: Scale at 80% capacity  
└─ Memory usage: Scale at 80%
```

```
Scaling behavior:  
└─ Scale up: +50% instances in 60s  
└─ Scale down: -10% instances in 5 minutes  
└─ Min: 10 instances, Max: 100 instances
```

Load Balancing:

```
Algorithm: Least connections + IP hash (for debugging)
```

```
Health checks:  
└─ HTTP GET /health every 10s  
└─ 2 consecutive failures → Remove from pool  
└─ 2 consecutive successes → Add back to pool  
└─ Drain connections before removal (30s grace period)
```

2. Vertical Scaling (Redis)

Redis Sizing:

Small (0–1K RPS):

- 2 GB RAM, 2 vCPU
- Single master + replica
- Cost: \$100/month

Medium (1K–10K RPS):

- 16 GB RAM, 4 vCPU
- 3 masters + 3 replicas (sharded)
- Cost: \$500/month

Large (10K–100K RPS):

- 64 GB RAM per node, 8 vCPU
- 10 masters + 10 replicas
- Cost: \$2,000/month

3. Data Partitioning (Redis Cluster)

Sharding Strategy:

Consistent Hashing:

`hash(user_id) % num_shards → shard_id`

Example with 3 shards:

`user_123 → hash → 42 → 42 % 3 = 0 → Shard 0`
`user_456 → hash → 87 → 87 % 3 = 0 → Shard 0`
`user_789 → hash → 14 → 14 % 3 = 2 → Shard 2`

Benefits:

- ✓ Evenly distributed load
- ✓ Easy to add shards
- ✓ Minimal data movement

Redis Cluster Mode:

Hash Slots: 16,384 slots

Distribution:

└─ Master 1: slots 0–5461
└─ Master 2: slots 5462–10922
└─ Master 3: slots 10923–16383

Auto-sharding:

Redis automatically routes keys to correct shard
Client libraries handle redirection

Performance Optimization Techniques

1. Connection Pooling

Problem: Creating new Redis connections is expensive (~10ms)

Solution: Connection Pool

- └─ Min connections: 10
- └─ Max connections: 50
- └─ Idle timeout: 60s
- └─ Connection reuse: ✓
- └─ Latency reduction: 10ms → 0.5ms

2. Pipelining

Without Pipelining:

```
GET key1 → 0.5ms  
GET key2 → 0.5ms  
GET key3 → 0.5ms  
Total: 1.5ms
```

With Pipelining:

```
PIPELINE {  
    GET key1  
    GET key2  
    GET key3  
}  
Total: 0.6ms (60% faster)
```

3. Caching Strategy

L1 Cache (In-Memory):

- └─ Rules cache: 10,000 rules
- └─ TTL: 5 minutes
- └─ Hit rate: 95%
- └─ Latency: 0.1ms
- └─ Eviction: LRU


L2 Cache (Redis):

- └─ Rules cache: 100,000 rules
- └─ TTL: 15 minutes
- └─ Hit rate: 99%
- └─ Latency: 0.5ms

- L3 (PostgreSQL):
 - Source of truth

Fault Tolerance & Reliability

High Availability Architecture



Recovery Objectives:

- RPO (Recovery Point Objective): <1 second
- RT0 (Recovery Time Objective): <30 seconds
- Data Loss: Accept <100 requests during failover

Failure Modes & Handling

1. Redis Unavailable

Detection:

- Health check every 1s
- 3 consecutive failures → Trigger failover
- Alert ops team

Response Options:

A. Fail Open (Recommended):

- ✓ Allow all requests
 - ✓ Log failures for audit
 - ✓ Resume normal operation when Redis recovers
- Risk: No rate limiting for ~30s

B. Fail Closed:

- ✗ Deny all requests
- ✗ Service outage

Benefit: Maintain strict limits

C. Degraded Mode:

- ✓ Use local in-memory counters
- ✓ Less accurate but functional
- ✓ Sync when Redis recovers

2. Configuration Service Down

Impact: Cannot load new rules or user tiers

Mitigation:

- └─ Rules cached in API gateways (5 min TTL)
- └─ Stale data acceptable for rate limiting
- └─ Continue with cached rules until service recovers
- └─ No immediate impact on rate limiting operations

3. Network Partition


Symptom: API gateways can't reach Redis


Decision Tree:

- └─ If latency > 50ms:
 - └─ Use cached counters (degraded mode)
- └─ If timeout (>1s):
 - └─ Fail open (allow requests)
- └─ Log all decisions for audit

Multi-Region Architecture

Global Deployment Strategy





Traffic Distribution:

- └─ US-WEST: 40% (California, Pacific)
- └─ US-EAST: 35% (East Coast, South America)
- └─ EU-WEST: 25% (Europe, Middle East, Africa)

Consistency Model Options

Option 1: Regional Independence (Recommended)

Strategy: Each region operates independently

Rate Limit Distribution:

- └─ Global limit: 1000 req/hour
- └─ US-WEST quota: 400 req/hour
- └─ US-EAST quota: 350 req/hour
- └─ EU-WEST quota: 250 req/hour

Pros:

- ✓ Zero cross-region latency
- ✓ Regions isolated from each other's failures
- ✓ Simple to implement

Cons:

- ✗ User can potentially use quota in multiple regions
- ✗ Over-limit by ~20% if user distributes requests

Acceptable because:

- Users typically stay in one region
- 20% over-limit acceptable per requirements
- Regional failures don't cascade

Option 2: Global Synchronization

Strategy: Central Redis for all regions

Architecture:

All regions → Central Redis (US-EAST)

Pros:

- ✓ Strict global limits
- ✓ Perfect accuracy

Cons:

- ✗ Cross-region latency (+50–200ms)
- ✗ Single point of failure
- ✗ Expensive (cross-region data transfer)

Use when:

- Strict accuracy required (financial APIs)
- Traffic is low (<1K RPS globally)
- Latency less critical

Option 3: Async Synchronization

Strategy: Sync counters every 1–5 seconds

Flow:

1. User hits US-WEST: 45 requests
2. Background job syncs to other regions
3. US-EAST Redis updated: 45 requests
4. User hits US-EAST: Sees 45 from counter

Pros:

- ✓ Better accuracy than regional independence
- ✓ Lower latency than global sync

Cons:

- ✗ Complex implementation
- ✗ Eventual consistency issues
- ✗ Race conditions during sync window

Region Failover Strategy

Normal Operation:

User (California) → US-WEST → Local Redis

US-WEST Failure:

User (California) → Route53 detects failure
→ Redirect to US-EAST
→ User quota starts fresh

Failover Implications:

- |— User quota reset in new region
- |— Can potentially double usage

- Acceptable trade-off for availability
- Monitor and alert on region failures

Recovery:

- US-WEST comes back online
- Gradual traffic shift (10% per minute)
- Monitor for issues
- Full traffic in ~10 minutes

Design Trade-offs

1. Accuracy vs Performance

Approach	Accuracy	Latency	Memory	Choice
Sliding Window Log	99.9%	2-5ms	High	High-value APIs
Sliding Window Counter	95-99%	1-2ms	Medium	Default
Fixed Window	90-95%	0.5ms	Low	MVP/Low-traffic
Token Bucket	95%	1ms	Low	Burst scenarios

Decision: Sliding Window Counter

- Sweet spot between accuracy and performance
- Handles 99% of use cases
- Can upgrade to Sliding Window Log for critical endpoints

2. Consistency vs Availability

CAP Theorem Applied to Rate Limiting:

Partition Tolerance (Required):

- Must work when networks fail
- Can't sacrifice this

Choose One:

Consistency (Strict Limits)

- Reject at exactly limit
- May deny valid requests on failure
- Lower availability

Trade-off: 99.9% availability

vs

Availability (Always Available)

- May allow 5-10% over-limit

- Always serve requests
- Higher availability

Trade-off: 99.99% availability

****Decision**:** Favor Availability

- Rate limiting is not security-critical
- 5% over-limit acceptable
- Better user experience

3. Centralized vs Distributed State

Centralized (Redis):

Pros:

- ✓ Single source of truth
- ✓ Accurate counts
- ✓ Simple logic

Cons:

- ✗ Network latency
- ✗ Single point of failure (mitigated with clustering)

Distributed (In-Memory):

Pros:

- ✓ Zero network latency
- ✓ No external dependencies

Cons:

- ✗ Inconsistent across servers
- ✗ Lost on restart
- ✗ Hard to implement correctly

****Decision**:** Centralized with caching

- Use Redis for counters (accuracy)
- Cache rules in-memory (performance)
- Best of both worlds

4. Pull vs Push for Rule Updates

Pull (Polling):

API Gateways poll PostgreSQL every 5 minutes

Pros:

- ✓ Simple implementation
- ✓ Predictable load

Cons:

- ✗ 5-minute delay for updates
- ✗ Unnecessary queries if no changes

Push (Event-Driven):

PostgreSQL → Kafka → API Gateways

Pros:

- ✓ Immediate updates
- ✓ No unnecessary queries

Cons:

- ✗ More complex
- ✗ Another service to manage (Kafka)

Decision: Pull for now, Push later

- 5-minute delay acceptable
- Simpler to implement and operate
- Can add push notifications later if needed

Real-World Case Studies

Case Study 1: Stripe API Rate Limiting

Scale: 100,000+ RPS globally

Architecture:

- └─ Algorithm: Token Bucket + Sliding Window
- └─ Multi-tier limits:
 - └─ Per API key: 100 reads/sec, 25 writes/sec
 - └─ Per endpoint: Different limits
 - └─ Test vs Production: Separate limits
- └─ Storage: Redis Cluster (15 nodes)
- └─ Regions: 5 global regions

Key Design Decisions:

1. **Separate read/write limits:** Writes more expensive, stricter limits
2. **Burst allowance:** Token bucket allows legitimate bursts
3. **Graceful degradation:** Fail open on Redis issues
4. **Rich error messages:** Include retry-after, remaining quota

Lessons Learned:

- Communicate limits clearly in documentation
- Provide SDK libraries that respect limits
- Monitor P95/P99 of limit usage (identify power users)
- Allow temporary limit increases for events

Case Study 2: GitHub API

Scale: 50,000+ RPS

Architecture:

```
└── Algorithm: Fixed Window (1 hour)
    └── Limits:
        ├── Unauthenticated: 60 req/hour per IP
        ├── Authenticated: 5,000 req/hour per user
        └── GraphQL: Point-based system
    └── Secondary limit: 100 req/minute (burst protection)
    └── Storage: Redis + PostgreSQL
```

Key Design Decisions:

1. **GraphQL complexity scoring:** Each query costs points based on complexity
2. **Conditional requests:** 304 Not Modified doesn't count
3. **GitHub Apps:** Higher limits (15,000 req/hour)
4. **Secondary rate limit:** Prevents rapid bursts

Lessons Learned:

- Simple hourly window sufficient for most cases
- Secondary burst limit catches abuse
- Different auth methods need different limits
- Conditional requests improve UX without cost

Case Study 3: Twitter API

Scale: 200,000+ RPS

Architecture:

```
└── Algorithm: Fixed Window (15 minutes)
    └── Multi-dimensional limits:
        ├── User context: Per user
        ├── App context: Per app
        └── Endpoint-specific: Different for each endpoint
    └── Storage: Distributed cache + Redis
    └── Regions: 4 global regions
```

Key Design Decisions:

1. **15-minute windows:** Balance between accuracy and UX
2. **Per-endpoint limits:** Tweet posting stricter than reading
3. **Multiple rate limit types:** User and app limits combined

4. **Premium tiers:** Different limits for different access levels

Lessons Learned:

- Clear documentation of limits crucial
- Rate limits as product differentiator (premium tiers)
- Window resets at fixed intervals (predictable for users)
- Must handle high-profile events (millions of users simultaneously)

Case Study 4: AWS API Gateway

Scale: 1,000,000+ RPS

Architecture:

- Algorithm: Token Bucket
- Hierarchical limits:
 - Account-level: 10,000 req/sec
 - API-level: Configurable
 - Stage-level: Per environment
 - Method-level: Per endpoint
- Burst: 5,000 requests
- Storage: Distributed in-memory + DynamoDB

Key Design Decisions:

1. **Token bucket for bursts:** AWS workloads often bursty
2. **Hierarchical limits:** Multiple layers of protection
3. **Per-stage configuration:** Dev/staging/prod different limits
4. **Usage plans:** Package limits with API keys

Lessons Learned:

- Burst capacity critical for serverless workloads
- Multiple limit tiers provide flexibility
- Integration with billing system (pay for higher limits)
- CloudWatch metrics for monitoring and alerting

Summary & Key Takeaways

Architecture Principles

1. Stateless Application Layer

- API gateways should be stateless
- Enables horizontal scaling
- Simplifies deployment and operations

2. Centralized State Management

- Use Redis for counters (fast, atomic)
- Use PostgreSQL for rules (durable, queryable)
- Cache aggressively for performance

3. Favor Availability Over Strict Consistency

- Accept 5% over-limit
- Fail open on errors
- Better UX than false rejections

4. Design for Failure

- Redis cluster with replication
- Multiple availability zones
- Graceful degradation strategies

Scaling Strategy

Growth Path:

Phase 1 (0–1K RPS): Single region, minimal setup

- └─ 3 API gateways
- └─ 1 Redis master + replica
- └─ Cost: \$500/month

Phase 2 (1K–10K RPS): Regional scaling

- └─ 10–20 API gateways
- └─ Redis cluster (3 masters + replicas)
- └─ Cost: \$5,000/month

Phase 3 (10K–100K RPS): Multi-region

- └─ 3 regions × 20 gateways each
- └─ Regional Redis clusters
- └─ Cross-region config replication
- └─ Cost: \$30,000/month

Phase 4 (100K+ RPS): Global scale

- └─ 5+ regions
- └─ CDN integration
- └─ Custom optimizations
- └─ Cost: \$100,000+/month

Decision Matrix

Choose Sliding Window Counter when:

- Building a new rate limiter
- Need balance of accuracy and performance

- Can tolerate ~5% over-limit
- This is the recommended default

Choose Sliding Window Log when:

- Need maximum accuracy
- Financial transactions or high-value operations
- Can afford higher memory usage

Choose Token Bucket when:

- Legitimate burst traffic expected
- Mobile apps, batch operations
- Variable request sizes

Choose Fixed Window when:

- MVP or prototype stage
- Very low traffic (<100 RPS)
- Simplicity more important than accuracy

Interview Success Tips

1. **Start with requirements** - Don't jump to solution
2. **Draw architecture diagram** - Visual communication is key
3. **Discuss trade-offs explicitly** - Show you understand nuances
4. **Consider failure scenarios** - What happens when things break?
5. **Estimate capacity** - Show you can do back-of-envelope math
6. **Reference real systems** - "Similar to how Stripe does it..."
7. **Think about operations** - Monitoring, alerting, debugging
8. **Iterate on design** - Start simple, add complexity as needed

Common Pitfalls to Avoid

- ✖ Over-engineering for day 1
 - ✖ Ignoring the CAP theorem
 - ✖ Forgetting about clock synchronization
 - ✖ Not planning for failure scenarios
 - ✖ Underestimating Redis performance
 - ✖ Not considering hot key problems
 - ✖ Forgetting about operational concerns
 - ✖ Focusing too much on code details vs architecture
-

Further Reading

Essential Resources

Rate Limiting Concepts:

- "Rate Limiting Strategies" - Redis documentation
- "Distributed Rate Limiting" - Martin Kleppmann
- "API Rate Limiting Best Practices" - Kong API Gateway

System Design:

- "Designing Data-Intensive Applications" - Martin Kleppmann
- "System Design Interview" - Alex Xu
- "Building Microservices" - Sam Newman

Real-World Implementations:

- Stripe API Rate Limiting - stripe.com/docs/rate-limits
- GitHub REST API - docs.github.com/en/rest/overview/resources-in-the-rest-api
- AWS API Gateway Throttling - AWS documentation

Redis & Distributed Systems:

- "Redis in Action" - Josiah Carlson
- "Designing Distributed Systems" - Brendan Burns
- "Site Reliability Engineering" - Google SRE Book

Document Version: 1.0 (HLD Focus)

Last Updated: 2024

Focus: High-Level Design, Architecture, Scalability

<task_progress>

- Create HLD-focused rate limiter design document
- Cover problem statement and requirements
- Design high-level architecture
- Explain core components
- Compare algorithms at conceptual level
- Detail distributed system challenges
- Cover scalability and performance
- Discuss fault tolerance
- Add multi-region architecture
- Explain design trade-offs
- Include real-world case studies
- Add summary and key takeaways

</task_progress>