

# System Design Decision Guide - When to Use What (HLD Focus)

---

## Table of Contents

1. [Introduction](#)
  2. [Load Balancer Decisions](#)
  3. [Database Decisions](#)
  4. [Caching Decisions](#)
  5. [Replication Decisions](#)
  6. [Sharding Decisions](#)
  7. [Consistency Decisions](#)
  8. [Communication Decisions](#)
  9. [Scaling Decisions](#)
  10. [Decision Trees](#)
- 

## Introduction

"System design is about making the **RIGHT** decisions for **YOUR** specific requirements."

This guide focuses on **decision-making** in HLD:

- When to use each approach
  - What questions to ask
  - Trade-offs for each choice
  - Real-world examples
  - No code (focus on architecture)
- 

## Load Balancer Decisions

Layer 4 vs Layer 7 Load Balancer

### When to Use Layer 4 (Transport Layer)

**Use When:**

- Need maximum performance (low latency critical)
- Simple routing (all requests go to any server)
- Non-HTTP traffic (database connections, WebSocket, custom protocols)
- Very high throughput (millions of requests/second)
- Cost-sensitive (cheaper hardware)

**Real Examples:**

- **Database connection pooling:** Route PostgreSQL connections
- **Gaming servers:** Route UDP game traffic
- **Video streaming:** Route RTMP streams
- **IoT:** Route MQTT messages from millions of devices

#### **Decision Factors:**

Traffic Type: Simple/Uniform → Layer 4  
 Traffic Type: HTTP with routing → Layer 7

Throughput: > 1M QPS → Layer 4  
 Throughput: < 1M QPS → Layer 7

Protocol: HTTP only → Layer 7  
 Protocol: Mixed (HTTP + WebSocket + DB) → Layer 4

### **When to Use Layer 7 (Application Layer)**

#### **Use When:**

- Need content-based routing (/api/users → User Service)
- Microservices architecture (route by URL path)
- Need SSL termination (decrypt once at LB)
- Want to inspect/modify HTTP headers
- A/B testing or canary deployments
- Need to aggregate responses from multiple services

#### **Real Examples:**

- **Uber:** Routes `/riders/*` to rider services, `/drivers/*` to driver services
- **Netflix:** Routes by device type (mobile vs TV vs web)
- **Amazon:** Routes `/products` to product service, `/cart` to cart service
- **Twitter:** Routes authenticated vs unauthenticated traffic differently

#### **Decision Factors:**

Architecture: Monolith → Layer 4  
 Architecture: Microservices → Layer 7

Routing Needs: IP/Port only → Layer 4  
 Routing Needs: URL/Headers → Layer 7

Security: Basic → Layer 4  
 Security: WAF, DDoS protection → Layer 7

---

## Load Balancing Algorithm Choice

### Round Robin

#### Use When:

- All requests take similar time
- All servers have same capacity
- Simple setup needed
- Starting point (simplest)

#### Avoid When:

- Mixed workload (uploads + API calls)
- Servers have different capacity
- Long-running requests

**Example:** Static website serving, API with uniform endpoints

---

### Least Connections

#### Use When:

- Mixed request types (fast + slow)
- Long-lived connections (WebSocket, HTTP/2)
- File uploads/downloads
- Varying request durations

#### Avoid When:

- All requests are similar (Round Robin simpler)
- Very high QPS (tracking connections has overhead)

**Example:** Netflix (streaming + API), file sharing services, chat applications

---

### IP Hash / Sticky Sessions

#### Use When:

- Server maintains session state in memory
- WebSocket connections
- Multi-step workflows
- Can't use external session store (legacy app)

#### Avoid When:

- Can use Redis for sessions (better!)
- Need even distribution
- Server failures cause data loss

**Better Alternative:** Always prefer Redis for sessions over sticky sessions

---

#### Interview Decision Framework:

Question: "What load balancing algorithm would you use?"

Answer Template:

"I'd use [ALGORITHM] because:

1. Traffic pattern: [uniform/mixed/long-lived]
2. Architecture: [monolith/microservices]
3. State management: [stateless/stateful]
4. Performance requirements: [throughput/latency]

Specifically:

- Layer 7 for microservices routing
- Least Connections for mixed workload
- Avoid sticky sessions by using Redis

Trade-off: [mention downside of chosen approach]"

---

## Database Decisions

SQL vs NoSQL

#### When to Use SQL (PostgreSQL, MySQL)

#### Use When:

- Data is structured and relationships matter
- Need ACID transactions (payments, banking, inventory)
- Complex queries with JOINs required

- Schema is stable and well-defined
- Data integrity critical (foreign keys, constraints)
- Strong consistency required
- Team familiar with SQL
- Data size < 10 TB (can fit in single instance)

### Real Examples:

- **Uber**: User accounts, payment transactions (need ACID)
- **Airbnb**: Booking system, payments (need transactions)
- **Instagram**: User profiles, relationships (structured data)
- **Stripe**: All financial data (strict consistency)

### When It Breaks Down:

- △ Write throughput > 10K QPS → Consider Cassandra
- △ Data size > 10 TB → Need sharding (complex)
- △ Schema changes frequently → Consider MongoDB
- △ Need geographic distribution → Consider Cassandra

## When to Use Cassandra (Wide-Column NoSQL)

### Use When:

- Write-heavy workload (> 10K write QPS)
- Time-series data (logs, events, metrics)
- Need linear scalability (add nodes = add capacity)
- Geographic distribution required
- High availability critical (no single point of failure)
- Data size > 100 TB
- OK with eventual consistency

### Real Examples:

- **Twitter**: 500M tweets/day (6K write QPS) → Cassandra
- **Netflix**: Viewing history (billions of events) → Cassandra
- **Instagram**: Activity feed (write-heavy) → Cassandra
- **Apple**: iMessage logs (billions of messages) → Cassandra

### When to Avoid:

- Need ACID transactions → Use PostgreSQL
- Complex queries with aggregations → Use PostgreSQL

- ✗ Data size < 1 TB → PostgreSQL simpler
- ✗ Need strong consistency → Use PostgreSQL

## When to Use MongoDB (Document Store)

### Use When:

- ✓ Schema is flexible/changing
- ✓ Nested documents (product catalog with varying fields)
- ✓ Rapid development (schema can evolve)
- ✓ Mobile app backend
- ✓ Content management system
- ✓ Need some transactions (v4.0+)

### Real Examples:

- **Uber**: Trip data (varying fields per trip type)
- **eBay**: Product catalog (each product different fields)
- **Forbes**: Content management system
- **Codecademy**: User progress tracking

### When to Avoid:

- ✗ Need multi-document transactions → PostgreSQL better
- ✗ Complex JOINs required → PostgreSQL
- ✗ Write-heavy (> 50K QPS) → Cassandra better

## When to Use Redis (Key-Value Store)

### Use When:

- ✓ Need caching layer (< 1ms latency)
- ✓ Session storage with TTL
- ✓ Real-time leaderboards
- ✓ Rate limiting counters
- ✓ Pub/Sub messaging
- ✓ Sorted timelines (Twitter/Instagram feed)

### Real Examples:

- **Twitter**: Timeline cache, rate limiting
- **Instagram**: Feed cache, session storage

- **Pinterest:** Board cache, follower lists
- **StackOverflow:** View counts, vote caching

### When to Avoid:

- ✗ Primary data store (not durable enough)
- ✗ Data > 512 GB per instance (need clustering)
- ✗ Complex queries → Use Elasticsearch
- ✗ Need ACID transactions → Use PostgreSQL

## Database Decision Tree

START: What are your requirements?

Data Structure?

- Structured + Relationships → SQL
  - Need transactions? YES → PostgreSQL/MySQL
  - Write QPS < 10K? YES → PostgreSQL/MySQL
  - Write QPS > 10K? → Consider Cassandra
- Flexible Schema → MongoDB
  - Need transactions? → MongoDB 4.0+
  - Write-heavy (> 50K QPS)? → Cassandra
- Time-Series + Write-Heavy → Cassandra
  - Logs, Events, Metrics → Cassandra
  - Need strong consistency? → PostgreSQL with partitioning
- Simple Key-Value → Redis (cache) or DynamoDB (primary)

Caching Layer?

- ALWAYS add Redis for read-heavy workloads

## Caching Decisions

### When to Add Caching

#### Add Cache When:

- ✓ Read-to-write ratio > 10:1
- ✓ Database CPU > 50%
- ✓ Query latency > 100ms
- ✓ Same data queried repeatedly
- ✓ Data doesn't change frequently

## **Don't Add Cache When:**

- Write-heavy workload
- Data changes constantly
- Each query unique (no cache hits)
- Database handling load fine
- Premature optimization

## **Real Decision Example:**

Scenario: Social media feed

Metrics:

- 100K read QPS, 1K write QPS (100:1 ratio)
- Same users refresh feed multiple times
- Database CPU at 80%

Decision:  ADD CACHE (Redis)

- Will handle 80K QPS (80% hit ratio)
- Reduce database to 20K QPS
- Database CPU drops to 16%

Cost: \$750/month cache vs \$12K/month more databases

## Cache-Aside vs Write-Through vs Write-Behind

### **Use Cache-Aside When:**

- Read-heavy (100:1 ratio)
- Data updated infrequently
- OK with slightly stale data (seconds)
- Most common pattern

**Example:** User profiles, product catalog, blog posts

### **Use Write-Through When:**

- Need strong consistency (cache = database)
- Can't serve stale data
- Banking, inventory, booking systems
- Critical data that must be accurate

---

**Example:** Bank balance, seat reservation, stock count

---

### Use Write-Behind When:

- Write-heavy workload
- OK to lose small amount of data
- Non-critical updates (likes, views, counters)
- Need sub-millisecond write latency

---

**Example:** Instagram likes, YouTube views, page analytics

---

### Interview Decision:

Question: "What caching strategy would you use for Twitter?"

Answer:

"I'd use different strategies for different data:

User Profiles (Cache-Aside):

- Read-heavy, updates rare
- TTL: 1 hour
- 90% hit ratio expected

Tweet Likes (Write-Behind):

- Write-heavy (50K likes/second)
- Batch to database every 5 seconds
- Acceptable if count slightly off

Account Balance (Write-Through):

- Critical data, must be accurate
  - Can't show wrong balance
  - Worth the latency trade-off
- "

---

## Replication Decisions

### Master-Slave vs Master-Master

#### Use Master-Slave When:

- Read-heavy workload (90%+ reads)
- Writes can go to single master

- Write QPS < 10K
- Acceptable to have replication lag
- Simpler to manage

### Scaling Pattern:

1 Master + 1 Slave: Handle 20K read QPS  
 1 Master + 5 Slaves: Handle 100K read QPS  
 1 Master + 10 Slaves: Handle 200K read QPS

Write capacity: Still limited to master!

### Real Examples:

- **Instagram**: 1 master + 12 slaves per region
- **GitHub**: MySQL with read replicas
- **Reddit**: PostgreSQL master-slave

### When It Breaks:

- △ Write QPS > 10K → Master becomes bottleneck → Need sharding
- △ Master fails → Downtime during failover (30–60 seconds)
- △ Replication lag > 1 second → Users see stale data

### Use Master-Master When:

- Need write scalability (> 10K write QPS)
- High availability critical (zero downtime)
- Multi-region writes required
- Can handle conflict resolution

### Scaling Pattern:

2 Masters: Handle 20K write QPS total  
 4 Masters: Handle 40K write QPS total

But: Complex conflict resolution needed!

### Real Examples:

- **Global services**: Need writes in multiple regions

- **CockroachDB:** Distributed SQL with multi-master
- **Cassandra:** All nodes are masters (no single master)

### When to Avoid:

- ✗ Simple read-heavy workload → Master-slave simpler
- ✗ Can't handle conflicts → Master-slave safer
- ✗ Small scale (< 10K write QPS) → Master-slave sufficient

### Interview Decision:

Question: "Master-slave or master-master replication?"

Answer Template:

"I'd choose master-slave because:

Workload Analysis:

- Read-heavy (200:1 ratio)
- Write QPS: 2K (single master can handle)
- Read QPS: 400K (need replicas)

Decision:

- 1 Master for writes
- 10 Slaves for reads (40K QPS each)
- Simpler than master-master
- Replication lag < 500ms acceptable for this use case

Would consider master-master if:

- Write QPS > 10K
  - Need zero-downtime writes
  - Multi-region writes required
- "

## Sharding Decisions

### When to Shard

#### Shard When:

- ✓ Single database can't handle load:
  - Write QPS > 10K
  - Data size > 10 TB
  - Query latency degrading
- ✓ Can't add more read replicas:

- Already have 10+ replicas
- Replication lag increasing

- Clear sharding key exists:
  - user\_id, region, tenant\_id
- Worth the complexity:
  - Scale justifies effort

#### Don't Shard When:

- Can add read replicas instead
- Can add cache instead
- Data < 1 TB
- Premature optimization
- No clear sharding key

### Sharding Strategy Selection

#### Use Hash-Based Sharding When:

- Need even distribution
- No hot partitions expected
- Access pattern: Lookup by ID
- Don't need range queries

Example: Twitter shards by tweet\_id

- Even distribution across 64 shards
- No celebrity hot partition
- Trade-off: Can't efficiently query "all tweets by user"

#### Use Range-Based Sharding When:

- Need range queries
- Time-series data
- Natural ranges exist

Example: Logs sharded by timestamp

- Q1 2024 → Shard 1
- Q2 2024 → Shard 2
- Efficient for "get logs from March"

Avoid: User shards by ID range (hot partitions!)

## Use Geographic Sharding When:

- Users primarily in specific regions
- Data sovereignty requirements (GDPR)
- Low latency critical
- Regions have independent data

Example: Uber trips sharded by region

- US trips → US shard (50ms latency)
- EU trips → EU shard (50ms latency)
- vs 150ms cross-continent

Trade-off: Uneven distribution (US has most traffic)

## Interview Decision:

Question: "How would you shard Twitter?"

Answer:

"I'd use hash-based sharding by tweet\_id because:

Reasoning:

1. Write-heavy (500M tweets/day = 6K QPS)
2. Need even distribution (avoid celebrity problem)
3. Primary access: Get tweet by ID
4. Can sacrifice: Query all tweets by user

Sharding Plan:

- 64 shards
- Each handles  $6K/64 = 94$  write QPS
- Each stores 4 TB (manageable)

Trade-offs:

- Can't efficiently get user timeline (scatter-gather)
  - Solution: Pre-compute timelines in Redis cache
  - Benefits: Even load, no hot shards
- "

## Consistency Decisions

### Strong vs Eventual Consistency

#### Use Strong Consistency When:

- Financial transactions (bank balance, payments)
- Inventory management (can't oversell)
- Booking systems (seats, hotels, tickets)
- User authentication (password changes)
- Incorrect data has serious consequences

### Systems for Strong Consistency:

- PostgreSQL, MySQL (with single master)
- MongoDB (with majority writes)
- DynamoDB (optional, costs more)

### Real Examples:

- **Banking**: Account balance MUST be accurate
  - **Ticketmaster**: Can't sell same seat twice
  - **Airbnb**: Can't double-book property
  - **Amazon Checkout**: Inventory count accurate
- 

### Use Eventual Consistency When:

- Social media feeds (stale OK for seconds)
- View counts, like counts (approximate OK)
- Product recommendations
- Search results
- Analytics dashboards
- High availability more important than accuracy

### Systems for Eventual Consistency:

- Cassandra (default)
- DynamoDB (default)
- DNS
- CDN caches

### Real Examples:

- **Twitter**: Feed shows tweets within 1-2 seconds
  - **Instagram**: Like count can be slightly off
  - **YouTube**: View count updates every few seconds
  - **Facebook**: News feed eventual consistency
- 

### Interview Decision:

Question: "Design a banking system"

Answer:

"I'd choose strong consistency because:

Requirements:

- Account balance must be accurate
- Can't show \$100 when actually \$0
- Better to show error than wrong balance

System Choice:

- PostgreSQL with ACID transactions
- Master-slave replication
- Read your own writes from master

Trade-offs:

- Slower than eventual consistency
- Less available during partitions
- Acceptable because correctness > speed

vs. Social Feed:

- Eventual consistency acceptable
- Availability > consistency
- Cassandra with AP model

"

## CDN Decisions

When to Use CDN

**Use CDN When:**

- ✓ Serving media files (images, videos, documents)
- ✓ Static assets (CSS, JS, fonts)
- ✓ Global user base
- ✓ High bandwidth costs
- ✓ Origin server overwhelmed

**Benefits Calculation:**

Without CDN:

- Latency: 150ms (cross-continent)
- Bandwidth: All traffic from origin
- Cost: \$0.09/GB egress from AWS

With CDN:

- Latency: 10ms (local edge)
- Bandwidth: 95% from CDN edge
- Cost: \$0.02/GB from CloudFront
- Savings: 78% bandwidth cost + 15x faster

#### Don't Use CDN When:

- ✗ All users in single region (CDN adds complexity)
- ✗ Dynamic content (personalized for each user)
- ✗ Content changes every second
- ✗ Small scale (< 10K users)

#### Push vs Pull CDN

##### Use Push CDN When:

- ✓ Content updates infrequently
- ✓ Small number of files
- ✓ Know what will be popular
- ✓ Want immediate global availability

**Example:** Marketing website, documentation site, product catalog

##### Use Pull CDN When:

- ✓ User-generated content (millions of files)
- ✓ Don't know what will be popular
- ✓ Content updates frequently
- ✓ Large catalog

**Example:** Instagram photos, YouTube videos, Twitter images

#### Interview Decision:

Question: "Design Instagram photo delivery"

Answer:

"I'd use Pull CDN (Akamai/CloudFront) because:

Requirements:

- 95M photos uploaded daily
- Don't know which will be popular
- User-generated content

#### Strategy:

- Store in S3 (origin)
- CDN pulls on first request
- Popular photos cached at edge
- Unpopular photos only in S3

#### Benefits:

- 95% cache hit ratio
- Origin bandwidth reduced 20x
- Latency: 150ms → 10ms globally

#### Trade-off:

- First viewer in region sees 150ms
- Acceptable for user-generated content
- "

## Kafka vs RabbitMQ vs SQS

### When to Use Kafka

#### Use When:

- High throughput (> 100K messages/second)
- Need message replay (debugging, new features)
- Event sourcing architecture
- Multiple consumers need same data
- Order matters (per partition)
- Long retention needed (days/weeks)

#### Real Examples:

- **Uber**: Ride events, surge pricing calculations
- **LinkedIn**: Activity streams, news feed generation
- **Netflix**: Viewing events, recommendation engine
- **Twitter**: Tweet fanout, analytics pipeline

### When to Use RabbitMQ

#### Use When:

- Need complex routing (topic exchanges, headers)
- Task queues with worker pools

- Lower throughput (< 50K msg/sec)
- Need message acknowledgments
- Priority queues

### Real Examples:

- **Instagram:** Background job processing
  - **Reddit:** Comment processing, moderation queue
  - **Task scheduling:** Celery with RabbitMQ backend
- 

### When to Use Amazon SQS

#### Use When:

- Want fully managed (no ops overhead)
- AWS-based architecture
- Don't need ordering
- Don't need replay
- Simple use case

### Real Examples:

- **Serverless applications:** Lambda triggers
  - **Microservices:** Decoupling in AWS
  - **Background jobs:** Image processing, email sending
- 

#### Decision Matrix:

Requirement	Choose
High throughput	→ Kafka
Need replay	→ Kafka
Event sourcing	→ Kafka
Simple task queue	→ RabbitMQ or SQS
AWS serverless	→ SQS
Complex routing	→ RabbitMQ
Managed service	→ SQS

## Communication Decisions

### REST vs GraphQL vs gRPC

#### Use REST When:

- ✓ Public API (most developers know REST)
- ✓ CRUD operations (standard HTTP verbs)
- ✓ Caching important (HTTP caching)
- ✓ Simple requirements
- ✓ Resource-oriented data

**Example:** Twitter API, GitHub API, Stripe API (all REST)

---

#### Use GraphQL When:

- ✓ Client needs flexibility (mobile apps)
- ✓ Multiple resources per request
- ✓ Avoid over-fetching/under-fetching
- ✓ Rapidly changing requirements
- ✓ BFF pattern (Backend for Frontend)

**Example:** Facebook, GitHub (alongside REST), Shopify

---

#### Use gRPC When:

- ✓ Internal microservice communication
- ✓ Need high performance (binary protocol)
- ✓ Bi-directional streaming
- ✓ Strong typing required
- ✓ Polyglot environment (multiple languages)

**Example:** Google's internal services, Uber's microservices

---

#### WebSocket vs Polling

##### Use WebSocket When:

- ✓ Real-time bi-directional communication
- ✓ Chat applications
- ✓ Live notifications
- ✓ Collaborative editing
- ✓ Gaming
- ✓ Stock tickers

**Example:** Slack messages, WhatsApp Web, Google Docs

---

**Scaling Challenge:** Each connection consumes memory

1M concurrent WebSocket connections:

- 1 MB per connection (buffers)
- Total: 1 TB memory needed
- Solution: Multiple WebSocket servers, sticky routing

**Use Polling (Long or Short) When:**

- WebSocket not supported (legacy systems)
- Firewall restrictions
- Simple occasional updates
- Don't need sub-second latency

**Example:** Legacy chat systems, simple notifications

**Efficiency:**

Short Polling: 3,600 requests/hour per user

Long Polling: 60 requests/hour per user

WebSocket: 1 connection for entire session

## Scaling Decisions by User Count

0 - 10,000 Users

**Architecture:**

[Users] → [Single Server] → [PostgreSQL]

Keep it simple!

**Decisions:**

- No load balancer yet
- No caching (database fast enough)
- No replication (can tolerate downtime)
- Vertical scaling only
- Managed database (RDS)

**Cost:** ~\$200/month

---

10,000 - 100,000 Users

**Architecture:**

```
[Users] → [Single Server (upgraded)] → [RDS PostgreSQL]
                                         → [Redis Cache]
```

**Decisions:**

- Vertical scaling (bigger server)
- Add Redis cache (80% hit ratio)
- Use managed database (RDS)
- No load balancer yet (single server OK)
- No replication yet

**Cost:** ~\$800/month

---

100,000 - 1M Users

**Architecture:**

```
[Users] → [Load Balancer]
          ↓
          [10 App Servers]
          ↓
          [Redis Cluster]
          ↓
          [PostgreSQL Master + 3 Slaves]
```

**Decisions:**

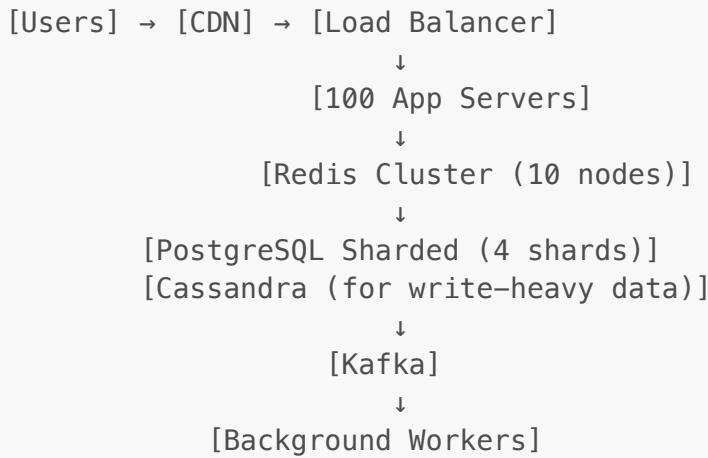
- Horizontal scaling (10 servers)
- Load balancer (Layer 7)
- Redis cluster (distributed cache)
- Database replication (master-slave)
- CDN for media files
- No sharding yet

**Cost:** ~\$5,000/month

---

1M - 10M Users

**Architecture:**



#### Decisions:

- Database sharding (PostgreSQL can't handle writes)
- Add Cassandra for write-heavy data
- Kafka for async processing
- Microservices architecture
- Multi-region deployment starts

**Cost:** ~\$50,000/month

---

10M - 100M Users

#### Architecture:

Multi-Region:

Region 1 (50% traffic):

[CDN] → [LB] → [500 Servers] → [Cache] → [DB Shards]

Region 2 (30% traffic):

[CDN] → [LB] → [300 Servers] → [Cache] → [DB Shards]

Region 3 (20% traffic):

[CDN] → [LB] → [200 Servers] → [Cache] → [DB Shards]

#### Decisions:

- Multi-region for global users
- Extensive sharding (32+ shards)
- Multiple NoSQL databases (Cassandra, DynamoDB)
- Kafka clusters for event streaming
- Separate analytics pipeline

**Cost:** ~\$500,000/month

---

---

## Decision Trees

### Decision Tree: Choose Database

START: What's your data?

Structured + Relationships?

- └ YES → Need Transactions?
  - └ YES → PostgreSQL/MySQL
  - └ NO → MongoDB (if schema flexible)

└ NO → What's your workload?

- └ Write-Heavy (> 10K QPS) → Cassandra
- └ Read-Heavy → MongoDB or PostgreSQL + Cache
- └ Key-Value only → Redis or DynamoDB
- └ Time-Series → Cassandra or InfluxDB

Special Needs:

- └ Full-Text Search? → Add Elasticsearch
- └ Graph Relationships? → Add Neo4j
- └ File Storage? → S3 or equivalent
- └ Analytics? → Redshift or BigQuery

---

### Decision Tree: Choose Caching Strategy

START: Do you need cache?

Read-to-Write Ratio?

- └ > 10:1 → YES, add cache
  - └ What data?
    - └ User profiles → Cache-Aside, TTL 1hr
    - └ Social feed → Cache-Aside, TTL 5min
    - └ Product catalog → Cache-Aside, TTL 1day
    - └ Session data → Write-Through, TTL 30min
- └ < 10:1 → Maybe not needed
  - └ Database handling load?
    - └ YES → Skip cache
    - └ NO → Add cache-aside
- └ Write-Heavy → Write-Behind
  - └ Examples: View counts, like counts

---

### Decision Tree: Choose Load Balancer Type

START: What's your architecture?

Microservices?

- |- YES → Layer 7
  - |- Route by URL path
- |- NO → Monolith
  - |- What protocol?
    - |- HTTP only → Layer 7 (for SSL termination)
    - |- Mixed protocols → Layer 4

Throughput?

- |- > 1M QPS → Layer 4 (faster)
- |- < 1M QPS → Layer 7 (more features)

Budget?

- |- High → Managed LB (AWS ALB, Google CLB)
- |- Low → Self-managed (NGINX, HAProxy)

## Quick Decision Reference

For Every System Design Question

### 1. Choose Database:

User Data, Transactions → PostgreSQL  
Posts, Feed (write-heavy) → Cassandra  
Search → Elasticsearch  
Cache → Redis  
Files → S3  
Analytics → Redshift

### 2. Add Caching:

Read-heavy (> 10:1) → Always add Redis  
TTLs:

- User profiles: 1 hour
- Feeds: 5 minutes
- Static data: 24 hours

### 3. Choose Replication:

Read-heavy → Master-Slave (add replicas)  
Write-heavy → Shard or Cassandra

Both heavy → Cassandra (masterless)

#### 4. Add CDN:

Media files → Always use CDN  
Global users → CDN mandatory  
Cost > \$1000/mo bandwidth → CDN pays for itself

#### 5. Message Queue:

Need async → Always add queue (Kafka/RabbitMQ)  
High throughput → Kafka  
Simple tasks → RabbitMQ or SQS

## Complete Interview Examples

Example 1: "Design Twitter" - Decision Walkthrough

### Step 1: Clarify Requirements

Scale: 400M DAU, 500M tweets/day  
Read-heavy: 200:1 ratio  
Latency: < 200ms for feed  
Availability: 99.99%

### Step 2: Database Decisions

User Data → PostgreSQL

Why:  
– Structured (user profiles, relationships)  
– Need ACID for account operations  
– Moderate write QPS (user updates)

Scaling:  
– Master-slave replication  
– 10 read replicas  
– Shard by user\_id when > 100M users

Tweets → Cassandra

Why:

- Write-heavy (6K QPS)
- Time-series (sorted by time)
- No ACID needed
- Need linear scalability

Scaling:

- 64 shards (Cassandra nodes)
- Each handles 94 write QPS
- Geographic distribution

### Step 3: Caching Decision

#### Timeline Cache → Redis Sorted Sets

Why:

- Read QPS: 463K (database can't handle)
- Perfect for sorted timeline
- 80% hit ratio expected

Configuration:

- Redis Cluster (20 nodes)
- 2 TB total cache
- TTL: 2 days

### Step 4: CDN Decision

#### Media → CloudFront + S3

Why:

- 40M images per day
- Global users
- Reduce origin bandwidth

Benefits:

- 95% cache hit ratio
- Latency: 150ms → 10ms
- Bandwidth cost: 20x cheaper

### Step 5: Async Processing

#### Feed Fanout → Kafka

Why:

- Don't block tweet posting
- Multiple consumers (notifications, search, analytics)

- Handle traffic spikes
- Can replay events

Architecture:

- Producer: Tweet Service
- Consumers: Feed Fanout, Notifications, Search Indexer

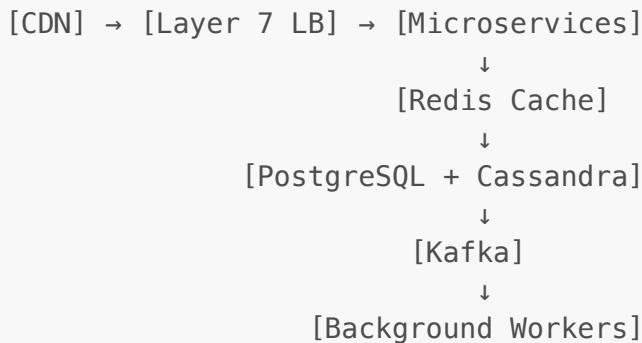
## Step 6: Load Balancer

### Layer 7 ALB

Why:

- Microservices (route by path)
- SSL termination
- Geographic routing
- Health checks

## Final Architecture Summary:




---

## Example 2: "Design Uber" - Decision Walkthrough

### Step 1: Key Decisions

#### Real-Time Locations → Redis Geo

Why:

- Need < 50ms latency
- Geo-spatial queries (find drivers within 5km)
- High write volume (location updates every 3 seconds)
- Small dataset (active drivers fit in memory)

Alternative considered:

- PostgreSQL with PostGIS → Too slow for real-time
- Cassandra → No geo-spatial queries
- Redis Geo → Perfect fit

## Trip History → Cassandra

Why:

- Time-series data
- Write-heavy (50M trips/day)
- Long-term storage (years)
- Geographic distribution

Alternative considered:

- PostgreSQL → Can't handle write volume
- MongoDB → Cassandra better for time-series

## Payments → PostgreSQL

Why:

- ACID transactions mandatory (can't lose money!)
- Strong consistency required
- Financial data (audit trail)

No alternatives acceptable for payments!

Must use SQL with ACID

## Surge Pricing → Kafka + Redis

Why Kafka:

- Real-time event streaming
- Multiple consumers (riders, drivers, analytics)
- High throughput

Why Redis:

- Fast aggregations (count requests per region)
- Real-time counters
- Sub-millisecond latency

## Example 3: "Design Instagram" - Decision Walkthrough

### Photo Storage → S3 + CloudFront CDN

Why S3:

- Durable (11 9's)
- Infinite scalability
- Cost-effective (\$0.023/GB)
- 95M photos/day = 19 TB/day

Why CDN:

- Global users
- 95% cache hit ratio
- Latency: 150ms → 10ms
- Bandwidth: 20x cheaper

Alternative considered:

- Store on app servers → Not scalable
- EBS volumes → Expensive, not durable
- S3 + CDN → Clear winner

## Feed Generation → Hybrid Fan-out

Decision: Hybrid (Push + Pull)

Normal users (< 10K followers):

- Fan-out on write (push to followers)
- Pre-compute in Redis
- Fast reads

Celebrities (> 10K followers):

- Fan-out on read (pull on demand)
- Avoid 10M writes per tweet
- Merge with pre-computed feed

Why hybrid:

- Solves celebrity problem
- Fast for everyone
- Optimal resource usage

## Photo Metadata → PostgreSQL + Cassandra

PostgreSQL for:

- User profiles
- Relationships

Cassandra for:

- Photo metadata (write-heavy)
- Activity feed
- Comments/likes

Why both:

- Use right tool for each job
- PostgreSQL: Relationships need SQL
- Cassandra: Feed needs write scalability

# Common Interview Questions

Q: "Your database is slow, what do you do?"

**Decision Process:**

## Step 1: Identify Bottleneck

Check metrics:

- CPU > 80% → Add cache or replica
- Disk I/O maxed → Add SSD or cache
- Memory low → Increase RAM or cache
- Connections maxed → Connection pooling

## Step 2: Quick Wins

1. Add indexes (10–100x speedup)
2. Add Redis cache (5–10x fewer DB queries)
3. Query optimization (avoid N+1)
4. Connection pooling

## Step 3: Architectural Changes

Read-heavy:

- Add read replicas (scale reads)

Write-heavy:

- Add cache (write-behind for non-critical)
- Consider Cassandra

Both heavy:

- Shard database
- Or move to Cassandra

Q: "How do you handle 1M concurrent WebSocket connections?"

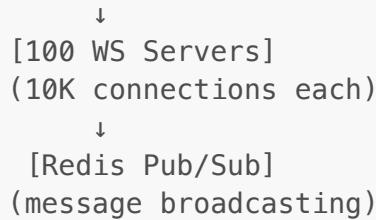
**Decision:**

Problem:  $1\text{M connections} \times 1\text{ MB} = 1\text{ TB memory}$

Solution: Distribute across servers

Architecture:

[Users] → [LB with Sticky Sessions]



**Why sticky:**

- WebSocket is stateful
- Must route to same server
- Use IP hash or session cookie

**Scaling:**

- Add servers → 10K connections each
- 100 servers → 1M connections

**Q: "When would you use Cassandra over PostgreSQL?"**

**Decision Framework:**

Choose Cassandra when:

1. Write throughput > 10K QPS  
Example: Twitter (6K QPS average, 60K peak)
2. Time-series data  
Example: Netflix viewing history, IoT sensor data
3. Geographic distribution  
Example: Global app needs low latency everywhere
4. High availability critical  
Example: Can't tolerate single point of failure
5. Data size > 100 TB  
Example: Instagram activity feed

Choose PostgreSQL when:

1. Need ACID transactions
2. Complex JOINs required
3. Data < 10 TB
4. Write QPS < 10K
5. Strong consistency required

**Q: "How do you prevent database from being overwhelmed?"**

**Multi-Layer Defense:**

## **Layer 1: Rate Limiting**

Prevent abuse at API Gateway

- 1000 requests/hour per user
- Protects all downstream systems

## **Layer 2: Caching**

Redis cache intercepts 80% of reads

- Database only sees 20%
- 5x load reduction

## **Layer 3: Connection Pooling**

Limit concurrent DB connections

- Pool of 50 connections
- Prevents connection exhaustion
- Graceful degradation under load

## **Layer 4: Read Replicas**

Distribute reads across slaves

- 10 replicas → 10x read capacity
- Master only handles writes

## **Layer 5: Sharding (Last Resort)**

When all else fails, shard

- 64 shards → 64x capacity
- But adds complexity

## **Trade-Off Analysis**

Every Decision Has Trade-Offs

### **Caching Trade-offs:**

Benefits:

- 10–100x faster reads

- Reduce database load 80%+
- Better user experience

Costs:

- Stale data possible
- Cache invalidation complexity
- Additional infrastructure
- More memory needed

### Sharding Trade-offs:

Benefits:

- Scale writes linearly
- Scale storage linearly
- No single database limit

Costs:

- Can't JOIN across shards
- Complex queries difficult
- Rebalancing expensive
- Application complexity

### CDN Trade-offs:

Benefits:

- 10–15x faster (global)
- 20x cheaper bandwidth
- Reduce origin load 95%

Costs:

- Cache invalidation delays
- Additional cost
- Complexity
- First request slow (miss)

### Microservices Trade-offs:

Benefits:

- Independent scaling
- Team autonomy
- Technology flexibility
- Fault isolation

Costs:

- Network latency between services
- Distributed debugging hard

- ✖ Operational complexity
- ✖ Data consistency challenges

## The Universal Decision Framework

For ANY Component Decision

### 1. REQUIREMENTS

What do we need?

- Scale (QPS, storage)
- Latency targets
- Consistency requirements
- Availability requirements

### 2. OPTIONS

What are the choices?

- List 2–3 alternatives
- Know what each offers

### 3. TRADE-OFFS

What do we sacrifice?

- Cost
- Complexity
- Performance
- Consistency

### 4. DECISION

What do we choose?

- Pick based on requirements
- Justify with trade-offs
- Mention alternatives considered

### 5. SCALE PLAN

How does this scale?

- What happens at 10x growth?
- When do we need to change?

## Key Principles

Start Simple, Scale When Needed

### Phase 1: MVP (< 100K users)

Keep it simple:

- Single server
- Managed database (RDS)

- No caching
- No load balancer

Why: Fast to build, cheap, good enough

## Phase 2: Growth (100K - 1M users)

Add essentials:

- Load balancer
- Horizontal scaling
- Redis cache
- Database replicas
- CDN for media

Why: Needed for scale, proven patterns

## Phase 3: Scale (1M - 10M users)

Add advanced:

- Database sharding
- Kafka for events
- Microservices
- Multi-region

Why: Single region/database can't handle load

## Phase 4: Hyper-Scale (> 10M users)

Optimize everything:

- Multiple NoSQL databases
- Extensive caching
- Global CDN
- Custom infrastructure

Why: At this scale, custom solutions pay off

## Default Technology Choices

**For 90% of interviews, default to:**

Load Balancer: Layer 7 (ALB, NGINX)

Database: PostgreSQL (users), Cassandra (events)

Cache: Redis

CDN: CloudFront / Akamai  
Message Queue: Kafka  
Search: Elasticsearch  
File Storage: S3  
Analytics: Redshift

These are safe, proven choices!

### Only deviate when you can justify:

- "I'd use X instead of Y because [specific reason]"
  - Always mention the trade-off
- 

## Interview Scoring

### What Interviewers Look For

#### In Decision-Making:

##### Strong Candidate :

- States assumptions clearly
- Considers multiple options
- Explains trade-offs
- Justifies with requirements
- Mentions when to change approach
- Uses real company examples

##### Weak Candidate :

- "I'd use MongoDB" (no justification)
  - Doesn't mention alternatives
  - No trade-off discussion
  - Over-engineers (Kafka for 100 users)
  - Under-engineers (PostgreSQL for 1B writes/day)
- 

## Final Checklist

### Before Finishing Any Design

#### Ask Yourself:

##### Scalability: How does this scale to 10x? 100x?

- "At 10x traffic, we'd add sharding"
-

- "At 100x, we'd move to Cassandra"

**Availability:** How do we handle failures?

- "Replication across 3 zones"
- "Automatic failover in 30 seconds"

**Consistency:** Strong or eventual?

- "Strong for payments (PostgreSQL)"
- "Eventual for feed (Cassandra)"

**Trade-offs:** What did we sacrifice?

- "Chose availability over consistency for feed"
- "Sharding sacrifices JOINs for write scalability"

**Bottlenecks:** What will break first?

- "Database will bottleneck at 10K write QPS"
- "Solution: Shard or move to Cassandra"

**Cost:** Is this cost-effective?

- "CDN saves \$10K/month in bandwidth"
- "Cache reduces database costs 80%"

## Summary: The Golden Rules

1. **Start with SQL** (PostgreSQL) unless you have specific reason not to
2. **Add Redis cache** for any read-heavy workload (> 10:1)
3. **Use CDN** for any media files and global users
4. **Layer 7 LB** for microservices, Layer 4 for performance
5. **Master-Slave** for read-heavy, shard for write-heavy
6. **Strong consistency** for money, eventual for social
7. **Kafka** for high throughput async, RabbitMQ for tasks
8. **Cassandra** for write-heavy time-series, PostgreSQL otherwise
9. **Always mention trade-offs** - nothing is free
10. **Scale incrementally** - don't over-engineer early

**Document Version:** 1.0

**Last Updated:** January 8, 2025

**For:** System Design HLD Interview Preparation

**Status:** Complete - Pure Decision-Making Guide

**Remember:** It's not about knowing every technology - it's about making the RIGHT choice for the SPECIFIC requirements!