

WhatsApp Messenger System Design - High-Level Design (HLD)

Table of Contents

1. Problem Statement
 2. Functional Requirements
 3. Non-Functional Requirements
 4. Capacity Estimation
 5. High-Level Architecture
 6. Core Components
 7. Database Design
 8. Message Delivery
 9. Deep Dives
 10. Scalability & Reliability
 11. Trade-offs & Alternatives
-

Problem Statement

Design a real-time messaging platform like WhatsApp that enables:

- One-on-one text messaging with delivery guarantees
- Group chat (up to 256 participants)
- Media sharing (images, videos, documents)
- Voice and video calls (out of scope for detailed design)
- Online/offline status indicators
- Read receipts (blue ticks)
- Message encryption (end-to-end)
- Cross-platform support (mobile, web, desktop)

Scale Requirements

- **2 billion daily active users (DAU)**
 - **100 billion messages per day**
 - **500 million active users simultaneously**
 - **4.5 billion group messages per day**
 - **Average message size: 100 bytes**
 - **Latency requirement: < 100ms for message delivery**
 - **99.99% availability**
-

Functional Requirements

Must Have (P0)

1. User Management

- User registration with phone number
- Profile management (name, status, profile picture)
- Contact list synchronization
- Block/unblock users

2. Messaging

- Send text messages (1:1)
- Receive messages in real-time
- Message delivery status:
 - Sent (single tick) ✓
 - Delivered (double tick) ✓✓
 - Read (blue tick) ✓✓ (blue)
- Message history (retrieve old messages)
- Delete messages (delete for me, delete for everyone)
- Forward messages
- Reply to specific messages (quoted)

3. Group Chat

- Create groups (up to 256 members)
- Add/remove participants
- Admin roles (only admins can change settings)
- Group name, description, icon
- See participant list
- Group delivery receipts (show who read)
- Leave group

4. Media Sharing

- Share images (JPEG, PNG, max 16MB)
- Share videos (MP4, max 16MB for sending)
- Share documents (PDF, DOC, max 100MB)
- Share voice messages
- Share location
- Share contacts

5. Status & Presence

- Online/offline status
- Last seen timestamp
- Typing indicators
- "Recording audio" indicator

6. Calls (Basic mention)

- Voice calls (1:1)
- Video calls (1:1, group)
- Call history

Nice to Have (P1)

- Stories/Status updates (24-hour expiry)
 - Message reactions (emoji)
 - Voice messages
 - Live location sharing
 - Disappearing messages
 - Archived chats
 - Starred messages
 - Message search
 - Two-factor authentication
 - Multi-device support
-

Non-Functional Requirements

Performance

- **Message delivery latency:** < 100ms (p95)
- **Message send latency:** < 50ms
- **Connection establishment:** < 1 second
- **Image upload:** < 3 seconds (5MB image)
- **Conversation load:** < 200ms

Scalability

- Support 2 billion DAU
- Handle 100 billion messages/day (~1.15M messages/second)
- Support 500M concurrent connections
- Handle message spikes during peak hours (New Year, events)
- Scale to 10x current load

Availability

- **99.99% uptime** (< 53 minutes downtime/year)
- Multi-region deployment
- No message loss (durability guarantee)
- Offline message queueing

Reliability

- **Message delivery guarantee:** At-least-once delivery
- **Message ordering:** Maintain order per conversation
- **No duplicate messages:** Idempotent delivery

- **Persistent storage:** Messages stored forever (unless deleted)

Security

- **End-to-end encryption** (E2EE) for all messages
- **Perfect forward secrecy** (compromised key doesn't decrypt past messages)
- **Server-side encryption** at rest
- **TLS** for all connections
- **2FA** for account access

Consistency

- **Eventual consistency** for last seen, online status
- **Strong consistency** for messages (no message loss)
- **Causal ordering** for conversation messages

Capacity Estimation

Traffic Estimates

```
Daily Active Users (DAU): 2B  
Messages per user per day: 50  
Total messages/day:  $2B \times 50 = 100B$ 
```

```
Messages/second (average):  $100B / 86,400 \approx 1.15M$  QPS  
Messages/second (peak - 3x):  $3.5M$  QPS
```

```
Group messages/day: 4.5B (10% have media)  
1:1 messages/day: 95.5B
```

```
Concurrent WebSocket connections: 500M
```

Storage Estimates

Text Messages:

```
Average message size: 100 bytes (text + metadata)  
Daily text messages: 100B  
Daily storage:  $100B \times 100 \text{ bytes} = 10 \text{ TB/day}$ 
```

```
Yearly storage:  $10 \text{ TB} \times 365 = 3.65 \text{ PB/year}$   
Storage for 5 years: 18.25 PB (text only)
```

Media Messages:

Media messages/day: 10B (10% of total)

Average sizes:

- Images: 500 KB
- Videos: 5 MB
- Voice messages: 50 KB
- Documents: 2 MB

Assumption: 60% images, 20% videos, 15% voice, 5% docs

Weighted average: $(0.6 \times 500\text{KB}) + (0.2 \times 5\text{MB}) + (0.15 \times 50\text{KB}) + (0.05 \times 2\text{MB})$

$$= 300\text{KB} + 1\text{MB} + 7.5\text{KB} + 100\text{KB} \approx 1.4 \text{ MB per media message}$$

Daily media storage: $10\text{B} \times 1.4 \text{ MB} = 14 \text{ PB/day}$

Yearly: $14 \text{ PB} \times 365 = 5,110 \text{ PB} = 5.1 \text{ EB/year}$

With compression (50% reduction): 2.55 EB/year

Total Storage:

Year 1: 2.6 EB (2.55 media + 0.05 text)

Year 5: 13 EB

Note: WhatsApp stores messages temporarily server-side

Permanent storage is on client devices

Server retention: 30 days for undelivered messages

Bandwidth Estimates

Incoming (message sends):

$10 \text{ TB text} + 14 \text{ PB media} = 14.01 \text{ PB/day}$

$14.01 \text{ PB} / 86,400 \text{ seconds} \approx 162 \text{ GB/second}$

Outgoing (message deliveries – 1:1 each message):

Similar to incoming: $\sim 162 \text{ GB/second}$

Peak (3x average): $\sim 486 \text{ GB/second}$

Memory Requirements (Caching)

Online user presence:

$500\text{M concurrent} \times 100 \text{ bytes} = 50 \text{ GB}$

Recent conversations (per user, last 50 messages):

$500\text{M users} \times 50 \text{ messages} \times 100 \text{ bytes} = 2.5 \text{ TB}$

Undelivered messages queue (24 hour buffer):
Assume 10% users offline at any time
 $200\text{M offline users} \times 100 \text{ messages} \times 100 \text{ bytes} = 2 \text{ TB}$

Connection metadata:
500M connections × 500 bytes = 250 GB

Total memory needed: ~5 TB
Distributed across 200+ Redis nodes = 25 GB per node

Server Estimates


WebSocket connections per server: 65K (C10K problem limit)
Servers needed for 500M connections: $500M / 65K \approx 8,000$ servers


With redundancy and headroom (2x): 16,000 servers
Organized into clusters by region


Cost estimation (AWS):

- WebSocket servers: $16K \times \$100/\text{month} = \$1.6\text{M}/\text{month}$
 - Database servers: $\$500\text{K}/\text{month}$
 - Storage (S3): $13 \text{ EB} \times \$0.023/\text{GB} = \$300\text{K}/\text{month}$
 - Bandwidth: $\$200\text{K}/\text{month}$
 - Total: $\sim \$2.6\text{M}/\text{month}$

High-Level Architecture







Key Architectural Decisions

1. WebSocket for Real-time Communication

- Persistent bidirectional connection
- Low latency (< 100ms)
- Push capability (server → client)

2. Cassandra for Message Storage

- Write-heavy workload (1.15M messages/sec)
- Time-series data (sorted by timestamp)
- Linear scalability
- No single point of failure

3. Redis for Presence & Session

- In-memory speed for online status
- Sub-millisecond lookups
- TTL support for automatic cleanup
- Pub/Sub for real-time updates

4. Kafka for Message Pipeline

- High throughput (1M+ messages/sec)

- Durability (no message loss)
- Replay capability (disaster recovery)
- Decouples producers and consumers

5. S3 for Media Storage

- Virtually unlimited storage
- 99.99999999% durability (11 nines)
- Integrated with CDN
- Lifecycle policies (auto-archival)

Core Components


1. WebSocket Gateway

Purpose: Maintain persistent connections with clients

Responsibilities:

- Establish WebSocket connections
- Handle heartbeats (keep-alive every 30 seconds)
- Route messages to/from clients
- Maintain connection metadata (user_id, device_id, connection_id)
- Handle reconnection logic
- Graceful degradation to HTTP long-polling

Architecture:



Connection Distribution:

User connects → Load Balancer → Available WebSocket Gateway

Gateway Selection:

- Geographic proximity (latency)
- Current connection count (load balancing)
- Server health status

Connection sticky: Yes (same user → same gateway if possible)

Benefit: Reduced state synchronization

Scaling:

- Each server: 65,000 concurrent connections (C10K problem solved)
- Total servers: 500M / 65K ≈ 8,000 servers
- Organized in clusters per region

State Management:

Problem: WebSocket is stateful, but we need to route messages

Solution: Redis-based session store

Key: "session:{user_id}"

Value: {

```
"gateway_server": "ws-gateway-42.us-east",
"connection_id": "conn_abc123",
"device_id": "device_xyz789",
"connected_at": 1704672000,
"last_heartbeat": 1704673000
```

}

TTL: 70 seconds (auto-cleanup if connection dies)

2. Message Service

Purpose: Core business logic for message handling

Responsibilities:

- Validate message content (length, format)
- Generate unique message IDs (Snowflake algorithm)
- Store messages in Cassandra
- Publish to Kafka for async processing
- Handle message deletion
- Manage message metadata

Message ID Generation (Snowflake):

64-bit Message ID structure:

1 bit unused	41 bits timestamp (ms)	10 bits server ID	12 bits sequence (4096/ms)
-----------------	------------------------------	-------------------------	----------------------------------

Benefits:

- Globally unique
- Sortable by time
- No coordination needed
- 4096 IDs per millisecond per server

Message Processing Pipeline:

1. Receive message from sender
2. Validate (length, format, permissions)
3. Generate message_id
4. Check if recipient online (Redis presence)
5. If online:
 - Push to recipient's gateway (via Redis Pub/Sub)
 - Update delivery status → SENT
6. If offline:
 - Queue in Redis undelivered messages
 - Send push notification
7. Store in Cassandra (async via Kafka)
8. Return ACK to sender with message_id

Total Latency Breakdown:

Validation: 1ms
Generate ID: 0.1ms
Presence check (Redis): 1ms
Push to recipient: 10–50ms (WebSocket)
Store in Cassandra (async): 0ms (non-blocking)
ACK to sender: 2ms

Total: ~15–55ms (well under 100ms target)

3. User Service

Purpose: Manage user accounts and profiles

Responsibilities:

- User registration (phone number verification)

- Authentication (JWT tokens)
- Profile management (name, status, picture)
- Contact list synchronization
- Block/unblock management
- Privacy settings

Registration Flow:

1. User enters phone number (+1234567890)
2. System sends 6-digit OTP via SMS (Twilio/SNS)
3. User enters OTP
4. System validates OTP (5-minute expiry)
5. Create user account:
 - Generate user_id (UUID)
 - Hash phone number for privacy
 - Create initial profile
6. Generate JWT tokens (access + refresh)
7. Return to client with tokens

Contact Sync:

Problem: User has 1000 contacts, need to check who's on WhatsApp

Naive approach: Query DB for each contact (1000 queries) X

Optimized approach (Bloom Filter):

1. Client sends hash of all phone numbers
2. Server checks against Bloom filter (probabilistic)
3. Possible matches: Query database (100 queries)
4. Return list of users on WhatsApp
5. Cache result on client (refresh daily)

Benefits:

- Reduces queries by 90%
- Fast membership testing
- Privacy-preserving (hashes, not plain numbers)

4. Group Service

Purpose: Manage group chats and memberships

Group Creation:

1. User creates group → Group Service
2. Generate group_id (UUID)
3. Add creator as admin
4. Invite initial members (send invites)

5. Store group metadata in Cassandra:
 - group_id, name, description, icon_url
 - created_at, created_by, admins[]
6. Create group_members table entry for each member
7. Return group_id to creator

Group Message Fanout:

Challenge: Send message to 256 group members efficiently

Solution: Async fanout via Kafka

Flow:

1. User sends message to group
2. Store message once in Cassandra
3. Publish to Kafka: "group.message.sent"
4. Fanout worker:
 - Fetch all group members (cached)
 - For each member:
 - a. Check if online (Redis)
 - b. If online: Push to WebSocket
 - c. If offline: Queue message + push notification
 - Update delivery receipts
5. Total time: ~100-200ms for 256 members

Group Metadata Caching:

Cache group members (Redis):

Key: "group:{group_id}:members"

Type: Set

Members: {user_id_1, user_id_2, ...}

TTL: 1 hour

Benefits:

- Fast member lookup ($O(1)$)
- Avoid DB query for every message
- Invalidate on member add/remove

5. Presence Service

Purpose: Track online/offline status of users

Online Status Flow:

User opens app:

1. Establish WebSocket connection

2. Send "presence.online" event
3. Update Redis: SET user:{user_id}:online 1 EX 70
4. Broadcast to user's contacts (Pub/Sub)

User closes app / disconnects:

1. Connection drops (detected by gateway)
2. Update Redis: DEL user:{user_id}:online
3. Set last_seen: SET user:{user_id}:last_seen {timestamp}
4. Broadcast offline status

Heartbeat (every 30 seconds):

1. Client sends ping
2. Server updates TTL: EXPIRE user:{user_id}:online 70
3. Server sends pong

Last Seen Timestamp:

Storage: Redis + MySQL

Redis (hot data):

Key: "user:{user_id}:last_seen"
 Value: Unix timestamp
 TTL: 7 days

MySQL (permanent):

UPDATE users SET last_seen_at = NOW() WHERE user_id = ?
 (Batched updates every 5 minutes to reduce DB load)

Privacy Controls:

User settings:

- Everyone: Show last seen to all
- My contacts: Show only to contacts
- Nobody: Don't show last seen

Implementation:

Before returning last_seen, check:

1. Requester's relationship with user
2. User's privacy setting
3. Return last_seen or "hidden" accordingly

6. Media Service

Purpose: Handle media uploads and processing

Upload Flow:

1. Client requests upload URL:
 POST /api/v1/media/upload/init
 Body: {"type": "image", "size": 5242880}
2. Server validates:
 - Size within limits
 - User has storage quota
 - Rate limit not exceeded
3. Generate pre-signed S3 URL:
 URL expires in 15 minutes
 Unique key: media/{user_id}/{timestamp}/{uuid}.jpg
4. Return upload URL to client
5. Client uploads directly to S3 (parallel, chunked)
6. Client confirms upload:
 POST /api/v1/media/upload/complete
 Body: {"media_id": "abc123", "s3_key": "media/..."}
7. Async processing (Lambda / Worker):
 - Generate thumbnails (150px, 400px)
 - Compress original (WebP for images, H.264 for video)
 - Extract metadata (dimensions, duration)
 - Virus scan
 - Store processed versions
8. Update media metadata in database
9. Media now ready to attach to messages

Media Storage Organization:

S3 Bucket Structure:

```

whatsapp-media-{region}/
  └── images/
    ├── original/{user_id}/{yyyy-mm}/{uuid}.jpg
    ├── large/{user_id}/{yyyy-mm}/{uuid}_800.webp
    └── thumb/{user_id}/{yyyy-mm}/{uuid}_150.webp
  └── videos/
    ├── original/{user_id}/{yyyy-mm}/{uuid}.mp4
    └── compressed/{user_id}/{yyyy-mm}/{uuid}_720p.mp4
  └── documents/
    └── {user_id}/{yyyy-mm}/{uuid}.{ext}
  └── voice/
    └── {user_id}/{yyyy-mm}/{uuid}.opus

```

Lifecycle Policy:

- Hot tier (SSD): 0-30 days

- Standard tier: 31–90 days
- Glacier: > 90 days (if user doesn't access)

CDN Integration:

CloudFront in front of S3:

- Cache popular media (24 hour TTL)
- Serve from nearest edge location
- ~90% cache hit ratio

Origin: S3 bucket

Edge Locations: 400+ globally

Latency: ~20ms (edge) vs 200ms (origin)

7. Notification Service

Purpose: Send push notifications when user offline

Notification Types:

1. New message (priority: HIGH)
2. Missed call (priority: HIGH)
3. Group mention (priority: MEDIUM)
4. Group message (priority: LOW if muted)

Push Notification Flow:

1. Message arrives, user offline
2. Check user's device tokens (cached in Redis)
3. Build notification payload:


```
{
        "title": "Alice",
        "body": "Hey, how are you?",
        "message_id": "msg_123",
        "conversation_id": "conv_456",
        "priority": "HIGH",
        "badge_count": 5 // Unread messages
      }
```
4. Send to push service:
 - iOS → APNs (Apple Push Notification Service)
 - Android → FCM (Firebase Cloud Messaging)
5. Handle delivery status:
 - Success: Log
 - Failure (token expired): Remove from cache
 - Retry: 3 attempts with exponential backoff

Optimization: Batching:

Don't send notification for every message in group chat

Strategy:

- Buffer messages for 3 seconds
- If multiple messages in group: Send one notification
"3 new messages in Family Group"
- Prevents notification spam

8. Delivery Worker

Purpose: Ensure message reaches recipient(s)

Delivery Guarantee: At-least-once delivery

Flow for 1:1 Message:

1. Message stored in Cassandra
2. Check recipient online status (Redis)
3. If online:
 - Get recipient's WebSocket gateway from Redis
 - Push message directly via that gateway
 - Wait for ACK from recipient
 - Update status: DELIVERED
4. If offline:
 - Store in Redis undelivered queue
 - Send push notification
 - Status remains: SENT
5. When recipient comes online:
 - Retrieve undelivered messages from queue
 - Deliver in order
 - Update status: DELIVERED
6. When recipient opens chat:
 - Send READ receipt
 - Update status: READ

Flow for Group Message:

1. Message stored once in Cassandra
2. Publish to Kafka: "group.message.new"
3. Fanout worker consumes event:
 - Fetch all 256 group members (cached)
 - For each member (parallel processing):
 - a. Skip sender
 - b. Check online status
 - c. If online: Push via WebSocket

- d. If offline: Queue + push notification
 - Track delivery: {member_1: DELIVERED, member_2: SENT, ...}
- 4. Update group message status
- 5. Show sender: "Delivered to 200/255"

Idempotency:

Problem: Network issues may cause duplicate deliveries

Solution: Message ID + De-duplication

Client side:

- Track received message IDs (last 1000)
- Ignore duplicates based on message_id
- Send ACK for both original and duplicate

Server side:

- Use message_id as idempotency key
- Cassandra: PRIMARY KEY handles duplicates
- Redis: SETNX for atomic operations

9. Analytics Service

Purpose: Track usage metrics and generate insights

Metrics Collected:

User Metrics:

- DAU, WAU, MAU
- Message sent per user
- Active time per session
- Retention rate (D1, D7, D30)


Message Metrics:

- Messages per second
- Average message length
- Media vs text ratio
- Group vs 1:1 ratio

Performance Metrics:

- Message delivery latency (p50, p95, p99)
- Connection success rate
- Heartbeat response time
- Error rates by type

Data Pipeline:



Database Design

1. Message Database (Cassandra)

Why Cassandra?

- Write-optimized (1.15M writes/second)
- Time-series data (messages sorted by time)
- Linear scalability (add nodes for more throughput)
- Multi-datacenter replication
- No single point of failure
- Tunable consistency

Schema Design:

Messages Table (1:1 conversations):

```

CREATE TABLE messages (
    conversation_id uuid,
    message_id timeuuid,
    sender_id uuid,
    receiver_id uuid,
    content text,
    message_type text, -- TEXT, IMAGE, VIDEO, DOCUMENT, VOICE
    media_url text,
    created_at timestamp,
    status text, -- SENT, DELIVERED, READ
    is_deleted boolean DEFAULT false,
    PRIMARY KEY ((conversation_id), message_id)
) WITH CLUSTERING ORDER BY (message_id DESC);

-- Index for user's all conversations
CREATE TABLE user_messages (
    user_id uuid,
    conversation_id uuid,
    message_id timeuuid,
    created_at timestamp,

```

```
PRIMARY KEY ((user_id), created_at, message_id)
) WITH CLUSTERING ORDER BY (created_at DESC);
```

Group Messages Table:

```
CREATE TABLE group_messages (
    group_id uuid,
    message_id timeuuid,
    sender_id uuid,
    content text,
    message_type text,
    media_url text,
    created_at timestamp,
    PRIMARY KEY ((group_id), message_id)
) WITH CLUSTERING ORDER BY (message_id DESC);

-- Delivery receipts for group messages
CREATE TABLE group_message_receipts (
    group_id uuid,
    message_id timeuuid,
    user_id uuid,
    status text, -- DELIVERED, READ
    timestamp timestamp,
    PRIMARY KEY ((group_id, message_id), user_id)
);
```

Partitioning Strategy:

Partition by conversation_id for 1:1 messages
Partition by group_id for group messages

Benefits:

- All messages in conversation co-located
- Efficient range queries (get last N messages)
- Natural hot-spot distribution

Replication Factor: 3 (across datacenters)
Consistency Level: QUORUM (2 of 3 must acknowledge)

2. User Database (MySQL)

Why MySQL?

- User data requires ACID transactions
- Complex queries (contact matching)
- Relatively small dataset (2B users × 2KB ≈ 4TB)
- Strong consistency for user profiles

Schema:

```
CREATE TABLE users (
    user_id CHAR(36) PRIMARY KEY,
    phone_number_hash CHAR(64) UNIQUE NOT NULL, -- SHA-256 hash
    username VARCHAR(50),
    display_name VARCHAR(100),
    status_message VARCHAR(139), -- Like Twitter
    profile_picture_url VARCHAR(500),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    last_seen_at TIMESTAMP,
    is_online BOOLEAN DEFAULT FALSE,
    privacy_last_seen ENUM('everyone', 'contacts', 'nobody') DEFAULT
    'everyone',
    privacy_profile_pic ENUM('everyone', 'contacts', 'nobody') DEFAULT
    'everyone',
    privacy_status ENUM('everyone', 'contacts', 'nobody') DEFAULT
    'everyone',
    INDEX idx_phone_hash (phone_number_hash),
    INDEX idx_last_seen (last_seen_at)
);

CREATE TABLE contacts (
    user_id CHAR(36),
    contact_user_id CHAR(36),
    contact_name VARCHAR(100), -- User's local name for contact
    added_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    PRIMARY KEY (user_id, contact_user_id),
    INDEX idx_contact (contact_user_id)
);

CREATE TABLE blocked_users (
    user_id CHAR(36),
    blocked_user_id CHAR(36),
    blocked_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    PRIMARY KEY (user_id, blocked_user_id)
);

CREATE TABLE device_tokens (
    user_id CHAR(36),
    device_id VARCHAR(100),
    platform ENUM('ios', 'android', 'web'),
    push_token VARCHAR(500), -- APNs/FCM token
    last_active_at TIMESTAMP,
    PRIMARY KEY (user_id, device_id)
);
```

Sharding Strategy:

Shard by user_id using consistent hashing

Number of shards: 100 (can scale to 1000)
Shard calculation: hash(user_id) % 100

Example:

```
user_id: "550e8400-e29b-41d4-a716-446655440000"
hash: 12345
shard: 12345 % 100 = 45
→ Database shard 45
```

3. Group Database (Cassandra)

Schema:

```
CREATE TABLE groups (
    group_id uuid PRIMARY KEY,
    name text,
    description text,
    icon_url text,
    created_by uuid,
    created_at timestamp,
    max_members int DEFAULT 256,
    member_count int
);

CREATE TABLE group_members (
    group_id uuid,
    user_id uuid,
    role text, -- ADMIN, MEMBER
    joined_at timestamp,
    added_by uuid,
    PRIMARY KEY ((group_id), user_id)
);

-- Index for user's groups
CREATE TABLE user_groups (
    user_id uuid,
    group_id uuid,
    joined_at timestamp,
    PRIMARY KEY ((user_id), joined_at, group_id)
) WITH CLUSTERING ORDER BY (joined_at DESC);
```

4. Conversation Metadata (Redis)

Purpose: Fast access to conversation data

Data Structures:

Recent Conversations (per user):

Key: "conversations:{user_id}"

Type: Sorted Set (ZSET)

Score: Last message timestamp

Value: conversation_id

```
ZADD conversations:user123 1704672000 conv-abc
```

```
ZREVRANGE conversations:user123 0 49 # Last 50 conversations
```

TTL: 30 days

Unread Count:

Key: "unread:{user_id}:{conversation_id}"

Type: Integer

Value: Count of unread messages

```
INCR unread:user123:conv-abc
```

```
GET unread:user123:conv-abc → 5
```

TTL: 30 days

Typing Indicator:

Key: "typing:{conversation_id}:{user_id}"

Type: String

Value: 1 (is typing)

```
SETEX typing:conv-abc:user123 5 1 # 5 second TTL
```

```
GET typing:conv-abc:user123
```

Message Delivery

End-to-End Message Flow

Scenario: Alice sends "Hello" to Bob

PHASE 1: Message Send (Alice's side)

T=0ms: Alice types "Hello" and hits send

T=1ms: Alice's app encrypts message (E2EE)

- Generate ephemeral key pair
- Encrypt content with Bob's public key
- Add digital signature

T=5ms: Send via WebSocket to gateway

Message payload: {

```
  "message_id": "temp_client_id_xyz",
  "receiver_id": "bob_uuid",
```

```
        "encrypted_content": "base64_encrypted_data",
        "message_type": "TEXT",
        "timestamp": 1704672000000
    }
```

T=10ms: Gateway forwards to Message Service

PHASE 2: Server Processing

T=11ms: Message Service processing

- Validate message format
- Check Alice not blocked by Bob
- Generate server message_id (Snowflake)
- Replace temp_client_id with server message_id

T=12ms: Check Bob's online status (Redis)

Result: Bob is ONLINE

T=15ms: Lookup Bob's WebSocket gateway (Redis)

Result: ws-gateway-15.us-west

T=18ms: Publish to Redis Pub/Sub channel

Channel: "messages:ws-gateway-15"

Message: {message_id, receiver_id: bob, ...}

T=20ms: Send ACK to Alice

ACK: {
 "temp_id": "temp_client_id_xyz",
 "message_id": "msg_server_id_abc",
 "status": "SENT" ✓
}

Alice's app: Show single tick ✓

T=22ms: Async Kafka publish (non-blocking)

Topic: "message.created"

Purpose: Persistence, analytics, audit

PHASE 3: Delivery to Bob

T=23ms: ws-gateway-15 receives message (Redis Pub/Sub)

T=25ms: Gateway pushes to Bob's WebSocket connection

T=30ms: Bob's app receives encrypted message

T=32ms: Bob's app decrypts message

- Use Bob's private key
- Verify Alice's signature

– Decrypt content → "Hello"

T=35ms: Bob's app sends DELIVERED ACK to gateway

T=37ms: Gateway forwards ACK to Message Service

T=40ms: Message Service updates status:

- Update Redis: message:msg_abc:status = DELIVERED
- Publish to Kafka: "message.delivered"

T=45ms: Gateway pushes status update to Alice

Alice's app: Show double tick √√

T=50ms: Bob opens chat, reads message

T=52ms: Bob's app sends READ receipt

T=55ms: Message Service updates status: READ

T=60ms: Alice receives READ receipt

Alice's app: Ticks turn blue √√ (blue)

PHASE 4: Persistence (Async, parallel with above)

T=25ms: Kafka consumer picks up "message.created" event

T=30ms: Write to Cassandra:

```
INSERT INTO messages (
    conversation_id, message_id, sender_id,
    receiver_id, content, created_at, status
) VALUES (...);
```

T=50ms: Cassandra write complete (3 replicas)

T=100ms: All async processes complete

TOTAL END-TO-END LATENCY: 60ms (Alice sees blue ticks)

Offline Message Handling

Scenario: Alice sends message, Bob is offline

Flow:

1. Alice sends message (T=0)
2. Server checks Bob status: OFFLINE
3. Server actions:

- a. Store in Cassandra
 - b. Queue in Redis: LPUSH undelivered:bob msg_id
 - c. Send push notification to Bob
 - d. Return SENT status to Alice (single tick ✓)
4. Bob comes online (T=3600s, 1 hour later)
 5. Connection established
 6. Server retrieves undelivered messages:
messages = LRANGE undelivered:bob 0 -1
→ [msg_abc, msg_def, msg_xyz]
7. Deliver messages in order:
For each message:
 - Fetch from Cassandra
 - Push to Bob's WebSocket
 - Wait for ACK
 - Update status: DELIVERED
 - Remove from queue: LREM undelivered:bob
8. Alice receives DELIVERED updates (ticks turn double ✓✓)
 9. Bob opens chat with Alice
 10. Bob's app sends READ receipts
 11. Alice's ticks turn blue ✓✓

Message Ordering Guarantee

Problem: Messages may arrive out of order due to:

- Network delays
- Retries
- Multiple paths

Solution: Sequence numbers + Client-side reordering

Server side:

- Use timeuuid for message_id (sortable by time)
- Messages naturally ordered in Cassandra

Client side:

- Buffer messages temporarily
- Sort by message_id (time-sortable)
- Display in correct order
- Handle gaps (missing sequence numbers)

Example:

```
Receive: msg_5, msg_3, msg_7, msg_4
Buffer and sort: msg_3, msg_4, msg_5, msg_7
Gap detected: msg_6 missing
Request msg_6 from server
Display complete sequence
```

Deep Dives

1. WebSocket Connection Management

Connection Lifecycle:

1. CONNECTING
 - TCP handshake
 - TLS negotiation
 - WebSocket upgrade

Duration: ~200–500ms

↓

2. AUTHENTICATING
 - Send JWT token
 - Server validates
 - Load user session

Duration: ~100ms

↓

3. CONNECTED
 - Register connection in Redis
 - Start heartbeat (every 30s)
 - Update presence: ONLINE
 - Deliver queued messages

Duration: Indefinite

↓

4. DISCONNECTED
 - Heartbeat timeout (60s) OR
 - Clean disconnect signal
 - Remove from Redis
 - Update presence: OFFLINE
 - Record last_seen

Heartbeat Mechanism:

Purpose: Detect dead connections quickly

Client → Server (every 30 seconds):
{
 "type": "PING",

```
        "timestamp": 1704672000
    }

Server → Client (immediate response):
{
    "type": "PONG",
    "timestamp": 1704672001,
    "server_time": 1704672001 // For time sync
}
```

Timeout handling:

- If no PING for 60 seconds: Mark connection dead
- Clean up Redis session
- Update presence to OFFLINE

Reconnection Strategy:

Network interruption detected:

Client side:

1. Detect connection drop (no PONG received)
2. Attempt reconnect with exponential backoff:
 - Try 1: Immediately
 - Try 2: After 1 second
 - Try 3: After 2 seconds
 - Try 4: After 4 seconds
 - Max backoff: 30 seconds
3. If reconnect succeeds:
 - Re-authenticate
 - Sync messages (get missed messages)
 - Resume normal operation

Server side:

1. Old connection cleanup (automatic via TTL)
2. New connection accepted
3. Return missed messages since last_ack

2. End-to-End Encryption (E2EE)

Signal Protocol (Used by WhatsApp):

Key Components:

1. Identity Key Pair:
 - Long-term key pair per user
 - Public key stored on server
 - Private key never leaves device

2. Pre-Keys:
 - One-time keys for initial handshake
 - User uploads 100 pre-keys to server
 - Server distributes to initiators

3. Session Keys:
 - Ephemeral keys for each session
 - Rotated regularly (every 1000 messages)
 - Provides perfect forward secrecy

Initial Handshake (Alice → Bob, first message):

1. Alice's device:
 - Fetch Bob's identity key (from server)
 - Fetch one of Bob's pre-keys (from server)
 - Derive shared secret using ECDH
 - Generate session key

2. Encrypt message:
 - Use session key + AES-256-GCM
 - Include metadata for Bob to derive same key

3. Send encrypted message + handshake data

4. Bob's device:
 - Receive message
 - Use handshake data + own private keys
 - Derive same shared secret
 - Decrypt message

5. Both devices now have session key
6. Subsequent messages use session key directly

Perfect Forward Secrecy:

Mechanism: Rotate keys frequently

After every 1000 messages or 1 week:

- Generate new ephemeral key
- Derive new session key
- Delete old session key

Result:

If attacker compromises current key,
they can't decrypt past messages
(old keys already deleted)

Server's Role in E2EE:

Server CANNOT read message content!

Server can see:

- ✓ Sender user_id
- ✓ Receiver user_id
- ✓ Timestamp
- ✓ Message type (text/image)
- ✓ Encrypted payload (can't decrypt)

Server cannot see:

- ✗ Message content
- ✗ Media content (encrypted before upload)
- ✗ Who user talks to (metadata encrypted too in latest version)

3. Group Chat Architecture

Group Limits:

WhatsApp Groups:

- Max members: 256 (1024 in newer versions)
- Max admins: Unlimited
- Messages: No limit

Technical constraints:

- 256 members = 256 delivery operations
- Must complete in < 200ms
- Need efficient fanout mechanism

Group Message Fanout Strategy:

Challenge: Send 1 message to 256 recipients efficiently

Solution: Parallel async fanout

Architecture:

Group Message Arrives




Store Once in Cassandra
(group_id + message_id as composite key)



```
Publish to Kafka  
Topic: "group.message.fanout"  
Partition by: group_id (keeps order)
```



```
Fanout Workers (20+ instances)  
Consume in parallel
```



```
Online Members  
(Push via WS)      Offline Members  
(Queue + Push)
```

Processing:

1. Fetch 256 members from cache (Redis SET)
2. Split into online (150) and offline (106)
3. Parallel delivery:
 - Online: Push to 150 WebSockets (parallel)
 - Offline: Queue for 106 users + push notifications
4. Track delivery: Store receipt per member
5. Update sender: "Delivered to 150/256"

Total time: ~100–150ms

Group Admin Features:

Admin Operations:

- Add/remove members
- Change group name/description/icon
- Promote/demote admins
- Delete messages (for everyone)
- Change group settings (who can message, edit info)

Implementation:

1. Check if user is admin (cached role)
2. If yes: Perform operation
3. If no: Reject with 403 Forbidden
4. Broadcast change to all members
5. Update cache

4. Multi-Device Support

Challenge: User has phone, tablet, desktop - sync messages across all

Solution: Linked Devices Architecture

Primary Device (Phone):

- Stores encryption keys
- Can add/remove linked devices
- Acts as "master" device

Linked Devices (Tablet, Desktop):

- Paired with primary device
- Get encryption keys from primary
- Receive messages independently
- Send messages through primary device bridge

Pairing Flow:

1. User opens WhatsApp Web on desktop
2. Desktop generates QR code containing:
 - Device public key
 - Random session token
3. User scans QR code with phone
4. Phone:
 - Sends device public key to server
 - Establishes secure channel with desktop
 - Shares encryption keys
5. Server:
 - Registers new device for user
 - Creates device_id
 - Links to user account
6. Desktop now receives all messages
7. Phone can revoke desktop access anytime

Message Sync:

When message arrives:

1. Server delivers to ALL online devices
2. Each device independently:
 - Receives encrypted message
 - Decrypts using shared keys
 - Displays message

When primary device receives message:

1. Decrypt on phone
2. Re-encrypt for each linked device
3. Server forwards to linked devices

Result: All devices stay in sync

5. Message Search

Challenge: Search through billions of messages per user

Local Search (Preferred):

Storage: SQLite database on device

- Messages stored locally
- Full-text search index
- Fast, private (no server query)

Query: `SELECT * FROM messages
WHERE conversation_id = ?
AND content LIKE '%keyword%'
ORDER BY created_at DESC
LIMIT 50;`

Benefits:

- ✓ Fast (local database)
- ✓ Private (search term not sent to server)
- ✓ Works offline
- ✓ No server load

Server-Side Search (Fallback):

Use case: Search across devices, or old messages not on device

Architecture:

Messages → Kafka → Indexer → Elasticsearch

Elasticsearch Index:

```
{  
  "message_id": "msg_abc",  
  "user_id": "user_123", // Who can search  
  "conversation_id": "conv_xyz",  
  "content_encrypted": "...", // Still encrypted!  
  "created_at": "2024-01-08T10:00:00Z",  
  "message_type": "text"  
}
```

Note: Content encrypted even in search index!

Search by metadata only (sender, date, type)
Full-text search requires client-side decryption

6. Read Receipts (Blue Ticks)

Privacy Design:

User can control read receipts:

- ON: Others see when you read (default)
- OFF: Others don't see, but you also can't see theirs

Implementation:

- Setting stored in user profile
- Enforced bidirectionally
- If Alice disables: Bob doesn't see Alice's reads, and Alice doesn't see Bob's reads

Group Read Receipts:

Show who read the message in group

Display:

- ✓✓ Delivered to 200/256
- ✓✓ (blue) Read by 150/256

Tap to see: List of members who read

- Alice (read 2 mins ago)
- Bob (read 5 mins ago)
- Charlie (delivered, not read)
- ...

Storage:

Cassandra group_message_receipts table
Track per-member delivery status

7. Typing Indicators

Requirements:

- Show "Alice is typing..." in real-time
- Don't spam (rate limit updates)
- Ephemeral (don't store permanently)

Implementation:

Client side:

User starts typing:

- After 1 second of typing: Send "typing_start"
- While typing: Send "typing_continue" every 3 seconds
- Stop typing: Send "typing_stop" OR wait 5 seconds

Server side:

Receive "typing_start":

1. Store in Redis: SETEX typing:conv_abc:user123 5 1
2. Broadcast to conversation participants:
 - 1:1: Send to other person
 - Group: Send to all members (max 256)
3. Auto-expire after 5 seconds (TTL)

Receive "typing_stop":

1. Delete from Redis: DEL typing:conv_abc:user123
2. Broadcast "stopped" to participants

Recipient side:

- Show indicator if Redis key exists
- Hide automatically after 5 seconds (TTL)

Optimization: Debouncing

Don't send on every keystroke

Client logic:

- Wait 1 second after first keystroke
- Then send "typing_start"
- Send "typing_continue" every 3 seconds
- Send "typing_stop" on submit or 5s inactivity

Result: Reduce network traffic by 95%

8. Message Deletion

Two Types:

Delete for Me:

- Delete locally on device
- Mark as deleted in local DB
- Server still has message
- Other participants unaffected

Flow:

1. User selects "Delete for me"

2. App marks message as deleted locally
3. Hide from UI
4. Optional: Tell server to not sync to new devices

Delete for Everyone:

- Delete for all conversation participants
- Time limit: 1 hour after sending (configurable)
- Leaves tombstone: "This message was deleted"

Flow:

1. User selects "Delete for everyone"
2. Check time limit (< 1 hour since sent)
3. Server updates message:
 - Set `is_deleted = TRUE`
 - Keep `message_id` (for ordering)
 - Clear content
4. Server broadcasts deletion event:
 - To all online participants: Delete immediately
 - To offline participants: Deliver deletion on next connect
5. All devices show: "This message was deleted"

Storage:

- Keep `message_id` to maintain conversation flow
- Clear content to save space
- Mark with tombstone

Scalability & Reliability

Horizontal Scaling

Stateless Services (Easy to scale):

- Message Service
- User Service
- Group Service
- API Gateway

Auto-Scaling Rules:

Scale Up Triggers:

- └ CPU > 70% for 5 minutes
- └ Memory > 80% for 5 minutes
- └ Request queue > 1000
- └ WebSocket connections > 80% capacity

Scale Down Triggers:

- └─ CPU < 30% for 15 minutes
- └─ Request queue < 100
- └─ Min instances: 100 per region (always)

Scaling speed:

- Scale up: +50% capacity in 2 minutes
- Scale down: -10% capacity in 10 minutes (gradual)

Stateful Services (Complex scaling):

WebSocket Gateways:

Challenge: Can't just kill connections

Graceful scaling:

Scale up: Simple, new servers accept new connections

Scale down:

1. Mark server for removal
2. Stop accepting new connections
3. Wait for existing connections to disconnect naturally
4. Or drain connections:
 - Send reconnect signal to clients
 - Clients reconnect to other gateways
5. After all connections drained: Remove server

Drain time: 5–10 minutes

Database Scaling

Cassandra Scaling:

Horizontal scaling (add nodes):

Current: 100 nodes

Each node: 2TB storage, 10K writes/sec

Total capacity: 200TB, 1M writes/sec

To double capacity:

Add 100 nodes

Cassandra automatically rebalances data

Time: 2–4 hours for rebalancing

No downtime during scaling

MySQL Scaling:

Vertical scaling (larger instances):

- Upgrade from db.r5.4xlarge → db.r5.8xlarge
- Requires brief downtime (5–10 minutes)
- Plan during maintenance window

Horizontal scaling (read replicas):

- Add read replicas for read queries
- Write still goes to master
- Eventual consistency for reads (< 1 second lag)

Current setup:

- 1 master (writes)
- 5 read replicas per region (reads)
- Cross-region replicas for DR

Redis Scaling:

Redis Cluster mode:

- 16,384 hash slots
- Distributed across nodes
- Add nodes to increase capacity

Current setup:

- 200 nodes × 25GB = 5TB total memory
- 3 replicas per shard
- Handle 10M operations/sec

Multi-Region Architecture

Global Deployment:

REGION: US-EAST (Primary)

- WebSocket Gateways: 3000
- Cassandra Cluster: 30 nodes
- MySQL: Master + 5 replicas
- Redis: 50 nodes
- Serves: North America, South America

REGION: EU-WEST (Secondary)

- WebSocket Gateways: 2500
- Cassandra Cluster: 25 nodes
- MySQL: Replica
- Redis: 40 nodes
- Serves: Europe, Middle East, Africa

```
REGION: AP-SOUTH (Secondary)
└── WebSocket Gateways: 2500
└── Cassandra Cluster: 25 nodes
└── MySQL: Replica
└── Redis: 40 nodes
└── Serves: Asia-Pacific
```

Cross-Region Replication:

- Cassandra: Automatic multi-datacenter replication
- MySQL: Async replication (< 1 second lag)
- Redis: Independent per region (eventual consistency OK)

User Routing:

Strategy: Route to nearest region

User in India:

1. DNS resolves to ap-south load balancer
2. Connect to ap-south WebSocket gateway
3. Messages stored in global Cassandra
4. Low latency (<50ms in-region)

Cross-region messaging:

Alice (US) → Bob (India):

1. Alice sends to us-east gateway
2. Stored in Cassandra (replicated to ap-south)
3. Bob's gateway (ap-south) reads from local Cassandra
4. Delivers to Bob

Total latency: ~100–150ms (acceptable)

High Availability

Failure Scenarios & Handling:

WebSocket Gateway Failure:

Detection: Health check failure (3 consecutive)

Impact: Users on that gateway disconnected

Mitigation:

1. Load balancer removes failed gateway

2. Clients detect disconnection
3. Clients reconnect to