

# Back-of-the-Envelope Calculations Guide for System Design Interviews

---

## Table of Contents

1. [Introduction](#)
  2. [Core Principles](#)
  3. [Essential Numbers to Memorize](#)
  4. [Step-by-Step Framework](#)
  5. [Traffic Estimation](#)
  6. [Storage Estimation](#)
  7. [Bandwidth Estimation](#)
  8. [Memory/Cache Estimation](#)
  9. [Real-World Examples](#)
  10. [Quick Reference Cheat Sheet](#)
- 

## Introduction

**"In system design interviews, being approximately right is better than being precisely wrong."**

Back-of-the-envelope calculations are rough estimates used to evaluate the feasibility of a system design. Interviewers want to see:

- Can you estimate scale and resource requirements?
- Do you understand trade-offs between different approaches?
- Can you identify potential bottlenecks?

**You don't need a calculator!** Simple multiplication and division with rough numbers is enough.

---

## Core Principles

### 1. Round Numbers Liberally

**Bad:** "We have 387,654,321 daily active users..."

**Good:** "We have ~400 million daily active users..."

**Why:** Easier to calculate, equally accurate for capacity planning

### 2. Use Powers of 10/2

**Instead of:** 1,234,567 bytes

**Use:** ~1 MB (1 million bytes)

**Instead of:** 987 requests

**Use:** ~1K requests (1,000)

### 3. State Your Assumptions Clearly

**Always say:**

- "Assuming 100 million daily active users..."
- "Let's assume average photo size is 200 KB..."
- "I'll estimate read-to-write ratio as 100:1..."

### 4. Work in Standard Units

**Storage:** Bytes → KB → MB → GB → TB → PB

**Time:** Seconds → Minutes → Hours → Days → Years

**Traffic:** QPS (Queries Per Second)

### 5. Show Your Work

**Don't say:** "We need 10 TB of storage"

**Do say:** "100M photos × 200 KB = 20 TB, minus 50% compression = 10 TB"

---

## Essential Numbers to Memorize

### Powers of Two (Storage)

$2^{10} = 1,024$	$\approx 1$ Thousand	= 1 KB
$2^{20} = 1,048,576$	$\approx 1$ Million	= 1 MB
$2^{30} = 1$ Billion	$\approx 1$ GB	
$2^{40} = 1$ Trillion	$\approx 1$ TB	
$2^{50} = 1$ Quadrillion	$\approx 1$ PB	

**Pro Tip:** Just remember "1024" and multiply!

- 1 KB = ~1,000 bytes
- 1 MB = ~1,000 KB = 1 million bytes
- 1 GB = ~1,000 MB = 1 billion bytes
- 1 TB = ~1,000 GB = 1 trillion bytes

### Time Conversions

1 minute = 60 seconds
1 hour = 3,600 seconds ( $60 \times 60$ )
1 day = 86,400 seconds ( $24 \times 3,600$ )
1 month = ~2.5 million seconds ( $30 \times 86,400$ )
1 year = ~31.5 million seconds ( $365 \times 86,400$ )

**Pro Tip:** Remember "86,400 seconds per day"

- Daily to QPS: Divide by 86,400 ( $\approx$  100K)
- Monthly to QPS: Divide by 2.5M

## Latency Numbers (Google's Famous List)

L1 cache reference:	0.5 ns
Branch mispredict:	5 ns
L2 cache reference:	7 ns
Main memory reference:	100 ns
Compress 1KB with Snappy:	10 $\mu$ s
Send 1 KB over 1 Gbps network:	10 $\mu$ s
Read 1 MB from memory:	250 $\mu$ s
Round trip within datacenter:	500 $\mu$ s
Read 1 MB from SSD:	1 ms
Disk seek:	10 ms
Read 1 MB from network:	10 ms
Read 1 MB from disk:	30 ms
Send packet CA→Netherlands→CA:	150 ms

### Takeaways:

- Memory is  $\sim$ 100x faster than SSD
- SSD is  $\sim$ 10x faster than HDD
- Network within datacenter is fast (0.5ms)
- Cross-continent is slow (150ms)

## Common Data Sizes

1 character = 1 byte (ASCII)  
 1 character = 2–4 bytes (Unicode/UTF-8)

User ID (integer): 8 bytes  
 Timestamp: 8 bytes  
 IP address (IPv4): 4 bytes  
 IP address (IPv6): 16 bytes

Tweet text (280 chars): ~280 bytes  
 Small image (thumbnail): ~50 KB  
 Regular image (compressed): ~200 KB  
 High-res image: ~2 MB  
 Short video (1 min): ~10 MB  
 HD video (1 min): ~50 MB

User profile: ~1–2 KB  
 Email: ~75 KB (average)  
 Web page: ~2 MB (with assets)

# Step-by-Step Framework

## Step 1: Clarify Scale Requirements

### Questions to Ask:

#### Users:

- How many daily active users (DAU)?
- How many monthly active users (MAU)?
- User growth rate?

#### Usage:

- How many posts/uploads per user per day?
- How many reads per user per day?
- Read-to-write ratio?

#### Data:

- Average size of user data?
- Average size of content (posts, photos, videos)?
- Data retention period?

## Step 2: Calculate Traffic

### Formula:

$$\text{QPS} = \text{Total Operations per Day} / 86,400 \text{ seconds}$$

#### Example:

$$\begin{aligned} & 100\text{M tweets per day} \\ & = 100,000,000 / 86,400 \\ & = 1,157 \text{ QPS (round to } \sim 1.2\text{K QPS)} \end{aligned}$$

$$\begin{aligned} & \text{Peak QPS} = \text{Average QPS} \times \text{Peak Factor} \\ & \text{Peak Factor typically: 2-5x} \\ & = 1.2\text{K} \times 3 = 3.6\text{K QPS (round to } \sim 4\text{K QPS)} \end{aligned}$$

## Step 3: Calculate Storage

### Formula:

$$\text{Total Storage} = \text{Number of Items} \times \text{Size per Item} \times \text{Time Period}$$

#### Example:

$$\begin{aligned} & 100\text{M photos per month} \times 200 \text{ KB} \times 12 \text{ months} \times 10 \text{ years} \\ & = 100\text{M} \times 200 \text{ KB} \times 120 \\ & = 2.4 \text{ PB} \end{aligned}$$

## Step 4: Calculate Bandwidth

**Formula:**

$$\text{Bandwidth} = \text{QPS} \times \text{Average Response Size}$$

Example:

Read QPS: 10K

Average response: 100 KB

$$\text{Bandwidth} = 10,000 \times 100 \text{ KB} = 1 \text{ GB/s}$$

## Step 5: Calculate Memory (Cache)

**Formula:**

$$\text{Cache Size} = \% \text{ of Data to Cache} \times \text{Data Size}$$

Example (80-20 rule):

Total data: 10 TB

$$\text{Cache 20\% hot data: } 10 \text{ TB} \times 0.2 = 2 \text{ TB}$$

$$\text{Distributed across 20 servers: } 2 \text{ TB} / 20 = 100 \text{ GB per server}$$

## Traffic Estimation

Calculation Method

**Given:** Daily Active Users (DAU) and actions per user

**Formula:**

$$\text{Total Daily Operations} = \text{DAU} \times \text{Operations per User}$$

$$\text{QPS} = \text{Total Daily Operations} / 86,400$$

$$\text{Peak QPS} = \text{QPS} \times \text{Peak Factor (2-5x)}$$

Example 1: Twitter

**Given:**

- 400M DAU
- Each user tweets 0.5 times per day (average)
- Each user reads 100 tweets per day (scrolling feed)

**Calculate:**

Write Traffic (Tweets):

- Tweets per day:  $400M \times 0.5 = 200M$
- Write QPS:  $200M / 86,400 \approx 2,315$  QPS
- Peak write QPS:  $2,315 \times 3 \approx 7K$  QPS

Read Traffic (Timeline):

- Reads per day:  $400M \times 100 = 40$  billion
- Read QPS:  $40B / 86,400 \approx 462K$  QPS
- Peak read QPS:  $462K \times 3 \approx 1.4M$  QPS

Read-to-Write Ratio:  $40B / 200M = 200:1$

**Interview Insight:** "This is read-heavy, so we need caching and read replicas"

Example 2: Instagram

**Given:**

- 500M DAU
- Each user uploads 0.1 photos per day (1 photo per 10 days)
- Each user views 50 photos per day

**Calculate:**

Write Traffic (Photos):

- Photos per day:  $500M \times 0.1 = 50M$
- Write QPS:  $50M / 86,400 \approx 578$  QPS
- Peak:  $\sim 1.7K$  QPS

Read Traffic (Feed):

- Views per day:  $500M \times 50 = 25$  billion
- Read QPS:  $25B / 86,400 \approx 289K$  QPS
- Peak:  $\sim 867K$  QPS

Read-to-Write Ratio:  $25B / 50M = 500:1$

**Interview Insight:** "Extremely read-heavy, CDN caching is critical"

Example 3: WhatsApp

**Given:**

- 2 billion MAU
- 50% daily active = 1B DAU
- Each user sends 40 messages per day
- Each user receives 40 messages per day

**Calculate:**

Messages per day:  $1B \times 40 = 40$  billion

Message writes:  $40B / 86,400 \approx 463K$  QPS

Peak:  $\sim 1.4M$  QPS

With reads (checking for new messages):

- Check every 10 seconds:  $1B \text{ users} \times 8,640 \text{ checks} = 8.64 \text{ trillion checks/day}$
- Check QPS:  $8.64T / 86,400 = 100M \text{ QPS}$  (too high!)
- Solution: WebSocket connections, push notifications

**Interview Insight:** "Need WebSocket for efficiency, not polling"

---

## Storage Estimation

Calculation Method

**Formula:**

$$\text{Storage} = \text{Number of Objects} \times \text{Size per Object} \times \text{Time Period}$$

Example 1: YouTube

**Given:**

- 500 hours of video uploaded per minute
- Average video size: 50 MB per minute
- Store for 10 years

**Calculate:**

Daily uploads:

- $500 \text{ hours/min} \times 60 \text{ min/hour} \times 24 \text{ hours} = 720,000 \text{ hours per day}$
- $720K \text{ hours} \times 60 \text{ min} \times 50 \text{ MB} = 2,160,000,000 \text{ MB per day}$
- $= 2,160 \text{ TB per day} = \sim 2 \text{ PB per day}$

Yearly:  $2 \text{ PB} \times 365 = 730 \text{ PB per year}$

10 years:  $730 \text{ PB} \times 10 = 7.3 \text{ exabytes (7,300 PB)}$

With multiple resolutions (360p, 720p, 1080p, 4K):

- 4 versions per video
- Total:  $7,300 \text{ PB} \times 4 = 29 \text{ exabytes}$

With compression and deduplication:

- Assume 30% savings
- Effective:  $29 \text{ EB} \times 0.7 = \sim 20 \text{ exabytes}$

**Interview Insight:** "This is why YouTube needs Google's infrastructure!"

### Example 2: Instagram

#### Given:

- 500M DAU
- 50M photos uploaded per day
- Average photo: 200 KB
- Store for 10 years

#### Calculate:

Daily storage:

$$\begin{aligned} - 50M \text{ photos} \times 200 \text{ KB} &= 10,000,000,000 \text{ KB} \\ - &= 10 \text{ TB per day} \end{aligned}$$

$$\text{Yearly: } 10 \text{ TB} \times 365 = 3.65 \text{ PB}$$

$$10 \text{ years: } 3.65 \text{ PB} \times 10 = 36.5 \text{ PB}$$

With thumbnails (3 sizes per photo):

- Original: 200 KB
- Large: 100 KB
- Thumbnail: 20 KB
- Total per photo: 320 KB

$$\text{Adjusted: } 50M \times 320 \text{ KB} = 16 \text{ TB per day}$$

$$10 \text{ years: } 16 \text{ TB} \times 365 \times 10 = 58 \text{ PB}$$

**Interview Insight:** "Need S3 or similar object storage, with lifecycle policies"

### Example 3: Twitter

#### Given:

- 400M DAU
- 200M tweets per day
- Average tweet: 280 characters = ~280 bytes
- Store for 5 years
- 20% of tweets have images (~500 KB each)

#### Calculate:

Text storage:

$$- 200M \text{ tweets} \times 280 \text{ bytes} = 56 \text{ GB per day}$$

Image storage:

- $200M \times 0.2 = 40M$  images per day
- $40M \times 500 \text{ KB} = 20 \text{ TB}$  per day

Total per day:  $56 \text{ GB} + 20 \text{ TB} \approx 20 \text{ TB}$   
 Yearly:  $20 \text{ TB} \times 365 = 7.3 \text{ PB}$   
 5 years:  $7.3 \text{ PB} \times 5 = 36.5 \text{ PB}$

With metadata (user\_id, timestamp, likes, etc.):

- Additional 100 bytes per tweet
- $200M \times 100 \text{ bytes} = 20 \text{ GB}$  per day
- Negligible compared to images

Total:  $\sim 37 \text{ PB}$  for 5 years

**Interview Insight:** "Text is cheap, media is expensive"

---

## Bandwidth Estimation

### Calculation Method

#### Formula:

$$\text{Bandwidth} = \text{QPS} \times \text{Average Object Size}$$

### Example 1: Instagram Feed

#### Given:

- 289K read QPS (from earlier calculation)
- Average feed response:  $50 \text{ photos} \times 200 \text{ KB thumbnails} = 10 \text{ MB}$

#### Calculate:

Incoming bandwidth (uploads):

$$- 578 \text{ QPS (writes)} \times 200 \text{ KB} = 115 \text{ MB/s}$$

Outgoing bandwidth (feed loads):

$$- 289K \text{ QPS} \times 10 \text{ MB} = 2,890 \text{ GB/s} = 2.8 \text{ TB/s}$$

Wait, that's too high! Let's optimize:

With CDN caching (95% cache hit ratio):

- Only 5% hits origin:  $2.8 \text{ TB/s} \times 0.05 = 140 \text{ GB/s}$
- CDN bandwidth:  $2.8 \text{ TB/s}$  (distributed globally)

Actual Instagram approach:

- Serve thumbnails first (150 KB instead of 200 KB)

- Lazy load full images
- Reduced to ~1 TB/s from CDN

**Interview Insight:** "CDN is essential, not optional"

Example 2: YouTube

**Given:**

- 1 billion video views per day
- Average video: 10 minutes at 720p = 100 MB

**Calculate:**

Views per second:  $1B / 86,400 \approx 11,574 \text{ QPS}$

Bandwidth:  $11,574 \times 100 \text{ MB} = 1,157 \text{ GB/s} = 1.1 \text{ TB/s}$

Peak (3x): 3.3 TB/s

With CDN (80% cache hit):

- Origin:  $3.3 \text{ TB/s} \times 0.2 = 660 \text{ GB/s}$
- CDN:  $3.3 \text{ TB/s} \times 0.8 = 2.6 \text{ TB/s}$

Optimization with adaptive bitrate:

- Serve lower quality for slow connections
- Average drops to 50 MB per view
- New bandwidth: 578 GB/s origin

## Memory/Cache Estimation

The 80-20 Rule (Pareto Principle)

**Key Insight:** 20% of data generates 80% of traffic

**Cache Strategy:** Cache the hot 20%, serve 80% of requests

Example 1: Twitter Timeline Cache

**Given:**

- 400M DAU
- Each timeline:  $50 \text{ tweets} \times 500 \text{ bytes} = 25 \text{ KB}$
- Cache 20% of active users

**Calculate:**

Users to cache:  $400M \times 0.2 = 80M$  users  
Cache size:  $80M \times 25\text{ KB} = 2,000\text{ GB} = 2\text{ TB}$

Distributed across servers:

- 20 cache servers  $\times$  100 GB RAM each = 2 TB total
- Cost: ~\$2,000/month (AWS ElastiCache)

Cache hit ratio achieved: ~80%

Database queries saved: 80% of 462K QPS = 370K QPS

**Interview Insight:** "Caching 20% of users saves 80% of database queries"

### Example 2: Netflix Thumbnails

**Given:**

- 10,000 videos in catalog
- Each thumbnail: 50 KB
- Cache most popular 1,000 videos (10%)

**Calculate:**

Cache size: 1,000 videos  $\times$  50 KB = 50 MB

With 50% cache hit ratio:

- 11,574 QPS  $\times$  0.5 = 5,787 QPS to origin
- Saves significant bandwidth and latency

Even full catalog cache:

- $10,000 \times 50\text{ KB} = 500\text{ MB}$
- Tiny! Cache everything!

**Interview Insight:** "Sometimes you can cache the entire dataset"

## Real-World Examples

### Example 1: Design Twitter

#### Step 1: Clarify Requirements

Assumptions:

- 400M DAU
- 200M tweets per day (0.5 tweets per user)
- Each user views 100 tweets per day
- Average tweet: 280 characters = 280 bytes

- 20% tweets have images (500 KB)
- Store for 5 years

## Step 2: Traffic Estimation

Writes (Tweets):

$$200M \text{ tweets/day} \div 86,400 = 2,315 \text{ QPS}$$

$$\text{Peak: } 2,315 \times 3 = \sim 7K \text{ QPS}$$

Reads (Timeline):

$$400M \text{ users} \times 100 \text{ tweets} = 40B \text{ reads/day}$$

$$40B \div 86,400 = 462,963 \text{ QPS} \approx 463K \text{ QPS}$$

$$\text{Peak: } 463K \times 3 = \sim 1.4M \text{ QPS}$$

$$\text{Read-to-Write Ratio: } 463K / 2.3K = 200:1$$

## Step 3: Storage Estimation

Daily storage:

- Text:  $200M \times 280 \text{ bytes} = 56 \text{ GB}$
- Images:  $200M \times 0.2 \times 500 \text{ KB} = 20 \text{ TB}$
- Total:  $\sim 20 \text{ TB/day}$

$$5 \text{ years: } 20 \text{ TB} \times 365 \times 5 = 36.5 \text{ PB}$$

With metadata and indexes (+20%):

$$\text{Total: } 36.5 \text{ PB} \times 1.2 = \sim 44 \text{ PB}$$

## Step 4: Bandwidth Estimation

Incoming:

- $2.3K \text{ QPS} \times 280 \text{ bytes} = 644 \text{ KB/s (text)}$
- Images:  $40M/\text{day} \div 86,400 = 463 \text{ QPS}$
- $463 \text{ QPS} \times 500 \text{ KB} = 231 \text{ MB/s}$
- Total incoming:  $\sim 232 \text{ MB/s}$

Outgoing:

- $463K \text{ QPS} \times 280 \text{ bytes} = 130 \text{ MB/s (text)}$
- Images:  $463K \times 0.2 \times 200 \text{ KB (thumbnails)} = 18.5 \text{ GB/s}$
- Total:  $\sim 19 \text{ GB/s}$

With CDN (95% cache hit):

- Origin:  $19 \text{ GB/s} \times 0.05 = 950 \text{ MB/s}$
- CDN:  $19 \text{ GB/s} \times 0.95 = 18 \text{ GB/s}$

## Step 5: Memory Estimation

Cache 20% of active user timelines:

- $400M \times 0.2 = 80M$  users
- Each timeline:  $50 \text{ tweets} \times 280 \text{ bytes} = 14 \text{ KB}$
- Total:  $80M \times 14 \text{ KB} = 1.1 \text{ TB}$

Cache servers needed:

- $1.1 \text{ TB} \div 100 \text{ GB per server} = 11 \text{ servers}$
- Round up to 20 servers for redundancy

## Summary for Interview:

Traffic: ~7K write, ~1.4M read (peak)

Storage: ~44 PB for 5 years

Bandwidth: ~19 GB/s (950 MB/s with CDN)

Cache: ~1.1 TB across 20 servers

Database: Cassandra (write-heavy), PostgreSQL (users), Redis (cache)

## Example 2: Design Instagram

### Step 1: Requirements

- 500M DAU
- 50M photos uploaded per day (0.1 per user)
- Each user views 50 photos per day
- Photo: 2 MB original, 200 KB feed, 20 KB thumbnail
- Store for 10 years

### Step 2: Traffic

Uploads:  $50M \div 86,400 = 578 \text{ QPS}$  (peak: 1.7K QPS)

Views:  $500M \times 50 = 25B/\text{day}$

Views QPS:  $25B \div 86,400 = 289K \text{ QPS}$  (peak: 867K QPS)

Ratio: 500:1 (extremely read-heavy)

### Step 3: Storage

Per photo:

- Original: 2 MB
- Feed size: 200 KB
- Thumbnail: 20 KB

- Total: 2.22 MB per photo

Daily:  $50M \times 2.22 \text{ MB} = 111 \text{ TB}$   
 Yearly:  $111 \text{ TB} \times 365 = 40.5 \text{ PB}$   
 10 years:  $405 \text{ PB}$

With compression (save 30%):  
 Actual:  $405 \text{ PB} \times 0.7 = \sim 283 \text{ PB}$

#### Step 4: Bandwidth

Upload:  $578 \text{ QPS} \times 2 \text{ MB} = 1.2 \text{ GB/s}$   
 Download:  $289K \text{ QPS} \times 200 \text{ KB} = 57.8 \text{ GB/s}$

With CDN (90% cache hit):  
 - Origin:  $57.8 \text{ GB/s} \times 0.1 = 5.8 \text{ GB/s}$   
 - CDN:  $57.8 \text{ GB/s} \times 0.9 = 52 \text{ GB/s}$

#### Step 5: Cache

Cache hot photos (20% of last 30 days):  
 -  $50M/\text{day} \times 30 \text{ days} = 1.5B \text{ photos}$   
 -  $1.5B \times 0.2 = 300M \text{ hot photos}$   
 -  $300M \times 200 \text{ KB} = 60 \text{ TB}$

Distributed:  $60 \text{ TB} \div 100 \text{ servers} = 600 \text{ GB each}$

### Example 3: Design WhatsApp

#### Step 1: Requirements

- 2B MAU, 1B DAU (50% active daily)
- Each user sends 40 messages/day
- Each user receives 40 messages/day
- Average message: 100 bytes
- Store for 1 year

#### Step 2: Traffic

Messages per day:  $1B \times 40 = 40B$   
 Message QPS:  $40B \div 86,400 = 462K \text{ QPS}$   
 Peak:  $462K \times 3 = 1.4M \text{ QPS}$

This is WRITE-heavy (send/receive are both writes)

### Step 3: Storage

Daily:  $40B \times 100 \text{ bytes} = 4 \text{ TB}$

Yearly:  $4 \text{ TB} \times 365 = 1.46 \text{ PB}$

With metadata (timestamp, read status, etc.):

- +50 bytes per message
- $40B \times 150 \text{ bytes} = 6 \text{ TB/day}$
- Yearly: 2.19 PB

With images/videos (10% of messages):

- $40B \times 0.1 = 4B$  media messages
- Average: 1 MB per media
- $4B \times 1 \text{ MB} = 4 \text{ PB/day}$
- Yearly:  $4 \text{ PB} \times 365 = 1,460 \text{ PB} = 1.46 \text{ exabytes}$

### Step 4: Bandwidth

Text:  $462K \text{ QPS} \times 100 \text{ bytes} = 46 \text{ MB/s}$

Media:  $462K \times 0.1 \times 1 \text{ MB} = 46 \text{ GB/s}$

Total: ~46 GB/s

Note: This is why WhatsApp compresses images aggressively!

With compression (80% reduction):

Actual:  $46 \text{ GB/s} \times 0.2 = \sim 9 \text{ GB/s}$

## Example 4: Design Uber

### Step 1: Requirements

- 50M daily rides
- Each ride: 1 driver, 1 rider
- Location updates every 3 seconds during ride
- Average ride: 15 minutes
- Store location history for 1 year

### Step 2: Traffic

Location Updates:

- Updates per ride:  $(15 \text{ min} \times 60 \text{ sec}) / 3 \text{ sec} = 300 \text{ updates}$

- Updates per day:  $50M \text{ rides} \times 300 \times 2 \text{ (driver + rider)} = 30B \text{ updates}$
- QPS:  $30B \div 86,400 = 347K \text{ QPS}$
- Peak (during rush hour, 3x):  $\sim 1M \text{ QPS}$

This is WRITE-HEAVY (real-time location tracking)

### Step 3: Storage

Per location update:

- User ID: 8 bytes
- Latitude: 8 bytes
- Longitude: 8 bytes
- Timestamp: 8 bytes
- Total: 32 bytes

Daily:  $30B \times 32 \text{ bytes} = 960 \text{ GB per day}$

Yearly:  $960 \text{ GB} \times 365 = 350 \text{ TB}$

With trip metadata (pickup, dropoff, fare, etc.):

- Additional 1 KB per trip
- $50M \times 1 \text{ KB} = 50 \text{ GB per day}$
- Yearly: 18 TB

Total:  $\sim 368 \text{ TB per year}$

### Step 4: Memory (Real-time Matching)

Active drivers at any time:

- Assume 5M drivers online (10% of daily drivers)
- Each driver location: 32 bytes
- Total:  $5M \times 32 \text{ bytes} = 160 \text{ MB}$

Cache all active drivers in memory:

- Easily fits in single Redis instance!
- Use Redis Geo data structure (GEOADD, GEORADIUS)

Active riders:

- 5M riders waiting (matching period)
- Total: 160 MB

Combined: 320 MB (tiny!)

## Quick Calculation Shortcuts

### Shortcut 1: Daily to QPS

**Formula:** Daily Requests  $\div$  100K  $\approx$  QPS

**Why:** 86,400  $\approx$  100,000 (close enough for estimates)

**Examples:**

$$10M \text{ requests/day} \div 100K = 100 \text{ QPS}$$

$$1B \text{ requests/day} \div 100K = 10K \text{ QPS}$$

$$100B \text{ requests/day} \div 100K = 1M \text{ QPS}$$

## Shortcut 2: Monthly to Daily

**Formula:** Monthly  $\div$  30 = Daily

**Example:**

$$3B \text{ requests/month} \div 30 = 100M \text{ requests/day}$$

$$100M \div 100K = 1K \text{ QPS}$$

## Shortcut 3: Storage Growth

**Formula:** Daily  $\times$  365  $\times$  Years = Total

**Example:**

$$10 \text{ TB/day} \times 365 \times 5 = 18,250 \text{ TB} = \sim 18 \text{ PB}$$

Round to 20 PB for safety margin

## Shortcut 4: Bandwidth from QPS

**Formula:** QPS  $\times$  Object Size

**Example:**

$$10K \text{ QPS} \times 100 \text{ KB} = 1,000,000 \text{ KB/s} = 1 \text{ GB/s}$$

Round to 1 GB/s for simplicity

## Shortcut 5: 80-20 Rule for Caching

**Formula:** Total Data  $\times$  0.2 = Cache Size

**Example:**

100 TB total data  
Cache:  $100 \text{ TB} \times 0.2 = 20 \text{ TB}$   
Achieves ~80% cache hit ratio

## Common Estimation Patterns

### Pattern 1: Social Media (Twitter, Instagram)

#### Characteristics:

- Read-heavy (100:1 to 1000:1 ratio)
- Small writes (text, metadata)
- Large reads (media files)

#### Formula:

$$\begin{aligned}\text{Write QPS} &= \text{DAU} \times \text{Posts per User} / 86,400 \\ \text{Read QPS} &= \text{DAU} \times \text{Feed Loads} \times \text{Items per Load} / 86,400 \\ \text{Storage} &= \text{Posts per Day} \times \text{Size} \times \text{Retention Period}\end{aligned}$$

### Pattern 2: Messaging (WhatsApp, Telegram)

#### Characteristics:

- Balanced read/write (1:1 ratio)
- Small messages (bytes to KB)
- High frequency (billions per day)

#### Formula:

$$\begin{aligned}\text{Message QPS} &= \text{DAU} \times \text{Messages per User} / 86,400 \\ \text{Storage} &= \text{Messages per Day} \times \text{Size} \times \text{Retention} \\ \text{Real-time Connections} &= \text{DAU} \text{ (for WebSocket)}\end{aligned}$$

### Pattern 3: Video Streaming (YouTube, Netflix)

#### Characteristics:

- Extremely read-heavy (1000:1)
- Large files (GB per video)
- CDN is mandatory

#### Formula:

Upload QPS = Uploads per Day / 86,400  
View QPS = Views per Day / 86,400  
Storage = Videos per Day × Size × Retention  
Bandwidth = View QPS × Bitrate / 8 (bits to bytes)

## Pattern 4: E-Commerce (Amazon, eBay)

### Characteristics:

- Read-heavy (browse > buy)
- Need ACID (transactions)
- Spike during sales events

### Formula:

Browse QPS = DAU × Page Views / 86,400  
Purchase QPS = Orders per Day / 86,400  
Storage = Products × Size + Orders × Size  
Peak Factor = 10x (during Black Friday)

## Interview Example Walkthroughs

### Walkthrough 1: "Design a News Feed"

**Interviewer:** "Design a news feed like Facebook. 1 billion users."

**You:**

"Let me estimate the scale:

#### Assumptions:

- 1B total users, 500M DAU (50% active)
- Each user posts 0.5 times per day
- Each user refreshes feed 10 times per day
- Each feed load shows 20 posts
- Average post: 1 KB (text + metadata)
- Store for 2 years

#### Traffic:

Posts:  $500M \times 0.5 = 250M/\text{day} \div 100K = 2.5K$  write QPS

Reads:  $500M \times 10 \times 20 = 100B$  reads/day  $\div 100K = 1M$  read QPS

Ratio: 400:1 (very read-heavy)

#### Storage:

Posts:  $250M/\text{day} \times 1 \text{ KB} = 250 \text{ GB}/\text{day}$

2 years:  $250 \text{ GB} \times 365 \times 2 = 182 \text{ TB}$

Round to 200 TB with indexes

Bandwidth:

Outgoing:  $1\text{M QPS} \times 1\text{ KB} = 1\text{ GB/s}$

Cache:

Cache 20% of 500M users' feeds: 100M users

Each feed:  $20\text{ posts} \times 1\text{ KB} = 20\text{ KB}$

Cache size:  $100\text{M} \times 20\text{ KB} = 2\text{ TB}$  across 20 servers

Database:

- PostgreSQL/MySQL for users (ACID)
  - Cassandra for posts (write-heavy, 2.5K QPS)
  - Redis for feed cache ( $1\text{M read QPS} \times 80\% = 800\text{K QPS}$  served from cache)
- "

## Walkthrough 2: "Design a URL Shortener"

**Interviewer:** "Design bit.ly. 100M URLs created per month."

**You:**

"Let me calculate the scale:

Assumptions:

- 100M new URLs per month
- 10B redirects per month (100:1 read-to-write)
- Average URL: 200 bytes
- Store for 10 years

Traffic:

Writes:  $100\text{M/month} \div 30 \div 100\text{K} = 33\text{ QPS}$  (round to 40 QPS)

Reads:  $10\text{B/month} \div 30 \div 100\text{K} = 3,333\text{ QPS}$  (round to 4K QPS)

Peak: 40 write QPS, 12K read QPS (3x)

Storage:

URLs:  $100\text{M} \times 12 \times 10 = 12\text{B URLs}$

Size:  $12\text{B} \times 200\text{ bytes} = 2.4\text{ TB}$

With analytics: +6 TB (assume 10 clicks per URL)

Total: ~9 TB over 10 years

Cache:

Cache hot URLs (20% of last 30 days):

- $100\text{M/month} = 3.3\text{M/day}$
- $3.3\text{M} \times 30\text{ days} = 100\text{M URLs}$
- $100\text{M} \times 0.2 = 20\text{M hot URLs}$
- $20\text{M} \times 200\text{ bytes} = 4\text{ GB}$

This is tiny! Cache everything recent!

Bandwidth:

Incoming: 40 QPS × 200 bytes = 8 KB/s  
Outgoing: 4K QPS × 200 bytes = 800 KB/s

Database:

- PostgreSQL (simple, ACID, 4K QPS easily handled)
- Redis cache (80% hit ratio, < 1ms latency)
- Cassandra for analytics (time-series clicks)
- "

### Walkthrough 3: "Design Dropbox"

**Interviewer:** "Design file sync service. 50M users."

**You:**

"Let me estimate:

Assumptions:

- 50M registered users, 10M DAU (20% active)
- Each user has 200 files (average)
- Average file: 1 MB
- Each user syncs 5 files per day

Traffic:

Syncs:  $10M \times 5 = 50M$  syncs/day  $\div 100K = 500$  QPS

Peak:  $500 \times 3 = 1.5K$  QPS

Storage:

Total files:  $50M \text{ users} \times 200 \text{ files} = 10 \text{ billion files}$

Total storage:  $10B \times 1 \text{ MB} = 10 \text{ PB}$

With 3 replicas (redundancy):

Total:  $10 \text{ PB} \times 3 = 30 \text{ PB}$

Bandwidth:

Upload/Download:  $500 \text{ QPS} \times 1 \text{ MB} = 500 \text{ MB/s}$

Peak: 1.5 GB/s

Metadata cache:

- File paths, versions, permissions
- $10B \text{ files} \times 1 \text{ KB metadata} = 10 \text{ TB}$
- Cache 20%: 2 TB across servers
- "

## Quick Reference Cheat Sheet

### Time Conversions

1 second	= 1,000 ms
1 minute	= 60 seconds
1 hour	= 3,600 seconds
1 day	= 86,400 seconds (~100K for rough estimates)
1 month	= 2,592,000 seconds (~2.5M)
1 year	= 31,536,000 seconds (~30M)

## Storage Conversions

1 KB	= 1,024 bytes	(~1 thousand)
1 MB	= 1,024 KB	(~1 million bytes)
1 GB	= 1,024 MB	(~1 billion bytes)
1 TB	= 1,024 GB	(~1 trillion bytes)
1 PB	= 1,024 TB	(~1 quadrillion bytes)
1 EB	= 1,024 PB	(~1 quintillion bytes)

## QPS Estimation Shortcuts

Daily Operations → QPS:

1M/day	= 10 QPS
10M/day	= 100 QPS
100M/day	= 1K QPS
1B/day	= 10K QPS
10B/day	= 100K QPS
100B/day	= 1M QPS

## Typical Ratios

Read-to-Write Ratios:

- Social media: 100:1 to 1000:1
- Messaging: 1:1
- E-commerce: 10:1 to 100:1
- Analytics: Write-heavy (opposite)

Cache Hit Ratios:

- Well-designed: 80–90%
- Average: 70–80%
- Poor: < 70%

Peak to Average:

- Normal: 2–3x
- Social (viral): 5–10x
- E-commerce (Black Friday): 10–20x

## Storage Per Object

User profile: 1–2 KB  
Tweet/Post: 280 bytes – 1 KB  
Email: 50–100 KB  
Small image: 50–100 KB  
Medium image: 200–500 KB  
Large image: 1–3 MB  
Short video (1 min): 10–50 MB  
HD video (1 min): 50–100 MB

## Interview Tips & Tricks

### Tip 1: Start with Round Numbers

#### Example:

Interviewer: "500 million daily active users"  
You: "So roughly 500M DAU, I'll round to 500M for calculations"

Interviewer: "Each user posts 2.3 times per day"  
You: "I'll approximate that as 2 posts per day"

**Why:** Makes math easier, shows you understand precision isn't critical

### Tip 2: Show Units in Every Step

#### Bad:

$$\begin{aligned} 100 \times 200 &= 20,000 \\ 20,000 \times 365 &= 7,300,000 \end{aligned}$$

#### Good:

$$\begin{aligned} 100M \text{ photos} \times 200 \text{ KB} &= 20,000,000,000 \text{ KB} = 20 \text{ TB per day} \\ 20 \text{ TB per day} \times 365 \text{ days} &= 7,300 \text{ TB} = 7.3 \text{ PB per year} \end{aligned}$$

### Tip 3: Sanity Check Your Numbers

#### After calculating, ask yourself:

- "Does this make sense?"

- "Is 1 PB reasonable for Instagram photos? Yes!"
- "Is 10 EB reasonable for Twitter text? No, that's too high!"

### Common Sanity Checks:

Twitter text storage should be < 100 TB (tiny)  
Instagram photo storage should be > 10 PB (huge)  
WhatsApp messages should be > 1 PB but < 10 PB  
YouTube videos should be > 1 EB (massive)

### Tip 4: Mention Optimizations

#### After calculations, always mention:

Storage optimization:  
"With compression, we can reduce this by 30-50%"

Bandwidth optimization:  
"With CDN caching at 90% hit ratio, origin only serves 10%"

Memory optimization:  
"Using the 80-20 rule, we only need to cache 20% of data"

### Tip 5: Compare to Real Systems

#### Say things like:

"This is similar to Twitter's scale"  
"Instagram actually uses ~100 PB of storage"  
"Netflix serves petabytes per day from CDN"  
"WhatsApp handles billions of messages"

## Common Mistakes to Avoid

### ✗ Mistake 1: Being Too Precise

**Bad:** "387,654,321 requests per day"

**Good:** "~400 million requests per day"

**Why:** Precision gives false confidence, rough numbers are sufficient

### ✗ Mistake 2: Forgetting Peak Traffic

**Bad:** "10K QPS average"

**Good:** "10K QPS average, 30K QPS peak (3x factor)"

**Why:** Systems must handle peak load, not average

### ✖ Mistake 3: Not Showing Units

**Bad:** "Storage is 100"

**Good:** "Storage is 100 TB"

**Why:** Interviewer doesn't know if you mean MB, GB, or TB

### ✖ Mistake 4: Ignoring Media

**Bad:** "Twitter stores 1 TB of tweets"

**Good:** "Twitter stores 50 GB of text + 10 PB of images"

**Why:** Media dominates storage, don't ignore it

### ✖ Mistake 5: Forgetting Replication

**Bad:** "Need 10 TB storage"

**Good:** "Need 10 TB  $\times$  3 replicas = 30 TB storage"

**Why:** Replication is standard for reliability

---

## Practice Problems

### Problem 1: Design Google Photos

**Given:**

- 1B users, 500M DAU
- Each user uploads 5 photos per day
- Each photo: 3 MB
- Each user views 50 photos per day
- Store for lifetime (20 years)

**Your turn!** Calculate:

1. Write QPS?
2. Read QPS?
3. Total storage for 20 years?
4. Bandwidth required?
5. Cache size?

► Click for Answer

**Traffic:**

- Uploads:  $500M \times 5 = 2.5B/\text{day} \div 100K = 25K$  write QPS
- Views:  $500M \times 50 = 25B/\text{day} \div 100K = 250K$  read QPS
- Peak: 75K write, 750K read QPS

**Storage:**

- Daily:  $2.5B \times 3 \text{ MB} = 7.5 \text{ PB}/\text{day}$
- Yearly:  $7.5 \text{ PB} \times 365 = 2,737 \text{ PB} = 2.7 \text{ EB}$
- 20 years:  $2.7 \text{ EB} \times 20 = 54 \text{ EB}$
- With compression (50%): 27 EB

**Bandwidth:**

- Upload:  $25K \text{ QPS} \times 3 \text{ MB} = 75 \text{ GB/s}$
- Download:  $250K \text{ QPS} \times 3 \text{ MB} = 750 \text{ GB/s}$
- With CDN (95% cache): 37.5 GB/s origin

**Cache:**

- Hot photos (20% of last 90 days):
- $2.5B \times 90 \times 0.2 = 45B$  photos
- $45B \times 3 \text{ MB} = 135 \text{ PB}$  (too large to cache originals!)
- Cache thumbnails instead:  $45B \times 200 \text{ KB} = 9 \text{ PB}$
- Still large! Cache metadata only:  $45B \times 1 \text{ KB} = 45 \text{ TB}$

## Problem 2: Design Spotify

**Given:**

- 400M users, 100M DAU
- Each user streams 10 songs per day
- Each song: 5 MB (compressed)
- Catalog: 100M songs
- Store catalog forever

**Calculate:** Traffic, Storage, Bandwidth, Cache

► Click for Answer

**Traffic:**

- Streams:  $100M \times 10 = 1B/\text{day} \div 100K = 10K$  QPS
- Peak: 30K QPS

**Storage:**

- Catalog:  $100M \text{ songs} \times 5 \text{ MB} = 500 \text{ TB}$
- With multiple qualities:  $500 \text{ TB} \times 3 = 1.5 \text{ PB}$

**Bandwidth:**

- Download:  $10K \text{ QPS} \times 5 \text{ MB} = 50 \text{ GB/s}$
- With CDN (90% cache): 5 GB/s origin

### Cache:

- Hot songs (20% of catalog):  $100M \times 0.2 = 20M$  songs
- $20M \times 5 \text{ MB} = 100 \text{ TB}$
- Distributed:  $100 \text{ TB} \div 100 \text{ servers} = 1 \text{ TB each}$

## The Universal Template

Use this template for ANY system design question:

### 1. TRAFFIC (QPS):

Write QPS = Writes per Day  $\div 100K$

Read QPS = Reads per Day  $\div 100K$

Peak QPS = Average  $\times 3$

Ratio = Reads / Writes

### 2. STORAGE:

Daily = Items per Day  $\times$  Size per Item

Total = Daily  $\times 365 \times$  Years

With Replication = Total  $\times 3$

### 3. BANDWIDTH:

Incoming = Write QPS  $\times$  Object Size

Outgoing = Read QPS  $\times$  Object Size

With CDN = Outgoing  $\times (1 - \text{Cache Hit Ratio})$

### 4. MEMORY (Cache):

Cache Size = Total Data  $\times 0.2$  (80–20 rule)

Servers = Cache Size  $\div 100 \text{ GB}$

Cache Hit Ratio = ~80%

### 5. DATABASE:

If Read-Heavy  $\rightarrow$  SQL + Read Replicas + Cache

If Write-Heavy  $\rightarrow$  Cassandra/NoSQL

If < 1 TB  $\rightarrow$  Single SQL instance

If > 10 TB  $\rightarrow$  Sharding or NoSQL

## Real Numbers from Production Systems

Twitter (Actual)

DAU: 400M

Tweets/day: 500M

QPS: ~6K write, ~300K read

Storage: ~100 PB

Bandwidth: ~1 TB/s (with CDN)

## Instagram (Actual)

DAU: 500M  
Photos/day: 95M  
QPS: ~1K write, ~500K read  
Storage: ~100 PB  
Cache: Redis Cluster (50+ nodes)

## WhatsApp (Actual)

DAU: 2B  
Messages/day: 100B  
QPS: ~1M write, ~1M read  
Storage: ~10 PB  
Erlang servers: 1000+

## YouTube (Actual)

Videos watched/day: 5B  
Upload: 500 hours/minute  
Storage: ~1 EB (yes, exabyte!)  
Bandwidth: ~1 PB/s globally  
CDN: Critical infrastructure

---

## Interview Scoring Rubric

### What Interviewers Look For:

#### **Structured Approach** (30%)

- Clear methodology
- Step-by-step calculations
- Organized presentation

#### **Reasonable Assumptions** (20%)

- States assumptions clearly
- Numbers are realistic
- Justifies choices

#### **Correct Math** (20%)

- Gets order of magnitude right
- Shows work

- Unit conversions correct

**Identifies Bottlenecks (15%)**

- "At 1M QPS, we need caching"
- "44 PB requires distributed storage"
- "This is write-heavy, need Cassandra"

**Mentions Optimizations (15%)**

- CDN for bandwidth
  - Caching for reads
  - Compression for storage
- 

## The Mental Math Trick

### Multiplication by Powers of 10

```

× 10:      Add one zero
× 100:     Add two zeros
× 1,000:   Add three zeros (1K)
× 1,000,000: Add six zeros (1M)
× 1,000,000,000: Add nine zeros (1B)

```

**Example:**

$$5 \text{ MB} \times 1,000 = 5,000 \text{ MB} = 5 \text{ GB}$$

$$100 \text{ KB} \times 1,000,000 = 100,000,000 \text{ KB} = 100 \text{ GB}$$

### Division by Large Numbers

```

÷ 100:      Remove two zeros
÷ 1,000:    Remove three zeros
÷ 86,400:   Divide by 100,000 (close enough)

```

**Example:**

$$10,000,000 \div 100,000 = 100$$

$$1,000,000,000 \div 100,000 = 10,000$$

---

## Advanced: Server Capacity Estimation

## Typical Server Specs

### **Small Instance (t3.medium):**

- 2 vCPU, 4 GB RAM
- Handles: ~500 QPS
- Cost: ~\$50/month

### **Medium Instance (m5.xlarge):**

- 4 vCPU, 16 GB RAM
- Handles: ~2,000 QPS
- Cost: ~\$150/month

### **Large Instance (m5.4xlarge):**

- 16 vCPU, 64 GB RAM
- Handles: ~10,000 QPS
- Cost: ~\$600/month

## Calculating Server Count

### **Formula:**

$$\text{Servers} = (\text{Peak QPS} / \text{QPS per Server}) \times \text{Safety Factor}$$

Example:

Peak: 30K QPS

Server capacity: 2K QPS

Safety factor: 1.5 (50% headroom)

$$\text{Servers} = (30,000 / 2,000) \times 1.5 = 22.5$$

Round up: 25 servers

## Database Server Capacity

### **PostgreSQL (single instance):**

- Reads: 10K-50K QPS
- Writes: 1K-10K QPS
- Storage: Up to 16 TB

### **Cassandra (per node):**

- Reads: 10K-50K QPS
- Writes: 10K-100K QPS
- Storage: Up to 8 TB per node

### **Redis (single instance):**

- Reads: 100K-1M QPS
  - Writes: 100K-1M QPS
  - Storage: Up to 512 GB RAM
- 

## Estimation Worksheet Template

**Copy this for every interview:**

SYSTEM: \_\_\_\_\_

### ASSUMPTIONS:

- DAU: \_\_\_\_\_
- Operations per user: \_\_\_\_\_
- Data size: \_\_\_\_\_
- Retention: \_\_\_\_\_

### TRAFFIC CALCULATIONS:

- Daily operations: DAU × ops = \_\_\_\_\_
- Average QPS: daily ÷ 100K = \_\_\_\_\_
- Peak QPS: avg × 3 = \_\_\_\_\_
- Read/Write ratio: \_\_\_\_\_

### STORAGE CALCULATIONS:

- Per item size: \_\_\_\_\_
- Daily: items/day × size = \_\_\_\_\_
- Yearly: daily × 365 = \_\_\_\_\_
- Total: yearly × years = \_\_\_\_\_
- With replication: total × 3 = \_\_\_\_\_

### BANDWIDTH CALCULATIONS:

- Incoming: write QPS × size = \_\_\_\_\_
- Outgoing: read QPS × size = \_\_\_\_\_
- With CDN: outgoing × (1 - hit ratio) = \_\_\_\_\_

### MEMORY/CACHE CALCULATIONS:

- Hot data (20%): total × 0.2 = \_\_\_\_\_
- Cache servers: cache ÷ 100 GB = \_\_\_\_\_
- Cache hit ratio: \_\_\_\_\_

### DATABASE CHOICE:

- Primary: \_\_\_\_\_ (because: \_\_\_\_\_)
- Cache: \_\_\_\_\_
- Analytics: \_\_\_\_\_

---

## Final Checklist

Before saying "I'm done with calculations", verify:

## Traffic Estimation

- Calculated write QPS
- Calculated read QPS
- Mentioned peak factor (2-3x)
- Identified if read-heavy or write-heavy

## Storage Estimation

- Calculated total storage
- Considered retention period
- Mentioned replication (3x)
- Separated data types (text vs media)

## Bandwidth Estimation

- Calculated incoming bandwidth
- Calculated outgoing bandwidth
- Mentioned CDN impact

## Memory Estimation

- Applied 80-20 rule
- Calculated cache size
- Determined number of servers

## Identified Bottlenecks

- Database: "At 10K write QPS, need sharding"
  - Network: "At 1 TB/s, need CDN"
  - Memory: "At 10 TB cache, need 100 servers"
- 

## Summary: The Golden Rules

1. **Round liberally** - 387M → 400M
  2. **State assumptions** - "Assuming X..."
  3. **Show your work** - Write calculations down
  4. **Use shortcuts** - Daily ÷ 100K = QPS
  5. **Apply 80-20 rule** - Cache 20%, serve 80%
  6. **Mention peak traffic** - Average × 3
  7. **Don't forget media** - Images dominate storage
  8. **Include replication** - × 3 for redundancy
  9. **Sanity check** - Does 1 EB for Twitter make sense? No!
  10. **Mention optimizations** - CDN, compression, caching
- 

## References

### Books

1. "Designing Data-Intensive Applications" - Martin Kleppmann
2. "System Design Interview" - Alex Xu

## Online Resources

1. **System Design Primer** - GitHub (donnemartin)
2. **Google's Latency Numbers** - Jeff Dean
3. **AWS Calculator** - For cost estimates
4. **High Scalability Blog** - Real-world numbers

## Videos

1. "System Design Course" - YouTube (Gaurav Sen)
2. "Back of Envelope Calculations" - YouTube (SystemDesignInterview)

## Appendix: Pre-Calculated Common Scenarios

### Scenario: 1 Million Users

```
Assumptions: 1M DAU, 10 operations/user/day
Traffic: 10M/day ÷ 100K = 100 QPS
Storage (1 KB per op): 10M × 1 KB = 10 GB/day = 3.6 TB/year
Cache: 1M users × 10 KB = 10 GB
Servers: 5–10 app servers
```

### Scenario: 10 Million Users

```
Assumptions: 10M DAU, 10 operations/user/day
Traffic: 100M/day ÷ 100K = 1K QPS
Storage: 100M × 1 KB = 100 GB/day = 36 TB/year
Cache: 10M × 10 KB = 100 GB
Servers: 10–20 app servers, single DB with replicas
```

### Scenario: 100 Million Users

```
Assumptions: 100M DAU, 10 operations/user/day
Traffic: 1B/day ÷ 100K = 10K QPS
Storage: 1B × 1 KB = 1 TB/day = 365 TB/year
Cache: 100M × 10 KB = 1 TB across 10 servers
Servers: 50–100 app servers, sharded DB or NoSQL
```

### Scenario: 1 Billion Users (Facebook Scale)

Assumptions: 1B DAU, 10 operations/user/day  
Traffic:  $10\text{B/day} \div 100\text{K} = 100\text{K QPS}$   
Storage:  $10\text{B} \times 1 \text{ KB} = 10 \text{ TB/day} = 3.6 \text{ PB/year}$   
Cache:  $200\text{M hot users} \times 10 \text{ KB} = 2 \text{ TB across 20 servers}$   
Servers: 500–1000 app servers, distributed NoSQL

---

**Document Version:** 1.0

**Last Updated:** January 8, 2025

**For:** System Design Interview Preparation

**Status:** Complete & Interview-Ready 

**Remember:** In interviews, being roughly right with clear reasoning beats being precisely wrong!