

Back-of-the-Envelope Calculations Guide for System Design Interviews

Table of Contents

1. [Introduction](#)
 2. [Core Principles](#)
 3. [Essential Numbers to Memorize](#)
 4. [Step-by-Step Framework](#)
 5. [Traffic Estimation](#)
 6. [Storage Estimation](#)
 7. [Bandwidth Estimation](#)
 8. [Memory/Cache Estimation](#)
 9. [Real-World Examples](#)
 10. [Quick Reference Cheat Sheet](#)
-

Introduction

"In system design interviews, being approximately right is better than being precisely wrong."

Back-of-the-envelope calculations are rough estimates used to evaluate the feasibility of a system design. Interviewers want to see:

- Can you estimate scale and resource requirements?
- Do you understand trade-offs between different approaches?
- Can you identify potential bottlenecks?

You don't need a calculator! Simple multiplication and division with rough numbers is enough.

Core Principles

1. Round Numbers Liberally

Bad: "We have 387,654,321 daily active users..."

Good: "We have ~400 million daily active users..."

Why: Easier to calculate, equally accurate for capacity planning

2. Use Powers of 10/2

Instead of: 1,234,567 bytes

Use: ~1 MB (1 million bytes)

Instead of: 987 requests

Use: ~1K requests (1,000)

3. State Your Assumptions Clearly

Always say:

- "Assuming 100 million daily active users..."
- "Let's assume average photo size is 200 KB..."
- "I'll estimate read-to-write ratio as 100:1..."

4. Work in Standard Units

Storage: Bytes → KB → MB → GB → TB → PB

Time: Seconds → Minutes → Hours → Days → Years

Traffic: QPS (Queries Per Second)

5. Show Your Work

Don't say: "We need 10 TB of storage"

Do say: "100M photos × 200 KB = 20 TB, minus 50% compression = 10 TB"

Essential Numbers to Memorize

Powers of Two (Storage)

$2^{10} = 1,024$	≈ 1 Thousand	= 1 KB
$2^{20} = 1,048,576$	≈ 1 Million	= 1 MB
$2^{30} = 1$ Billion	≈ 1 GB	
$2^{40} = 1$ Trillion	≈ 1 TB	
$2^{50} = 1$ Quadrillion	≈ 1 PB	

Pro Tip: Just remember "1024" and multiply!

- 1 KB = ~1,000 bytes
- 1 MB = ~1,000 KB = 1 million bytes
- 1 GB = ~1,000 MB = 1 billion bytes
- 1 TB = ~1,000 GB = 1 trillion bytes

Time Conversions

1 minute = 60 seconds
1 hour = 3,600 seconds (60×60)
1 day = 86,400 seconds ($24 \times 3,600$)
1 month = ~2.5 million seconds ($30 \times 86,400$)
1 year = ~31.5 million seconds ($365 \times 86,400$)

Pro Tip: Remember "86,400 seconds per day"

- Daily to QPS: Divide by 86,400 (\approx 100K)
- Monthly to QPS: Divide by 2.5M

Latency Numbers (Google's Famous List)

L1 cache reference:	0.5 ns
Branch mispredict:	5 ns
L2 cache reference:	7 ns
Main memory reference:	100 ns
Compress 1KB with Snappy:	10 μ s
Send 1 KB over 1 Gbps network:	10 μ s
Read 1 MB from memory:	250 μ s
Round trip within datacenter:	500 μ s
Read 1 MB from SSD:	1 ms
Disk seek:	10 ms
Read 1 MB from network:	10 ms
Read 1 MB from disk:	30 ms
Send packet CA→Netherlands→CA:	150 ms

Takeaways:

- Memory is \sim 100x faster than SSD
- SSD is \sim 10x faster than HDD
- Network within datacenter is fast (0.5ms)
- Cross-continent is slow (150ms)

Common Data Sizes

1 character = 1 byte (ASCII)
 1 character = 2–4 bytes (Unicode/UTF-8)

User ID (integer): 8 bytes

Timestamp: 8 bytes

IP address (IPv4): 4 bytes

IP address (IPv6): 16 bytes

Tweet text (280 chars): \sim 280 bytes

Small image (thumbnail): \sim 50 KB

Regular image (compressed): \sim 200 KB

High-res image: \sim 2 MB

Short video (1 min): \sim 10 MB

HD video (1 min): \sim 50 MB

User profile: \sim 1–2 KB

Email: \sim 75 KB (average)

Web page: \sim 2 MB (with assets)

Step-by-Step Framework

Step 1: Clarify Scale Requirements

Questions to Ask:

Users:

- How many daily active users (DAU)?
- How many monthly active users (MAU)?
- User growth rate?

Usage:

- How many posts/uploads per user per day?
- How many reads per user per day?
- Read-to-write ratio?

Data:

- Average size of user data?
- Average size of content (posts, photos, videos)?
- Data retention period?

Step 2: Calculate Traffic

Formula:

$$\text{QPS} = \text{Total Operations per Day} / 86,400 \text{ seconds}$$

Example:

$$\begin{aligned} & 100\text{M tweets per day} \\ & = 100,000,000 / 86,400 \\ & = 1,157 \text{ QPS (round to } \sim 1.2\text{K QPS)} \end{aligned}$$

$$\begin{aligned} & \text{Peak QPS} = \text{Average QPS} \times \text{Peak Factor} \\ & \text{Peak Factor typically: 2-5x} \\ & = 1.2\text{K} \times 3 = 3.6\text{K QPS (round to } \sim 4\text{K QPS)} \end{aligned}$$

Step 3: Calculate Storage

Formula:

$$\text{Total Storage} = \text{Number of Items} \times \text{Size per Item} \times \text{Time Period}$$

Example:

$$\begin{aligned} & 100\text{M photos per month} \times 200 \text{ KB} \times 12 \text{ months} \times 10 \text{ years} \\ & = 100\text{M} \times 200 \text{ KB} \times 120 \\ & = 2.4 \text{ PB} \end{aligned}$$

Step 4: Calculate Bandwidth

Formula:

$$\text{Bandwidth} = \text{QPS} \times \text{Average Response Size}$$

Example:

Read QPS: 10K

Average response: 100 KB

$$\text{Bandwidth} = 10,000 \times 100 \text{ KB} = 1 \text{ GB/s}$$

Step 5: Calculate Memory (Cache)

Formula:

$$\text{Cache Size} = \% \text{ of Data to Cache} \times \text{Data Size}$$

Example (80-20 rule):

Total data: 10 TB

$$\text{Cache 20\% hot data: } 10 \text{ TB} \times 0.2 = 2 \text{ TB}$$

$$\text{Distributed across 20 servers: } 2 \text{ TB} / 20 = 100 \text{ GB per server}$$

Traffic Estimation

Calculation Method

Given: Daily Active Users (DAU) and actions per user

Formula:

$$\text{Total Daily Operations} = \text{DAU} \times \text{Operations per User}$$

$$\text{QPS} = \text{Total Daily Operations} / 86,400$$

$$\text{Peak QPS} = \text{QPS} \times \text{Peak Factor (2-5x)}$$

Example 1: Twitter

Given:

- 400M DAU
- Each user tweets 0.5 times per day (average)
- Each user reads 100 tweets per day (scrolling feed)

Calculate:

Write Traffic (Tweets):

- Tweets per day: $400M \times 0.5 = 200M$
- Write QPS: $200M / 86,400 \approx 2,315$ QPS
- Peak write QPS: $2,315 \times 3 \approx 7K$ QPS

Read Traffic (Timeline):

- Reads per day: $400M \times 100 = 40$ billion
- Read QPS: $40B / 86,400 \approx 462K$ QPS
- Peak read QPS: $462K \times 3 \approx 1.4M$ QPS

Read-to-Write Ratio: $40B / 200M = 200:1$

Interview Insight: "This is read-heavy, so we need caching and read replicas"

Example 2: Instagram

Given:

- 500M DAU
- Each user uploads 0.1 photos per day (1 photo per 10 days)
- Each user views 50 photos per day

Calculate:

Write Traffic (Photos):

- Photos per day: $500M \times 0.1 = 50M$
- Write QPS: $50M / 86,400 \approx 578$ QPS
- Peak: $\sim 1.7K$ QPS

Read Traffic (Feed):

- Views per day: $500M \times 50 = 25$ billion
- Read QPS: $25B / 86,400 \approx 289K$ QPS
- Peak: $\sim 867K$ QPS

Read-to-Write Ratio: $25B / 50M = 500:1$

Interview Insight: "Extremely read-heavy, CDN caching is critical"

Example 3: WhatsApp

Given:

- 2 billion MAU
- 50% daily active = 1B DAU
- Each user sends 40 messages per day
- Each user receives 40 messages per day

Calculate:

Messages per day: $1B \times 40 = 40$ billion

Message writes: $40B / 86,400 \approx 463K$ QPS

Peak: $\sim 1.4M$ QPS

With reads (checking for new messages):

- Check every 10 seconds: $1B \text{ users} \times 8,640 \text{ checks} = 8.64 \text{ trillion checks/day}$
- Check QPS: $8.64T / 86,400 = 100M$ QPS (too high!)
- Solution: WebSocket connections, push notifications

Interview Insight: "Need WebSocket for efficiency, not polling"

Storage Estimation

Calculation Method

Formula:

$$\text{Storage} = \text{Number of Objects} \times \text{Size per Object} \times \text{Time Period}$$

Example 1: YouTube

Given:

- 500 hours of video uploaded per minute
- Average video size: 50 MB per minute
- Store for 10 years

Calculate:

Daily uploads:

- $500 \text{ hours/min} \times 60 \text{ min/hour} \times 24 \text{ hours} = 720,000 \text{ hours per day}$
- $720K \text{ hours} \times 60 \text{ min} \times 50 \text{ MB} = 2,160,000,000 \text{ MB per day}$
- $= 2,160 \text{ TB per day} = \sim 2 \text{ PB per day}$

Yearly: $2 \text{ PB} \times 365 = 730 \text{ PB per year}$

10 years: $730 \text{ PB} \times 10 = 7.3 \text{ exabytes (7,300 PB)}$

With multiple resolutions (360p, 720p, 1080p, 4K):

- 4 versions per video
- Total: $7,300 \text{ PB} \times 4 = 29 \text{ exabytes}$

With compression and deduplication:

- Assume 30% savings
- Effective: $29 \text{ EB} \times 0.7 = \sim 20 \text{ exabytes}$

Interview Insight: "This is why YouTube needs Google's infrastructure!"

Example 2: Instagram

Given:

- 500M DAU
- 50M photos uploaded per day
- Average photo: 200 KB
- Store for 10 years

Calculate:

Daily storage:

$$\begin{aligned} - 50M \text{ photos} \times 200 \text{ KB} &= 10,000,000,000 \text{ KB} \\ - &= 10 \text{ TB per day} \end{aligned}$$

$$\text{Yearly: } 10 \text{ TB} \times 365 = 3.65 \text{ PB}$$

$$10 \text{ years: } 3.65 \text{ PB} \times 10 = 36.5 \text{ PB}$$

With thumbnails (3 sizes per photo):

- Original: 200 KB
- Large: 100 KB
- Thumbnail: 20 KB
- Total per photo: 320 KB

$$\text{Adjusted: } 50M \times 320 \text{ KB} = 16 \text{ TB per day}$$

$$10 \text{ years: } 16 \text{ TB} \times 365 \times 10 = 58 \text{ PB}$$

Interview Insight: "Need S3 or similar object storage, with lifecycle policies"

Example 3: Twitter

Given:

- 400M DAU
- 200M tweets per day
- Average tweet: 280 characters = ~280 bytes
- Store for 5 years
- 20% of tweets have images (~500 KB each)

Calculate:

Text storage:

$$- 200M \text{ tweets} \times 280 \text{ bytes} = 56 \text{ GB per day}$$

Image storage:

- $200M \times 0.2 = 40M$ images per day
- $40M \times 500 \text{ KB} = 20 \text{ TB}$ per day

Total per day: $56 \text{ GB} + 20 \text{ TB} \approx 20 \text{ TB}$
 Yearly: $20 \text{ TB} \times 365 = 7.3 \text{ PB}$
 5 years: $7.3 \text{ PB} \times 5 = 36.5 \text{ PB}$

With metadata (user_id, timestamp, likes, etc.):

- Additional 100 bytes per tweet
- $200M \times 100 \text{ bytes} = 20 \text{ GB}$ per day
- Negligible compared to images

Total: $\sim 37 \text{ PB}$ for 5 years

Interview Insight: "Text is cheap, media is expensive"

Bandwidth Estimation

Calculation Method

Formula:

$$\text{Bandwidth} = \text{QPS} \times \text{Average Object Size}$$

Example 1: Instagram Feed

Given:

- 289K read QPS (from earlier calculation)
- Average feed response: $50 \text{ photos} \times 200 \text{ KB thumbnails} = 10 \text{ MB}$

Calculate:

Incoming bandwidth (uploads):

$$- 578 \text{ QPS (writes)} \times 200 \text{ KB} = 115 \text{ MB/s}$$

Outgoing bandwidth (feed loads):

$$- 289K \text{ QPS} \times 10 \text{ MB} = 2,890 \text{ GB/s} = 2.8 \text{ TB/s}$$

Wait, that's too high! Let's optimize:

With CDN caching (95% cache hit ratio):

- Only 5% hits origin: $2.8 \text{ TB/s} \times 0.05 = 140 \text{ GB/s}$
- CDN bandwidth: 2.8 TB/s (distributed globally)

Actual Instagram approach:

- Serve thumbnails first (150 KB instead of 200 KB)

- Lazy load full images
- Reduced to ~1 TB/s from CDN

Interview Insight: "CDN is essential, not optional"

Example 2: YouTube

Given:

- 1 billion video views per day
- Average video: 10 minutes at 720p = 100 MB

Calculate:

Views per second: $1B / 86,400 \approx 11,574$ QPS

Bandwidth: $11,574 \times 100$ MB = 1,157 GB/s = 1.1 TB/s

Peak (3x): 3.3 TB/s

With CDN (80% cache hit):

- Origin: $3.3 \text{ TB/s} \times 0.2 = 660 \text{ GB/s}$
- CDN: $3.3 \text{ TB/s} \times 0.8 = 2.6 \text{ TB/s}$

Optimization with adaptive bitrate:

- Serve lower quality for slow connections
- Average drops to 50 MB per view
- New bandwidth: 578 GB/s origin

Memory/Cache Estimation

The 80-20 Rule (Pareto Principle)

Key Insight: 20% of data generates 80% of traffic

Cache Strategy: Cache the hot 20%, serve 80% of requests

Example 1: Twitter Timeline Cache

Given:

- 400M DAU
- Each timeline: 50 tweets \times 500 bytes = 25 KB
- Cache 20% of active users

Calculate:

Users to cache: $400M \times 0.2 = 80M$ users
Cache size: $80M \times 25\text{ KB} = 2,000\text{ GB} = 2\text{ TB}$

Distributed across servers:

- 20 cache servers \times 100 GB RAM each = 2 TB total
- Cost: ~\$2,000/month (AWS ElastiCache)

Cache hit ratio achieved: ~80%

Database queries saved: 80% of 462K QPS = 370K QPS

Interview Insight: "Caching 20% of users saves 80% of database queries"

Example 2: Netflix Thumbnails

Given:

- 10,000 videos in catalog
- Each thumbnail: 50 KB
- Cache most popular 1,000 videos (10%)

Calculate:

Cache size: 1,000 videos \times 50 KB = 50 MB

With 50% cache hit ratio:

- 11,574 QPS \times 0.5 = 5,787 QPS to origin
- Saves significant bandwidth and latency

Even full catalog cache:

- $10,000 \times 50\text{ KB} = 500\text{ MB}$
- Tiny! Cache everything!

Interview Insight: "Sometimes you can cache the entire dataset"

Real-World Examples

Example 1: Design Twitter

Step 1: Clarify Requirements

Assumptions:

- 400M DAU
- 200M tweets per day (0.5 tweets per user)
- Each user views 100 tweets per day
- Average tweet: 280 characters = 280 bytes

- 20% tweets have images (500 KB)
- Store for 5 years

Step 2: Traffic Estimation

Writes (Tweets):

$$200M \text{ tweets/day} \div 86,400 = 2,315 \text{ QPS}$$

$$\text{Peak: } 2,315 \times 3 = \sim 7K \text{ QPS}$$

Reads (Timeline):

$$400M \text{ users} \times 100 \text{ tweets} = 40B \text{ reads/day}$$

$$40B \div 86,400 = 462,963 \text{ QPS} \approx 463K \text{ QPS}$$

$$\text{Peak: } 463K \times 3 = \sim 1.4M \text{ QPS}$$

$$\text{Read-to-Write Ratio: } 463K / 2.3K = 200:1$$

Step 3: Storage Estimation

Daily storage:

- Text: $200M \times 280 \text{ bytes} = 56 \text{ GB}$
- Images: $200M \times 0.2 \times 500 \text{ KB} = 20 \text{ TB}$
- Total: $\sim 20 \text{ TB/day}$

$$5 \text{ years: } 20 \text{ TB} \times 365 \times 5 = 36.5 \text{ PB}$$

With metadata and indexes (+20%):

$$\text{Total: } 36.5 \text{ PB} \times 1.2 = \sim 44 \text{ PB}$$

Step 4: Bandwidth Estimation

Incoming:

- $2.3K \text{ QPS} \times 280 \text{ bytes} = 644 \text{ KB/s (text)}$
- Images: $40M/\text{day} \div 86,400 = 463 \text{ QPS}$
- $463 \text{ QPS} \times 500 \text{ KB} = 231 \text{ MB/s}$
- Total incoming: $\sim 232 \text{ MB/s}$

Outgoing:

- $463K \text{ QPS} \times 280 \text{ bytes} = 130 \text{ MB/s (text)}$
- Images: $463K \times 0.2 \times 200 \text{ KB (thumbnails)} = 18.5 \text{ GB/s}$
- Total: $\sim 19 \text{ GB/s}$

With CDN (95% cache hit):

- Origin: $19 \text{ GB/s} \times 0.05 = 950 \text{ MB/s}$
- CDN: $19 \text{ GB/s} \times 0.95 = 18 \text{ GB/s}$

Step 5: Memory Estimation

Cache 20% of active user timelines:

- $400M \times 0.2 = 80M$ users
- Each timeline: $50 \text{ tweets} \times 280 \text{ bytes} = 14 \text{ KB}$
- Total: $80M \times 14 \text{ KB} = 1.1 \text{ TB}$

Cache servers needed:

- $1.1 \text{ TB} \div 100 \text{ GB per server} = 11 \text{ servers}$
- Round up to 20 servers for redundancy

Summary for Interview:

Traffic: ~7K write, ~1.4M read (peak)

Storage: ~44 PB for 5 years

Bandwidth: ~19 GB/s (950 MB/s with CDN)

Cache: ~1.1 TB across 20 servers

Database: Cassandra (write-heavy), PostgreSQL (users), Redis (cache)

Example 2: Design Instagram

Step 1: Requirements

- 500M DAU
- 50M photos uploaded per day (0.1 per user)
- Each user views 50 photos per day
- Photo: 2 MB original, 200 KB feed, 20 KB thumbnail
- Store for 10 years

Step 2: Traffic

Uploads: $50M \div 86,400 = 578 \text{ QPS}$ (peak: 1.7K QPS)

Views: $500M \times 50 = 25B/\text{day}$

Views QPS: $25B \div 86,400 = 289K \text{ QPS}$ (peak: 867K QPS)

Ratio: 500:1 (extremely read-heavy)

Step 3: Storage

Per photo:

- Original: 2 MB
- Feed size: 200 KB
- Thumbnail: 20 KB

- Total: 2.22 MB per photo

Daily: $50M \times 2.22 \text{ MB} = 111 \text{ TB}$
 Yearly: $111 \text{ TB} \times 365 = 40.5 \text{ PB}$
 10 years: 405 PB

With compression (save 30%):
 Actual: $405 \text{ PB} \times 0.7 = \sim 283 \text{ PB}$

Step 4: Bandwidth

Upload: $578 \text{ QPS} \times 2 \text{ MB} = 1.2 \text{ GB/s}$
 Download: $289K \text{ QPS} \times 200 \text{ KB} = 57.8 \text{ GB/s}$

With CDN (90% cache hit):
 - Origin: $57.8 \text{ GB/s} \times 0.1 = 5.8 \text{ GB/s}$
 - CDN: $57.8 \text{ GB/s} \times 0.9 = 52 \text{ GB/s}$

Step 5: Cache

Cache hot photos (20% of last 30 days):
 - $50M/\text{day} \times 30 \text{ days} = 1.5B \text{ photos}$
 - $1.5B \times 0.2 = 300M \text{ hot photos}$
 - $300M \times 200 \text{ KB} = 60 \text{ TB}$

Distributed: $60 \text{ TB} \div 100 \text{ servers} = 600 \text{ GB each}$

Example 3: Design WhatsApp

Step 1: Requirements

- 2B MAU, 1B DAU (50% active daily)
- Each user sends 40 messages/day
- Each user receives 40 messages/day
- Average message: 100 bytes
- Store for 1 year

Step 2: Traffic

Messages per day: $1B \times 40 = 40B$
 Message QPS: $40B \div 86,400 = 462K \text{ QPS}$
 Peak: $462K \times 3 = 1.4M \text{ QPS}$

This is WRITE-heavy (send/receive are both writes)

Step 3: Storage

Daily: $40B \times 100 \text{ bytes} = 4 \text{ TB}$

Yearly: $4 \text{ TB} \times 365 = 1.46 \text{ PB}$

With metadata (timestamp, read status, etc.):

- +50 bytes per message
- $40B \times 150 \text{ bytes} = 6 \text{ TB/day}$
- Yearly: 2.19 PB

With images/videos (10% of messages):

- $40B \times 0.1 = 4B$ media messages
- Average: 1 MB per media
- $4B \times 1 \text{ MB} = 4 \text{ PB/day}$
- Yearly: $4 \text{ PB} \times 365 = 1,460 \text{ PB} = 1.46 \text{ exabytes}$

Step 4: Bandwidth

Text: $462K \text{ QPS} \times 100 \text{ bytes} = 46 \text{ MB/s}$

Media: $462K \times 0.1 \times 1 \text{ MB} = 46 \text{ GB/s}$

Total: ~46 GB/s

Note: This is why WhatsApp compresses images aggressively!

With compression (80% reduction):

Actual: $46 \text{ GB/s} \times 0.2 = \sim 9 \text{ GB/s}$

Example 4: Design Uber

Step 1: Requirements

- 50M daily rides
- Each ride: 1 driver, 1 rider
- Location updates every 3 seconds during ride
- Average ride: 15 minutes
- Store location history for 1 year

Step 2: Traffic

Location Updates:

- Updates per ride: $(15 \text{ min} \times 60 \text{ sec}) / 3 \text{ sec} = 300 \text{ updates}$

- Updates per day: $50M \text{ rides} \times 300 \times 2 \text{ (driver + rider)} = 30B \text{ updates}$
- QPS: $30B \div 86,400 = 347K \text{ QPS}$
- Peak (during rush hour, 3x): $\sim 1M \text{ QPS}$

This is WRITE-HEAVY (real-time location tracking)

Step 3: Storage

Per location update:

- User ID: 8 bytes
- Latitude: 8 bytes
- Longitude: 8 bytes
- Timestamp: 8 bytes
- Total: 32 bytes

Daily: $30B \times 32 \text{ bytes} = 960 \text{ GB per day}$

Yearly: $960 \text{ GB} \times 365 = 350 \text{ TB}$

With trip metadata (pickup, dropoff, fare, etc.):

- Additional 1 KB per trip
- $50M \times 1 \text{ KB} = 50 \text{ GB per day}$
- Yearly: 18 TB

Total: $\sim 368 \text{ TB per year}$

Step 4: Memory (Real-time Matching)

Active drivers at any time:

- Assume 5M drivers online (10% of daily drivers)
- Each driver location: 32 bytes
- Total: $5M \times 32 \text{ bytes} = 160 \text{ MB}$

Cache all active drivers in memory:

- Easily fits in single Redis instance!
- Use Redis Geo data structure (GEOADD, GEORADIUS)

Active riders:

- 5M riders waiting (matching period)
- Total: 160 MB

Combined: 320 MB (tiny!)

Quick Calculation Shortcuts

Shortcut 1: Daily to QPS

Formula: Daily Requests \div 100K \approx QPS

Why: 86,400 \approx 100,000 (close enough for estimates)

Examples:

$$10M \text{ requests/day} \div 100K = 100 \text{ QPS}$$

$$1B \text{ requests/day} \div 100K = 10K \text{ QPS}$$

$$100B \text{ requests/day} \div 100K = 1M \text{ QPS}$$

Shortcut 2: Monthly to Daily

Formula: Monthly \div 30 = Daily

Example:

$$3B \text{ requests/month} \div 30 = 100M \text{ requests/day}$$

$$100M \div 100K = 1K \text{ QPS}$$

Shortcut 3: Storage Growth

Formula: Daily \times 365 \times Years = Total

Example:

$$10 \text{ TB/day} \times 365 \times 5 = 18,250 \text{ TB} = \sim 18 \text{ PB}$$

Round to 20 PB for safety margin

Shortcut 4: Bandwidth from QPS

Formula: QPS \times Object Size

Example:

$$10K \text{ QPS} \times 100 \text{ KB} = 1,000,000 \text{ KB/s} = 1 \text{ GB/s}$$

Round to 1 GB/s for simplicity

Shortcut 5: 80-20 Rule for Caching

Formula: Total Data \times 0.2 = Cache Size

Example:

100 TB total data
Cache: $100 \text{ TB} \times 0.2 = 20 \text{ TB}$
Achieves ~80% cache hit ratio

Common Estimation Patterns

Pattern 1: Social Media (Twitter, Instagram)

Characteristics:

- Read-heavy (100:1 to 1000:1 ratio)
- Small writes (text, metadata)
- Large reads (media files)

Formula:

$$\begin{aligned}\text{Write QPS} &= \text{DAU} \times \text{Posts per User} / 86,400 \\ \text{Read QPS} &= \text{DAU} \times \text{Feed Loads} \times \text{Items per Load} / 86,400 \\ \text{Storage} &= \text{Posts per Day} \times \text{Size} \times \text{Retention Period}\end{aligned}$$

Pattern 2: Messaging (WhatsApp, Telegram)

Characteristics:

- Balanced read/write (1:1 ratio)
- Small messages (bytes to KB)
- High frequency (billions per day)

Formula:

$$\begin{aligned}\text{Message QPS} &= \text{DAU} \times \text{Messages per User} / 86,400 \\ \text{Storage} &= \text{Messages per Day} \times \text{Size} \times \text{Retention} \\ \text{Real-time Connections} &= \text{DAU} \text{ (for WebSocket)}\end{aligned}$$

Pattern 3: Video Streaming (YouTube, Netflix)

Characteristics:

- Extremely read-heavy (1000:1)
- Large files (GB per video)
- CDN is mandatory

Formula:

Upload QPS = Uploads per Day / 86,400
View QPS = Views per Day / 86,400
Storage = Videos per Day × Size × Retention
Bandwidth = View QPS × Bitrate / 8 (bits to bytes)

Pattern 4: E-Commerce (Amazon, eBay)

Characteristics:

- Read-heavy (browse > buy)
- Need ACID (transactions)
- Spike during sales events

Formula:

Browse QPS = DAU × Page Views / 86,400
Purchase QPS = Orders per Day / 86,400
Storage = Products × Size + Orders × Size
Peak Factor = 10x (during Black Friday)

Interview Example Walkthroughs

Walkthrough 1: "Design a News Feed"

Interviewer: "Design a news feed like Facebook. 1 billion users."

You:

"Let me estimate the scale:

Assumptions:

- 1B total users, 500M DAU (50% active)
- Each user posts 0.5 times per day
- Each user refreshes feed 10 times per day
- Each feed load shows 20 posts
- Average post: 1 KB (text + metadata)
- Store for 2 years

Traffic:

Posts: $500M \times 0.5 = 250M/\text{day} \div 100K = 2.5K$ write QPS

Reads: $500M \times 10 \times 20 = 100B$ reads/day $\div 100K = 1M$ read QPS

Ratio: 400:1 (very read-heavy)

Storage:

Posts: $250M/\text{day} \times 1 \text{ KB} = 250 \text{ GB}/\text{day}$

2 years: $250 \text{ GB} \times 365 \times 2 = 182 \text{ TB}$

Round to 200 TB with indexes

Bandwidth:

Outgoing: $1\text{M QPS} \times 1\text{ KB} = 1\text{ GB/s}$

Cache:

Cache 20% of 500M users' feeds: 100M users

Each feed: $20\text{ posts} \times 1\text{ KB} = 20\text{ KB}$

Cache size: $100\text{M} \times 20\text{ KB} = 2\text{ TB}$ across 20 servers

Database:

- PostgreSQL/MySQL for users (ACID)
 - Cassandra for posts (write-heavy, 2.5K QPS)
 - Redis for feed cache ($1\text{M read QPS} \times 80\% = 800\text{K QPS}$ served from cache)
- "

Walkthrough 2: "Design a URL Shortener"

Interviewer: "Design bit.ly. 100M URLs created per month."

You:

"Let me calculate the scale:

Assumptions:

- 100M new URLs per month
- 10B redirects per month (100:1 read-to-write)
- Average URL: 200 bytes
- Store for 10 years

Traffic:

Writes: $100\text{M/month} \div 30 \div 100\text{K} = 33\text{ QPS}$ (round to 40 QPS)

Reads: $10\text{B/month} \div 30 \div 100\text{K} = 3,333\text{ QPS}$ (round to 4K QPS)

Peak: 40 write QPS, 12K read QPS (3x)

Storage:

URLs: $100\text{M} \times 12 \times 10 = 12\text{B URLs}$

Size: $12\text{B} \times 200\text{ bytes} = 2.4\text{ TB}$

With analytics: +6 TB (assume 10 clicks per URL)

Total: ~9 TB over 10 years

Cache:

Cache hot URLs (20% of last 30 days):

- $100\text{M/month} = 3.3\text{M/day}$
- $3.3\text{M} \times 30\text{ days} = 100\text{M URLs}$
- $100\text{M} \times 0.2 = 20\text{M hot URLs}$
- $20\text{M} \times 200\text{ bytes} = 4\text{ GB}$

This is tiny! Cache everything recent!

Bandwidth:

Incoming: 40 QPS × 200 bytes = 8 KB/s
Outgoing: 4K QPS × 200 bytes = 800 KB/s

Database:

- PostgreSQL (simple, ACID, 4K QPS easily handled)
- Redis cache (80% hit ratio, < 1ms latency)
- Cassandra for analytics (time-series clicks)
- "

Walkthrough 3: "Design Dropbox"

Interviewer: "Design file sync service. 50M users."

You:

"Let me estimate:

Assumptions:

- 50M registered users, 10M DAU (20% active)
- Each user has 200 files (average)
- Average file: 1 MB
- Each user syncs 5 files per day

Traffic:

Syncs: $10M \times 5 = 50M$ syncs/day $\div 100K = 500$ QPS

Peak: $500 \times 3 = 1.5K$ QPS

Storage:

Total files: $50M \text{ users} \times 200 \text{ files} = 10 \text{ billion files}$

Total storage: $10B \times 1 \text{ MB} = 10 \text{ PB}$

With 3 replicas (redundancy):

Total: $10 \text{ PB} \times 3 = 30 \text{ PB}$

Bandwidth:

Upload/Download: $500 \text{ QPS} \times 1 \text{ MB} = 500 \text{ MB/s}$

Peak: 1.5 GB/s

Metadata cache:

- File paths, versions, permissions
- $10B \text{ files} \times 1 \text{ KB metadata} = 10 \text{ TB}$
- Cache 20%: 2 TB across servers
- "

Quick Reference Cheat Sheet

Time Conversions

1 second	= 1,000 ms
1 minute	= 60 seconds
1 hour	= 3,600 seconds
1 day	= 86,400 seconds (~100K for rough estimates)
1 month	= 2,592,000 seconds (~2.5M)
1 year	= 31,536,000 seconds (~30M)

Storage Conversions

1 KB	= 1,024 bytes	(~1 thousand)
1 MB	= 1,024 KB	(~1 million bytes)
1 GB	= 1,024 MB	(~1 billion bytes)
1 TB	= 1,024 GB	(~1 trillion bytes)
1 PB	= 1,024 TB	(~1 quadrillion bytes)
1 EB	= 1,024 PB	(~1 quintillion bytes)

QPS Estimation Shortcuts

Daily Operations → QPS:

1M/day	= 10 QPS
10M/day	= 100 QPS
100M/day	= 1K QPS
1B/day	= 10K QPS
10B/day	= 100K QPS
100B/day	= 1M QPS

Typical Ratios

Read-to-Write Ratios:

- Social media: 100:1 to 1000:1
- Messaging: 1:1
- E-commerce: 10:1 to 100:1
- Analytics: Write-heavy (opposite)

Cache Hit Ratios:

- Well-designed: 80–90%
- Average: 70–80%
- Poor: < 70%

Peak to Average:

- Normal: 2–3x
- Social (viral): 5–10x
- E-commerce (Black Friday): 10–20x

Storage Per Object

User profile: 1–2 KB
Tweet/Post: 280 bytes – 1 KB
Email: 50–100 KB
Small image: 50–100 KB
Medium image: 200–500 KB
Large image: 1–3 MB
Short video (1 min): 10–50 MB
HD video (1 min): 50–100 MB

Interview Tips & Tricks

Tip 1: Start with Round Numbers

Example:

Interviewer: "500 million daily active users"
You: "So roughly 500M DAU, I'll round to 500M for calculations"

Interviewer: "Each user posts 2.3 times per day"
You: "I'll approximate that as 2 posts per day"

Why: Makes math easier, shows you understand precision isn't critical

Tip 2: Show Units in Every Step

Bad:

$$\begin{aligned} 100 \times 200 &= 20,000 \\ 20,000 \times 365 &= 7,300,000 \end{aligned}$$

Good:

$$\begin{aligned} 100M \text{ photos} \times 200 \text{ KB} &= 20,000,000,000 \text{ KB} = 20 \text{ TB per day} \\ 20 \text{ TB per day} \times 365 \text{ days} &= 7,300 \text{ TB} = 7.3 \text{ PB per year} \end{aligned}$$

Tip 3: Sanity Check Your Numbers

After calculating, ask yourself:

- "Does this make sense?"

- "Is 1 PB reasonable for Instagram photos? Yes!"
- "Is 10 EB reasonable for Twitter text? No, that's too high!"

Common Sanity Checks:

Twitter text storage should be < 100 TB (tiny)
 Instagram photo storage should be > 10 PB (huge)
 WhatsApp messages should be > 1 PB but < 10 PB
 YouTube videos should be > 1 EB (massive)

Tip 4: Mention Optimizations

After calculations, always mention:

Storage optimization:
 "With compression, we can reduce this by 30-50%"

Bandwidth optimization:
 "With CDN caching at 90% hit ratio, origin only serves 10%"

Memory optimization:
 "Using the 80-20 rule, we only need to cache 20% of data"

Tip 5: Compare to Real Systems

Say things like:

"This is similar to Twitter's scale"
 "Instagram actually uses ~100 PB of storage"
 "Netflix serves petabytes per day from CDN"
 "WhatsApp handles billions of messages"

Common Mistakes to Avoid

✗ Mistake 1: Being Too Precise

Bad: "387,654,321 requests per day"

Good: "~400 million requests per day"

Why: Precision gives false confidence, rough numbers are sufficient

✗ Mistake 2: Forgetting Peak Traffic

Bad: "10K QPS average"

Good: "10K QPS average, 30K QPS peak (3x factor)"

Why: Systems must handle peak load, not average

✖ Mistake 3: Not Showing Units

Bad: "Storage is 100"

Good: "Storage is 100 TB"

Why: Interviewer doesn't know if you mean MB, GB, or TB

✖ Mistake 4: Ignoring Media

Bad: "Twitter stores 1 TB of tweets"

Good: "Twitter stores 50 GB of text + 10 PB of images"

Why: Media dominates storage, don't ignore it

✖ Mistake 5: Forgetting Replication

Bad: "Need 10 TB storage"

Good: "Need 10 TB \times 3 replicas = 30 TB storage"

Why: Replication is standard for reliability

Practice Problems

Problem 1: Design Google Photos

Given:

- 1B users, 500M DAU
- Each user uploads 5 photos per day
- Each photo: 3 MB
- Each user views 50 photos per day
- Store for lifetime (20 years)

Your turn! Calculate:

1. Write QPS?
2. Read QPS?
3. Total storage for 20 years?
4. Bandwidth required?
5. Cache size?

► Click for Answer

Traffic:

- Uploads: $500M \times 5 = 2.5B/\text{day} \div 100K = 25K$ write QPS
- Views: $500M \times 50 = 25B/\text{day} \div 100K = 250K$ read QPS
- Peak: 75K write, 750K read QPS

Storage:

- Daily: $2.5B \times 3 \text{ MB} = 7.5 \text{ PB}/\text{day}$
- Yearly: $7.5 \text{ PB} \times 365 = 2,737 \text{ PB} = 2.7 \text{ EB}$
- 20 years: $2.7 \text{ EB} \times 20 = 54 \text{ EB}$
- With compression (50%): 27 EB

Bandwidth:

- Upload: $25K \text{ QPS} \times 3 \text{ MB} = 75 \text{ GB/s}$
- Download: $250K \text{ QPS} \times 3 \text{ MB} = 750 \text{ GB/s}$
- With CDN (95% cache): 37.5 GB/s origin

Cache:

- Hot photos (20% of last 90 days):
- $2.5B \times 90 \times 0.2 = 45B$ photos
- $45B \times 3 \text{ MB} = 135 \text{ PB}$ (too large to cache originals!)
- Cache thumbnails instead: $45B \times 200 \text{ KB} = 9 \text{ PB}$
- Still large! Cache metadata only: $45B \times 1 \text{ KB} = 45 \text{ TB}$

Problem 2: Design Spotify**Given:**

- 400M users, 100M DAU
- Each user streams 10 songs per day
- Each song: 5 MB (compressed)
- Catalog: 100M songs
- Store catalog forever

Calculate: Traffic, Storage, Bandwidth, Cache

► Click for Answer

Traffic:

- Streams: $100M \times 10 = 1B/\text{day} \div 100K = 10K$ QPS
- Peak: 30K QPS

Storage:

- Catalog: $100M \text{ songs} \times 5 \text{ MB} = 500 \text{ TB}$
- With multiple qualities: $500 \text{ TB} \times 3 = 1.5 \text{ PB}$

Bandwidth:

- Download: $10K \text{ QPS} \times 5 \text{ MB} = 50 \text{ GB/s}$
- With CDN (90% cache): 5 GB/s origin

Cache:

- Hot songs (20% of catalog): $100M \times 0.2 = 20M$ songs
- $20M \times 5 \text{ MB} = 100 \text{ TB}$
- Distributed: $100 \text{ TB} \div 100 \text{ servers} = 1 \text{ TB each}$

The Universal Template

Use this template for ANY system design question:

1. TRAFFIC (QPS):

Write QPS = Writes per Day $\div 100K$

Read QPS = Reads per Day $\div 100K$

Peak QPS = Average $\times 3$

Ratio = Reads / Writes

2. STORAGE:

Daily = Items per Day \times Size per Item

Total = Daily $\times 365 \times$ Years

With Replication = Total $\times 3$

3. BANDWIDTH:

Incoming = Write QPS \times Object Size

Outgoing = Read QPS \times Object Size

With CDN = Outgoing $\times (1 - \text{Cache Hit Ratio})$

4. MEMORY (Cache):

Cache Size = Total Data $\times 0.2$ (80–20 rule)

Servers = Cache Size $\div 100 \text{ GB}$

Cache Hit Ratio = ~80%

5. DATABASE:

If Read-Heavy \rightarrow SQL + Read Replicas + Cache

If Write-Heavy \rightarrow Cassandra/NoSQL

If < 1 TB \rightarrow Single SQL instance

If > 10 TB \rightarrow Sharding or NoSQL

Real Numbers from Production Systems

Twitter (Actual)

DAU: 400M

Tweets/day: 500M

QPS: ~6K write, ~300K read

Storage: ~100 PB

Bandwidth: ~1 TB/s (with CDN)

Instagram (Actual)

DAU: 500M
Photos/day: 95M
QPS: ~1K write, ~500K read
Storage: ~100 PB
Cache: Redis Cluster (50+ nodes)

WhatsApp (Actual)

DAU: 2B
Messages/day: 100B
QPS: ~1M write, ~1M read
Storage: ~10 PB
Erlang servers: 1000+

YouTube (Actual)

Videos watched/day: 5B
Upload: 500 hours/minute
Storage: ~1 EB (yes, exabyte!)
Bandwidth: ~1 PB/s globally
CDN: Critical infrastructure

Technology Selection Based on Scale Estimates

Database Selection Framework

PostgreSQL/MySQL - When Numbers Justify It

Use PostgreSQL when your calculations show:

Storage: < 10 TB
Write QPS: < 10,000
Read QPS: < 100,000 (with 5-10 read replicas)
Need: ACID transactions, complex queries, JOINs

Real Examples from This Guide:

TinyURL:

- ✓ Storage: 8 TB over 10 years (< 10 TB threshold)
- ✓ Write QPS: 40 (far below 10K limit)
- ✓ Read QPS: 4K (easily handled with 2-3 replicas)
- ✓ Needs: ACID for URL creation (prevent duplicates)

Justification:

"PostgreSQL handles our scale easily. With 8 TB storage, we're well under the 10–16 TB single instance limit. Our 40 write QPS is 0.4% of PostgreSQL's 10K capacity. With 3 read replicas, we can serve 30K QPS, but we only need 4K. PostgreSQL provides ACID guarantees we need for URL creation."

Configuration:

- 1 master (r5.2xlarge): 8 vCPU, 64 GB RAM
- 3 read replicas for HA
- Cost: ~\$800/month

Pastebin:

- ✓ Storage: 21 TB over 5 years
- ✓ Write QPS: 115
- ✓ Read QPS: 12K
- ✓ Features: TTL (auto-delete expired pastes)

Justification:

"21 TB fits in sharded PostgreSQL (3 shards × 7 TB each) or single instance with cleanup. 115 write QPS is 1% of capacity. 12K read QPS handled by 5 replicas. PostgreSQL's TTL support (DELETE WHERE created_at < NOW() - INTERVAL '24 hours') perfectly matches our expire requirement."

Configuration:

- Option 1: 1 master + 5 read replicas (~\$1,500/month)
- Option 2: 3-shard setup (~\$2,000/month) for future growth

Cassandra/NoSQL - When Numbers Demand Scale

Use Cassandra when your calculations show:

Storage: > 10 TB (need horizontal scaling)
 Write QPS: > 10,000 (need distributed writes)
 OR
 Storage: > 100 TB (PostgreSQL sharding too complex)
 Data Pattern: Time-series, append-heavy
 Consistency: Eventual acceptable

Real Examples from This Guide:

Twitter Tweets:

- ✗ Storage: 44 PB (4,400x PostgreSQL limit!)
- ✓ Write QPS: 7,000 (approaches PostgreSQL limit)
- ✓ Pattern: Time-series (tweets sorted by timestamp)
- ✓ Consistency: Eventual OK (tweets appear within 1 sec)

Justification:

"44 PB storage requires distributed database. PostgreSQL would need $44,000 \text{ GB} \div 16 \text{ TB} = 2,750$ shards – impossible to manage. Cassandra scales linearly: $44 \text{ PB} \div 2 \text{ TB per node} = 22,000$ nodes. Each node adds 100 write QPS, giving us 2.2M QPS capacity for our 7K need. Cassandra's write-optimized LSM tree perfect for tweet insertion pattern."

Configuration:

- 22,000 nodes (in reality, fewer due to data lifecycle)
- Each: i3.2xlarge (8 vCPU, 61 GB RAM, 1.9 TB NVMe)
- Replication factor: 3
- Cost: ~\$300K/month (at scale pricing)

WhatsApp Messages:

- ✗ Storage: 1.46 EB = 1,460 PB (impossible for PostgreSQL)
- ✓ Write QPS: 1.4 million (140x PostgreSQL limit!)
- ✓ Pattern: Append-only messages
- ✓ Consistency: Eventual OK for message delivery

Justification:

"1.4M write QPS requires massive distribution. PostgreSQL max 10K writes = would need 140 masters! Cassandra: $1.4M \div 100 \text{ per node} = 14,000$ nodes minimum. Storage: $1.46 \text{ EB} \div 2 \text{ TB} = 750,000$ nodes. In practice, messages deleted after 30 days, so 50K nodes at ~30 PB. Cassandra's peer-to-peer architecture enables this scale."

Configuration:

- 50,000 nodes (with 30-day retention)
- Distributed globally (20 regions × 2,500 nodes)
- Cost: ~\$700K/month

Uber Location Updates:

- ✗ Storage: 2 PB (requires 125 PostgreSQL shards)
- ✓ Write QPS: 1 million (100x PostgreSQL limit)
- ✓ Pattern: Constant location writes every 3 seconds

- ✓ Query: Time-range lookups (trip replay)

Justification:

"1M location writes/sec impossible for PostgreSQL (10K max). Cassandra:
1M ÷ 1K per node = 1,000 nodes. Storage: 2 PB ÷ 2 TB = 1,000 nodes.
Perfect match! Cassandra's write path (MemTable → SSTable) optimized
for constant writes. Time-series partition key enables efficient trip
replay queries."

Configuration:

- 1,000 nodes globally
- Partition key: (driver_id, date)
- Clustering key: timestamp
- Cost: ~\$140K/month

MongoDB - The Middle Ground

Use MongoDB when your calculations show:

Storage: 1–100 TB (too big for single PostgreSQL, not big enough for Cassandra)
Write QPS: 1,000–50,000 (moderate)
Schema: Flexible/varying structure
Need: Document queries, indexes

Reddit Example:

- ✓ Storage: 250 TB (would need 16 PostgreSQL shards)
- ✓ Write QPS: 30,000 (3x PostgreSQL limit)
- ✓ Schema: Posts vary (text, images, videos, polls)
- ✓ Queries: Complex (find by subreddit, user, time)

Justification:

"250 TB too large for single PostgreSQL (need 16 shards). But 30K write
QPS doesn't justify Cassandra's complexity. MongoDB offers middle
ground:
16 shards × 15 TB each. Flexible schema handles varying post types.
Built-in sharding simpler than manual PostgreSQL sharding."

Configuration:

- 16 shards × 3 replica sets = 48 nodes
- Each: r5.2xlarge (8 vCPU, 64 GB RAM)
- Cost: ~\$9,000/month

Cache Technology Selection

Redis - The Default Choice

Use Redis (almost always) when:

Cache size: Any (MB to TB)
QPS: Up to 1M per instance
Need: Data structures, persistence, Pub/Sub, TTL

Redis Sizing by Numbers:

Cache < 1 GB:

- Single Redis instance (r5.large: 16 GB RAM)
- Handles: 100K QPS
- Cost: ~\$140/month
- Examples: Rate Limiter (400 MB), Parking Lot (183 MB)

Cache 1–100 GB:

- 2–5 Redis instances (HA cluster)
- Handles: 200–500K QPS
- Cost: ~\$280–700/month
- Examples: Reddit (20 GB), Autocomplete (10 GB)

Cache 100 GB – 1 TB:

- Redis Cluster (10–20 nodes)
- Handles: 1–2M QPS
- Cost: ~\$1,400–2,800/month
- Examples: Twitter timelines (1.1 TB), Instagram (60 TB needs 600 nodes)

Cache > 1 TB:

- Redis Cluster (50–1000 nodes)
- Handles: 5–50M QPS
- Cost: ~\$7,000–140,000/month
- Examples: Instagram hot photos (60 TB), LinkedIn (7.68 TB)

Real Examples with Justifications:

TinyURL Cache (2.5 GB):

Calculations show:

- Cache size: 20M hot URLs × 125 bytes = 2.5 GB
- QPS: 12K reads (with 95% cache hit)
- Latency requirement: < 10ms

Redis Configuration:

- Single instance: cache.r5.large (13.5 GB RAM)
- Headroom: 11 GB free (440% capacity)

- Performance: 100K QPS capacity vs 12K need (8x headroom)
- Latency: 0.5–1ms average (meets < 10ms requirement)

Cost Justification:

- Redis: \$140/month
- Alternative (Memcached): \$130/month (save \$10)
- Benefit of Redis: Persistence, TTL, sorted sets
- Decision: Redis (minimal cost difference, more features)

Twitter Timeline Cache (1.1 TB):

Calculations show:

- Cache size: 80M users × 14 KB timeline = 1.1 TB
- QPS: 1.4M reads
- Cache hit requirement: 80%+ (save DB from 1.1M QPS)

Redis Cluster Configuration:

- 20 nodes × 64 GB RAM = 1.28 TB capacity
- Per node: cache.r5.4xlarge (52 GB usable)
- Each handles: 70K QPS
- Total capacity: 1.4M QPS ✓

Cost Analysis:

- Redis Cluster: \$5,600/month
- Saves: Database capacity for 1.1M QPS
- Alternative DB scaling: \$50K/month (20x more expensive!)
- ROI: Cache pays for itself 9x over

Configuration Details:

- Partition: 16,384 hash slots across 20 nodes
- Replication: Each node has 1 replica (40 nodes total)
- Eviction: LRU (Least Recently Used)
- Total cost: \$11,200/month (with replicas)

Rate Limiter Cache (400 MB):

Calculations show:

- Active counters: 10M users × 20 bytes = 200 MB
- With sliding window: 200 MB × 2 = 400 MB
- QPS: 36K checks/sec
- Latency requirement: < 1ms (critical!)

Redis Configuration:

- Single instance: cache.r5.large (13.5 GB)
- Memory usage: 400 MB = 3% of capacity
- Operations: INCR, EXPIRE (atomic, fast)
- Performance: Sub-millisecond latency ✓

Why Redis over Memcached:

- Need atomic INCR (prevent race conditions)
- Need TTL per key (auto-cleanup windows)
- Memcached lacks both features
- Cost: Same (\$140/month)
- Decision: Redis (required features, same cost)

Memcached - Rare Use Case

Only use Memcached when:

- Pure key-value (no data structures needed)
- Multi-threading benefit significant (>1M QPS per instance)
- No persistence needed
- No TTL per-key needed

Realistic scenario: < 1% of systems need Memcached over Redis

CDN Selection Based on Bandwidth Numbers

When CDN is Mandatory

Use CDN when calculations show:

Bandwidth: > 1 GB/s outgoing
Content: Static (images, videos, CSS, JS)
Users: Global distribution
Pattern: Read-heavy (can cache)

Cost-Benefit Analysis:

Instagram Example (57 GB/s bandwidth):

Without CDN (Origin Only):

- Bandwidth: $57 \text{ GB/s} \times \$0.08/\text{GB} = \$4.56/\text{sec}$
- Per hour: $\$4.56/\text{sec} \times 3600 \text{ sec} = \$16,416$
- Per month: $\$16,416 \times 30 \text{ days} = \11.82 million

With CDN (90% cache hit):

- Origin: $5.7 \text{ GB/s} \times \$0.08/\text{GB} = \$0.456/\text{sec} = \$328\text{K}/\text{month}$
- CDN: $51.3 \text{ GB/s} \times \$0.02/\text{GB} = \$1.026/\text{sec} = \$739\text{K}/\text{month}$
- Total: $\$1.067 \text{ million/month}$

Savings: $\$11.82\text{M} - \$1.067\text{M} = \$10.75\text{M}/\text{month}$ (91% cost reduction!)

Additional Benefits:

- Latency: 200ms (origin) → 20ms (edge) = 10x faster
- Origin load: 57 GB/s → 5.7 GB/s = 90% reduction
- Global reach: 400+ edge locations vs single origin
- DDoS protection: Included

Decision: CDN absolutely mandatory, not optional

Netflix Example (8.3 TB/s bandwidth):

Without CDN:

- Cost: $8,300 \text{ GB/s} \times \$0.08 = \$664/\text{sec} = \$1.9 \text{ billion/month}$
- Impossible! Can't serve from origin.

With Custom CDN (95% from edge):

- Origin: $415 \text{ GB/s} \times \$0.08 = \$33.2/\text{sec} = \86M/month
- CDN (Open Connect): Custom infrastructure
- Netflix deploys servers in ISP facilities (free transit)
- Estimated cost: \$10M/month (vs \$1.9B without CDN)

Savings: Must use CDN to be economically viable

Netflix's Approach:

- Open Connect: Custom CDN appliances
- Pre-populate hot content (new releases)
- 90% of traffic served within ISP network (no internet transit)
- Result: Streams 1+ billion hours/day profitably

When CDN is Optional

Small bandwidth (< 100 MB/s):

Rate Limiter (36 KB/s):

- Cost without CDN: $\$0.0000288/\text{sec} = \$75/\text{month}$
- Cost with CDN: \$50/month (minimum) + setup complexity
- Decision: Skip CDN, not worth it

Parking Lot (1 MB/s):

- Cost without CDN: $\$0.08/\text{sec} = \$2,100/\text{month}$
- Cost with CDN: \$400/month
- Savings: \$1,700/month
- But: Local system (single facility), CDN adds latency
- Decision: Skip CDN for local-only service

Message Queue Selection

Kafka - For High Throughput

Use Kafka when calculations show:

Events: > 10K/sec
Need: Durability (no message loss)
Need: Replay capability
Pattern: Multiple consumers

Real Examples:

Twitter Fanout (7K tweets/sec):

Requirements:
- 7K tweet events/sec
- Each tweet → fanout to avg 500 followers
- Total: $7K \times 500 = 3.5M$ fanout operations/sec
- Need durability (can't lose tweets)
- Need ordering per user

Kafka Configuration:
- 10 brokers (each handles 100K msg/sec)
- $3.5M \div 100K = 35$ brokers minimum
- With replication (factor 3): 105 brokers
- Cost: ~\$15,000/month

Why Kafka over alternatives:
- RabbitMQ: Max 50K/sec/node (need 70 nodes, similar cost)
- SQS: $\$0.40/\text{million} = 3.5M \times 86,400 \times 30 = \$3.6M/\text{month}$ (240x more!)
- Redis Pub/Sub: No durability (lose messages on crash)

Decision: Kafka for durability + performance at scale

WhatsApp Messages (1.4M/sec):

Requirements:
- 1.4M message events/sec
- Must guarantee delivery
- Need message ordering per conversation
- Global distribution

Kafka Configuration:
- 100 brokers globally (20 regions × 5 brokers)
- Each region: 70K msg/sec
- Replication: 3x within region
- Cost: ~\$140K/month

Alternative (AWS SQS):

- Cost: $1.4M \times 86,400 \times 30 \times \$0.40/\text{million} = \$1.45M/\text{month}$
- 10x more expensive!
- Kafka justified by cost savings alone

SQS/Cloud Queue - For Moderate Load

Use SQS when:

Events: < 10K/sec

Pattern: Simple queue (no replay needed)

Need: Managed service (no ops overhead)

Cost Comparison:

Notification System (36K events/sec):

Kafka:

- 5 brokers $\times \$140/\text{month} = \$700/\text{month}$
- Operational overhead: $20 \text{ hours/month} \times \$100/\text{hr} = \$2,000$
- Total: \$2,700/month

SQS:

- $36K \times 86,400 \times 30 = 93 \text{ billion requests/month}$
- Cost: $93B \times \$0.40/\text{million} = \$37,200/\text{month}$
- No operational overhead
- Total: \$37,200/month

Decision: Kafka (13x cheaper including ops cost)

Low-volume system (100 events/sec):

Kafka:

- Minimum 3 brokers: \$420/month
- Ops overhead: \$2,000/month
- Total: \$2,420/month

SQS:

- $100 \times 86,400 \times 30 = 259 \text{ million requests/month}$
- Cost: $259M \times \$0.40/\text{million} = \$104/\text{month}$
- Total: \$104/month

Decision: SQS (23x cheaper for low volume)

Object Storage Selection

S3/Cloud Storage - For Media Files

Use S3 when calculations show:

Storage: > 100 GB of media files
Type: Images, videos, documents
Need: 11-nines durability
Access: Random (not sequential)

Cost Analysis with Lifecycle Policies:

Instagram Photos (283 PB):

Storage Tiers:

- Hot (< 30 days): $50M \times 30 \text{ days} = 1.5B \text{ photos}$
Size: $1.5B \times 200 \text{ KB} = 300 \text{ TB}$
Cost: $300 \text{ TB} \times \$23/\text{TB} = \$6,900/\text{month}$ (S3 Standard)
- Warm (30–90 days): $50M \times 60 \text{ days} = 3B \text{ photos}$
Size: $3B \times 200 \text{ KB} = 600 \text{ TB}$
Cost: $600 \text{ TB} \times \$12.50/\text{TB} = \$7,500/\text{month}$ (S3 IA)
- Cold (> 90 days): $50M \times 3,560 \text{ days} = 178B \text{ photos}$
Size: $178B \times 200 \text{ KB} = 35.6 \text{ PB}$
Cost: $35.6 \text{ PB} \times \$4/\text{TB} = \$142,400/\text{month}$ (S3 Glacier)

Total monthly cost: \$156,800/month

Total storage: 283 PB

Without lifecycle policies:

- All in Standard: $283 \text{ PB} \times \$23/\text{TB} = \$6,509,000/\text{month}$
- Savings: \$6.35M/month (97% reduction!)

Justification:

"98% of photo views are < 90 days old. By using lifecycle policies, we move 126 PB to warm and 35.6 PB to cold storage. This reduces costs by 97% while maintaining instant access to hot photos."

Netflix Catalog (300 TB):

All videos in S3 Standard:

- $300 \text{ TB} \times \$23/\text{TB} = \$6,900/\text{month}$
- Need instant access for streaming
- Can't use lifecycle (all content frequently accessed)

Additional copies:

- 4 resolutions \times 300 TB = 1.2 PB
- Cost: 1.2 PB \times \$23/TB = \$27,600/month

CDN pre-population:

- Popular content (20%): 240 GB at each edge
- 100 edges \times 240 GB = 24 TB
- Cost: Included in CDN edge cache

Total S3 cost: \$27,600/month for origin storage

Reasonable for \$30B/year revenue company

Complete Technology Stack Examples

Small Scale (< 1M users, < 10K QPS)

Example: Parking Lot System

Numbers:

- DAU: 5M vehicles
- QPS: 519 peak
- Storage: 11 TB

Technology Stack:

- ✓ Database: PostgreSQL (single instance)
Why: 11 TB fits, 519 QPS trivial
Instance: db.r5.xlarge (4 vCPU, 32 GB)
Cost: \$350/month

- ✓ Cache: Redis (single instance)
Why: 183 MB cache
Instance: cache.r5.large (13.5 GB)
Cost: \$140/month

- ✓ Web Servers: 5 instances
Why: 519 QPS \div 100 per server = 5.19
Instance: t3.medium (2 vCPU, 4 GB) \times 5
Cost: \$250/month

- ✗ CDN: Not needed
Why: 1 MB/s bandwidth, local facility only

- ✗ Message Queue: Not needed
Why: Synchronous flow sufficient

Total Cost: \$740/month

Scaling headroom: 10x (can handle 5K QPS)

Medium Scale (10-100M users, 10-100K QPS)

Example: Reddit

Numbers:

- DAU: 50M
- Read QPS: 88K peak
- Write QPS: 30K peak
- Storage: 250 TB

Technology Stack:

- ✓ Database: MongoDB (sharded)
 - Why: 250 TB too large for single PostgreSQL
30K write QPS exceeds single instance
 - Configuration: 16 shards × 3 replicas = 48 nodes
 - Instance: r5.2xlarge per node
 - Cost: \$9,000/month
 - ✓ Cache: Redis Cluster
 - Why: 20 GB cache, 88K read QPS
 - Configuration: 5 nodes
 - Instance: cache.r5.xlarge per node
 - Cost: \$700/month
 - ✓ Search: Elasticsearch
 - Why: Full-text search on posts/comments
 - Configuration: 10 nodes
 - Cost: \$2,000/month
 - ✓ CDN: CloudFront
 - Why: Image content, reduce latency
 - Bandwidth: 5 GB/s × 30% CDN-worthy = 1.5 GB/s
 - Cost: \$1,500/month
 - ✓ Queue: Kafka
 - Why: Vote aggregation, comment fanout
 - Configuration: 5 brokers
 - Cost: \$700/month
 - ✓ Web Servers: 50 instances
 - Why: 88K QPS ÷ 2K per server = 44 servers
 - Instance: c5.2xlarge × 50
 - Cost: \$6,000/month
- Total Cost: \$19,900/month
User cost: \$0.40 per user per month
Revenue needed: \$0.50+ per user for profitability

Large Scale (100M-1B users, 100K-1M QPS)

Example: Twitter

Numbers:

- DAU: 400M
- Read QPS: 1.4M peak
- Write QPS: 7K peak
- Storage: 44 PB

Technology Stack:

- ✓ Database (Tweets): Cassandra
 - Why: 44 PB storage, write-optimized
 - Configuration: 22,000 nodes (in reality ~5K with lifecycle)
 - Instance: i3.2xlarge per node
 - Cost: \$300K/month (negotiated volume pricing)
- ✓ Database (Users): PostgreSQL
 - Why: 800 GB, need ACID
 - Configuration: 1 master + 10 read replicas
 - Instance: db.r5.4xlarge
 - Cost: \$2,500/month
- ✓ Cache: Redis Cluster
 - Why: 1.1 TB timeline cache
 - Configuration: 20 master + 20 replica = 40 nodes
 - Instance: cache.r5.4xlarge per node
 - Cost: \$11,200/month
- ✓ Search: Elasticsearch
 - Why: Tweet search, trending topics
 - Configuration: 100 nodes
 - Cost: \$20,000/month
- ✓ CDN: CloudFront
 - Why: 19 GB/s bandwidth for media
 - Configuration: Global distribution
 - Cost: \$50,000/month (with 95% cache hit)
- ✓ Queue: Kafka
 - Why: Tweet fanout, timeline generation
 - Configuration: 50 brokers globally
 - Cost: \$7,000/month
- ✓ Web Servers: 300 instances
 - Why: $1.4M \text{ QPS} \div 5K \text{ per server} = 280$
 - Instance: c5.4xlarge $\times 300$
 - Cost: \$75,000/month
- ✓ Object Storage (S3)
 - Why: Image/video storage (20 PB with lifecycle)
 - Cost: \$50,000/month

Total Cost: \$515,700/month = \$6.2M/year

DAU: 400M
Cost per DAU: \$1.29/month
Revenue needed: \$2+/user/month for profitability
Twitter actual revenue: \$5+/user (ads + premium)

Extreme Scale (1B+ users, 1M+ QPS)

Example: WhatsApp

Numbers:

- DAU: 1B
- Message QPS: 1.4M peak
- Storage: 30 PB (with 30-day retention)
- Connections: 500M concurrent WebSocket

Technology Stack:

- ✓ Database: Cassandra
 - Why: 1.4M write QPS, 30 PB storage
 - Configuration: 50,000 nodes globally
 - Instance: i3.2xlarge per node
 - Cost: \$700K/month
- ✓ Cache: Redis Cluster
 - Why: 2.5 TB (presence, sessions, queues)
 - Configuration: 200 nodes
 - Cost: \$28,000/month
- ✓ WebSocket Servers: 10,000 instances
 - Why: 500M connections ÷ 50K per server
 - Instance: c5.xlarge × 10,000
 - Cost: \$150,000/month
- ✓ Queue: Kafka
 - Why: 1.4M message events/sec
 - Configuration: 100 brokers (20 regions)
 - Cost: \$140,000/month
- ✓ Object Storage: S3
 - Why: Media files (aggressive lifecycle)
 - Storage: 1 PB (after compression + lifecycle)
 - Cost: \$23,000/month
- ✓ Media Processing: GPU instances
 - Why: Image/video compression
 - Configuration: 1,000 instances
 - Cost: \$200,000/month

Total Cost: \$1.241 million/month = \$14.9M/year
DAU: 1B
Cost per DAU: \$0.01/month (!!)

Revenue: Free service (Meta subsidizes)

Actual monetization: Data for ad targeting across Meta products

Technology Decision Checklist

Before Recommending Technology, Verify:

Database Selection:

- What is storage requirement? (< 10 TB = PostgreSQL possible)
- What is write QPS? (< 10K = PostgreSQL possible)
- What is read QPS? (< 100K = PostgreSQL + replicas possible)
- Need ACID? (Yes = prefer PostgreSQL)
- Time-series data? (Yes = consider Cassandra)
- Storage > 100 TB? (Yes = Cassandra/DynamoDB)

Cache Selection:

- What is cache size? (Size Redis cluster accordingly)
- What is QPS? (1 Redis = 100K QPS)
- Need data structures? (Yes = Redis, No = Memcached option)
- Need persistence? (Yes = Redis, No = Memcached option)
- Need Pub/Sub? (Yes = Redis)
- Calculate cost: $\text{size} \div 64 \text{ GB} \times \$280/\text{month}$

CDN Decision:

- What is bandwidth? (< 100 MB/s = optional, > 1 GB/s = mandatory)
- Content type? (Static = yes, Dynamic = no)
- Global users? (Yes = CDN beneficial)
- Calculate savings: Compare origin vs CDN cost
- Current: $X \text{ GB/s} \times \$0.08 = \Y/month
- With CDN (90% hit): $X \times 0.1 \times \$0.08 + X \times 0.9 \times \$0.02 = \$Z/\text{month}$

Message Queue:

- What is event rate? (< 1K/sec = SQS, > 10K/sec = Kafka)
- Need durability? (Yes = Kafka, No = Redis Pub/Sub)
- Need replay? (Yes = Kafka)
- Calculate cost: $\text{Events/month} \times \$0.40/\text{million (SQS)} \text{ vs nodes} \times \140 (Kafka)

Interview Scoring Rubric

What Interviewers Look For:

Structured Approach (30%)

- Clear methodology
- Step-by-step calculations
- Organized presentation

Reasonable Assumptions (20%)

- States assumptions clearly
- Numbers are realistic
- Justifies choices

Correct Math (20%)

- Gets order of magnitude right
- Shows work
- Unit conversions correct

Identifies Bottlenecks (15%)

- "At 1M QPS, we need caching"
- "44 PB requires distributed storage"
- "This is write-heavy, need Cassandra"

Mentions Optimizations (15%)

- CDN for bandwidth
- Caching for reads
- Compression for storage

Technology Justification (NEW - Critical!)

- Explains WHY specific tech chosen
- Uses numbers to justify decisions
- Compares alternatives with costs
- Shows scalability path

Example Good Answer:

"Based on our calculations:

- 44 PB storage requires distributed DB (PostgreSQL max 16 TB)
- 7K write QPS approaches single instance limit

The Mental Math Trick

Multiplication by Powers of 10

- × 10: Add one zero
- × 100: Add two zeros
- × 1,000: Add three zeros (1K)
- × 1,000,000: Add six zeros (1M)
- × 1,000,000,000: Add nine zeros (1B)

Example:

$$5 \text{ MB} \times 1,000 = 5,000 \text{ MB} = 5 \text{ GB}$$

$$100 \text{ KB} \times 1,000,000 = 100,000,000 \text{ KB} = 100 \text{ GB}$$

Division by Large Numbers

- ÷ 100: Remove two zeros
- ÷ 1,000: Remove three zeros
- ÷ 86,400: Divide by 100,000 (close enough)

Example:

$$10,000,000 \div 100,000 = 100$$

$$1,000,000,000 \div 100,000 = 10,000$$

Advanced: Server Capacity Estimation

Typical Server Specs

Small Instance (t3.medium):

- 2 vCPU, 4 GB RAM
- Handles: ~500 QPS
- Cost: ~\$50/month

Medium Instance (m5.xlarge):

- 4 vCPU, 16 GB RAM
- Handles: ~2,000 QPS

- Cost: ~\$150/month

Large Instance (m5.4xlarge):

- 16 vCPU, 64 GB RAM
- Handles: ~10,000 QPS
- Cost: ~\$600/month

Calculating Server Count

Formula:

$$\text{Servers} = (\text{Peak QPS} / \text{QPS per Server}) \times \text{Safety Factor}$$

Example:

Peak: 30K QPS

Server capacity: 2K QPS

Safety factor: 1.5 (50% headroom)

$$\text{Servers} = (30,000 / 2,000) \times 1.5 = 22.5$$

Round up: 25 servers

Database Server Capacity

PostgreSQL (single instance):

- Reads: 10K–50K QPS
- Writes: 1K–10K QPS
- Storage: Up to 16 TB

Cassandra (per node):

- Reads: 10K–50K QPS
- Writes: 10K–100K QPS
- Storage: Up to 8 TB per node

Redis (single instance):

- Reads: 100K–1M QPS
- Writes: 100K–1M QPS
- Storage: Up to 512 GB RAM

Estimation Worksheet Template

Copy this for every interview:

SYSTEM: _____

ASSUMPTIONS:

- DAU: _____
- Operations per user: _____
- Data size: _____
- Retention: _____

TRAFFIC CALCULATIONS:

- Daily operations: DAU × ops = _____
- Average QPS: daily ÷ 100K = _____
- Peak QPS: avg × 3 = _____
- Read/Write ratio: _____

STORAGE CALCULATIONS:

- Per item size: _____
- Daily: items/day × size = _____
- Yearly: daily × 365 = _____
- Total: yearly × years = _____
- With replication: total × 3 = _____

BANDWIDTH CALCULATIONS:

- Incoming: write QPS × size = _____
- Outgoing: read QPS × size = _____
- With CDN: outgoing × (1 - hit ratio) = _____

MEMORY/CACHE CALCULATIONS:

- Hot data (20%): total × 0.2 = _____
- Cache servers: cache ÷ 100 GB = _____
- Cache hit ratio: _____

DATABASE CHOICE:

- Primary: _____ (because: _____)
- Cache: _____
- Analytics: _____

Final Checklist

Before saying "I'm done with calculations", verify:

- **Traffic Estimation**
 - Calculated write QPS
 - Calculated read QPS
 - Mentioned peak factor (2-3x)
 - Identified if read-heavy or write-heavy

****Storage Estimation****

- Calculated total storage
- Considered retention period
- Mentioned replication (3x)
- Separated data types (text vs media)

****Bandwidth Estimation****

- Calculated incoming bandwidth
- Calculated outgoing bandwidth
- Mentioned CDN impact

****Memory Estimation****

- Applied 80-20 rule
- Calculated cache size
- Determined number of servers

****Identified Bottlenecks****

- Database: "At 10K write QPS, need sharding"
- Network: "At 1 TB/s, need CDN"
- Memory: "At 10 TB cache, need 100 servers"

Summary: The Golden Rules

1. ****Round liberally**** – 387M → 400M
2. ****State assumptions**** – "Assuming X..."
3. ****Show your work**** – Write calculations down
4. ****Use shortcuts**** – Daily ÷ 100K = QPS
5. ****Apply 80-20 rule**** – Cache 20%, serve 80%
6. ****Mention peak traffic**** – Average × 3
7. ****Don't forget media**** – Images dominate storage
8. ****Include replication**** – × 3 for redundancy
9. ****Sanity check**** – Does 1 EB for Twitter make sense? No!
10. ****Mention optimizations**** – CDN, compression, caching

References

Books

1. ****"Designing Data-Intensive Applications"**** – Martin Kleppmann
2. ****"System Design Interview"**** – Alex Xu

Online Resources

1. ****System Design Primer**** – GitHub ([donnemartin](#))
2. ****Google's Latency Numbers**** – Jeff Dean
3. ****AWS Calculator**** – For cost estimates
4. ****High Scalability Blog**** – Real-world numbers

Videos

1. ****"System Design Course"**** – YouTube (Gaurav Sen)

2. **"Back of Envelope Calculations"** – YouTube (SystemDesignInterview)

Appendix: Pre-Calculated Common Scenarios

Scenario: 1 Million Users

Assumptions: 1M DAU, 10 operations/user/day

Traffic: $10M/\text{day} \div 100K = 100 \text{ QPS}$

Storage (1 KB per op): $10M \times 1 \text{ KB} = 10 \text{ GB/day} = 3.6 \text{ TB/year}$

Cache: $1M \text{ users} \times 10 \text{ KB} = 10 \text{ GB}$

Servers: 5-10 app servers

Scenario: 10 Million Users

Assumptions: 10M DAU, 10 operations/user/day

Traffic: $100M/\text{day} \div 100K = 1K \text{ QPS}$

Storage: $100M \times 1 \text{ KB} = 100 \text{ GB/day} = 36 \text{ TB/year}$

Cache: $10M \times 10 \text{ KB} = 100 \text{ GB}$

Servers: 10-20 app servers, single DB with replicas

Scenario: 100 Million Users

Assumptions: 100M DAU, 10 operations/user/day

Traffic: $1B/\text{day} \div 100K = 10K \text{ QPS}$

Storage: $1B \times 1 \text{ KB} = 1 \text{ TB/day} = 365 \text{ TB/year}$

Cache: $100M \times 10 \text{ KB} = 1 \text{ TB}$ across 10 servers

Servers: 50-100 app servers, sharded DB or NoSQL

Scenario: 1 Billion Users (Facebook Scale)

Assumptions: 1B DAU, 10 operations/user/day

Traffic: $10B/\text{day} \div 100K = 100K \text{ QPS}$

Storage: $10B \times 1 \text{ KB} = 10 \text{ TB/day} = 3.6 \text{ PB/year}$

Cache: $200M \text{ hot users} \times 10 \text{ KB} = 2 \text{ TB}$ across 20 servers

Servers: 500-1000 app servers, distributed NoSQL

Top 15 HLD Questions: Detailed Calculations

1. Design TinyURL / URL Shortener ★★★★

Requirements:

- 100M new URLs per month
- 100:1 read-to-write ratio
- URLs never expire
- Short URL: 7 characters (a-zA-Z0-9) = $62^7 = 3.5$ trillion combinations

Traffic Calculation:

Writes (URL creation):

100M URLs/month ÷ 30 days = 3.3M URLs/day

3.3M ÷ 86,400 = 38 QPS (round to 40 QPS)

Peak: $40 \times 3 = 120$ QPS

Reads (Redirects):

100:1 ratio → 10B redirects/month

10B ÷ 30 = 333M redirects/day

333M ÷ 86,400 = 3,854 QPS (round to 4K QPS)

Peak: $4K \times 3 = 12K$ QPS

Storage Calculation:

Per URL entry:

- Original URL: 100 bytes (average)
- Short URL: 7 bytes
- Created timestamp: 8 bytes
- User ID: 8 bytes
- Total: ~125 bytes per entry

Over 10 years:

- 100M/month × 12 × 10 = 12 billion URLs
- $12B \times 125$ bytes = 1.5 TB
- With indexes: $1.5\text{ TB} \times 1.5 = 2.25$ TB

Analytics (click tracking):

- Assume 10 clicks per URL average
- 12B URLs \times 10 clicks = 120B click events
- Each event: 50 bytes (URL_id, timestamp, IP, user_agent)
- 120B \times 50 bytes = 6 TB
- Total with analytics: 2.25 TB + 6 TB = 8.25 TB

Cache Calculation:

80-20 Rule: 20% URLs get 80% traffic

Hot URLs (recent + popular):

- Recent: Last 30 days = 100M URLs
- Apply 80-20: 100M \times 0.2 = 20M hot URLs
- Size: 20M \times 125 bytes = 2.5 GB

Decision: Cache ALL recent URLs! Only 2.5 GB needed

Single Redis instance can handle entire cache

Server Estimate:

Database:

- Single PostgreSQL instance handles 4K read QPS easily
- With 5 read replicas: 20K+ QPS capacity
- Well within requirements

Cache:

- 1 Redis instance (16 GB) - plenty of headroom
- Expected cache hit rate: 95%
- Only 5% (200 QPS) hits database

Application servers:

- 4K peak QPS \div 500 QPS per server = 8 servers
- With redundancy: 15-20 servers

Interview Summary:

"TinyURL is actually a simple system at this scale:

- Traffic: 4K read QPS, 40 write QPS (manageable)
- Storage: 8 TB over 10 years (tiny!)
- Cache: 2.5 GB (fits in single Redis)
- Database: Single PostgreSQL + read replicas sufficient
- The key optimization is caching hot URLs for 95%+ hit rate"

2. Design Netflix / Video Streaming ★★★

Requirements:

- 200M subscribers, 100M DAU (50% daily active)
- Each user watches 2 hours per day
- Average video: 1 hour duration
- Multiple qualities: 360p, 720p, 1080p, 4K
- Store catalog forever

Traffic Calculation:

Video plays per day:

$$100M \text{ users} \times (2 \text{ hours} / 1 \text{ hour per video}) = 200M \text{ video starts/day}$$

$$200M \div 86,400 = 2,314 \text{ QPS} (\text{round to } 2.5\text{K QPS})$$

$$\text{Peak: } 2.5\text{K} \times 3 = 7.5\text{K QPS}$$

Streaming QPS (continuous):

Active streamers at any moment:

- $100M \text{ DAU} \times 2 \text{ hours} / 24 \text{ hours} = 8.3M \text{ concurrent streams}$
- Each stream = continuous data flow (not traditional QPS)

API calls (browse, search, metadata):

- Each user: 20 API calls per session
- $100M \times 20 = 2B \text{ calls/day}$
- $2B \div 86,400 = 23K \text{ QPS}$
- Peak: 70K QPS

Storage Calculation:

Catalog size:

- Assume 10,000 movies/shows
- Average content length: 90 minutes

Per video, multiple qualities:

- 360p: 500 MB (0.5 GB)
- 720p: 1.5 GB
- 1080p: 3 GB
- 4K: 8 GB
- Total per video: 13 GB

Catalog storage:

- $10,000 \text{ videos} \times 13 \text{ GB} = 130 \text{ TB}$

New content per month:

- Assume 100 new titles per month
- $100 \times 13 \text{ GB} = 1.3 \text{ TB per month}$
- Yearly: $1.3 \text{ TB} \times 12 = 15.6 \text{ TB}$

10-year catalog: $130 \text{ TB} + (15.6 \text{ TB} \times 10) = 286 \text{ TB}$

User data (watch history, preferences):

- $200\text{M users} \times 10 \text{ KB} = 2 \text{ TB} \text{ (negligible)}$

Total: ~300 TB for catalog

Bandwidth Calculation:

Concurrent streams: 8.3M

Bitrate by quality:

- 360p: 1 Mbps = 125 KB/s
- 720p: 3 Mbps = 375 KB/s
- 1080p: 5 Mbps = 625 KB/s
- 4K: 25 Mbps = 3.125 MB/s

Average (assume 40% HD, 40% SD, 20% 4K):

- $(0.4 \times 625 \text{ KB/s}) + (0.4 \times 375 \text{ KB/s}) + (0.2 \times 3.125 \text{ MB/s})$
- $= 250 + 150 + 625 = 1,025 \text{ KB/s} \approx 1 \text{ MB/s per stream}$

Total bandwidth:

- $8.3 \text{M streams} \times 1 \text{ MB/s} = 8.3 \text{ TB/s}$

With CDN (Netflix uses 95% CDN):

- Origin: $8.3 \text{ TB/s} \times 0.05 = 415 \text{ GB/s}$
- CDN: $8.3 \text{ TB/s} \times 0.95 = 7.8 \text{ TB/s}$

CDN edge locations: 100+ globally

Per edge: $7.8 \text{ TB/s} \div 100 = 78 \text{ GB/s}$

****Cache Calculation**:**

Popular content (80-20 rule):

- $10,000 \text{ videos} \times 0.2 = 2,000 \text{ hot videos}$
- $2,000 \times 13 \text{ GB} = 26 \text{ TB}$

CDN cache per edge:

- Distribute 26 TB across 100 edges
- $26 \text{ TB} \div 100 = 260 \text{ GB per edge}$
- Manageable!

Metadata cache (Redis):

- $10,000 \text{ videos} \times 10 \text{ KB} = 100 \text{ MB}$
- Cache ALL metadata! Tiny!

****Interview Summary**:**

"Netflix's key challenge is bandwidth, not storage:

- Traffic: 8.3 TB/s global streaming bandwidth
- Storage: 300 TB catalog (manageable)
- CDN: Absolutely critical - serves 95% of traffic
- Cache: Popular content (20%) at edge locations
- Hot content (new releases) pre-pushed to all edges"

3. Design Uber / Lyft ⭐⭐⭐

****Requirements**:**

- 50M daily rides
- 100M registered users (50% are drivers)
- Location updates every 3 seconds while active
- Average ride: 15 minutes
- Store trip data for 5 years

****Traffic Calculation**:**

Active drivers at peak:

- Assume 10% of drivers active: $50M \times 0.1 = 5M$ drivers
- Peak hour (10% of daily rides): $50M \times 0.1 = 5M$ rides/hour

Location updates (real-time):

Per ride: $15 \text{ min} \div 3 \text{ sec} = 300$ updates per person

Daily updates: $50M \text{ rides} \times 300 \times 2 \text{ (driver + rider)} = 30B$ updates

Update QPS: $30B \div 86,400 = 347K$ QPS

Peak: $347K \times 3 = 1M$ QPS

Ride matching requests:

$50M \text{ rides/day} \div 86,400 = 578$ QPS

Each match query searches nearby drivers

Database queries: $578 \times 10 \text{ (search iterations)} = 5.7K$ QPS

Trip requests (ETAs, pricing):

- Users check prices multiple times
- $50M \text{ rides} \times 5 \text{ price checks} = 250M$ requests/day
- $250M \div 86,400 = 2,893$ QPS (round to 3K QPS)

****Storage Calculation**:**

Location history:

- $30B \text{ updates/day} \times 32 \text{ bytes (lat, long, timestamp, user_id)} = 960 \text{ GB/day}$
- Yearly: $960 \text{ GB} \times 365 = 350 \text{ TB}$
- 5 years: $350 \text{ TB} \times 5 = 1.75 \text{ PB}$

Trip data:

- $50M \text{ trips/day} \times 2 \text{ KB (pickup, dropoff, fare, route)} = 100 \text{ GB/day}$
- Yearly: $100 \text{ GB} \times 365 = 36 \text{ TB}$
- 5 years: $36 \text{ TB} \times 5 = 180 \text{ TB}$

Driver/Rider profiles:

- $100M \text{ users} \times 5 \text{ KB} = 500 \text{ GB}$ (negligible)

Total: $1.75 \text{ PB} + 180 \text{ TB} = 1.93 \text{ PB} \approx 2 \text{ PB}$

****Memory Calculation (Critical for Real-time)**:**

Active drivers (for matching):

- $5M \text{ drivers} \times 100 \text{ bytes} (\text{location} + \text{metadata}) = 500 \text{ MB}$
- Fits in single Redis instance!

Active ride requests:

- During peak: $100K \text{ concurrent searches}$
- $100K \times 200 \text{ bytes} = 20 \text{ MB}$

Surge pricing data:

- Store per geohash (city divided into grids)
- $1M \text{ geohashes} \times 50 \text{ bytes} = 50 \text{ MB}$

Route cache:

- Popular routes: $1M \text{ routes} \times 500 \text{ bytes} = 500 \text{ MB}$

Total in-memory: $500 \text{ MB} + 20 \text{ MB} + 50 \text{ MB} + 500 \text{ MB} = 1.07 \text{ GB}$

Single Redis instance handles everything!

****Server Estimate**:**

Application servers:

- $1M \text{ peak QPS} \div 1K \text{ QPS per server} = 1,000 \text{ servers}$
- With redundancy: 1,500 servers globally

Database:

- Cassandra for location history (write-heavy, 1M QPS)
- PostgreSQL for trips (ACID needed for payments)
- Redis for real-time matching (sub-second)

****Interview Summary**:**

"Uber's main challenge is real-time location processing:

- Traffic: 1M location updates/sec peak
- Storage: 2 PB for 5 years (manageable)
- Memory: 1 GB for real-time data (surprisingly small!)
- Key: Redis Geo commands (GEOADD, GEORADIUS) for O(log N) nearby search
- Database: Cassandra for high write throughput"

4. Design Amazon / E-commerce ★★★★

Requirements:

- 300M customers, 50M DAU (active shoppers)
- 1M products in catalog
- 1M orders per day
- Each user views 50 product pages per day
- Average product page: 500 KB (images, reviews)
- Store order history forever

Traffic Calculation:

Product views:

$$50\text{M users} \times 50 \text{ pages} = 2.5\text{B page views/day}$$

$$2.5\text{B} \div 86,400 = 28,935 \text{ QPS} (\text{round to } 30\text{K QPS})$$

Peak (Black Friday, 10x): 300K QPS

Product searches:

$$50\text{M users} \times 10 \text{ searches} = 500\text{M searches/day}$$

$$500\text{M} \div 86,400 = 5,787 \text{ QPS} (\text{round to } 6\text{K QPS})$$

Peak: 60K QPS

Orders (writes):

$$1\text{M orders/day} \div 86,400 = 11.5 \text{ QPS} (\text{round to } 12 \text{ QPS})$$

Peak: 120 QPS (easily handled)

Checkout process (critical path):

$$1\text{M checkouts/day} \div 86,400 = 11.5 \text{ QPS}$$

Must be ACID-compliant (payment processing)

Inventory checks (frequent):

- Every add-to-cart checks inventory
- $50\text{M users} \times 5 \text{ add-to-cart} = 250\text{M checks/day}$
- $250\text{M} \div 86,400 = 2,893 \text{ QPS}$ (round to 3K QPS)

****Storage Calculation**:**

Product catalog:

- $1\text{M products} \times 50 \text{ KB (details, specs)} = 50 \text{ GB}$
- Product images: $1\text{M} \times 2 \text{ MB} = 2 \text{ TB}$
- Total catalog: 2.05 TB (small!)

Order history:

- $1\text{M orders/day} \times 5 \text{ KB (order details, items, customer)} = 5 \text{ GB/day}$
- Yearly: $5 \text{ GB} \times 365 = 1.825 \text{ TB}$
- 10 years: 18.25 TB

Product reviews:

- Assume 10% of orders leave review
- $1\text{M} \times 0.1 = 100\text{K reviews/day}$
- $100\text{K} \times 1 \text{ KB} = 100 \text{ MB/day}$
- 10 years: $100 \text{ MB} \times 365 \times 10 = 365 \text{ GB}$

User data:

- $300\text{M users} \times 2 \text{ KB} = 600 \text{ GB}$

Total storage: $2 \text{ TB} + 18 \text{ TB} + 0.4 \text{ TB} + 0.6 \text{ TB} = 21 \text{ TB}$

With product images at multiple sizes:

- Original: 2 TB
- Large (product page): 1 TB
- Thumbnail: 200 GB
- Additional: 3.2 TB

Grand total: 25 TB (surprisingly small!)

****Cache Calculation**:**

Hot products (80-20 rule):

- $1\text{M products} \times 0.2 = 200\text{K hot products}$

- $200K \times 50 \text{ KB} = 10 \text{ GB}$ (metadata)
- Images served from CDN

Recent orders (last 90 days):

- $1\text{M/day} \times 90 = 90\text{M}$ orders
- $90\text{M} \times 5 \text{ KB} = 450 \text{ GB}$

Shopping carts (active sessions):

- $5\text{M active shoppers} \times 20 \text{ KB} = 100 \text{ GB}$

Total cache: $10 \text{ GB} + 450 \text{ GB} + 100 \text{ GB} = 560 \text{ GB}$

Across 10 Redis nodes: 56 GB each

****Bandwidth Calculation**:**

Product page loads:

$30\text{K QPS} \times 500 \text{ KB} = 15 \text{ GB/s}$

With CDN caching (90% hit rate):

- Origin: $15 \text{ GB/s} \times 0.1 = 1.5 \text{ GB/s}$
- CDN: $15 \text{ GB/s} \times 0.9 = 13.5 \text{ GB/s}$

Product images dominate bandwidth, not product data

****Interview Summary**:**

"E-commerce is read-heavy with critical write path:

- Traffic: 30K product view QPS, only 12 order QPS
- Storage: 25 TB total (orders + catalog)
- Cache: 560 GB for hot products + carts
- Critical: Inventory management (3K QPS checks)
- Database: PostgreSQL for orders (ACID), Redis for cart/inventory
- Need: Strong consistency for checkout, eventual OK for catalog"

5. Design Reddit / News Aggregator ★★

****Requirements**:**

- 50M DAU
- 500K posts per day (0.01 posts per user)
- Each user views 100 posts per day
- Each post: 2 KB (text, metadata)
- Comments: 5M per day (avg 10 per post)
- Store for 10 years

****Traffic Calculation**:**

Post creation:

$$500\text{K posts/day} \div 86,400 = 5.7 \text{ QPS} (\text{round to } 6 \text{ QPS})$$

Peak: 18 QPS

Comment creation:

$$5\text{M comments/day} \div 86,400 = 57 \text{ QPS}$$

Peak: 171 QPS

Post views:

$$50\text{M} \times 100 = 5\text{B views/day}$$

$$5\text{B} \div 86,400 = 57,870 \text{ QPS} (\text{round to } 58\text{K QPS})$$

Peak: 174K QPS

Vote operations (upvote/downvote):

- Assume each user votes 50 times/day
- $50\text{M} \times 50 = 2.5\text{B votes/day}$
- $2.5\text{B} \div 86,400 = 28,935 \text{ QPS} (\text{round to } 30\text{K QPS})$
- Peak: 90K QPS

Total read QPS: 58K (views) + 30K (votes) = 88K QPS

Total write QPS: 6 (posts) + 57 (comments) + 30K (votes) = 30K QPS

****Storage Calculation**:**

Posts:

- $500\text{K/day} \times 2 \text{ KB} = 1 \text{ GB/day}$
- 10 years: $1 \text{ GB} \times 365 \times 10 = 3.65 \text{ TB}$

Comments:

- $5\text{M/day} \times 1 \text{ KB} = 5 \text{ GB/day}$
- 10 years: $5 \text{ GB} \times 365 \times 10 = 18.25 \text{ TB}$

Votes:

- Store user_id + post_id + vote_type
- $2.5\text{B votes/day} \times 25\text{ bytes} = 62.5\text{ GB/day}$
- 10 years: $62.5\text{ GB} \times 365 \times 10 = 228\text{ TB}$

User data:

- $50\text{M} \times 5\text{ KB} = 250\text{ GB}$

Total: $3.65\text{ TB} + 18.25\text{ TB} + 228\text{ TB} + 0.25\text{ TB} = 250\text{ TB}$

Cache Calculation:

Front page posts (hot):

- Top 1000 posts (refreshed every 30 min)
- $1000 \times 2\text{ KB} = 2\text{ MB}$ (tiny!)
- Cache entire front page

Hot posts (trending):

- Top 10K posts $\times 2\text{ KB} = 20\text{ MB}$

Comment trees (popular threads):

- Top 1000 threads $\times 100\text{ comments} \times 1\text{ KB} = 100\text{ MB}$

User karma/profile:

- $10\text{M active} \times 2\text{ KB} = 20\text{ GB}$

Total cache: $20\text{ GB} + 100\text{ MB} + 20\text{ MB} = \sim 20\text{ GB}$

Single Redis instance sufficient!

Server Estimate:

Read servers:

$88\text{K QPS} \div 2\text{K QPS per server} = 44\text{ servers}$

With redundancy: 70 servers

Write servers:

$30\text{K QPS} \div 2\text{K QPS per server} = 15\text{ servers}$

With redundancy: 25 servers

Database:

- PostgreSQL for posts/comments (small dataset)
- Redis for vote counters (high write QPS)
- Elasticsearch for search

****Interview Summary**:**

"Reddit is interesting - high vote QPS but small storage:

- Traffic: 88K read QPS, 30K write QPS (balanced)
- Storage: 250 TB over 10 years (modest)
- Cache: 20 GB (front page + hot threads)
- Challenge: Vote aggregation in real-time
- Solution: Redis counters + periodic batch updates"

6. Design Rate Limiter ★★★

****Requirements**:**

- 1B API calls per day
- Rate limit: 100 requests per minute per user
- 10M active API users
- Track per user, per API key, per IP

****Traffic Calculation**:**

API calls:

$$1\text{B calls/day} \div 86,400 = 11,574 \text{ QPS} \text{ (round to 12K QPS)}$$

$$\text{Peak: } 12\text{K} \times 3 = 36\text{K QPS}$$

Rate limit checks:

- Every API call needs rate limit check
- Same as API traffic: 12K QPS, peak 36K QPS
- Must add < 10ms latency (or API too slow)

****Storage Calculation**:**

Counter storage (in-memory):

Per user counters:

- user_id: 8 bytes
- counter: 4 bytes
- window_timestamp: 8 bytes
- Total: 20 bytes per counter

Active users in last minute:

- 10M users \times 20 bytes = 200 MB
- With multiple time windows (sliding window):
- Current + Previous window: $200\text{ MB} \times 2 = 400\text{ MB}$

Total memory: 400 MB (fits easily in Redis!)

Historical data (analytics):

- Store violations for analysis
- 1% violation rate: $1\text{B} \times 0.01 = 10\text{M violations/day}$
- $10\text{M} \times 100\text{ bytes} = 1\text{ GB/day}$
- Yearly: 365 GB

Rule configuration:

- 10K rate limit rules \times 500 bytes = 5 MB
- Cache ALL rules in memory

Memory Calculation:

Redis for counters:

- 400 MB active data
- Single Redis instance: 16 GB
- Plenty of headroom for growth

Rules cache (application memory):

- 5 MB per server
- Loaded at startup
- Refresh every 5 minutes

Server Estimate:

Rate limiter must be FAST:

- Target: < 1ms overhead
- Redis GET + INCR: ~0.5ms
- Network: ~0.5ms
- Total: ~1ms ✓

Servers needed:

- $36K \text{ peak QPS} \div 10K \text{ QPS per server} = 3.6 \text{ servers}$
- With redundancy: 10 servers globally
- Each server: Connection pool to Redis

Redis cluster:

- 5 nodes for redundancy
- Each handles 7.2K QPS (easy for Redis)

****Interview Summary**:**

"Rate limiter is simple but critical:

- Traffic: 36K peak QPS (must be fast!)
- Storage: 400 MB in-memory (Redis)
- Latency: < 1ms overhead required
- Algorithm: Sliding window counter or Token bucket
- Key: Use Redis atomic operations (INCR)
- Scale: Single Redis instance handles millions of users"

7. Design Notification System ★★

****Requirements**:**

- 1B users, 100M DAU
- Each user receives 10 notifications per day
- Notification types: Push, Email, SMS
- 80% push, 15% email, 5% SMS
- Response time: < 100ms to accept notification
- Delivery: Best effort (OK to miss some)

****Traffic Calculation**:**

Total notifications:

$100M \text{ users} \times 10 = 1B \text{ notifications/day}$

$1B \div 86,400 = 11,574 \text{ QPS} (\text{round to } 12\text{K QPS})$

Peak: $12\text{K} \times 3 = 36\text{K QPS}$

By type:

- Push: $1B \times 0.8 = 800M/\text{day} = 9.2\text{K QPS}$
- Email: $1B \times 0.15 = 150M/\text{day} = 1.7\text{K QPS}$
- SMS: $1B \times 0.05 = 50M/\text{day} = 578 \text{ QPS}$

Fan-out scenario:

- User posts → notify 1000 followers
- 1M posts/day = 1B notifications
- Fanout QPS: Same 12K QPS

Storage Calculation:

Notification logs:

- $1B \text{ notifications/day} \times 500 \text{ bytes} = 500 \text{ GB/day}$
- 90-day retention: $500 \text{ GB} \times 90 = 45 \text{ TB}$

User preferences:

- $1B \text{ users} \times 2 \text{ KB} (\text{preferences per notification type}) = 2 \text{ TB}$

Device tokens:

- $1B \text{ users} \times 300 \text{ bytes (APNs/FCM tokens)} = 300 \text{ GB}$

Templates:

- $1000 \text{ templates} \times 10 \text{ KB} = 10 \text{ MB} (\text{negligible})$

Total: $45 \text{ TB} + 2 \text{ TB} + 0.3 \text{ TB} = 47 \text{ TB}$

Memory Calculation:

Cache user preferences:

- $100M \text{ DAU} \times 2 \text{ KB} = 200 \text{ GB}$
- Cache hot users ($20M$) $\times 2 \text{ KB} = 40 \text{ GB}$

Cache device tokens:

- $100M \text{ DAU} \times 300 \text{ bytes} = 30 \text{ GB}$

Notification queue (pending):

- Buffer 5 minutes of notifications
- $12K \text{ QPS} \times 300 \text{ seconds} = 3.6M \text{ notifications}$
- $3.6M \times 500 \text{ bytes} = 1.8 \text{ GB}$

Total cache: $40 \text{ GB} + 30 \text{ GB} + 2 \text{ GB} = 72 \text{ GB}$

Across 5 Redis nodes: 15 GB each

****Server Estimate**:**

Notification workers:

- $36K \text{ peak QPS} \div 5K \text{ QPS per worker} = 7.2 \text{ workers}$
- With redundancy: 15 workers

Push notification service:

- APNs connection pool: 100 connections
- FCM connection pool: 100 connections
- Each handles 1K QPS

Email service:

- Use SES/SendGrid
- 1.7K QPS (52M emails/day)
- Batched sending

SMS service:

- Use Twilio
- 578 QPS (50M SMS/day)
- Most expensive channel

****Interview Summary**:**

"Notification system needs reliability and fanout:

- Traffic: 36K peak QPS (manageable)
- Storage: 47 TB (mostly logs)
- Challenge: Fan-out (1 post → 1000 notifications)
- Solution: Queue-based async processing (Kafka)
- Priority queue: High (transactions) > Medium (social) > Low (marketing)

- Deduplication: Merge similar notifications"

8. Design Google Search Autocomplete ★★

Requirements:

- 5B searches per day (Google scale)
- Autocomplete: Suggest after 3+ characters
- Suggest top 10 completions
- Update suggestions hourly
- Multi-language support

Traffic Calculation:

Autocomplete requests:

- Each search has ~5 autocomplete requests (typing)
- 5B searches \times 5 = 25B autocomplete requests/day
- $25B \div 86,400 = 289,351$ QPS (round to 300K QPS)
- Peak: $300K \times 3 = 900K$ QPS

This is READ-ONLY and latency-critical!

Must respond in < 100ms (ideally < 50ms)

Storage Calculation:

Search query logs:

- 5B searches/day \times 50 bytes (query text + metadata) = 250 GB/day
- 30-day retention: $250 \text{ GB} \times 30 = 7.5 \text{ TB}$

Trie/Prefix data structure:

- Store popular queries (top 10M)
- Each node: ~100 bytes
- Total nodes: ~100M (branching factor)
- $100M \times 100 \text{ bytes} = 10 \text{ GB}$

Suggestions cache:

- Pre-compute top 10 suggestions for popular prefixes
- 1M popular prefixes \times 1 KB (10 suggestions) = 1 GB

Total: 7.5 TB (logs) + 10 GB (trie) + 1 GB (cache) = 7.5 TB

Memory Calculation (Critical):

Must be in-memory for speed!

Trie structure: 10 GB

- Fits in single server memory
- Replicate across regions: 10 GB \times 5 regions = 50 GB total

Top 10M queries cache:

- 10M \times 100 bytes = 1 GB

Per-language tries:

- 20 languages \times 10 GB = 200 GB
- Distribute across 20 servers: 10 GB each

Total: 200 GB distributed

Server Estimate:

Query servers:

- 900K peak QPS \div 10K QPS per server = 90 servers
- With redundancy: 150 servers globally

Per region (5 regions):

- 150 \div 5 = 30 servers per region

Each server:

- Loads trie into memory (10 GB)
- Serves 10K QPS
- < 10ms latency

Interview Summary:

"Autocomplete is latency-critical:

- Traffic: 900K peak QPS (very high!)
- Storage: 10 GB trie (must fit in memory)
- Latency: < 50ms (in-memory trie essential)
- Solution: Trie data structure in memory
- Optimization: Pre-compute popular suggestions
- Global: Replicate trie to all regions
- Update: Rebuild trie hourly from logs"

9. Design Web Crawler (Google Bot) ★★

Requirements:

- Crawl 1 billion web pages
- Re-crawl every 7 days (freshness)
- Average page: 500 KB
- Follow 20 links per page
- Respect robots.txt and rate limits (1 req/sec per domain)

Traffic Calculation:

Pages to crawl per day:

$$1B \text{ pages} \div 7 \text{ days} = 143M \text{ pages/day}$$

$$143M \div 86,400 = 1,655 \text{ QPS} (\text{round to } 1.7K \text{ QPS})$$

Peak: 5K QPS

URLs discovered per day:

$$143M \text{ pages} \times 20 \text{ links} = 2.86B \text{ URLs}$$

Need to check if already crawled: 2.86B lookups/day

$$2.86B \div 86,400 = 33K \text{ lookup QPS}$$

DNS lookups:

- Assume 100K unique domains
- Each page needs DNS: 1.7K QPS DNS lookups
- With caching (90% hit): 170 QPS actual DNS queries

Storage Calculation:

Page content:

- $143M \text{ pages/day} \times 500 \text{ KB} = 71.5 \text{ TB/day}$
- With 7-day rotation: $71.5 \text{ TB} \times 7 = 500 \text{ TB}$
- Only need latest snapshot: 500 TB

URL frontier (to be crawled):

- $2.86B \text{ URLs} \times 200 \text{ bytes} = 572 \text{ GB}$
- Keep in Redis for fast access

Already crawled URLs (Bloom filter):

- 1B URLs tracked
- Bloom filter: ~1.2 GB (1% false positive rate)

Extracted metadata:

- $1B \text{ pages} \times 10 \text{ KB} (\text{title, meta, links}) = 10 \text{ TB}$

Total: $500 \text{ TB} + 0.57 \text{ TB} + 10 \text{ TB} = 511 \text{ TB}$

Memory Calculation:

URL frontier (priority queue):

- 100M URLs ready to crawl
- $100M \times 200 \text{ bytes} = 20 \text{ GB}$
- Redis sorted set (score = priority)

DNS cache:

- $100K \text{ domains} \times 100 \text{ bytes} = 10 \text{ MB}$

Bloom filter (visited URLs):

- 1B URLs: 1.2 GB in memory

Robots.txt cache:

- $100K \text{ domains} \times 10 \text{ KB} = 1 \text{ GB}$

Total: $20 \text{ GB} + 0.01 \text{ GB} + 1.2 \text{ GB} + 1 \text{ GB} = 22 \text{ GB}$

Fits in 3-4 servers

Server Estimate:

Crawler workers:

- 1.7K pages/sec to crawl
- Each worker: 10 concurrent requests = 10 pages/sec
- Workers needed: $1,700 \div 10 = 170$ workers
- With redundancy: 250 workers

Rate limiting consideration:

- 1 req/sec per domain
- 100K domains = 100K req/sec theoretical max
- Our 1.7K QPS is well within limit

Interview Summary:

"Web crawler needs politeness and de-duplication:

- Traffic: 1.7K crawl QPS, 33K URL dedup QPS
- Storage: 511 TB (page snapshots rotate every 7 days)
- Memory: 22 GB (frontier + bloom filter + DNS cache)
- Challenge: Avoid re-crawling, respect rate limits
- Solution: Bloom filter for visited URLs, robots.txt caching
- Politeness: Max 1 req/sec per domain"

10. Design Pastebin ★★

Requirements:

- 10M DAU
- Each user creates 1 paste per day
- Each paste: 10 KB (code snippet)
- Paste expiry: 24 hours (default)
- Each paste viewed 100 times (read-heavy)
- Store active pastes (< 30 days old)

Traffic Calculation:

Writes (create paste):

$$10M \text{ creates/day} \div 86,400 = 115 \text{ QPS}$$

Peak: 345 QPS

Reads (view paste):

$10M \text{ pastes} \times 100 \text{ views} = 1B \text{ views/day}$

$1B \div 86,400 = 11,574 \text{ QPS} (\text{round to } 12\text{K QPS})$

Peak: 36K QPS

Ratio: 100:1 (read-heavy)

Storage Calculation:

Active pastes (30-day retention):

- $10M/\text{day} \times 30 \text{ days} = 300M \text{ active pastes}$
- $300M \times 10 \text{ KB} = 3 \text{ TB}$

With automatic expiry:

- 24-hour pastes: $10M \times 10 \text{ KB} = 100 \text{ GB}$
- 1-week pastes: $70M \times 10 \text{ KB} = 700 \text{ GB}$
- Permanent pastes (10%): $1M/\text{day} \times 10 \text{ KB} = 10 \text{ GB/day}$
Over 5 years: $10 \text{ GB} \times 365 \times 5 = 18 \text{ TB}$

Total: 3 TB (active) + 18 TB (permanent) = 21 TB

Cache Calculation:

Hot pastes (recent + popular):

- Last 24 hours: 10M pastes
- Apply 80-20: $10M \times 0.2 = 2M \text{ hot pastes}$
- $2M \times 10 \text{ KB} = 20 \text{ GB}$

Cache ALL recent pastes!

- $10M \times 10 \text{ KB} = 100 \text{ GB}$
- Distributed: 10 servers $\times 10 \text{ GB each}$

Cache hit rate: 95%+ (recent pastes very popular)

Bandwidth Calculation:

Incoming: 115 QPS \times 10 KB = 1.15 MB/s

Outgoing: 12K QPS \times 10 KB = 120 MB/s

With CDN (optional for code):

- Origin: 120 MB/s \times 0.1 = 12 MB/s
- CDN: 120 MB/s \times 0.9 = 108 MB/s

Interview Summary:

"Pastebin is surprisingly simple:

- Traffic: 12K read QPS, 115 write QPS
- Storage: 21 TB over 5 years
- Cache: 100 GB (recent pastes)
- TTL: Critical feature (auto-cleanup)
- Database: PostgreSQL sufficient (small dataset)
- Optimization: Cache recent pastes (95%+ hit rate)"

11. Design LinkedIn / Professional Network ★★

Requirements:

- 800M users, 300M DAU
- Each user views 20 profiles per day
- Each user updates profile 0.1 times per day
- Each user sees 50 feed posts per day
- Profile size: 20 KB (resume, skills, experience)
- Connection graph: 500 connections average per user

Traffic Calculation:

Profile views:

$$300M \times 20 = 6B \text{ views/day}$$

$$6B \div 86,400 = 69,444 \text{ QPS} (\text{round to } 70\text{K QPS})$$

Peak: 210K QPS

Profile updates:

$$300M \times 0.1 = 30M \text{ updates/day}$$

$30M \div 86,400 = 347 \text{ QPS}$

Peak: 1K QPS

Feed views:

$300M \times 50 = 15B \text{ posts viewed/day}$

$15B \div 86,400 = 173,611 \text{ QPS} (\text{round to } 174\text{K QPS})$

Peak: 522K QPS

Connection requests:

- Assume 2 new connections per user per day
- $300M \times 2 = 600M \text{ requests/day}$
- $600M \div 86,400 = 6,944 \text{ QPS} (\text{round to } 7\text{K QPS})$

Storage Calculation:

User profiles:

- $800M \text{ users} \times 20 \text{ KB} = 16 \text{ TB}$

Connections (social graph):

- $800M \text{ users} \times 500 \text{ connections} = 400B \text{ edges}$
- Each edge: 16 bytes (`user_id + connection_id`)
- $400B \times 16 \text{ bytes} = 6.4 \text{ TB}$

Posts/Updates:

- $30M \text{ posts/day} \times 2 \text{ KB} = 60 \text{ GB/day}$
- 5 years: $60 \text{ GB} \times 365 \times 5 = 109 \text{ TB}$

Messages:

- Similar to LinkedIn messaging
- Assume $100M \text{ messages/day} \times 500 \text{ bytes} = 50 \text{ GB/day}$
- 2 years: $50 \text{ GB} \times 365 \times 2 = 36 \text{ TB}$

Total: $16 \text{ TB} + 6.4 \text{ TB} + 109 \text{ TB} + 36 \text{ TB} = 167 \text{ TB}$

Cache Calculation:

Hot profiles (20%):

- $300M \text{ DAU} \times 0.2 = 60M \text{ profiles}$
- $60M \times 20 \text{ KB} = 1.2 \text{ TB}$

- Distributed: 20 servers \times 60 GB each

Connection graph:

- Cache 60M users' connections
- $60M \times 500 \times 16 \text{ bytes} = 480 \text{ GB}$

Feed cache:

- $60M \text{ users} \times 50 \text{ posts} \times 2 \text{ KB} = 6 \text{ TB}$

Total: $1.2 \text{ TB} + 0.48 \text{ TB} + 6 \text{ TB} = 7.68 \text{ TB}$

Across 80 servers: ~100 GB each

****Interview Summary**:**

"LinkedIn is profile-heavy with social graph:

- Traffic: 174K feed QPS, 70K profile QPS
- Storage: 167 TB (profiles + connections + content)
- Cache: 7.68 TB (profiles + connections + feeds)
- Challenge: Connection graph queries (2nd/3rd degree)
- Database: Neo4j/graph DB for connections, PostgreSQL for profiles
- Optimization: Denormalize feed data, cache aggressively"

12. Design Discord / Slack ★★

****Requirements**:**

- 150M DAU
- 10M active servers/workspaces
- Average: 1000 messages per server per day
- Each user in 5 servers average
- Real-time messaging
- Store for 90 days (free), forever (premium)

****Traffic Calculation**:**

Messages per day:

$10M \text{ servers} \times 1000 = 10B \text{ messages/day}$

$10B \div 86,400 = 115,740$ QPS (round to 116K QPS)

Peak: 348K QPS

User presence updates:

- 150M users \times online/offline changes
- Assume 10 status changes per user per day
- $150M \times 10 = 1.5B$ updates/day
- $1.5B \div 86,400 = 17,361$ QPS (round to 17K QPS)

Typing indicators:

- Active: 15M users typing at any moment
- Update every 3 seconds
- $15M \div 3 = 5M$ updates/sec = 5M QPS (very high!)
- Optimization: Throttle to 1 update per 3 seconds = 17K QPS

Storage Calculation:

Messages:

- $10B/day \times 200$ bytes = 2 TB/day
- 90 days: $2 \text{ TB} \times 90 = 180 \text{ TB}$

Voice/Video data (separate calculation):

- Assume 10M minutes of voice per day
- $10M \text{ min} \times 1 \text{ MB/min} = 10 \text{ TB/day}$
- 90 days: $10 \text{ TB} \times 90 = 900 \text{ TB}$

File uploads:

- Assume 100M files/day \times 1 MB avg = 100 TB/day
- 90 days: $100 \text{ TB} \times 90 = 9 \text{ PB}$

User data:

- $150M \text{ users} \times 5 \text{ KB} = 750 \text{ GB}$

Server metadata:

- $10M \text{ servers} \times 10 \text{ KB} = 100 \text{ GB}$

Total: $180 \text{ TB} + 900 \text{ TB} + 9 \text{ PB} + 0.75 \text{ TB} = 10 \text{ PB}$

Memory Calculation:

Online users (presence):

- 150M concurrent \times 100 bytes = 15 GB

Active servers (hot):

- 1M active servers \times 1 MB (recent messages) = 1 TB

Typing indicators (ephemeral):

- 15M users \times 50 bytes = 750 MB

Message queue (undelivered):

- Buffer 10 minutes: 116K QPS \times 600 sec = 70M messages
- 70M \times 200 bytes = 14 GB

Total: 15 GB + 1 TB + 0.75 GB + 14 GB = 1.03 TB

Across 20 Redis nodes: 51 GB each

Interview Summary:

"Discord/Slack is real-time group messaging at scale:

- Traffic: 348K message QPS peak
- Storage: 10 PB (90-day retention)
- Memory: 1 TB (presence + messages + queue)
- Challenge: Real-time delivery to 1000s in server
- Solution: WebSocket + fanout via Redis Pub/Sub
- Database: Cassandra for messages, PostgreSQL for users/servers"

13. Design Zoom / Video Conferencing ★★★★

Requirements:

- 300M DAU
- 10M concurrent meetings at peak
- Average meeting: 4 participants, 30 minutes
- Video: 720p (1.5 Mbps), Audio: 64 Kbps
- Store recordings: 10% of meetings, for 1 year

Traffic Calculation:

Concurrent connections:

- $10M \text{ meetings} \times 4 \text{ participants} = 40M \text{ concurrent connections}$
- Real-time streams (not traditional QPS)

Signaling QPS (WebRTC setup):

- $10M \text{ meetings/day} \div 86,400 = 115 \text{ meetings/sec}$
- Each meeting: $4 \text{ participants} \times 10 \text{ signaling messages} = 40 \text{ messages}$
- $115 \times 40 = 4,600 \text{ signaling QPS}$

Media relay QPS:

- Not all peer-to-peer, some need relay
- Assume 30% need relay: $40M \times 0.3 = 12M \text{ relayed connections}$
- Each sending/receiving: 24M streams total

Bandwidth Calculation (Critical):

Per participant bandwidth:

- Video upload: 1.5 Mbps = 187.5 KB/s
- Audio upload: 64 Kbps = 8 KB/s
- Total upload: ~195 KB/s
- Video downloads: 1.5 Mbps \times 3 (other participants) = 4.5 Mbps = 562.5 KB/s
- Audio downloads: 64 Kbps \times 3 = 192 Kbps = 24 KB/s
- Total download: ~586 KB/s

For 40M concurrent participants:

- Upload bandwidth: $40M \times 195 \text{ KB/s} = 7.8 \text{ TB/s}$
- Download bandwidth: $40M \times 586 \text{ KB/s} = 23.4 \text{ TB/s}$

This is HUGE! Need optimization:

With P2P (peer-to-peer):

- 70% connections don't hit servers
- Server bandwidth: $7.8 \text{ TB/s} \times 0.3 = 2.3 \text{ TB/s}$

With SFU (Selective Forwarding Unit) optimization:

- Server receives once, forwards to multiple
- Reduction: ~50%
- Server bandwidth: ~1.2 TB/s

Storage Calculation:

Meeting recordings:

- 10% of 10M meetings/day = 1M recordings/day
- Each: $30 \text{ min} \times 4 \text{ participants} \times 1.5 \text{ Mbps} (\text{video}) + 64 \text{ Kbps} (\text{audio})$
- Per recording: $30 \text{ min} \times 4 \times (1.5 \text{ Mbps} \times 60 \text{ sec} / 8) = 30 \times 4 \times 11.25 \text{ MB} = 1.35 \text{ GB}$
- Daily: $1\text{M} \times 1.35 \text{ GB} = 1.35 \text{ PB/day}$
- Yearly: $1.35 \text{ PB} \times 365 = 492 \text{ PB}$

With compression (30% reduction):

- Actual: $492 \text{ PB} \times 0.7 = 344 \text{ PB/year}$

Metadata:

- 10M meetings/day $\times 2 \text{ KB} = 20 \text{ GB/day}$
- Yearly: 7.3 TB (negligible)

Total: 344 PB per year

Interview Summary:

"Zoom's challenge is real-time bandwidth:

- Connections: 40M concurrent streams
- Bandwidth: 1.2 TB/s with SFU optimization
- Storage: 344 PB/year for recordings
- Critical: Low latency (< 200ms) for good UX
- Solution: Geographic distribution, P2P when possible, SFU for groups
- Media servers: Distributed globally (50+ regions)"

14. Design Ticketmaster / Event Booking ★★★

Requirements:

- 100M users, 10M DAU
- 10K events per day
- Average event: 10K tickets
- Flash sales: 100K users compete for 10K tickets in 1 minute
- Booking flow must prevent overselling

****Traffic Calculation** (Critical – Spike Pattern):**

Normal browsing:

- $10M \text{ users} \times 10 \text{ page views} = 100M \text{ views/day}$
- $100M \div 86,400 = 1,157 \text{ QPS}$ (round to 1.2K QPS)

Flash sale (worst case):

- 100K users trying to book in 1 minute
- Each user: 10 attempts = 1M booking requests in 60 seconds
- Spike QPS: $1M \div 60 = 16,666 \text{ QPS}$ (round to 17K QPS)
- This is 14x normal load!

Ticket inventory checks:

- Before every booking: Check seat available
- $17K \text{ booking attempts/sec} = 17K \text{ inventory checks/sec}$
- Must be ACID-compliant (no double booking)

****Storage Calculation**:**

Events:

- $10K \text{ events/day} \times 10 \text{ KB} = 100 \text{ MB/day}$
- 2 years: $100 \text{ MB} \times 365 \times 2 = 73 \text{ GB}$

Tickets:

- $10K \text{ events/day} \times 10K \text{ tickets} = 100M \text{ tickets/day}$
- Each ticket: 500 bytes (seat, price, status)
- $100M \times 500 \text{ bytes} = 50 \text{ GB/day}$
- 2 years: $50 \text{ GB} \times 365 \times 2 = 36 \text{ TB}$

Bookings:

- Assume 80% tickets sold
- $100M \times 0.8 = 80M \text{ bookings/day}$

- Each booking: 2 KB (user, ticket, payment)
- $80M \times 2 \text{ KB} = 160 \text{ GB/day}$
- 5 years: $160 \text{ GB} \times 365 \times 5 = 292 \text{ TB}$

Total: $0.07 \text{ TB} + 36 \text{ TB} + 292 \text{ TB} = 328 \text{ TB}$

Memory Calculation (Critical for Flash Sales):

Ticket inventory (must be in-memory for speed):

- Current active events: 10K events
- $10K \times 10K \text{ tickets} = 100M \text{ tickets}$
- Each: 100 bytes (seat_id, status, price)
- $100M \times 100 \text{ bytes} = 10 \text{ GB}$

Booking locks (prevent double-booking):

- During flash sale: 17K concurrent booking attempts
- Each lock: 50 bytes
- $17K \times 50 \text{ bytes} = 850 \text{ KB} (\text{tiny!})$

User sessions:

- $100K \text{ active users} \times 10 \text{ KB} = 1 \text{ GB}$

Total: $10 \text{ GB} + 0.85 \text{ MB} + 1 \text{ GB} = 11 \text{ GB}$

Single Redis instance handles it!

Interview Summary:

"Ticketmaster needs to handle extreme spikes:

- Traffic: 17K QPS during flash sales (vs 1.2K normal)
- Storage: 328 TB (manageable)
- Memory: 11 GB in-memory inventory (critical!)
- Challenge: Race condition (overselling)
- Solution: Redis distributed locks + optimistic locking
- Database: PostgreSQL with row-level locks for ACID
- Queue: Virtual waiting room to smooth spike"

15. Design Parking Lot System ★

Requirements:

- 10K parking spots per location
- 100 locations (malls, airports)
- 5M vehicles per day
- Average stay: 2 hours
- Real-time availability

Traffic Calculation:

Entry/exit events:

- $5M \text{ vehicles} \times 2 \text{ events (entry + exit)} = 10M \text{ events/day}$
- $10M \div 86,400 = 115 \text{ QPS}$
- Peak (rush hour, 3x): 345 QPS

Availability checks:

- Drivers check before arriving
- $5M \text{ drivers} \times 3 \text{ checks} = 15M \text{ checks/day}$
- $15M \div 86,400 = 173 \text{ QPS}$
- Peak: 519 QPS

Display updates:

- $100 \text{ locations} \times 10 \text{ displays} = 1000 \text{ displays}$
- Poll every 5 seconds = $1000 \div 5 = 200 \text{ QPS}$
- Always-on, no peak variation

Storage Calculation:

Parking spots:

- $100 \text{ locations} \times 10K \text{ spots} = 1M \text{ spots total}$
- Each spot: 200 bytes (spot_id, location, type, floor)
- $1M \times 200 \text{ bytes} = 200 \text{ MB (tiny!)}$

Parking events:

- $10M \text{ events/day} \times 500 \text{ bytes} = 5 \text{ GB/day}$
- 1 year retention: $5 \text{ GB} \times 365 = 1.825 \text{ TB}$

Payment records:

- 5M transactions/day \times 1 KB = 5 GB/day
- 5 years: 5 GB \times 365 \times 5 = 9 TB

Vehicle data:

- 5M unique vehicles \times 500 bytes = 2.5 GB

Total: 0.2 GB + 1.8 TB + 9 TB + 2.5 GB = 11 TB

****Memory Calculation**:**

Real-time availability:

- 1M spots \times 50 bytes (spot_id, status, occupied_since) = 50 MB
- Cache ALL spots! Tiny!

Active sessions:

- Peak concurrent vehicles: 5M \times (2 hours / 24 hours) = 416K
- 416K \times 200 bytes = 83 MB

Recent transactions (for display):

- Last 1000 per location \times 100 locations = 100K transactions
- 100K \times 500 bytes = 50 MB

Total: 50 MB + 83 MB + 50 MB = 183 MB

Single Redis instance more than sufficient!

****Interview Summary**:**

"Parking lot is small scale but real-time critical:

- Traffic: 519 QPS peak (very manageable)
- Storage: 11 TB over 5 years
- Memory: 183 MB (everything fits in cache!)
- Challenge: Prevent double-booking of spots
- Solution: Redis atomic operations (SETNX for locks)
- Display: WebSocket for real-time updates to screens
- Database: PostgreSQL sufficient (small dataset, ACID needed)"

Top HLD Questions Summary Table

System	DAU	QPS	Storage	Cache	Pattern
TinyURL	N/A	4K R	8 TB	2.5 GB	Read-heavy
Twitter	400M	1.4M R	44 PB	1.1 TB	Read-heavy
Instagram	500M	867K R	58 PB	60 TB	Read-heavy
WhatsApp	1B	1.4M RW	1.46 EB	2.5 TB	Balanced
Uber	50M	1M W	2 PB	1 GB	Write-heavy
Netflix	100M	70K API	300 TB	26 TB	Streaming
Amazon	50M	300K R	25 TB	560 GB	Read-heavy
Reddit	50M	88K R	250 TB	20 GB	Read-heavy
Rate Limiter	N/A	36K	400 MB	400 MB	Fast
Notifications	100M	36K	47 TB	72 GB	Fanout
Autocomplete	N/A	900K	10 GB	10 GB	Latency
Web Crawler	N/A	1.7K W	511 TB	22 GB	Write-heavy
Pastebin	10M	12K R	21 TB	100 GB	Read-heavy
LinkedIn	300M	522K R	167 TB	7.68 TB	Read-heavy
Discord/Slack	150M	348K	10 PB	1 TB	Real-time
Zoom	300M	40M Con	344 PB/y	N/A	Streaming
Ticketmaster	10M	17K	328 TB	11 GB	Spike
Parking Lot	5M	519	11 TB	183 MB	Real-time

Legend: R = Read, W = Write, RW = Balanced, Con = Connections

Technology Selection Based on Numbers

Decision Framework: Choose Database

When to Use PostgreSQL/MySQL (Relational)

Use when numbers show:

- ✓ Storage < 10 TB (fits in single instance)
- ✓ Write QPS < 10K (single master handles it)
- ✓ Read QPS < 100K (with 5-10 read replicas)

- ✓ Need ACID transactions
- ✓ Complex queries with JOINs

"Our calculations show:

- TinyURL: 8 TB storage, 40 write QPS → PostgreSQL ✓
- Pastebin: 21 TB storage, 115 write QPS → PostgreSQL ✓
- Amazon orders: ACID needed for payments → PostgreSQL ✓
- Parking Lot: 11 TB, 345 write QPS → PostgreSQL ✓

****Justification Template:****

"Our calculations show:

- Storage: X TB (< 10 TB threshold)
- Write QPS: Y (< 10K single master limit)
- Read QPS: Z (< 100K with replicas)
- Need: ACID for transactions

Therefore, PostgreSQL is sufficient and simpler than distributed DB"

When to Use Cassandra (NoSQL, Write-Heavy)

****Use when numbers show:****

- ✓ Write QPS > 10K (need distributed writes)
- ✓ Storage > 10 TB (need horizontal scaling)
- ✓ Time-series data (sorted by timestamp)
- ✓ Can accept eventual consistency
- ✓ Need linear scalability

"Our calculations show:

- Twitter tweets: 44 PB storage, 7K write QPS → Cassandra ✓
- WhatsApp messages: 1.46 EB, 1.4M write QPS → Cassandra ✓
- Uber locations: 2 PB, 1M write QPS → Cassandra ✓
- Discord messages: 10 PB, 348K write QPS → Cassandra ✓

****Justification Template:****

"Our calculations show:

- Write QPS: 1M/sec (far exceeds single DB limit of 10K)
- Storage: 2 PB (requires distributed storage)
- Data type: Time-series (naturally sorted)
Therefore, need Cassandra for:
- Linear write scalability (each node adds 10K write QPS)
- Horizontal storage scaling (each node adds 2 TB)
- No single point of failure"

When to Use MongoDB (NoSQL, Flexible)

Use when numbers show:

- ✓ Write QPS: 1K-50K (moderate)
- ✓ Storage: 1-100 TB
- ✓ Schema flexibility needed
- ✓ Read QPS < 100K
- ✓ Document-oriented data

Examples from above:

- Reddit posts: 250 TB, 6 write QPS → Could use MongoDB
- LinkedIn profiles: 167 TB, moderate complexity → MongoDB option

Justification Template:

"With 10K write QPS and 50 TB storage:

- Too large for single PostgreSQL (> 10 TB)
- Not enough write load for Cassandra
- Schema varies (profiles have different fields)
Therefore, MongoDB provides:
 - Horizontal scaling (sharding)
 - Schema flexibility
 - Simpler than Cassandra for this scale"

Decision Framework: Choose Cache

When to Use Redis

Use when numbers show:

- ✓ Cache size: Any (MB to TB)
- ✓ Need data structures (Lists, Sets, Sorted Sets)
- ✓ Need persistence
- ✓ Need Pub/Sub
- ✓ Operations QPS: Up to 1M per instance

Use Redis for (from examples):

- TinyURL: 2.5 GB cache → Single Redis ✓
- Twitter timelines: 1.1 TB → Redis Cluster (20 nodes) ✓
- Rate limiter: 400 MB counters → Redis (atomic INCR) ✓
- Uber: 1 GB geolocation → Redis (GEO commands) ✓
- All systems: Redis is the default choice

Justification by Numbers:

Cache < 1 GB:

"Cache is 500 MB, single Redis instance (16 GB) handles it easily with 15GB headroom for growth"

Cache 1-100 GB:

"Cache is 20 GB, use 2-3 Redis instances for redundancy

Each handles 100K+ QPS"

Cache 100 GB - 10 TB:

"Cache is 1.1 TB, use Redis Cluster with 20 nodes

Each node: 55 GB, handles 50K QPS

Total capacity: 1M QPS"

Cache > 10 TB:

"Cache is 60 TB, use Redis Cluster with 100 nodes

Each node: 600 GB, distribute evenly

Total capacity: 5M QPS"

When to Use Memcached

Use when numbers show:

- ✓ Pure key-value (no complex data structures)
- ✓ No persistence needed
- ✓ Multi-threaded benefit needed
- ✓ Extremely high QPS (> 1M per instance)

Rare in interviews - Redis almost always better choice

Only use Memcached:

"If calculations show > 1M QPS per cache instance and pure key-value storage, Memcached's multi-threading gives 10-20% performance edge. Otherwise, use Redis."

Decision Framework: When to Use CDN

Use CDN when numbers show:

- ✓ Bandwidth > 1 GB/s (CDN essential)
- ✓ Static content (images, videos, CSS, JS)
- ✓ Global users (serve from edge)
- ✓ Can accept caching (TTL-based)

Cost/Benefit Analysis:

Without CDN:

- Bandwidth: $50 \text{ GB/s} \times \$0.08/\text{GB} = \$4,000/\text{hour} = \$2.9\text{M}/\text{month}$

With CDN (90% cache hit):

- Origin: $5 \text{ GB/s} \times \$0.08/\text{GB} = \$400/\text{hour} = \$288\text{K}/\text{month}$
- CDN: $45 \text{ GB/s} \times \$0.02/\text{GB} = \$900/\text{hour} = \$648\text{K}/\text{month}$
- Total: $\$936\text{K}/\text{month}$ (saves $\$2\text{M}/\text{month} = 68\%$ savings!)

Use CDN for:

- Instagram: 57 GB/s → CDN mandatory
- Netflix: 8.3 TB/s → CDN absolutely critical
- Twitter: 19 GB/s → CDN essential
- Amazon: 15 GB/s → CDN required

Don't need CDN:

- Rate Limiter: 36 KB/s → Too small
- Parking Lot: 1 MB/s → Not worth it

Justification Template:

"Our bandwidth calculation shows X GB/s outgoing.

At this scale:

- Without CDN: \$Y/month in bandwidth costs

- With CDN (90% cache hit): \$Z/month
- Savings: \$(Y-Z)/month

Additionally:

- Latency improvement: 200ms → 20ms (edge vs origin)
- Origin load reduction: 90%
- Global reach: 400+ edge locations

Therefore, CDN is essential, not optional."

Decision Framework: Message Queue

Use Kafka when numbers show:

- ✓ Event throughput > 10K events/sec
- ✓ Need durability (no message loss)
- ✓ Need replay capability
- ✓ Multiple consumers for same events

Examples:

- Twitter: 7K tweet events/sec → Kafka ✓
- WhatsApp: 1.4M message events/sec → Kafka ✓
- Discord: 348K message events/sec → Kafka ✓
- Notifications: 36K events/sec → Kafka ✓

Justification Template:

"With 100K events/second:

- Kafka throughput: 1M+ events/sec per broker
- Need: $100K \div 100K = 1$ broker minimum
- With redundancy: 3 brokers (replication factor 3)
- Cost: ~\$500/month

Compared to alternatives:

- RabbitMQ: Max ~50K/sec per node (need 2+ nodes)
- SQS: \$0.40 per million = \$3.5K/month (more expensive)
- Redis Pub/Sub: No persistence, lose messages on failure

Therefore, Kafka for durability + throughput at scale"

```
#### Decision Framework: Object Storage
```

```
**Use S3/Cloud Storage when:**
```

- ✓ Media files (images, videos, documents)
- ✓ Storage > 1 TB
- ✓ Unstructured data
- ✓ Need durability (11 nines)

Cost Analysis:

S3 Standard:

- \$0.023/GB/month
- 100 TB = \$2,300/month
- 1 PB = \$23,000/month
- 1 EB = \$23M/month

S3 with lifecycle:

- Hot (< 30 days): Standard
- Warm (30-90 days): Infrequent Access (\$0.0125/GB)
- Cold (> 90 days): Glacier (\$0.004/GB)
- Savings: ~60%

Examples:

- Instagram: 58 PB → S3 with lifecycle
- Netflix: 300 TB catalog → S3 Standard
- WhatsApp: 1.46 EB → S3 with aggressive lifecycle

```
## Technology Decision Matrix
```

```
### Database Selection Matrix
```

Criteria	Recommendation
Write QPS < 1K Single PostgreSQL	
Write QPS 1K-10K PostgreSQL + Write Sharding	

Write QPS > 10K Cassandra / DynamoDB
Read QPS < 10K Single PostgreSQL
Read QPS 10K-100K PostgreSQL + 5-10 Read Replicas
Read QPS > 100K Add Redis Cache (80%+ hit rate)
Storage < 1 TB Single PostgreSQL Instance
Storage 1-10 TB PostgreSQL with partitioning
Storage 10-100 TB Sharded PostgreSQL OR Cassandra
Storage > 100 TB Cassandra / DynamoDB
Need ACID PostgreSQL (always)
Need Transactions PostgreSQL (always)
Eventual Consistency OK Cassandra (better performance)
Time-series data Cassandra / TimescaleDB
Social graph Neo4j / Cassandra wide columns
Document store MongoDB / Elasticsearch
Key-value Redis / DynamoDB

Cache Selection Matrix

Cache Size Configuration

< 1 GB Single Redis (16 GB instance)
1-100 GB 2-5 Redis instances (HA)
100 GB - 1 TB Redis Cluster (10-20 nodes)
1-10 TB Redis Cluster (50-100 nodes)
> 10 TB Redis Cluster (100+ nodes)
QPS < 10K Single Redis
QPS 10K-100K Redis with 2-3 replicas
QPS 100K-1M Redis Cluster (10+ nodes)
QPS > 1M Redis Cluster (50+ nodes)
TTL required Redis (native support)
Complex data structures Redis (Lists, Sets, Sorted Sets)

Pub/Sub needed	Redis (built-in)
Persistence needed	Redis (RDB + AOF)
Pure key-value only	Memcached (slightly faster)

↓

Complete Technology Recommendations by System

TinyURL

Numbers: 4K read QPS, 40 write QPS, 8 TB storage

Tech Stack:

✓ Database: PostgreSQL

Why: 8 TB fits in single instance, 4K QPS easily handled

✓ Cache: Single Redis (16 GB)

Why: 2.5 GB cache fits easily, 4K QPS trivial for Redis

✓ Web Servers: 10-15 servers

Why: 12K peak QPS ÷ 1K per server = 12 servers

✗ CDN: Not needed

Why: 800 KB/s bandwidth is tiny

✗ Message Queue: Not needed

Why: Synchronous flow sufficient at this scale

Twitter

Numbers: 1.4M read QPS, 7K write QPS, 44 PB storage

Tech Stack:

✓ Database: Cassandra for tweets

Why: 44 PB requires distributed storage, 7K write QPS

Each node: 2 TB storage, 100 write QPS

Nodes needed: 44 PB ÷ 2 TB = 22,000 nodes

Write capacity: 22K nodes × 100 = 2.2M QPS ✓

✓ Database: PostgreSQL for users

Why: 400M users × 2 KB = 800 GB (fits in single instance)

Need ACID for user operations

✓ Cache: Redis Cluster (20 nodes)

Why: $1.1 \text{ TB cache} \div 20 = 55 \text{ GB per node}$

Handles 1M read QPS (50K per node)

✓ CDN: CloudFront

Why: 19 GB/s bandwidth, reduces to 950 MB/s origin (95% savings)

✓ Message Queue: Kafka

Why: 7K tweet events/sec, need fanout to millions

✓ Search: Elasticsearch

Why: Full-text search on 44 PB of tweets

Web Servers: 200-300 servers ($1.4 \text{ M QPS} \div 5 \text{ K per server}$)

WhatsApp

Numbers: 1.4M message QPS, 1.46 EB storage

Tech Stack:

✓ Database: Cassandra for messages

Why: 1.4M write QPS requires massive distribution

1.46 EB requires huge horizontal scaling

Nodes needed: $1.46 \text{ EB} \div 2 \text{ TB} = 750,000 \text{ nodes!}$

(In reality: Messages deleted after 30 days, much less)

✓ Database: PostgreSQL for users

Why: $2 \text{ B users} \times 1 \text{ KB} = 2 \text{ TB}$ (manageable)

Need ACID for authentication

✓ Cache: Redis Cluster (200 nodes)

Why: $2.5 \text{ TB} \div 200 = 12.5 \text{ GB per node}$

Handles presence, sessions, message queues

✓ WebSocket Servers: 8,000-16,000

Why: $500 \text{ M connections} \div 65 \text{ K per server} = 7,692 \text{ servers}$

✓ Message Queue: Kafka

Why: 1.4M events/sec, need reliable delivery

Brokers needed: $1.4 \text{ M} \div 100 \text{ K per broker} = 14 \text{ brokers}$

✓ Object Storage: S3

Why: 1.46 EB of media files

Cost: ~\$300K/month with lifecycle policies

✗ CDN: Limited use

Why: Messages encrypted, can't cache

Only for profile pictures, minimal benefit

Netflix

Numbers: 8.3 TB/s bandwidth, 300 TB catalog

Tech Stack:

✓ Database: PostgreSQL for metadata

Why: 10K videos \times 10 KB = 100 MB (tiny!)

70K API QPS handled with 10 read replicas

✓ Object Storage: S3 for videos

Why: 300 TB catalog, need 11-nines durability

Cost: 300 TB \times \$0.023 = \$7K/month (cheap!)

✓ CDN: Custom CDN (Open Connect)

Why: 8.3 TB/s bandwidth is massive!

95% from CDN = 7.8 TB/s from edge

Origin only: 415 GB/s

CDN absolutely critical - without it, impossible

✓ Cache: Redis for metadata

Why: 100 MB metadata, cache everything

Single Redis instance sufficient

✓ Cache: 26 TB at CDN edges

Why: Hot content (20% of catalog)

Distributed: 260 GB per edge \times 100 edges

Web Servers: 100-150 (70K API QPS \div 500 per server)

Uber

Numbers: 1M location update QPS, 2 PB storage, 1 GB cache

Tech Stack:

✓ Database: Cassandra for locations

Why: 1M write QPS (extreme!)

2 PB storage requires distribution

Nodes: 2 PB \div 2 TB = 1,000 nodes

Write capacity: 1,000 \times 1K = 1M QPS ✓

✓ Database: PostgreSQL for trips/payments

Why: 12 QPS trips, need ACID for payments

180 TB over 5 years (manageable with partitioning)

✓ Cache: Single Redis (16 GB)

Why: Only 1 GB needed for real-time matching

Redis GEO commands for nearby search

Handles 100K QPS easily

✓ Message Queue: Kafka

Why: 1M location events/sec

Brokers: $1M \div 100K = 10$ brokers minimum

Web Servers: 1,500 servers ($1M \text{ QPS} \div 700 \text{ per server}$)

Rate Limiter

Numbers: 36K QPS, 400 MB storage

Tech Stack:

✓ Cache: Redis (primary storage!)

Why: 400 MB fits in single instance (16 GB)

Need atomic operations (INCR)

Need TTL support (auto-cleanup)

Sub-millisecond latency required

Redis handles 100K-1M QPS easily

✓ Database: PostgreSQL for rules

Why: $10K \text{ rules} \times 500 \text{ bytes} = 5 \text{ MB}$ (tiny!)

Infrequent updates

Need ACID for rule changes

✗ Message Queue: Not needed

Why: Synchronous flow, no async processing

Application Servers: 5-10 ($36K \text{ QPS} \div 5K \text{ per server}$)

Critical: Rate limiter adds < 10ms latency

Redis achieves < 1ms, perfect!

Scaling Thresholds & Technology Transitions

When to Shard Database

Threshold Indicators:

Single PostgreSQL limits:

- Write QPS: 10K
- Read QPS: 100K (with replicas)
- Storage: 10-16 TB

When to shard:

- ✓ Write QPS > 10K
- ✓ Storage > 10 TB
- ✓ Single table > 100M rows

Sharding Examples:

Instagram (500M users):

- User table: 500M rows
- Shard by user_id (hash-based)
- 10 shards \times 50M users each
- Each shard: $50M \times 2\text{ KB} = 100\text{ GB}$ ✓

Twitter (400M users):

- Users: 400M rows \rightarrow Shard by user_id (5 shards)
- Tweets: 200B rows \rightarrow Use Cassandra (can't shard PostgreSQL to this scale)

When to Add Read Replicas

Threshold: Read QPS > 10K

Rule of thumb:

- Master: 10K write QPS
- Each replica: +10K read QPS

Examples:

TinyURL (4K read QPS):

Replicas: $4K \div 10K = 0.4 \rightarrow 1\text{ master} + 1\text{ replica}$ sufficient

Amazon (300K read QPS):

Replicas: $300K \div 10K = 30$ replicas

But with 90% cache hit rate:

Actual DB QPS: $300K \times 0.1 = 30K$

Replicas needed: $30K \div 10K = 3$ replicas

Configuration: 1 master + 5 read replicas (with headroom)

When to Move to NoSQL

Decision Tree:

Storage > 100 TB AND Write QPS > 10K?

→ Cassandra (distributed, write-optimized)

Storage < 100 TB BUT Write QPS > 50K?

→ Cassandra (write throughput critical)

Storage > 10 TB AND Schema flexible?

→ MongoDB (easier than sharded PostgreSQL)

Need ACID no matter what?

→ Stay with PostgreSQL, accept complexity

→ Use sharding if needed

→ Examples: Banking, E-commerce orders

Time-series data > 1 TB?

→ Cassandra or TimescaleDB

→ Examples: Metrics, logs, sensor data

Cost-Based Technology Decisions

Database Cost Comparison

Scenario: 100 TB storage, 50K QPS

PostgreSQL (Sharded):

- 10 shards × r5.4xlarge = \$2,500/month
- Complex sharding logic
- Total: ~\$3,000/month

Cassandra (Managed):

- 50 nodes × \$100/month = \$5,000/month
- Simpler operations
- Better scalability
- Total: ~\$5,000/month

MongoDB Atlas:

- Cluster: ~\$4,000/month
- Middle ground

Decision: If budget tight and can manage sharding → PostgreSQL

If need simplicity and scale → Cassandra

Cache Cost Comparison

Scenario: 1 TB cache needed

Redis Cluster (Self-managed):

- 20 nodes × r5.large (\$50/month) = \$1,000/month
- Operational overhead
- Full control

ElastiCache (Managed Redis):

- 20 × cache.r5.large = \$2,000/month
- No operational overhead
- Auto-failover

Memcached:

- Slightly cheaper (~\$1,800/month)
- Fewer features

Decision: Use managed Redis unless budget critical

Interview Justification Examples

Example 1: Justify Cassandra for Twitter

"Let me justify why Cassandra for tweets:

Our calculations show:

1. Write QPS: 7K (peak)
 - Single PostgreSQL limit: 10K

- Close to limit, but manageable

2. Storage: 44 PB

- Single PostgreSQL limit: 16 TB
- Would need: $44 \text{ PB} \div 16 \text{ TB} = 2,750 \text{ shards!}$
- Cassandra: $44 \text{ PB} \div 2 \text{ TB} = 22\text{K nodes}$ (manageable)

3. Data pattern: Time-series

- Tweets naturally sorted by time
- Cassandra optimized for time-series
- PostgreSQL requires complex partitioning

4. Consistency: Eventual OK

- Tweet appears in timeline within 1 second acceptable
- Don't need immediate consistency
- Cassandra's eventual consistency works

5. Availability: Must be 99.99%

- Cassandra: No single point of failure
- PostgreSQL: Master failure = downtime

Therefore, Cassandra is the right choice for Twitter's tweet storage."

Summary: The Golden Rules

1. **Round liberally** – 387M → 400M
2. **State assumptions** – "Assuming X..."
3. **Show your work** – Write calculations down
4. **Use shortcuts** – Daily $\div 100\text{K} = \text{QPS}$
5. **Apply 80-20 rule** – Cache 20%, serve 80%
6. **Mention peak traffic** – Average $\times 3$
7. **Don't forget media** – Images dominate storage
8. **Include replication** – $\times 3$ for redundancy
9. **Sanity check** – Does 1 EB for Twitter make sense? No!
10. **Mention optimizations** – CDN, compression, caching
11. **Justify technology choices** – Use numbers to back decisions
12. **Compare alternatives** – Show cost/performance tradeoffs
13. **Scale path** – Explain how system grows

Final Interview Checklist

****Before presenting your design, verify:****

- **Calculations Complete****
- [] Traffic (QPS) calculated for reads and writes
 - [] Storage estimated with retention period
 - [] Bandwidth computed with CDN impact
 - [] Cache sized using 80-20 rule
 - [] Peak traffic considered (2-3x)

- **Technology Justified****
- [] Database choice explained with numbers
 - [] Cache strategy supported by calculations
 - [] CDN decision backed by cost analysis
 - [] Message queue justified by throughput
 - [] Each technology linked to specific requirements

- **Presentation Quality****
- [] Used round numbers throughout
 - [] Showed all units in calculations
 - [] Stated assumptions clearly
 - [] Sanity-checked results
 - [] Compared to real-world systems

References

Books

1. **"Designing Data-Intensive Applications"** – Martin Kleppmann
2. **"System Design Interview"** – Alex Xu (Volume 1 & 2)
3. **"Database Internals"** – Alex Petrov

Online Resources

1. **System Design Primer** – GitHub ([donnemartin](#))
2. **Google's Latency Numbers** – Jeff Dean
3. **AWS Calculator** – For cost estimates
4. **High Scalability Blog** – Real-world numbers

Videos

1. **"System Design Course"** – YouTube (Gaurav Sen)
2. **"Back of Envelope Calculations"** – YouTube (SystemDesignInterview)
3. **"Scaling to Millions"** – YouTube (Engineering with Utsav)

****END OF GUIDE****

This guide provides a comprehensive framework for back-of-the-envelope calculations and technology selection in system design interviews. Use it as a reference, practice the examples, and adapt the templates to your specific interview questions.

****Good luck with your system design interviews!****

