

# Non-Functional Requirements Guide for System Design Interviews

---

Latency, Throughput, Availability, Scalability, and More

## Table of Contents

1. Introduction
  2. Performance Requirements
  3. Scalability Requirements
  4. Availability & Reliability
  5. Consistency Requirements
  6. Security Requirements
  7. Maintainability & Operability
  8. Cost & Resource Constraints
  9. Compliance & Legal
  10. NFR Trade-offs
  11. How to Discuss NFRs in Interviews
  12. Common Interview Scenarios
  13. Measuring & Monitoring
  14. Interview Tips
- 

## Introduction

**Non-Functional Requirements (NFRs)** define HOW a system should behave, as opposed to functional requirements which define WHAT a system should do.

### Why NFRs Matter

NFRs are critical in system design interviews because they:

- **Drive Architecture Decisions:** Availability needs affect redundancy
- **Define Success Criteria:** What does "fast" mean?
- **Create Trade-offs:** Performance vs cost, consistency vs availability
- **Show Experience:** Production-minded thinking

### Functional vs Non-Functional

Functional Requirement (WHAT) :  
"Users should be able to send messages to each other"

Non-Functional Requirements (HOW) :  
– Messages delivered in < 100ms (Latency)  
– System handles 10K messages/sec (Throughput)  
– System available 99.99% of time (Availability)

- Scales to 1M concurrent users (Scalability)
- Messages never lost (Reliability)
- End-to-end encrypted (Security)

## Main Categories of NFRs

1. PERFORMANCE
  - Latency
  - Throughput
  - Response Time
2. SCALABILITY
  - Horizontal scaling
  - Vertical scaling
  - Elasticity
3. RELIABILITY
  - Availability
  - Fault Tolerance
  - Durability
4. CONSISTENCY
  - Strong consistency
  - Eventual consistency
  - Causal consistency
5. SECURITY
  - Authentication
  - Authorization
  - Encryption
6. OPERABILITY
  - Maintainability
  - Observability
  - Deployability
7. COST
  - Infrastructure costs
  - Operational costs
  - Development costs

---

## Performance Requirements

### 1. Latency

**Definition:** Time between request initiation and response receipt

#### Types of Latency

### Network Latency:

- Same datacenter: 1–2ms
- Cross-datacenter (same region): 5–10ms
- Cross-region (US East to West): 60–80ms
- Cross-continent (US to Europe): 100–150ms
- Cross-continent (US to Asia): 150–300ms

### Application Latency:

- Cache hit (Redis): <1ms
- Cache hit (Memcached): 1–5ms
- Database query (indexed): 5–10ms
- Database query (unindexed): 100–1000ms
- API call (internal): 10–50ms
- API call (external): 100–500ms

### End-to-End Latency Components:

Total = Network + Processing + Queue + Database

### Example (Web Request):

- DNS lookup: 20ms
- TCP handshake: 50ms
- TLS handshake: 100ms
- Request send: 10ms
- Server processing: 50ms
- Database query: 10ms
- Response send: 10ms

Total: 250ms (first request)

Subsequent: 70ms (connection reused)

## Latency Requirements by Use Case

Use Case	Acceptable Latency	Excellent Latency
Real-time Chat	< 1s	< 100ms
Social Media Feed	< 2s	< 500ms
Search Results	< 1s	< 200ms
E-commerce Checkout	< 3s	< 1s
Video Streaming	< 5s start	< 2s start
Gaming (action)	< 100ms	< 50ms
Trading Platform	< 50ms	< 10ms

File Upload	< 10s	< 5s
Dashboard Load	< 3s	< 1s
Email Delivery	< 5min	< 1min

## Optimizing Latency

### Strategies:

1. Caching
  - In-memory cache (Redis): <1ms
  - CDN edge cache: 10–50ms
  - Browser cache: 0ms (no request)
2. Geographic Distribution
  - Multiple regions: Reduces network latency
  - CDN edge locations: Serve from nearest
  - Database replicas: Read from nearest
3. Connection Optimization
  - Keep-alive connections
  - HTTP/2 multiplexing
  - Connection pooling
  - Reduce handshakes
4. Database Optimization
  - Indexes on query columns
  - Query optimization
  - Read replicas
  - Database caching
5. Async Processing
  - Return immediately
  - Process in background
  - Notify when complete

### Example:

Requirement: Page load < 1s

#### Without Optimization:

- Server: 500ms
  - Database: 300ms
  - Rendering: 200ms
- Total: 1000ms (at limit!)

#### With Optimization:

- Cache hit: 5ms (instead of 300ms DB)
- CDN for static: 20ms (instead of 100ms)

- Async non-critical: 0ms wait
- Total: 225ms (4x better!)

## 2. Throughput

**Definition:** Amount of work performed per unit time (requests/sec, transactions/sec, MB/sec)

### Throughput Benchmarks

#### Database Systems:

- MySQL: 1K–10K queries/sec (single instance)
- PostgreSQL: 5K–15K queries/sec
- MongoDB: 10K–50K operations/sec
- Cassandra: 100K–1M writes/sec (cluster)
- Redis: 100K–1M operations/sec

#### Message Queues:

- RabbitMQ: 20K–50K msgs/sec (single node)
- Kafka: 1M+ msgs/sec (cluster)
- AWS SQS: 300K msgs/sec (standard)

#### Web Servers:

- Node.js: 5K–10K req/sec (single core)
- Nginx: 50K–100K req/sec (reverse proxy)
- Java/Spring: 10K–20K req/sec

#### CDN:

- Cloudflare: Millions of req/sec (global)
- CloudFront: Millions of req/sec

### Calculating Required Throughput

Example: Social Media Platform

Daily Active Users (DAU): 10M  
 Average actions per user: 50/day  
 Peak factor: 3x average

#### Average RPS:

$$\begin{aligned}
 &= (10M * 50) / (24 * 3600) \\
 &= 500M / 86400 \\
 &= 5,787 \text{ requests/sec}
 \end{aligned}$$

#### Peak RPS:

$$\begin{aligned}
 &= 5,787 * 3 \\
 &= 17,361 \text{ requests/sec}
 \end{aligned}$$

Round up for safety: 20,000 RPS

Server Capacity Planning:

- Each server: 1000 RPS
- Servers needed:  $20,000 / 1000 = 20$  servers
- Add 50% buffer: 30 servers
- For 99.99% availability: Deploy across 3 AZs
- Per AZ: 10 servers

Database Capacity:

- 60% reads, 40% writes
- Reads: 12,000 RPS
- Writes: 8,000 RPS
- Read replicas:  $12,000 / 5,000 = 3$  replicas
- Primary for writes: 8,000 RPS (within limits)

## Optimizing Throughput

Strategies:

1. Horizontal Scaling
  - Add more servers
  - Distribute load
  - Load balancer
2. Caching
  - Reduce database load
  - Increase cache hit ratio
  - Tiered caching (L1, L2, L3)
3. Async Processing
  - Queue heavy operations
  - Process in background
  - Increase request capacity
4. Connection Pooling
  - Reuse connections
  - Reduce overhead
  - Optimize pool size
5. Batch Operations
  - Combine multiple operations
  - Reduce round trips
  - Bulk inserts/updates
6. Database Optimization
  - Sharding
  - Read replicas
  - Query optimization
  - Denormalization

- Throughput vs Latency Trade-off:
- Batching: ↑ Throughput, ↑ Latency
  - Parallel: ↑ Throughput, ↓ Latency
  - Caching: ↑ Throughput, ↓ Latency

### 3. Response Time

**Definition:** Time to complete an operation (includes processing + waiting)

#### Percentile-based SLA

Why Percentiles Matter:

Average can be misleading:

- 99 requests: 10ms
- 1 request: 1000ms
- Average: 19.9ms (looks good!)
- But 1% of users have terrible experience

Better: Use Percentiles

- P50 (Median): 10ms – half of requests
- P95: 15ms – 95% of requests
- P99: 50ms – 99% of requests
- P99.9: 200ms – 99.9% of requests

Example SLA:

"API response time:

- P50: < 100ms
- P95: < 500ms
- P99: < 1000ms"

Tail Latency Problem:

- P50: 50ms (great!)
- P99: 2000ms (terrible!)
- Means 1% of users wait 2 seconds
- At scale: 1% = 10,000 users!

Focus on: P95 and P99, not just average

#### Response Time Requirements

Interactive vs Batch:

Interactive (User waiting):

- Target: < 200ms (feels instant)
- Acceptable: < 1s (slight delay)

- Maximum: < 3s (noticeable delay)
- Beyond 3s: Show progress indicator

#### Batch Processing:

- Minutes to hours acceptable
- Show progress updates
- Allow cancellation
- Notify on completion

#### API Response Time by Type:


API Type	Target
GET (cached)	< 50ms
GET (uncached)	< 200ms
POST (simple)	< 500ms
POST (complex)	< 2s
Search query	< 500ms
Report generate	Async

## Scalability Requirements

### 1. Types of Scalability

#### Horizontal Scaling (Scale Out):

Add more servers to handle load




#### Pros:

- ✓ Linear scaling (usually)
- ✓ Fault tolerance (redundancy)
- ✓ Cost-effective
- ✓ Cloud-friendly

#### Cons:

- ✗ Complexity (distributed state)
- ✗ Load balancing needed
- ✗ Data consistency challenges

Vertical Scaling (Scale Up):  
Upgrade to more powerful server



Pros:

- Simpler (no distribution)
- No data consistency issues
- Less operational overhead

Cons:

- Limited by hardware
- Single point of failure
- Expensive at scale
- Downtime for upgrades

Best Practice: Design for horizontal scaling

## 2. Scalability Metrics

Key Metrics:

1. User Scalability
  - Current: 100K users
  - Target: 10M users (100x)
  - Growth rate: 2x per year
2. Data Scalability
  - Current: 1TB
  - Growth: 100GB/month
  - 5-year projection: 7TB
3. Traffic Scalability
  - Current: 5K RPS
  - Peak: 15K RPS (3x)
  - Black Friday: 50K RPS (10x)
4. Geographic Scalability
  - Current: US only
  - Phase 2: Add Europe
  - Phase 3: Global

Scaling Patterns:


- Read-heavy: Add read replicas, caching
- Write-heavy: Sharding, partitioning

- Compute-heavy: Horizontal scaling, async processing
- Storage-heavy: Distributed storage, tiered storage

### 3. Elasticity

**Definition:** Ability to automatically scale up/down based on load

**Auto-Scaling Strategy:**



**Scale-Out Triggers:**

- CPU usage > 70% for 5 minutes
- Request queue depth > 1000
- Response time P95 > 500ms
- Active connections > 8000

**Scale-In Triggers:**

- CPU usage < 30% for 15 minutes
- Request queue depth < 100
- After scale-out, load normalized

**Scaling Rules:**

- Minimum instances: 3 (HA)
- Maximum instances: 100 (cost limit)
- Scale-out: Add 30% capacity
- Scale-in: Remove 10% capacity
- Cooldown: 5 minutes between actions

**Benefits:**

- Cost optimization
- Handle traffic spikes

- Maintain performance
- Automatic management

## Availability & Reliability

### 1. Availability (Uptime)

**Definition:** Percentage of time system is operational and accessible

#### SLA Levels

SLA	Downtime/Year	Downtime/Mo	Common Use
90%	36.5 days	3 days	Internal tools
95%	18.25 days	1.5 days	Non-critical
99%	3.65 days	7.2 hours	Standard
99.9%	8.76 hours	43.2 min	High priority
99.95%	4.38 hours	21.6 min	Critical
99.99%	52.56 min	4.32 min	Mission-crit
99.999%	5.26 min	25.9 sec	Extreme HA

#### Cost of Nines:

Each additional "9" roughly doubles cost

- 99%: Baseline
- 99.9%: 2-3x cost
- 99.99%: 5-10x cost
- 99.999%: 20-50x cost

## Achieving High Availability

#### Strategies:

##### 1. Redundancy (Multiple Instances)



If S1 fails, S2 and S3 continue

##### 2. Geographic Distribution (Multi-Region)

US-East


US-West

Europe

If US-East fails, traffic to US-West

3. Health Checks & Failover
  - Active monitoring
  - Automatic failover
  - Traffic rerouting
  - Self-healing
4. Fault Isolation
  - Bulkheads (isolate failures)
  - Circuit breakers
  - Graceful degradation
  - Fail independently
5. Zero-Downtime Deployments
  - Blue-green deployment
  - Rolling updates
  - Canary releases
  - Feature flags

Multi-AZ Deployment:



Availability Calculation:

Single AZ: 99.5% availability  
Multi-AZ (3):  $1 - (0.005)^3 = 99.9999875\%$   
(Assumes independent failures)

## 2. Reliability (MTBF & MTTR)

Key Metrics:

MTBF (Mean Time Between Failures):

- Average time system runs before failure
- Higher is better
- Example: 720 hours (30 days)

MTTR (Mean Time To Repair):

- Average time to fix when failure occurs
- Lower is better

- Example: 1 hour

$\text{Availability} = \text{MTBF} / (\text{MTBF} + \text{MTTR})$

Example:  $720 / (720 + 1) = 99.86\%$

Reliability Patterns:

#### 1. Redundancy

- Active-Active: Both handle traffic
- Active-Passive: Standby for failover
- N+1: N needed, 1 spare

#### 2. Replication

- Synchronous: Consistent, slower
- Asynchronous: Faster, may lag
- Multi-master: Write anywhere

#### 3. Backup & Recovery

- Regular backups (hourly, daily)
- Point-in-time recovery
- Cross-region replication
- Disaster recovery plan

#### 4. Chaos Engineering

- Intentionally inject failures
- Test recovery procedures
- Improve resilience
- Netflix Chaos Monkey

### 3. Durability

Definition: Data is not lost once written

Durability Levels:

Storage Type	Durability	Use Case
Memory (Redis)	None/Low	Cache, temp
Single disk	95%	Development
RAID	99.9%	Production
Replicated (3x)	99.999%	Critical data
S3	99.99999999% (11 nines)	Archive

Achieving Durability:

1. Replication
  - Write to multiple copies
  - Acknowledge after N replicas
  - Quorum:  $W+R > N$

Example (3 replicas):

- Write to 2 replicas before ack
- Read from 2 replicas for consistency
- W=2, R=2, N=3:  $W+R > N$  (consistent)

2. Write-Ahead Logging (WAL)

- Log before applying
- Replay on recovery
- Ensures no data loss

3. Snapshots

- Periodic full copies
- Point-in-time recovery
- Backup strategy

4. Cross-Region Replication

- Protect against regional failure
- Higher latency
- Disaster recovery

## Consistency Requirements

### 1. CAP Theorem

Pick 2 of 3:

Consistency (C):


All nodes see same data at same time


Availability (A):

Every request receives response (success/failure)

Partition Tolerance (P):

System continues despite network partition





In distributed systems: P is mandatory (network always fails)  
 Real choice: CP vs AP

CP (Consistency + Partition Tolerance):

- Strong consistency
- May refuse requests during partition
- Examples: Banking, inventory
- Systems: MongoDB, HBase, Zookeeper

AP (Availability + Partition Tolerance):

- Always available
- May return stale data during partition
- Examples: Social media, recommendations
- Systems: Cassandra, DynamoDB, Riak

## 2. Consistency Models

**Strong Consistency:**

- Read always returns latest write
- Linearizable
- Highest guarantee

**Timeline:**

Write(X=1) → Complete  
 Read(X) → Returns 1 ✓

Example: Bank account balance

**Eventual Consistency:**

- Reads may return stale data
- Eventually all reads return latest
- Highest availability

**Timeline:**

Write(X=1) → Complete  
 Read(X) → Returns 0 (stale)  
 ...wait...  
 Read(X) → Returns 1 ✓ (caught up)

Example: Social media likes count

**Causal Consistency:**

- Causally related operations ordered

- Concurrent operations may be unordered
- Middle ground

**Timeline:**

User A: Post → Like own post (causal order maintained)

User B: Like post (concurrent, may see before post!)

**Read-Your-Own-Writes:**

- User sees their own updates immediately
- Others may see stale data temporarily

**Session Consistency:**

- Within session: consistent
- Across sessions: eventual

**Monotonic Reads:**

- Once you read X=1, never read X=0
- Time doesn't go backward

**Consistency Trade-offs:**

Strong → High latency, lower availability

Eventual → Low latency, high availability

### 3. ACID vs BASE

**ACID (Traditional Databases):**

**Atomicity:** All or nothing

- Transaction succeeds completely or fails completely

**Consistency:** Valid state

- Data adheres to rules and constraints

**Isolation:** Transactions independent

- Concurrent transactions don't interfere

**Durability:** Permanent writes

- Committed data survives crashes

**Use Case:** Financial transactions, inventory

**BASE (Modern Distributed Systems):**

**Basically Available:**

- System remains available
- May return stale data

**Soft State:**

- State may change over time
- Even without input (due to eventual consistency)

### Eventual Consistency:

- Will become consistent
- Given enough time

Use Case: Social media, recommendations, analytics

### Comparison:

Property	ACID	BASE
Consistency	Immediate	Eventual
Availability	May block	Always
Scalability	Limited	Excellent
Complexity	Simple	Complex

## Security Requirements

### 1. Authentication & Authorization

#### Authentication (Who are you?):

- Username/password
- OAuth 2.0
- JWT tokens
- Multi-factor (MFA)
- Biometric
- SSO (Single Sign-On)

#### Authorization (What can you do?):

- Role-Based Access Control (RBAC)
- Attribute-Based Access Control (ABAC)
- Fine-grained permissions
- Resource-level access

#### Security Levels:

Level	Requirements
Public	No auth, rate limiting
Authenticated	Login required
Sensitive	Auth + encryption
Financial	Auth + MFA + audit

Healthcare	HIPAA compliance + encryption
------------	-------------------------------

## 2. Data Protection

Encryption:

In Transit:

- TLS/SSL (HTTPS)
- Minimum TLS 1.2
- Certificate management
- Perfect forward secrecy

At Rest:

- Database encryption
- Disk encryption
- Key management (KMS)
- Field-level encryption

End-to-End:

- Only sender and recipient can read
- Server cannot decrypt
- Use case: WhatsApp, Signal

Data Classification:

Level	Protection Requirements
Public	None
Internal	Access control
Confidential	Encryption + access control
Restricted	Encryption + MFA + audit

PII (Personally Identifiable Information):

- Name, email, phone: Encrypted
- SSN, credit cards: Tokenized
- Passwords: Hashed (bcrypt, Argon2)
- Audit logs: Track all access

## 3. Security Best Practices

Defense in Depth (Multiple Layers):

#### Layer 1: Network Security

- Firewall
- DDoS protection
- VPC/Private network
- Security groups

#### Layer 2: Application Security

- Input validation
- SQL injection prevention
- XSS prevention
- CSRF tokens
- Rate limiting

#### Layer 3: Data Security

- Encryption
- Access control
- Data masking
- Audit logging

#### Layer 4: Monitoring

- Intrusion detection
- Anomaly detection
- Security alerts
- Incident response

#### Common Security Requirements:

- OWASP Top 10 compliance
- Regular security audits
- Penetration testing
- Vulnerability scanning
- Security patches
- Incident response plan

## Maintainability & Operability

### 1. Observability

#### Three Pillars of Observability:

1. Logging
  - Application logs
  - Access logs
  - Error logs
  - Audit logs

#### Log Levels:

- ERROR: System failures
- WARN: Potential issues
- INFO: Important events

- DEBUG: Detailed information

## 2. Metrics

- System metrics (CPU, memory, disk)
- Application metrics (RPS, latency)
- Business metrics (orders, revenue)
- Custom metrics

Key Metrics:

- Request rate
- Error rate
- Duration (P50, P95, P99)
- Saturation (resource usage)

## 3. Tracing

- Distributed tracing
- Request flow
- Performance bottlenecks
- Dependency mapping

Observability Stack:

Logs → ELK (Elasticsearch, Logstash, Kibana)  
Metrics → Prometheus + Grafana  
Tracing → Jaeger, Zipkin  
All-in-one → Datadog, New Relic

## 2. Maintainability

Code Maintainability:

- Clean code principles
- Documentation
- Code reviews
- Automated testing
- CI/CD pipelines

System Maintainability:

- Modular architecture
- Loose coupling
- Clear interfaces
- Version control
- Configuration management

Deployment Practices:

- Infrastructure as Code (Terraform)
- Automated deployments
- Rollback capability
- Feature flags
- Canary releases

Documentation Requirements:

- Architecture diagrams
- API documentation
- Runbooks
- Troubleshooting guides
- Post-mortems

## Cost & Resource Constraints

### 1. Infrastructure Costs

Cloud Cost Components:

Compute:

- EC2/VM instances
- Container orchestration
- Serverless functions
- Auto-scaling costs

Storage:

- Database storage
- Object storage (S3)
- Block storage
- Backup storage

Network:

- Data transfer out
- Cross-region transfer
- CDN bandwidth
- Load balancer

Cost Optimization Strategies:

#### 1. Right-sizing

- Monitor actual usage
- Scale down over-provisioned
- Use appropriate instance types

#### 2. Reserved Instances

- 1-year: 30–40% savings
- 3-year: 50–60% savings
- Trade flexibility for cost

#### 3. Spot Instances

- 70–90% savings
- For fault-tolerant workloads
- Batch processing
- Can be interrupted

#### 4. Caching

- Reduce database queries
- Lower data transfer
- Cheaper cache vs compute

#### Example Cost Analysis:

Requirement: 100K RPS, 99.9% availability

##### Option 1: On-Demand

- 50 servers @ \$100/mo = \$5,000/mo
- Database: \$2,000/mo
- Total: \$7,000/mo

##### Option 2: Reserved + Caching

- 30 reserved servers @ \$60/mo = \$1,800/mo
- 20 spot instances @ \$30/mo = \$600/mo
- Redis cache: \$500/mo
- Database: \$1,500/mo
- Total: \$4,400/mo (37% savings)

## Compliance & Legal

### 1. Data Privacy Regulations

#### GDPR (Europe):

- User consent for data collection
- Right to be forgotten
- Data portability
- Data breach notification (72 hours)
- Privacy by design

#### CCPA (California):

- Right to know what data collected
- Right to delete data
- Right to opt-out of data sale
- Non-discrimination

#### HIPAA (Healthcare – US):

- Protected Health Information (PHI)
- Access controls
- Encryption
- Audit trails
- Business Associate Agreements

#### PCI DSS (Payment Cards):

- Never store CVV
- Tokenization
- Encryption
- Network segmentation
- Regular audits

- Data Residency Requirements:**
- EU data must stay in EU
  - China data must stay in China
  - Russia localization law
  - Regional compliance

## 2. Compliance Requirements

**Audit Requirements:**

- Log all access to sensitive data
- Retain logs for X years
- Tamper-proof logs
- Regular audits

**Retention Policies:**

- Financial records: 7 years
- Medical records: 6 years
- Customer data: Until deletion request
- Logs: 1-3 years

**Data Deletion:**

- Soft delete vs hard delete
- Cascading deletes
- Backup retention
- Right to be forgotten compliance

## NFR Trade-offs

### 1. Performance vs Cost

**Trade-off Scenario:**

Requirement: P99 latency < 100ms

**Option A: Over-provision (Safe)**

- 100 servers
- 50% average utilization
- Cost: \$10,000/mo
- P99: 50ms ✓
- Waste: 50% capacity

**Option B: Right-sized (Efficient)**

- 60 servers
- 80% average utilization
- Cost: \$6,000/mo
- P99: 95ms ✓

- Risk: Spikes may breach SLA

Option C: Auto-scaling (Balanced)

- 40–80 servers (dynamic)
- 70% average utilization
- Cost: \$7,000/mo
- P99: 80ms ✓
- Complexity: Higher

Decision Framework:

- Critical system → Option A
- Cost-sensitive → Option B
- Variable load → Option C

## 2. Consistency vs Availability

Trade-off Scenario:

E-commerce Inventory System

Strong Consistency (CP):

```
Check inventory → Lock row
If available → Decrement
If not → Reject order
Other requests wait
```

Pros:

- Never oversell
- Accurate inventory
- Simple logic

Cons:

- Lower availability
- Higher latency
- Blocking

Eventual Consistency (AP):

```
Check inventory (may be stale)
Optimistically reserve
Async validation
Compensate if oversold
```

Pros:

- Higher availability
- Lower latency
- Better UX

Cons:

- ✗ May oversell
- ✗ Need compensation
- ✗ Complex logic

Best Practice: Hybrid

- Hot items: Strong consistency
- Regular items: Eventual consistency
- Based on stock level

### 3. Latency vs Throughput

Trade-off Example:

Message Processing System

Low Latency (Individual):

- Process each message immediately
- No batching
- Latency: 10ms per message
- Throughput: 100 msgs/sec

High Throughput (Batch):

- Batch 100 messages
- Process together
- Latency: 1000ms per message
- Throughput: 10,000 msgs/sec

Comparison:

Approach	Latency	Throughput
Individual	10ms	100/sec
Batch (10)	100ms	1,000/sec
Batch (100)	1000ms	10,000/sec

Decision Factors:

- Real-time needs → Low latency
- Analytics/ETL → High throughput
- Hybrid: Batch non-critical, immediate for critical

## How to Discuss NFRs in Interviews

Interview Framework

## STEP 1: Always Ask About NFRs (First 5 minutes)

### Essential Questions:

"Before we dive into the design, let me clarify the non-functional requirements:

#### Performance:

- What's our latency requirement?
- Expected throughput (requests/sec)?
- Any specific response time SLAs?

#### Scale:

- How many users (concurrent, daily active)?
- Data volume (current and projected)?
- Geographic distribution?

#### Reliability:

- Availability requirement (99.9%, 99.99%)?
- Can we tolerate data loss?
- Disaster recovery needs?

#### Consistency:

- Strong or eventual consistency acceptable?
- Read-your-own-writes needed?
- How stale can data be?

#### Other:

- Security/compliance requirements?
- Budget constraints?
- Any specific regulatory needs?"

This shows: Professional maturity, production thinking

## How to Present NFRs

### TEMPLATE:

"Based on the requirements, here are the key NFRs I'll design for:

### PERFORMANCE:

- Latency: P95 < 200ms, P99 < 500ms
- Throughput: 50K requests/sec peak
- Response time: < 1s for page loads

### SCALABILITY:

- Support 10M DAU
- Scale to 100K concurrent users
- Handle 5x traffic spikes

#### AVAILABILITY:

- 99.99% uptime (52 min downtime/year)
- Multi-AZ deployment
- < 1 min failover time

#### CONSISTENCY:

- Eventual consistency acceptable
- Read-your-own-writes for user data
- Strong consistency for financial transactions

#### SECURITY:

- Authentication required
- Data encrypted in transit and at rest
- PCI-DSS compliant for payments

These NFRs will guide my architecture decisions..."

Then explain HOW each NFR affects design:

"For 99.99% availability, I'll deploy across 3 AZs with..."

"For P99 < 500ms, I'll implement caching with..."

## Common Interview Scenarios

### Scenario 1: URL Shortener

#### FUNCTIONAL:

- Shorten long URLs
- Redirect to original URL
- Track click analytics

#### NON-FUNCTIONAL REQUIREMENTS TO DISCUSS:

Q: "What's the expected throughput?"

A: "100M new URLs/month, 10B redirects/month"

#### Calculation:

- Writes:  $100M / (30 * 86400) = \sim 40$  writes/sec
- Reads:  $10B / (30 * 86400) = \sim 4K$  reads/sec
- Read:Write ratio = 100:1 (read-heavy!)

#### Architecture Implications:

- Use caching aggressively (Redis)
- CDN for redirect endpoint
- Read replicas for database
- Cache hit ratio target: 95%

Q: "What about latency?"

A: "Redirects should be < 100ms"

#### Architecture Implications:

- Cache popular URLs (80/20 rule)
- CDN reduces network latency
- Database indexes on short\_url
- Monitor cache hit ratio

Q: "Availability requirements?"

A: "99.9% available (43 min downtime/month)"

#### Architecture Implications:

- Multi-AZ deployment
- Database replication
- Health checks and auto-failover
- No single points of failure

Q: "Consistency needs?"

A: "Eventual consistency OK for analytics,  
Strong consistency for URL creation"

#### Architecture Implications:

- Async analytics processing
- ACID database for URL storage
- Separate read/write paths

## Scenario 2: Twitter-like Feed

### NON-FUNCTIONAL REQUIREMENTS:

#### SCALE:

- 500M DAU
- 100M concurrent users
- 500M tweets/day
- 5B timeline reads/day

#### Calculations:

- Writes:  $500M / 86400 = 5,787$  tweets/sec
- Reads:  $5B / 86400 = 57,870$  reads/sec
- Peak: 3x = 17K writes/sec, 174K reads/sec

#### LATENCY:

- Feed load: < 1s
- Tweet post: < 500ms
- Notifications: < 2s

#### AVAILABILITY:

- 99.99% uptime
- No data loss for tweets
- Graceful degradation for non-critical features

#### CONSISTENCY:

- Eventual consistency for feed (stale < 5s OK)
- Strong consistency for tweet storage
- Eventually consistent follower counts

#### Architecture Decisions Based on NFRs:

##### For High Read Throughput:

- Redis cache for timelines (hot data)
- CDN for media assets
- Read replicas (10+)
- Fanout-on-write for active users
- Fanout-on-read for inactive users

##### For Low Latency:

- Pre-computed timelines in cache
- Geographic distribution
- Async processing for analytics
- CDN for static assets

##### For High Availability:

- Multi-region active-active
- Database replication across regions
- Cache replication
- Graceful degradation (show cached data if DB down)

##### For Eventual Consistency:

- Async fanout to followers
- Timeline updates propagate within 5s
- Follower counts updated async
- Strong consistency for tweet storage only

## Measuring & Monitoring

### Key Metrics to Track

#### 1. Golden Signals (Google SRE):

##### Latency:

- Request duration
- P50, P95, P99, P99.9
- By endpoint/service
- Historical trends

##### Traffic:

- Requests per second
- Bandwidth usage
- Active connections
- Request patterns

#### Errors:

- Error rate (%)
- Error types
- Failed requests
- 4xx vs 5xx errors

#### Saturation:

- CPU usage
- Memory usage
- Disk I/O
- Network bandwidth
- Queue depth

## 2. Business Metrics:

#### User Engagement:

- Daily/Monthly active users
- Session duration
- Feature usage
- Conversion rates

#### System Health:

- Availability (uptime %)
- Success rate
- Cache hit ratio
- Database connection pool

#### Financial:

- Infrastructure costs
- Cost per request
- Revenue per user
- ROI metrics

## Monitoring Strategy

#### Levels of Monitoring:

##### Level 1: Infrastructure

- Server health
- Network connectivity
- Disk space
- Resource utilization

##### Level 2: Application

- Request/response metrics
- Error rates
- API performance
- Queue depths

### Level 3: Business

- User actions
- Revenue metrics
- Feature adoption
- Customer satisfaction

### Alert Priorities:

#### P0 (Critical):

- System down
- Data loss
- Security breach

Response: Immediate (page on-call)

#### P1 (High):

- Degraded performance
- Error rate spike
- SLA breach

Response: Within 15 minutes

#### P2 (Medium):

- Resource warnings
- Slow queries
- Cache misses

Response: Within 1 hour

#### P3 (Low):

- Optimization opportunities
- Long-term trends
- Capacity planning

Response: Next business day

## Interview Tips

### How to Ace NFR Discussion

#### DEMONSTRATE PRODUCTION EXPERIENCE:

##### 1. Start with NFRs

"Before designing, I need to understand:  
- What's our latency requirement?  
- Expected scale?  
- Availability needs?"

##### 2. Use Specific Numbers

X "System should be fast"  
✓ "P99 latency < 200ms"  
  
X "Handle lots of users"

- "Scale to 1M concurrent, 10M DAU"
- "Highly available"
- "99.99% uptime, multi-AZ deployment"

### 3. Explain Trade-offs

"For 99.99% availability, I'll deploy across 3 AZs, which increases cost by 50% but provides redundancy. Alternatively, 99.9% with 2 AZs would be cheaper but allows more downtime."

### 4. Calculate Capacity

"With 10M DAU doing 50 actions each:

- Average: 5,787 RPS
- Peak (3x): 17,361 RPS
- Servers needed (1K RPS each): 20
- With buffer (1.5x): 30 servers"

### 5. Discuss Monitoring

"I'll monitor:

- Latency (P50, P95, P99)
- Error rate
- Throughput
- Resource utilization

Alert if P99 > 500ms or error rate > 1%"

## Common NFR Questions

Q: "How would you ensure 99.99% availability?"

Strong Answer:

"99.99% allows 52 minutes downtime per year. I'll achieve this through:

1. Multi-AZ deployment (3 availability zones)
2. Redundant components (no single point of failure)
3. Health checks with automatic failover (<1 min)
4. Zero-downtime deployments (rolling updates)
5. Chaos testing to verify resilience

This gives theoretical availability of 99.9995%, providing buffer for unexpected issues."

Q: "How do you handle 10x traffic spike?"

Strong Answer:

"Prepare for spikes through:

1. Auto-scaling (scale out to 10x capacity in 5 min)
2. Aggressive caching (reduces backend load 80%)
3. Rate limiting (protect system from overload)
4. Queue-based processing (smooth load)
5. Graceful degradation (disable non-critical features)
6. Pre-warming during known events (Black Friday)"

Q: "Strong consistency vs eventual consistency?"

Strong Answer:

"Depends on use case. For financial transactions like payments,

strong consistency is mandatory – can't have duplicate charges.

For social media likes, eventual consistency is fine – few seconds delay acceptable. I'd use:

– Strong: Financial, inventory, user auth

– Eventual: Feeds, likes, view counts, recommendations

Trade-off: Strong adds latency but prevents inconsistency."

## Quick Reference

### NFR Checklist for Interviews

#### ALWAYS ASK ABOUT:

- Latency requirements (ms, seconds?)
- Throughput needs (requests/sec, data/sec)
- Scale (users, data volume, geographic)
- Availability SLA (99.9%, 99.99%?)
- Consistency needs (strong, eventual?)
- Data durability (can we lose data?)
- Security requirements (auth, encryption?)
- Compliance (GDPR, HIPAA, PCI?)
- Budget constraints
- Growth projections

#### DURING DESIGN, ADDRESS:

- How you'll achieve latency targets
- Capacity calculations for throughput
- Scaling strategy (horizontal/vertical)
- High availability architecture
- Consistency model choice
- Security measures
- Monitoring and alerting
- Cost optimization

#### SHOW DEPTH WITH:

- Specific numbers (not just "fast")
- Percentiles (P95, P99, not just average)
- Calculations (capacity planning)
- Trade-off discussions
- Alternative approaches
- Real-world constraints

### Common NFR Values

## Typical Interview Numbers:

### Users:

- Small: 10K–100K users
- Medium: 1M–10M users
- Large: 100M+ users

### Throughput:

- Low: 100–1K RPS
- Medium: 10K–100K RPS
- High: 1M+ RPS

### Latency:

- Real-time: < 50ms
- Interactive: < 200ms
- Standard: < 1s
- Batch: Minutes–hours

### Availability:

- Standard: 99.9%
- High: 99.99%
- Critical: 99.999%

### Data:

- Small: GB range
- Medium: TB range
- Large: PB range

---

## Conclusion

Non-functional requirements are as important as functional requirements in system design. They drive architectural decisions and demonstrate your production experience.

### Key Principles

#### 1. Always Clarify NFRs First

- Don't assume
- Ask specific questions
- Get concrete numbers

#### 2. Use NFRs to Drive Design

- High availability → Multi-AZ
- Low latency → Caching, CDN
- High throughput → Horizontal scaling
- Strong consistency → ACID database

#### 3. Discuss Trade-offs

- Performance vs Cost
- Consistency vs Availability
- Latency vs Throughput
- Simplicity vs Features

#### 4. Be Specific

- Use numbers (99.99%, <100ms, 10K RPS)
- Use percentiles (P95, P99)
- Calculate capacity
- Estimate costs

#### 5. Show Production Thinking

- Monitoring and alerting
- Failure scenarios
- Operational concerns
- Cost awareness

### Interview Success Pattern

1. Requirements: "What are the NFRs?"
2. Calculations: "With 10M users, that's X RPS..."
3. Architecture: "To achieve 99.99% availability, I'll..."
4. Trade-offs: "We could also do Y, but the trade-off is..."
5. Monitoring: "I'll track these metrics..."

This demonstrates: Systematic thinking, production experience, depth of knowledge, ability to make engineering trade-offs

### Remember

**NFRs are not afterthoughts** - they're fundamental to system design. Strong candidates:

- Ask about NFRs upfront
- Use them to justify decisions
- Discuss trade-offs explicitly
- Show production awareness
- Demonstrate depth with numbers

Good luck with your interviews!

## Additional Resources

### Books

- "Site Reliability Engineering" by Google

- "Designing Data-Intensive Applications" by Martin Kleppmann
- "The Art of Scalability" by Martin L. Abbott

## Online Resources

- Google SRE Book (free online)
- AWS Well-Architected Framework
- Azure Architecture Best Practices
- High Scalability Blog

## Key Concepts to Study

- CAP Theorem
- PACELC Theorem
- SLA/SLO/SLI
- The Four Golden Signals
- Percentile-based monitoring
- Cost optimization strategies