

System Design Terminology Guide - Essential Concepts for Scale

Table of Contents

1. [Introduction](#)
 2. [Scalability Concepts](#)
 3. [Reliability & Availability](#)
 4. [Performance Concepts](#)
 5. [Database Concepts](#)
 6. [Caching Strategies](#)
 7. [Distributed Systems](#)
 8. [Load Balancing](#)
 9. [Data Partitioning](#)
 10. [Consistency Models](#)
 11. [Architecture Patterns](#)
 12. [Quick Reference](#)
-

Introduction

"Knowing the vocabulary is half the battle in system design interviews."

This guide covers every important term you need to know to discuss systems at scale. For each term:

- Clear definition
 - When to use it
 - Real-world examples
 - Interview context
-

Scalability Concepts

1. Scalability

Definition: A system's ability to handle increased load by adding resources

Two Types:

Vertical Scaling (Scale Up):

- Add more power to existing machine
 - Bigger CPU, more RAM, faster disk
 - Example: Upgrade from 8GB to 32GB RAM
 - **Limit:** Single machine ceiling (~1TB RAM)
 - **Pros:** Simple, no code changes
 - **Cons:** Expensive, single point of failure
-

Horizontal Scaling (Scale Out):

- Add more machines
- Distribute load across servers
- Example: 1 server → 10 servers → 100 servers
- **Limit:** Nearly unlimited
- **Pros:** Cost-effective, fault-tolerant
- **Cons:** Complex, requires distributed architecture

Interview Usage:

"We'd start with vertical scaling for simplicity, then move to horizontal scaling when traffic exceeds 10K QPS"

2. Load Balancer

Definition: Distributes incoming traffic across multiple servers

Types:

Layer 4 (Transport Layer):

- Routes based on IP address and port
- Fast, simple
- No visibility into application data
- Example: TCP/UDP load balancing

Layer 7 (Application Layer):

- Routes based on HTTP headers, URLs, cookies
- Can route /api/users to user service, /api/posts to post service
- Slower but more flexible
- Example: NGINX, HAProxy, AWS ALB

Algorithms:

- **Round Robin:** Distribute sequentially (server1, server2, server3, server1...)
- **Least Connections:** Route to server with fewest active connections
- **Least Response Time:** Route to fastest responding server
- **IP Hash:** Same client always routes to same server (sticky sessions)

Real Examples:

- **Netflix:** AWS ELB distributing traffic across thousands of instances
- **Twitter:** HAProxy for internal service routing

Interview Usage:

"I'd use Layer 7 load balancer to route API requests to appropriate microservices and handle SSL termination"

3. CDN (Content Delivery Network)

Definition: Globally distributed network of servers that cache content close to users

How It Works:

User in Tokyo → Requests image
→ CDN edge server in Tokyo (cache hit)
→ Returns image in 10ms

vs. without CDN:

User in Tokyo → Origin server in California
→ Returns image in 150ms (15x slower!)

Types:

Push CDN:

- You upload content to CDN manually
- Good for: Infrequently changing content
- Example: Uploading product catalog

Pull CDN:

- CDN pulls content from origin when first requested
- Good for: Frequently changing content
- Example: Social media posts

Real Examples:

- **Netflix:** Uses AWS CloudFront for streaming
- **Instagram:** Akamai for photo delivery
- **Twitter:** Fastly for image caching

Interview Usage:

"For media-heavy applications like Instagram, CDN is mandatory, not optional. It reduces origin bandwidth by 90%+ and improves latency from 150ms to 10ms"

4. Caching

Definition: Store frequently accessed data in fast storage (RAM) to avoid slow disk/database reads

Cache Levels:

Browser Cache → CDN Cache → Server Cache → Database Cache
(client-side) (edge) (application) (built-in)

Cache Strategies:

Cache-Aside (Lazy Loading):

1. Check cache
2. If miss → Query database → Store in cache
3. If hit → Return immediately

Use for: Read-heavy workloads

Example: User profiles, product catalog

Write-Through:

1. Write to cache
2. Cache synchronously writes to database
3. Return

Use for: Strong consistency needed

Example: Financial transactions

Write-Behind (Write-Back):

1. Write to cache
2. Return immediately
3. Cache asynchronously writes to database later

Use for: High write throughput

Example: Analytics events, logs

Cache Eviction Policies:

- **LRU** (Least Recently Used): Remove oldest accessed item
- **LFU** (Least Frequently Used): Remove least accessed item
- **FIFO** (First In First Out): Remove oldest item
- **TTL** (Time To Live): Auto-expire after time period

Real Examples:

- **Facebook:** Memcached for caching everything (photos, profiles, etc.)
- **Twitter:** Redis for timeline caching
- **Reddit:** Uses Memcached extensively

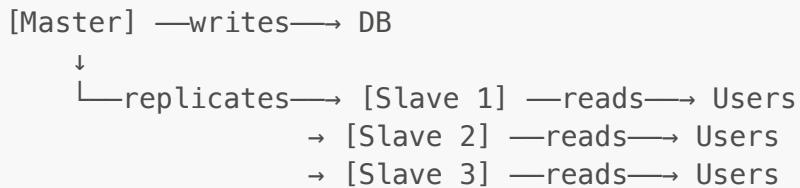
Interview Usage:

"I'd implement cache-aside pattern with Redis. This would reduce database load by 80% by caching hot user profiles and timelines"

5. Database Replication

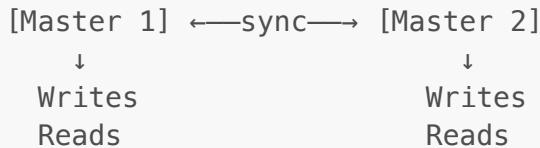
Definition: Copy data across multiple database servers

Master-Slave Replication:



Pros: Scales reads (add more slaves)
Cons: Doesn't scale writes, replication lag

Master-Master Replication:



Pros: Scales reads AND writes, high availability
Cons: Conflict resolution needed, complex

Replication Lag:

- Time for data to propagate from master to slave
- Typically milliseconds to seconds
- Can cause stale reads

Real Examples:

- **Instagram:** 12 PostgreSQL replicas in different zones
- **GitHub:** MySQL with multiple read replicas

Interview Usage:

"For read-heavy workload, I'd use master-slave replication with 5 read replicas to handle 50K read QPS"

6. Sharding (Horizontal Partitioning)

Definition: Split data across multiple databases, each holding a subset

Sharding Strategies:

Range-Based Sharding:

Users 1-1M	→ Shard 1
Users 1M-2M	→ Shard 2
Users 2M-3M	→ Shard 3

Pros: Simple to implement
Cons: Uneven distribution (hotspots)

Hash-Based Sharding:

```
Shard = hash(user_id) % num_shards

User 12345 → hash(12345) % 4 = 1 → Shard 1
User 67890 → hash(67890) % 4 = 2 → Shard 2

Pros: Even distribution
Cons: Hard to add/remove shards
```

Geographic Sharding:

```
US users → US shard
EU users → EU shard
Asia users → Asia shard

Pros: Low latency (data close to users)
Cons: Uneven distribution
```

Consistent Hashing:

```
Improves on hash-based sharding
When adding/removing nodes, only K/n keys need to be remapped
(K = total keys, n = number of nodes)

Used by: Cassandra, DynamoDB, Redis Cluster
```

Trade-offs:

- ✓ Scales both reads and writes
- ✗ Can't JOIN across shards
- ✗ Complex queries become difficult
- ✗ Rebalancing is expensive

Real Examples:

- **Instagram:** Shards PostgreSQL by user_id

- **Twitter**: Shards by tweet_id
- **Uber**: Geographic sharding for low latency

Interview Usage:

"When single PostgreSQL instance can't handle 10K+ write QPS, I'd shard by user_id using consistent hashing to distribute load evenly"

7. Federation (Vertical Partitioning)

Definition: Split database by function, not by rows

Example:

Single DB:

[Database] → Users, Products, Orders, Reviews

Federated:

[Users DB] → User tables only
[Products DB] → Product tables only
[Orders DB] → Order tables only

Pros:

- Smaller databases (fit in memory)
- Better cache locality
- Parallel writes

Cons:

- Can't JOIN across databases
- Application complexity

When to Use:

- Before sharding (simpler)
- Clear functional boundaries
- Different access patterns per function

Real Examples:

- **Amazon**: Separate databases for users, orders, inventory
- **Netflix**: Separate DBs for user data, viewing history, recommendations

Interview Usage:

"As first scaling step, I'd federate into Users DB and Content DB before implementing sharding"

Reliability & Availability

8. High Availability (HA)

Definition: System remains operational despite failures

Measured in "Nines":

99%	(2 nines) = 3.65 days downtime/year
99.9%	(3 nines) = 8.76 hours downtime/year
99.99%	(4 nines) = 52.6 minutes downtime/year
99.999%	(5 nines) = 5.26 minutes downtime/year
99.9999%	(6 nines) = 31.5 seconds downtime/year

Most companies target: 99.9% – 99.99%

Achieving HA:

- **Redundancy:** Multiple copies of everything
- **Failover:** Automatic switch to backup
- **Load Balancing:** Distribute across multiple servers
- **Health Checks:** Detect and route around failures
- **Circuit Breaker:** Fail fast instead of hanging

Real Examples:

- **AWS:** 99.99% SLA for most services
- **Google:** 99.95% SLA for Search
- **Netflix:** Designed for failure (Chaos Monkey)

Interview Usage:

"To achieve 99.99% availability (4 nines), I'd deploy across 3 availability zones with automatic failover"

9. Fault Tolerance

Definition: System continues operating even when components fail

Techniques:

Redundancy:

- N+1: One extra backup
- N+2: Two extra backups
- Example: 10 servers needed + 2 backup = 12 servers

Replication:

- Data copied to multiple nodes
- Typically 3 replicas
- Example: Cassandra writes to 3 nodes

Graceful Degradation:

- Core features work even if some components fail
- Example: Twitter timeline works even if trending topics fail

Real Examples:

- **Netflix**: If recommendation engine fails, show popular content
- **Amazon**: If product reviews fail, still show product page

Interview Usage:

"For fault tolerance, I'd replicate data across 3 nodes and implement graceful degradation where search failure doesn't break core features"

10. CAP Theorem

Definition: In distributed system, you can only have 2 of 3:

C - Consistency: All nodes see same data simultaneously

A - Availability: Every request gets a response (success or failure)

P - Partition Tolerance: System works despite network failures

Since network partitions WILL happen, you must choose:

CP (Consistency + Partition Tolerance):

- Sacrifice availability during partition
- Example: **MongoDB**, PostgreSQL, Redis
- Use for: Banking, inventory (can't show wrong balance)

AP (Availability + Partition Tolerance):

- Sacrifice consistency during partition
- Example: **Cassandra**, DynamoDB, Couchbase
- Use for: Social feeds, analytics (stale data acceptable)

Real Examples:

- **Bank Account**: CP (must show correct balance)
- **Twitter Feed**: AP (stale tweets acceptable)
- **Shopping Cart**: AP (eventual consistency OK)

Interview Usage:

"For user authentication, I'd choose CP (strong consistency) because showing wrong account balance is unacceptable. For the social feed, AP is fine since eventual consistency is acceptable"

11. Consistency Models

Definition: How quickly changes propagate in distributed system

Strong Consistency:

Write completes → All subsequent reads see the write
Example: Bank transactions

Timeline:
T1: Write($x = 5$)
T2: Read() → Returns 5 (guaranteed)

Eventual Consistency:

Write completes → Reads EVENTUALLY see the write (may take time)
Example: DNS updates, social media feeds

Timeline:
T1: Write($x = 5$)
T2: Read() → Might return old value (4)
T3: Read() → Returns 5 (eventually)

Causal Consistency:

If operation A causes operation B, everyone sees A before B
Example: Reply always appears after original post

Timeline:
T1: Post created
T2: Reply to post
T3: All users see post before reply (causality preserved)

Real Examples:

- **Strong:** PostgreSQL, MySQL, DynamoDB (optional)
- **Eventual:** Cassandra, DynamoDB (default), DNS
- **Causal:** Some distributed databases

Interview Usage:

"User profiles need strong consistency, but for the news feed, eventual consistency is acceptable since users won't notice if a post appears 1-2 seconds later"

Performance Concepts

12. Latency

Definition: Time taken to complete an operation

P50, P95, P99 Percentiles:

P50 (Median): 50% of requests faster than this

P95: 95% of requests faster than this

P99: 99% of requests faster than this

Example:

If P99 = 100ms, then 99% of requests complete in < 100ms

1% of requests take > 100ms (tail latency)

Why P99 matters:

- User with 100 parallel requests
- Probability all succeed in P50 time: $0.5^{100} \approx 0\%$ (practically 0)
- At least one hits P99: Very high probability

Target Latencies:

Cache read: < 1ms

Database read: < 10ms

API response: < 100ms

Page load: < 1 second

Video start: < 2 seconds

Interview Usage:

"Target p99 latency of 100ms for API responses, which means 99% of users get response in < 100ms"

13. Throughput

Definition: Number of operations per unit time (usually QPS)

vs. Latency:

- **Latency:** How long ONE request takes
- **Throughput:** How many requests in given time

Example:

System A: 10ms latency, 1,000 QPS throughput

System B: 20ms latency, 10,000 QPS throughput

System B is slower per request but handles more requests!

Interview Usage:

"System needs to handle 100K QPS throughput with p99 latency under 200ms"

14. Bottleneck

Definition: Component that limits overall system performance

Common Bottlenecks:

- **Database:** Too many queries, slow disk I/O
- **Network:** Bandwidth limit reached
- **CPU:** Complex computations, encryption
- **Memory:** Cache full, OOM errors
- **Disk I/O:** Database writes, log writes

How to Identify:

- Monitor CPU, memory, disk, network
- Find component at 80-100% utilization
- That's your bottleneck!

Interview Usage:

"At 10K write QPS, the database becomes the bottleneck. I'd add write-optimized database like Cassandra"

Database Concepts

15. ACID

Definition: Properties of reliable database transactions

A - Atomicity: Transaction is all-or-nothing

Transfer \$100:

1. Deduct \$100 from Account A
2. Add \$100 to Account B

If step 2 fails, step 1 rolls back
Either both happen or neither happens

C - Consistency: Database remains in valid state

If Account A - \$100, then total money in system unchanged
Constraints always satisfied

I - Isolation: Concurrent transactions don't interfere

Transaction 1: Read balance, transfer money
Transaction 2: Read balance, transfer money

Results same as if run sequentially

D - Durability: Committed data persists

After transaction commits, data survives crashes
Written to disk, not just memory

Databases with ACID:

- PostgreSQL, MySQL, Oracle, SQL Server

Interview Usage:

"For payment processing, I need ACID guarantees, so I'd use PostgreSQL"

16. BASE

Definition: Properties of NoSQL databases (alternative to ACID)

BA - Basically Available: System always responds (might be stale)

S - Soft State: State may change without input (replication happening)

E - Eventual Consistency: Eventually all nodes will be consistent

Example:

You update your Twitter bio:

- Write succeeds immediately (Basically Available)
- Some replicas still show old bio (Soft State)
- After 1 second, all replicas show new bio (Eventual Consistency)

Databases with BASE:

- Cassandra, DynamoDB, MongoDB (default), Couchbase

Interview Usage:

"For social media feed, BASE model is acceptable since users can tolerate seeing slightly stale data"

17. Normalization vs Denormalization

Normalization: Organize data to reduce redundancy

Users Table:

user_id	name
1	Alice

2 | Bob

Orders Table:

order_id	user_id	product
100	1	Laptop
101	1	Mouse

No data duplication

Need JOIN to get user name with order

Denormalization: Duplicate data to avoid JOINS

Orders Table:

order_id	user_id	user_name	product
100	1	Alice	Laptop
101	1	Alice	Mouse

Data duplicated (Alice appears twice)

No JOIN needed, faster reads

Trade-offs:

- **Normalized:** Slower reads (JOINS), easier writes, no redundancy
- **Denormalized:** Faster reads, complex writes (update many places), redundancy

Real Examples:

- **Facebook:** Denormalizes for timeline performance
- **Twitter:** Denormalizes user data in tweets

Interview Usage:

"I'd denormalize by storing user name with each tweet to avoid JOINS when loading timeline, trading write complexity for read speed"

18. Indexing

Definition: Data structure that improves query speed

How It Works:

Without Index:

SELECT * FROM users WHERE email = 'alice@example.com'
→ Scans entire table (O(n))
→ 1 billion users = very slow!

With Index on email:

→ B-tree lookup (O(log n))

- 1 billion users = ~30 comparisons
- Sub-millisecond query

Types:

- **Primary Index:** On primary key (automatic)
- **Secondary Index:** On other columns
- **Composite Index:** On multiple columns
- **Unique Index:** Enforces uniqueness

Trade-offs:

- Faster reads (10-1000x)
- Slower writes (must update index)
- More storage (index size)

Interview Usage:

"I'd create index on (user_id, created_at) for efficient timeline queries sorted by time"

Distributed Systems

19. Distributed System

Definition: Multiple computers working together as single system

Challenges:

- **Network is unreliable:** Packets drop, partitions happen
- **Latency:** Cross-datacenter calls take 50-150ms
- **No global clock:** Can't know exact order of events
- **Partial failures:** Some nodes work, some don't

Benefits:

- **Scalability:** Add more machines
- **Fault Tolerance:** One machine fails, others continue
- **Geographic Distribution:** Serve users globally

Real Examples:

- **Google:** Thousands of datacenters worldwide
- **Netflix:** Distributed across AWS regions
- **Facebook:** Multiple datacenters, cross-datacenter replication

Interview Usage:

"In distributed system, we must handle network partitions, so per CAP theorem, I'd choose between consistency and availability"

20. Consensus

Definition: Multiple nodes agreeing on a single value

Why Needed:

- Leader election (who's the master?)
- Configuration management (what's the current state?)
- Distributed locks (who has access?)

Algorithms:

- **Paxos:** Complex but proven
- **Raft:** Simpler, easier to understand
- **Zab:** Used by Zookeeper

Technologies:

- **Zookeeper:** Apache's consensus service
- **etcd:** Kubernetes uses this
- **Consul:** HashiCorp's service discovery

Interview Usage:

"For distributed counter in URL shortener, I'd use Zookeeper for consensus to ensure unique IDs"

21. Message Queue

Definition: Asynchronous communication between services

How It Works:

```
[Producer] → [Message Queue] → [Consumer]  
                           (Kafka)
```

Producer: Sends messages without waiting

Queue: Stores messages durably

Consumer: Processes messages at own pace

Use Cases:

- **Decouple services:** Producer doesn't depend on consumer
- **Handle spikes:** Queue buffers during traffic surge
- **Async processing:** Don't block user request
- **Retry logic:** Reprocess failed messages

Popular Queues:

- **Kafka:** High throughput, replay capability
- **RabbitMQ:** Flexible routing, AMQP protocol

- **Amazon SQS:** Managed, simple
- **Redis:** Simple, fast (but can lose messages)

Real Examples:

- **Uber:** Kafka for trip events, surge pricing
- **LinkedIn:** Kafka for activity streams
- **Instagram:** RabbitMQ for feed fanout

Interview Usage:

"I'd use Kafka message queue for asynchronous feed fanout, so posting a tweet doesn't block while updating millions of followers' timelines"

22. Microservices

Definition: Application as collection of small, independent services

vs. Monolith:

```
Monolith:  
[Single Application]  
├── User Module  
└── Post Module  
└── Payment Module  
└── Notification Module  
  
Microservices:  
[User Service] ↔ [Post Service]  
         ↓           ↓  
[Payment Service] ↔ [Notification Service]
```

Benefits:

- **Independent scaling:** Scale hot services
- **Independent deployment:** Deploy without affecting others
- **Technology flexibility:** Each service can use different tech
- **Team autonomy:** Teams own services

Challenges:

- **Network overhead:** HTTP calls between services
- **Distributed debugging:** Trace requests across services
- **Data consistency:** Each service has own database
- **Operational complexity:** Deploy/monitor many services

Real Examples:

- **Netflix:** 1000+ microservices

- **Uber:** 2000+ microservices
- **Amazon:** Service-oriented architecture

Interview Usage:

"I'd use microservices architecture with User Service, Post Service, and Feed Service, each scaling independently"

23. API Gateway

Definition: Single entry point for all client requests

Responsibilities:

- **Authentication:** Validate JWT tokens
- **Rate Limiting:** Prevent abuse
- **Request Routing:** Route to appropriate microservice
- **Response Aggregation:** Combine multiple service calls
- **Protocol Translation:** REST to gRPC
- **SSL Termination:** Handle HTTPS

Technologies:

- Kong, AWS API Gateway, Apigee, Express Gateway

Interview Usage:

"API Gateway handles authentication and rate limiting before routing requests to microservices"

Load Balancing

24. Sticky Sessions (Session Affinity)

Definition: Route same user to same server

How It Works:

```
User Alice → Server 1 (first request)
User Alice → Server 1 (all subsequent requests)
User Bob → Server 2 (first request)
User Bob → Server 2 (all subsequent requests)
```

When Needed:

- Server stores session in memory
- WebSocket connections
- File uploads (maintain state)

Problems:

- Uneven distribution (hot users)
- If server dies, lose session
- Better: Store session in Redis (shared)

Interview Usage:

"I'd avoid sticky sessions by storing sessions in Redis, making servers stateless for easier scaling"

25. Health Checks

Definition: Automatically detect and remove unhealthy servers

How It Works:

```
Load Balancer sends request every 10 seconds:  
GET /health  
  
If server responds 200 OK → Healthy  
If timeout or error → Unhealthy → Remove from pool
```

Health Endpoint:

```
@app.route('/health')
def health_check():
    # Check database connection
    if not db.ping():
        return "Unhealthy", 503

    # Check cache connection
    if not redis.ping():
        return "Unhealthy", 503

    return "Healthy", 200
```

Interview Usage:

"Load balancer performs health checks every 10 seconds, automatically removing unhealthy instances"

Caching Strategies

26. Cache Invalidation

Definition: Removing stale data from cache

Problem: "There are only two hard things in Computer Science: cache invalidation and naming things" - Phil Karlton

Strategies:

TTL (Time To Live):

```
SET user:123 "data" EX 3600 # Expire in 1 hour  
Pros: Simple, automatic  
Cons: Data might be stale before expiry
```

Write-Through Invalidation:

When data updated:
1. Update database
2. Update cache
Pros: Cache always fresh
Cons: Extra write to cache

Write-Invalidate:

When data updated:
1. Update database
2. Delete from cache
3. Next read will reload from DB
Pros: Simpler than write-through
Cons: Cache miss after every write

Interview Usage:

"I'd use TTL-based expiration for user profiles (1 hour TTL) and active invalidation for critical updates like password changes"

27. Cache Stampede (Thundering Herd)

Definition: Many requests hit database simultaneously when cache expires

Problem:

```
Popular cache entry expires  
1000 requests arrive simultaneously  
All 1000 miss cache  
All 1000 query database → Database overwhelmed!
```

Solution:

```
def get_with_lock(key):  
    value = cache.get(key)
```

```

if value:
    return value

# Acquire lock
if cache.set_nx(f"lock:{key}", "1", ex=10):
    # Only one request queries DB
    value = db.query(key)
    cache.set(key, value, ex=3600)
    cache.delete(f"lock:{key}")
    return value
else:
    # Others wait and retry
    time.sleep(0.1)
    return get_with_lock(key) # Retry

```

Interview Usage:

"To prevent cache stampede on popular items, I'd use distributed locks so only one request regenerates cache"

Data Partitioning

28. Consistent Hashing

Definition: Hash function that minimizes key redistribution when nodes added/removed

Problem with Simple Hashing:

hash(key) % N

Add/remove server → All keys need to be remapped!

Example: 4 servers → 5 servers

Only 1/5 of keys map to same server

4/5 need to move = 80% of data!

Consistent Hashing Solution:

Place servers on hash ring

Place keys on hash ring

Key goes to next server clockwise

Add/remove server → Only K/N keys remapped
(K = total keys, N = servers)

Example: 4 servers → 5 servers

Only 1/5 of keys move = 20% of data

Virtual Nodes:

Each physical server has multiple positions on ring
Better distribution, handles heterogeneous servers

Large server: 1000 virtual nodes
Small server: 100 virtual nodes

Used By:

- **Cassandra:** Data distribution
- **DynamoDB:** Partitioning
- **Redis Cluster:** Sharding
- **CDNs:** Server selection

Interview Usage:

"I'd use consistent hashing for cache distribution so adding cache servers doesn't cause massive key remapping"

29. Hot Partition (Hotspot)

Definition: One partition/shard receiving disproportionate traffic

Problem:

Celebrity user with 10M followers
All their data in one shard
Shard overloaded while others idle

Example:

Shard 1 (celebrity): 50K QPS
Shard 2 (normal users): 1K QPS
Shard 3 (normal users): 1K QPS

Solutions:

- **Partition by content, not user:** Tweet ID instead of user ID
- **Replicate hot data:** Multiple copies of celebrity data
- **Cache aggressively:** Serve hot data from cache

Real Examples:

- **Twitter:** Handles celebrity tweets differently (fan-out on read)
- **Instagram:** Caches celebrity profiles heavily

Interview Usage:

"To avoid hot partition with celebrity users, I'd shard by tweet_id not user_id, and cache celebrity data"

aggressively"

Architecture Patterns

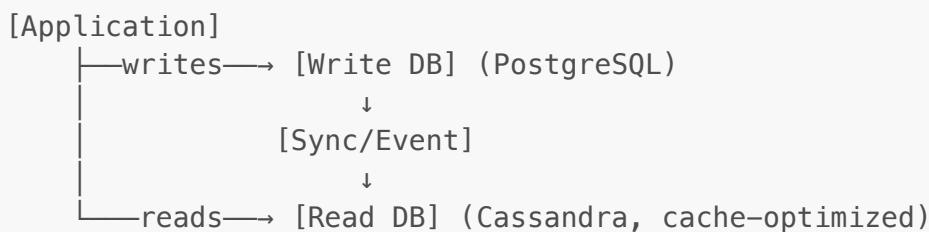
30. CQRS (Command Query Responsibility Segregation)

Definition: Separate read and write data models

Traditional:

```
[Application] ↔ [Single Database]  
(reads & writes)
```

CQRS:



Benefits:

- **Optimize separately:** Write DB for writes, Read DB for reads
- **Independent scaling:** Scale reads without affecting writes
- **Different models:** Denormalized for reads, normalized for writes

Use Cases:

- Social media (write tweets to Cassandra, read from Redis cache)
- E-commerce (write orders to PostgreSQL, read from MongoDB)
- Analytics (write events to Kafka, read from Redshift)

Real Examples:

- **Netflix:** Writes to Cassandra, reads from various caches
- **Amazon:** Order writes to Oracle, inventory reads from cache

Interview Usage:

"I'd implement CQRS with writes to PostgreSQL and reads from Redis cache for optimal performance"

31. Event Sourcing

Definition: Store all changes as sequence of events, not final state

Traditional (State-based):

```
User Account:  
balance = $100
```

```
After deposit $50:  
balance = $150 (old value lost)
```

Event Sourcing:

```
Event Log:  
1. Account created, balance = $100  
2. Deposited $50, balance = $150  
3. Withdrew $20, balance = $130
```

```
Current state = Replay all events  
Can rebuild state at any point in time
```

Benefits:

- **Audit trail:** Complete history
- **Time travel:** See state at any point
- **Debugging:** Replay events to reproduce bugs
- **Event replay:** Reprocess events with new logic

Combined with CQRS:

```
Commands → Event Store (Kafka) → Read Models  
(Multiple DBs)
```

Real Examples:

- **Banking:** Transaction history as events
- **E-commerce:** Order state changes as events
- **Git:** Commits are events!

Interview Usage:

"For audit requirements, I'd use event sourcing where all order state changes are stored as events in Kafka"

32. Circuit Breaker

Definition: Prevent cascading failures by failing fast

States:

CLOSED (Normal):
Requests pass through
If failures exceed threshold → OPEN

OPEN (Failing):
Reject requests immediately (fail fast)
After timeout → HALF_OPEN

HALF_OPEN (Testing):
Allow few requests through
If succeed → CLOSED
If fail → OPEN

Example:

```
class CircuitBreaker:  
    def __init__(self, failure_threshold=5, timeout=60):  
        self.failures = 0  
        self.threshold = failure_threshold  
        self.state = "CLOSED"  
        self.timeout = timeout  
  
    def call(self, func):  
        if self.state == "OPEN":  
            raise Exception("Circuit breaker OPEN")  
  
        try:  
            result = func()  
            self.failures = 0 # Reset on success  
            return result  
        except Exception:  
            self.failures += 1  
            if self.failures >= self.threshold:  
                self.state = "OPEN"  
                schedule_timeout(self.timeout)  
            raise
```

Benefits:

- Prevent cascading failures
- Give failing service time to recover
- Better error messages

Real Examples:

- **Netflix Hystrix**: Circuit breaker library

- **Uber:** Protects against downstream service failures

Interview Usage:

"I'd implement circuit breaker pattern so if recommendation service fails, we fail fast instead of timing out"

33. Rate Limiting

Definition: Limit number of requests per time period

Algorithms:

Token Bucket:

Bucket holds N tokens
Every request consumes 1 token
Refill R tokens per second
If no tokens → Reject request

Example:

Capacity: 100 tokens
Refill: 10 tokens/second
Allows burst of 100, then 10/second sustained

Leaky Bucket:

Requests enter bucket
Process at fixed rate
If bucket full → Reject

Example:

Process 10 requests/second, no burst allowed

Fixed Window:

100 requests per hour
Count resets at hour boundary

Problem: 100 at 12:59, 100 at 13:00 = 200 in 1 minute!

Sliding Window:

100 requests per rolling 60-minute window
Counts requests in last hour from current time

Real Examples:

- **Twitter API:** 300 requests per 15 minutes
- **GitHub API:** 5000 requests per hour
- **Stripe:** Prevents payment spam

Interview Usage:

"I'd implement token bucket algorithm in Redis for rate limiting - 100 requests per hour per user"

34. Idempotency

Definition: Operation can be applied multiple times with same result

Idempotent:

```
SET balance = 100
Call once: balance = 100
Call twice: balance = 100 (same result)

DELETE user_id = 5
Call once: User deleted
Call twice: User already deleted (same result, no error)
```

Not Idempotent:

```
INCREASE balance by 10
Call once: balance = 110
Call twice: balance = 120 (different result!)
```

Why Important:

- Network failures require retries
- Can safely retry idempotent operations
- Non-idempotent needs deduplication

Making Operations Idempotent:

```
# Use unique request ID
def process_payment(payment_id, amount):
    if redis.exists(f"processed:{payment_id}"):
        return "Already processed"

    # Process payment
    charge_card(amount)
```

```
# Mark as processed
redis.set(f"processed:{payment_id}", "1", ex=86400)
return "Success"
```

Interview Usage:

"For payment API, I'd make it idempotent using request IDs so retries don't charge customer twice"

35. Bloom Filter

Definition: Space-efficient probabilistic data structure to test set membership

How It Works:

```
Use k hash functions
Set k bits in bit array
```

Check membership:

- If all k bits are 1 → Probably in set
- If any bit is 0 → Definitely not in set

False positive: Possible (says yes when actually no)

False negative: Impossible (never says no when actually yes)

Use Cases:

- **Check if username exists** (before DB query)
- **Check if URL is malicious** (before fetching)
- **Weak password detection** (check against common passwords)
- **Cache key existence** (avoid cache misses)

Space Savings:

```
1 billion URLs
100 bytes each = 100 GB
```

```
Bloom filter (1% false positive):
1 billion URLs = ~1.2 GB (83x smaller!)
```

Real Examples:

- **Google Chrome:** Check malicious URLs
- **Medium:** Recommend articles you haven't read
- **Bitcoin:** Check if transaction exists

Interview Usage:

"I'd use Bloom filter to quickly check if shortened URL exists before querying database, saving 99% of queries for non-existent URLs"

36. Geohashing

Definition: Encode geographic coordinates into short string

How It Works:

```
Latitude, Longitude → Short string  
(37.7749, -122.4194) → "9q8yy" (San Francisco)
```

Nearby locations have similar prefixes:

```
9q8yy1 ← Close  
9q8yy2 ← Close  
9q9xyz ← Far
```

Use Cases:

- **Uber:** Find nearby drivers
- **Yelp:** Find nearby restaurants
- **Pokemon Go:** Spawn nearby Pokemon
- **Tinder:** Find matches in area

Benefits:

- Simple range query for nearby items
- Works with standard databases
- Variable precision (longer = more precise)

Interview Usage:

"For Uber driver matching, I'd use geohashing to efficiently query drivers within 5km radius using prefix matching"

37. WebSocket

Definition: Persistent bidirectional connection between client and server

vs. HTTP:

```
HTTP (Request–Response):  
Client → Request → Server  
Client ← Response ← Server  
(Connection closes)
```

```
WebSocket (Persistent):
```

Client ←→ Server
(Connection stays open, messages flow both ways)

Use Cases:

- **Chat applications:** Real-time messaging
- **Live updates:** Stock prices, sports scores
- **Collaborative editing:** Google Docs
- **Gaming:** Multiplayer games
- **Notifications:** Push notifications

Benefits:

- Low latency (no connection overhead)
- Server can push to client
- Efficient (one connection, many messages)

Challenges:

- Maintain many connections (memory)
- Load balancing with sticky sessions
- Scaling WebSocket servers

Real Examples:

- **Slack:** Chat messages
- **WhatsApp Web:** Message sync
- **Trading platforms:** Live prices

Interview Usage:

"For real-time chat, I'd use WebSocket for bidirectional communication instead of polling, reducing load by 100x"

38. Long Polling

Definition: Client requests, server holds connection until data available

How It Works:

1. Client sends request
2. Server holds connection open (30–60 seconds)
3. When data available OR timeout:
 - Server responds
 - Client immediately sends new request

vs. Short Polling:

Short Polling (wasteful):
Check every second → 3600 requests/hour
99% are "no new data"

Long Polling (efficient):
Hold connection → 60 requests/hour
Only responds when data available

Use Cases:

- Chat applications (before WebSocket)
- Notification systems
- Real-time updates

Interview Usage:

"Before WebSocket, I'd use long polling to reduce polling overhead from 3600 requests/hour to 60 requests/hour"

39. Server-Sent Events (SSE)

Definition: Server pushes updates to client over HTTP

vs. WebSocket:

WebSocket: Bidirectional
SSE: Unidirectional (server → client only)

SSE is simpler when only server needs to push

Use Cases:

- Live sports scores
- Stock price updates
- News feed updates
- Server status monitoring

Interview Usage:

"For live stock prices where client only receives updates, I'd use SSE instead of WebSocket for simplicity"

40. Reverse Proxy

Definition: Server that sits in front of web servers and forwards requests

vs. Forward Proxy:

Forward Proxy (client-side):

Client → Proxy → Internet

Hides client identity

Reverse Proxy (server-side):

Client → Proxy → Backend Servers

Hides server identity

Benefits:

- **Load balancing:** Distribute requests
- **SSL termination:** Handle HTTPS
- **Caching:** Cache responses
- **Compression:** Gzip responses
- **Security:** Hide backend servers

Technologies:

- NGINX, HAProxy, Apache

Interview Usage:

"I'd use NGINX as reverse proxy for SSL termination and response caching"

41. Fan-out

Definition: One message triggers multiple parallel operations

Example (Twitter):

User with 1M followers tweets

Fan-out on Write (Push):

Tweet posted → Write to 1M follower timelines

Fan-out on Read (Pull):

Tweet posted → Store once

Follower requests timeline → Fetch tweets from all following

Hybrid Approach:

Normal users (< 100K followers): Fan-out on write

Celebrities (> 100K followers): Fan-out on read

Real Examples:

- **Twitter**: Hybrid fan-out for timeline
- **Instagram**: Push notifications fan out to followers
- **Facebook**: News feed fan-out

Interview Usage:

"For timeline generation, I'd use hybrid fan-out: push for normal users, pull for celebrities to avoid the celebrity problem"

42. Back Pressure

Definition: Slow down producer when consumer can't keep up

Problem:

```
Producer: 10K messages/second  
Consumer: 1K messages/second  
Queue grows indefinitely → Memory exhausted
```

Solution:

```
When queue size > threshold:  
- Reject new messages (HTTP 503)  
- Client retries with exponential backoff  
- Or: Producer blocks until space available
```

Interview Usage:

"If Kafka queue backs up, I'd apply back pressure by rejecting new events and signaling clients to slow down"

43. Bulkhead Pattern

Definition: Isolate resources to prevent total system failure

Example:

```
Without Bulkhead:  
[Shared Thread Pool – 100 threads]  
Used by: API, Background Jobs, Admin
```

If background jobs consume all 100 threads →
API requests blocked!

```
With Bulkhead:  
[API Thread Pool – 70 threads]  
[Background Jobs – 20 threads]
```

[Admin – 10 threads]

Background jobs can't starve API

Real Examples:

- **Netflix:** Separate thread pools per service
- **Uber:** Isolate critical paths

Interview Usage:

"I'd use bulkhead pattern with separate thread pools for critical vs non-critical operations to prevent resource starvation"

44. Service Discovery

Definition: Automatically detect services in distributed system

Problem:

Microservices architecture:
How does Service A find Service B?
IPs change, instances scale up/down

Solution:

[Service Registry] (Consul/etc)
↓
Stores: service-name → [IP1, IP2, IP3]
↓
Services register on startup
Services query registry to find others

Technologies:

- **Consul:** HashiCorp's service discovery
- **etc:** Used by Kubernetes
- **Zookeeper:** Apache's coordinator
- **Eureka:** Netflix's service registry

Interview Usage:

"For microservices, I'd use Consul for service discovery so services can find each other dynamically"

45. Connection Pooling

Definition: Reuse database connections instead of creating new ones

Problem:

Creating new connection:

- TCP handshake
- Authentication
- Takes 10–100ms

$1000 \text{ requests/second} \times 50\text{ms} = 50 \text{ seconds of wasted time!}$

Solution:

Connection Pool:

- Pre-create 20–50 connections
- Reuse connections
- Return to pool after use

1000 requests reuse same 20 connections

Benefits:

- Faster requests (no connection overhead)
- Limit concurrent connections to database
- Handle connection failures gracefully

Interview Usage:

"I'd use connection pool with 50 connections to avoid overhead of creating new database connections on every request"

46. Monitoring & Observability

Monitoring: Collect metrics about system health

Observability: Understand system state from outputs

Three Pillars:

1. Metrics (numbers):

- Request rate (QPS)
- Error rate (%)
- Latency (p50, p95, p99)
- CPU, memory, disk usage

Tools: Prometheus, Datadog, CloudWatch

2. Logs (events):

- Application logs
- Error logs
- Access logs
- Audit logs

Tools: ELK Stack (Elasticsearch, Logstash, Kibana)

3. Traces (request flow):

[Request] → [API Gateway] → [User Service] → [Database]
5ms 10ms 20ms 30ms

Total: 65ms latency
Bottleneck: Database (30ms)

Tools: Jaeger, Zipkin, AWS X-Ray

Interview Usage:

"I'd use Prometheus for metrics, ELK for logs, and Jaeger for distributed tracing to identify bottlenecks"

47. Chaos Engineering

Definition: Deliberately inject failures to test system resilience

Netflix's Chaos Monkey:

Randomly terminates production instances
Forces engineers to design for failure
Reveals weaknesses before customers do

Types of Chaos:

- **Chaos Monkey:** Kill random instances
- **Chaos Gorilla:** Take down entire availability zone
- **Latency Monkey:** Inject artificial delays
- **Chaos Kong:** Take down entire region

Interview Usage:

"Netflix's Chaos Engineering approach ensures system handles failures gracefully by regularly testing fault tolerance in production"

48. Blue-Green Deployment

Definition: Zero-downtime deployment strategy

How It Works:

```
Blue (Current): v1.0 serving 100% traffic
Green (New):     v2.0 deployed, serving 0% traffic
```

```
Test Green → Switch traffic → Green serves 100%
If problems → Switch back to Blue instantly
```

Benefits:

- Zero downtime
- Instant rollback
- Test in production before full deployment

Interview Usage:

"I'd use blue-green deployment so we can deploy new version without downtime and instantly rollback if issues arise"

49. Canary Deployment

Definition: Gradual rollout to subset of users

How It Works:

```
Step 1: Deploy to 5% of servers
Step 2: Monitor metrics
Step 3: If good → 25% → 50% → 100%
Step 4: If bad → Rollback 5%
```

Benefits:

- Catch bugs early
- Limited blast radius
- Real user testing

Interview Usage:

"I'd use canary deployment, rolling out to 5% of users first to catch issues before affecting everyone"

50. Data Lake

Definition: Centralized repository for raw data (structured + unstructured)

vs. Data Warehouse:

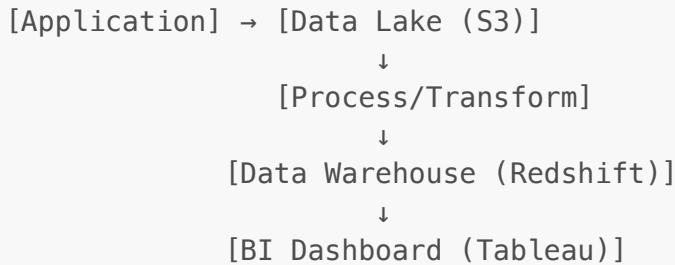
Data Lake:

- Raw, unprocessed data
- Schema-on-read
- Cheap storage (S3)
- For: Data science, exploration

Data Warehouse:

- Processed, structured data
- Schema-on-write
- Expensive (Redshift)
- For: Business intelligence

Example Flow:



Interview Usage:

"I'd store raw events in S3 data lake, then process into Redshift warehouse for analytics"

Quick Reference

Most Important Terms for Interviews

Must Know (Use in every interview):

1. Horizontal vs Vertical Scaling
2. Load Balancer (Layer 4 vs Layer 7)
3. Caching (Cache-aside pattern)
4. Database Replication (Master-Slave)
5. Sharding
6. CAP Theorem (CP vs AP)
7. CDN
8. Message Queue (Kafka)
9. Microservices
10. Rate Limiting

Should Know (Mention when relevant):

11. Consistent Hashing

12. CQRS
13. Event Sourcing
14. Circuit Breaker
15. WebSocket
16. Idempotency
17. Fan-out (Push vs Pull)
18. Bloom Filter
19. ACID vs BASE
20. Denormalization

Nice to Know (For senior roles):

21. Back Pressure
 22. Bulkhead Pattern
 23. Chaos Engineering
 24. Geohashing
 25. Service Discovery
-

Interview Cheat Sheet

When to Mention Each Term

"Design Twitter":

- Sharding (by tweet_id)
- Fan-out (hybrid push-pull)
- Redis Cache (timelines)
- Cassandra (write-heavy tweets)
- Message Queue (async processing)
- CDN (media delivery)

"Design Instagram":

- S3 + CDN (photo storage)
- Cassandra (photo metadata)
- Redis Cache (feed)
- Fan-out (posts to followers)
- Geohashing (location tags)

"Design Uber":

- Geohashing (driver matching)
- WebSocket (real-time location)
- ACID (payment transactions)
- Redis Geo (active drivers)
- Sharding (geographic)

"Design URL Shortener":

- Consistent Hashing (distributed counter)
- Base62 Encoding (short codes)
- Redis Cache (hot URLs)
- Rate Limiting (prevent abuse)
- Bloom Filter (check existence)

"Design WhatsApp":

- WebSocket (real-time messaging)
 - End-to-end encryption
 - Cassandra (message storage)
 - Push Notifications
 - CAP Theorem (AP for availability)
-

Terminology Combinations

Powerful Phrases for Interviews

"To achieve **high availability** with **fault tolerance**, I'd use **master-slave replication** across **3 availability zones** with **automatic failover**"

"For **scalability**, I'd start with **vertical scaling**, then move to **horizontal scaling** with **load balancer** and **database sharding** when traffic exceeds **10K QPS**"

"To handle **read-heavy workload**, I'd implement **cache-aside** pattern with **Redis**, achieving **80% cache hit ratio** and reducing **database load** by **5x**"

"For **write-heavy tweets (500M per day)**, I'd use **Cassandra** with **consistent hashing** for **linear scalability** and **geographic distribution**"

"To prevent **cascading failures**, I'd implement **circuit breaker** pattern and **bulkhead isolation** with **graceful degradation** for non-critical features"

Common Interview Questions & Terms to Use

Q: "How would you scale this system?"

Answer Template:

"I'd scale through:

1. Horizontal Scaling: Add more application servers behind load balancer
2. Database Replication: Master-slave with 5 read replicas
3. Caching: Redis with 80-20 rule, achieving 80% cache hit ratio
4. CDN: For static assets and media files
5. Sharding: When single DB can't handle writes (>10K QPS)
6. Message Queue: For async processing (Kafka)

7. Microservices: Split into independently scalable services
"

Q: "How would you ensure high availability?"

Answer Template:

"I'd ensure HA through:

1. Multi-Region Deployment: Deploy across 3 regions
 2. Redundancy: 3 replicas of all data
 3. Load Balancing: Active-active with health checks
 4. Failover: Automatic with < 30 second RTO
 5. Circuit Breaker: Prevent cascading failures
 6. Graceful Degradation: Core features work even if some components fail
 7. Monitoring: Prometheus + PagerDuty for alerts
- "

Q: "How would you improve performance?"

Answer Template:

"I'd improve performance by:

1. Caching: Redis for hot data (< 1ms latency)
 2. CDN: For media files (10x faster)
 3. Indexing: Database indexes for common queries
 4. Denormalization: Avoid expensive JOINs
 5. Connection Pooling: Reuse database connections
 6. Async Processing: Message queue for non-blocking operations
 7. Compression: Gzip API responses (70% smaller)
- "

Final Checklist

Before finishing any system design answer, mention:

- Scalability:** How system scales (horizontal, sharding, caching)
 - Availability:** How to achieve 99.99% (replication, failover)
 - Performance:** Latency targets (< 100ms p99)
 - Consistency:** Strong vs eventual (CAP theorem choice)
 - Monitoring:** How to detect issues (metrics, logs, traces)
 - Security:** Rate limiting, authentication, encryption
-

Summary: The Power Words

Use these terms to sound knowledgeable:

Scalability: Horizontal scaling, sharding, load balancer, CDN

Performance: Caching, indexing, denormalization, CDN

Reliability: Replication, failover, circuit breaker, graceful degradation

Consistency: ACID, BASE, CAP theorem, eventual consistency

Architecture: Microservices, CQRS, event sourcing, API gateway

Data: Sharding, consistent hashing, replication, federation

Document Version: 1.0

Last Updated: January 8, 2025

For: System Design Interview Preparation

Status: Complete & Interview-Ready 